

DATA_MODEL

Data Model

Revision: 1 Last modified: 2026-06-06T00:00:00Z

qBittorrent-Fixed has **no relational database**. All persistent state lives in qBittorrent itself (handled by the upstream image) and in a handful of JSON files under `config/`. The runtime data model is therefore a set of Python dataclasses (for internal state) and Pydantic models (for API request / response bodies).

This document enumerates them, their fields, and their lifecycle. Every class below was verified against the source tree — if you edit the model, update this document in the same PR.

Overview

```
classDiagram
    class SearchMetadata {
        +str search_id
        +str query
        +str category
        +datetime started_at
        +datetime completed_at
        +int total_results
        +int merged_results
        +list trackers_searched
        +list errors
        +str status
    }
    class SearchResult {
        +str name
        +str link
        +str size
        +int seeds
        +int leechers
        +str engine_url
    }
```

```

        +str desc_link
        +str pub_date
        +str tracker
        +str category
        +bool freeleech
    }
    class MergedResult {
        +CanonicalIdentity canonical_identity
        +list original_results
        +int total_seeds
        +int total_leechers
        +QualityTier best_quality
        +list download_urls
        +datetime created_at
    }
    class CanonicalIdentity {
        +str infohash
        +str title
        +int year
        +ContentType content_type
        +int season
        +int episode
        +str resolution
        +str codec
        +str group
        +str metadata_source
    }
    class TrackerSource {
        +str name
        +str url
        +bool enabled
        +datetime last_checked
        +str health_status
        +str scrape_url
    }
    class HookEvent {
        +HookEventType event_type
        +datetime timestamp
        +str search_id
        +str download_id
        +dict data
    }
    class HookConfig {
        +str name
        +HookEventType event
        +str script_path
        +bool enabled
        +int timeout
        +dict environment
    }
    class ScheduledSearch {
        +str id

```

```

+str name
+str query
+str category
+int interval_minutes
+bool enabled
+datetime last_run
+datetime next_run
+ScheduleStatus status
+int results_count
+str error_message
}
MergedResult --> CanonicalIdentity
MergedResult --> SearchResult : original_results
SearchMetadata --> TrackerSource : trackers_searched (names)

```

Dataclasses (internal state)

All dataclasses live in `download-proxy/src/merge_service/`.

SearchMetadata

File: `search.py:150`

Lifecycle: created in `SearchOrchestrator.search()`, stored in `_active_searches[search_id]`, returned from the API once the status is completed or failed.

Field	Type	Meaning
<code>search_id</code>	<code>str</code>	UUID generated at search start
<code>query</code>	<code>str</code>	Original search query
<code>category</code>	<code>str</code>	Category filter; default "all"
<code>started_at</code>	<code>datetime</code>	UTC, auto-populated
<code>completed_at</code>	<code>Optional[datetime]</code>	UTC, set when status leaves "running"
<code>total_results</code>	<code>int</code>	Sum of raw per-tracker counts
<code>merged_results</code>	<code>int</code>	Count after deduplication
<code>trackers_searched</code>	<code>List[str]</code>	Tracker names that were attempted

Field	Type	Meaning
errors	List[str]	"<tracker>: <message>" strings
status	str	"running", "completed", "failed"
tracker_stats	Dict[str, TrackerSearchStat]	Per-tracker run-time diagnostics; serialised as a sorted list in to_dict()

TrackerSearchStat

File: search.py:~150

Lifecycle: one instance per enabled tracker, created synchronously in `SearchOrchestrator.start_search()` (so the first HTTP response already lists every tracker that will be hit) and mutated in place by

`_run_search()` as the tracker transitions

pending → running → success | empty | error | timeout | cancelled.

Exposed through `SearchResponse.tracker_stats` on all three search endpoints and streamed via SSE as `tracker_started`

(pending→running) and `tracker_completed` (any terminal flip) events.

Field	Type	Meaning
name	str	Tracker identifier (e.g. "rutracker")
tracker_url	str	Base URL of the tracker
status	str	"pending", "running", "success", "empty", "error", "timeout", "cancelled"
results_count	int	Raw row count for this tracker's run
started_at	Optional[datetime]	UTC; set on pending→running
completed_at	Optional[datetime]	UTC; set on terminal flip
duration_ms	Optional[int]	Wall-clock milliseconds

Field	Type	Meaning
error	Optional[str]	Exception message, when status == "error" \ "timeout"
error_type	Optional[str]	Exception class name ("RuntimeError", "TimeoutError", ...)
authenticated	bool	True when a cached session exists or credentials env-vars are set
attempt	int	Retry counter (currently always 1; reserved for future retries)
http_status	Optional[int]	HTTP status from the plugin when surfaced
category	str	The category filter value used for this run
query	str	The user query
notes	Dict[str, Any]	Free-form diagnostics a plugin can stash for the dashboard dialog

SearchResult

File: search.py:82

Lifecycle: produced by each tracker's search method, stored in `_tracker_results[search_id][tracker_name]`, aggregated into `MergedResult` by the deduplicator.

Field	Type	Meaning
name	str	Torrent title
link	str	Download URL (.torrent or magnet)
size	str	

Field	Type	Meaning
		Human-readable size ("1.2 GB")
seeds	int	Current seeders; negative values clamped to 0
leechers	int	Current leechers
engine_url	str	Base URL of the source tracker
desc_link	Optional[str]	Permalink to the description page
pub_date	Optional[str]	Publication date as returned by the tracker
tracker	Optional[str]	Tracker identifier (e.g. "rutracker", "iptorrents")
category	Optional[str]	Category string
freeleech	bool	IPorrents-specific; drives [free] display suffix

`to_dict()` adds `tracker_display` ("`<tracker> [free]`" when `freeleech`, else "`<tracker>`").

MergedResult

File: `search.py:119`

Lifecycle: produced by `Deduplicator.merge_results()`, stored in `_last_merged_results[search_id]`, serialised into `SearchResultResponse`.

Field	Type	Meaning
canonical_identity	CanonicalIdentity	Normalised fingerprint used for dedup
original_results	List[SearchResult]	All source rows that collapsed into this merge
total_seeds	int	Sum of seeds across sources

Field	Type	Meaning
total_leechers	int	Sum of leechers across sources
best_quality	Optional[QualityTier]	Computed from sources; filled post-dedup
download_urls	List[str]	Unique URLs from the sources
created_at	datetime	UTC, auto-populated

CanonicalIdentity

File: search.py:54

Lifecycle: computed by the deduplicator when normalising a result; the stable key used for matching across trackers.

Field	Type	Meaning
infohash	Optional[str]	BitTorrent infohash (when the source exposes it)
title	Optional[str]	Normalised title
year	Optional[int]	Release year (strict \b(19\ 20)\d{2}\b match)
content_type	Optional[ContentType]	Enum: movie, tv, anime, music, audiobook, game, software, ebook, other, unknown
season	Optional[int]	TV season number
episode	Optional[int]	TV episode number
resolution	Optional[str]	"1080p", "4K", ...
codec	Optional[str]	"x264", "HEVC", ...
group	Optional[str]	Release group
metadata_source	Optional[str]	"omdb", "tmdb", "imdb", ...

TrackerSource

File: search.py:34

Lifecycle: built on demand by

SearchOrchestrator._get_enabled_trackers(); not persisted.

Field	Type	Meaning
name	str	Tracker identifier
url	str	Base URL
enabled	bool	Whether to search
last_checked	Optional[datetime]	Last health probe
health_status	str	"healthy", "degraded", "offline", "unknown"
scrape_url	Optional[str]	BEP 48 scrape endpoint

ContentType, QualityTier

File: search.py:11, 24

Enums used as the canonical vocabulary for content classification and quality tiers. Values exactly as they appear in the source:

- ContentType: movie, tv, anime, music, audiobook, game, software, ebook, other, unknown.
- QualityTier: sd, hd, full_hd, uhd_4k, uhd_8k, unknown.

MatchResult (deduplicator)

File: deduplicator.py:31

Intermediate result from a match attempt between two SearchResults.

Field	Type	Meaning
is_match	bool	Are they the same torrent?
confidence	float	0.0–1.0 similarity score
tier	int	1 (strongest, e.g. infohash) — 4

Field	Type	Meaning
		(weakest, e.g. fuzzy title)
reason	str	Human-readable reason used in logs

MetadataResult (enricher)

File: enricher.py:21

Result of an external metadata lookup (OMDb / TMDb / TVMaze / AniList / MusicBrainz / OpenLibrary).

Field	Type	Meaning
source	str	Provider name ("OMDb", "TMDb", ...)
title	str	Canonical title
year	Optional[int]	Release year
content_type	Optional[str]	"movie", "tv", "music", "book"
imdb_id	Optional[str]	—
tmdb_id	Optional[str]	—
anilist_id	Optional[str]	—
musicbrainz_id	Optional[str]	—
openlibrary_id	Optional[str]	—
poster_url	Optional[str]	—
overview	Optional[str]	Plot or blurb
genres	List[str]	Defaults to [] via __post_init__

ScrapeResult, TrackerStatus (validator)

File: validator.py:30, 40

Field	Type	Meaning
tracker	str	Tracker identifier
status	TrackerStatus	Enum: healthy, degraded, offline, unknown

Field	Type	Meaning
seeds	int	From scrape response
leechers	int	From scrape response
complete	int	From scrape response
error	Optional[str]	Message if scrape failed
scrape_time_ms	int	Wall-clock scrape duration

Hook types (merge_service/hooks.py)

HookEventType — enum of eight event types: search_start, search_progress, search_complete, download_start, download_progress, download_complete, merge_complete, validation_complete.

HookEvent (file: hooks.py:38)

Field	Type	Meaning
event_type	HookEventType	Which event fired
timestamp	datetime	Auto-populated UTC
search_id	Optional[str]	Correlation ID
download_id	Optional[str]	Correlation ID
data	Dict[str, Any]	Event-specific payload

HookConfig (file: hooks.py:58)

Field	Type	Meaning
name	str	Display name
event	HookEventType	Which event to listen on
script_path	str	Executable path; validated via os.path.exists
enabled	bool	Default True
timeout	int	Seconds, default 30
environment	Dict[str, str]	

Field	Type	Meaning
		Extra env vars injected into the script

Hooks persist at `/config/download-proxy/hooks.json` (see `api/hooks.py`).

ScheduledSearch, ScheduleStatus

File: `scheduler.py`:22, 32

Lifecycle: loaded from `/config/merge-service/scheduling.json`, kept in `Scheduler._scheduled_searches` during runtime.

Field	Type	Meaning
<code>id</code>	<code>str</code>	UUID
<code>name</code>	<code>str</code>	Display name
<code>query</code>	<code>str</code>	Query to re-run
<code>category</code>	<code>str</code>	Default "all"
<code>interval_minutes</code>	<code>int</code>	Default 60
<code>enabled</code>	<code>bool</code>	—
<code>last_run</code>	<code>Optional[datetime]</code>	—
<code>next_run</code>	<code>Optional[datetime]</code>	—
<code>status</code>	<code>ScheduleStatus</code>	active, paused, completed, failed
<code>results_count</code>	<code>int</code>	Last run's merged count
<code>error_message</code>	<code>Optional[str]</code>	Last error

Pydantic models (API surface)

Request / response bodies under `download-proxy/src/api/`.

api/routes.py

- `SearchRequest` — `{query, category, limit (1..100), enable_metadata, validate_trackers, sort_by, sort_order}`.

- SearchResultResponse — {name, size, seeds, leechers, download_urls, quality?, content_type?, desc_link?, tracker?, sources, metadata?, freeleech}.
- SearchResponse — {search_id, query, status, results, total_results, merged_results, trackers_searched, started_at, completed_at?}.
- DownloadRequest — {result_id, download_urls}.
- QBitLoginRequest (inner, line 369) — {username, password, remember?}.
- MagnetRequest (inner, line 675) — magnet-generation request.

api/hooks.py

- HookCreateRequest — {name, event, script_path, enabled, timeout (1..300), environment}.
- HookResponse — {hook_id, name, event, script_path, enabled, timeout, created_at}.

api/auth.py

- CaptchaLoginRequest — {cap_sid, cap_code_field, captcha_text, captcha_token}.
- CookieLoginRequest — {cookie_string}.

Note: The in-memory CAPTCHA challenge is **not** a distinct dataclass. It is stored as an anonymous dict in `_pending_captchas`: `dict = {}` at `api/auth.py:24`. Verified: no `CaptchaChallenge` class exists in the repo.

api/scheduler.py

- ScheduleCreateRequest — fields for creating a ScheduledSearch.
- ScheduleUpdateRequest — partial update payload.

JSON persistence

Path	Owner	Structure
/config/download-proxy/hooks.json	api/hooks.py	Array of HookConfig dicts
/config/download-proxy/qbittorrent_creds.json	api/routes.py	{username, password}

Path	Owner	Structure
/config/merge-service/ scheduling.json	merge_service/ scheduler.py	Array of ScheduledSearch dicts

Why no relational DB?

- The merge service is a stateless orchestrator over tracker APIs; durable state would duplicate what qBittorrent already tracks.
- Adding Postgres / SQLite would violate constitution Principle I (Container-First — the product is exactly two containers).
- Query surface is tiny (CRUD over 2–3 JSON files); the cognitive cost of a schema layer is higher than the benefit.

Gotchas

- All datetimes are UTC; the `to_dict()` serialisers call `.isoformat()`.
- `SearchMetadata.errors` is a plain string list. The tracker name is embedded as "`<name>: <msg>`", not in a separate field.
- `SearchResult.freeleech` matters only for IPTorrents; other trackers leave it `False`. The deduplicator uses this to forbid merging a non-freeleech IPTorrents row with any other source.
- `MergedResult.best_quality` is filled **post-dedup** by the enricher, not during initial construction.
- `_pending_captchas` lacks a TTL — this is a known leak, tracked in [CONCURRENCY.md](#).