Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Fixed (→ Fixed.md)	9
Bug	In progress	2
Bug	Operator-blocked	1
Bug	Queued	11
Bug	Ready for testing	1
Task	Completed (→ Fixed.md)	1
Task	In progress	9
Task	Operator-blocked	3
Task	Queued	19
Task	Ready for testing	2
TOTAL		**58**

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-ho SOCKS routing (containers pkg/egress
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download-proxy + qBitTorrent-go + Jackett + compose env-forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Aut commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-chann ledger — ongoing coverage-escape bac
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully abser from the mandated test-type matrix

ATM ID	Type	Status	Severity	Description
BOB-077	Task	Operator-blocked	High	RD2-10: Identify second host running t Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-cor mechanism through §11.4.234 dedicate commit/push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code review of the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workabl items for the four Lava-porting finding (closes GA-05)
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/ merge-service v1.0.0 readiness ledger (closes GA-10)
BOB-087	Task	Completed (→ Fixed.md)	High	RD2-20: Wire docs_chain / commit-sea sync hook per §11.4.106(F) so DB writ cannot land without MD mirror
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Tracked-Items + Status Documents tal row-completeness (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exercising all three start.sh subcomma end-to-end against real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + veri rutracker ReDoS regex bounds deploye container (closes runtime half of GA-12
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py wit §11.4.85 fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos for Go-side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage f exposed download-proxy/merge endpo (canonical impl of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one commit (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is –profile go parity still release goal? (gates RW-10..13) — OPERATOR-DECISION
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — tunnel 127.0.0.1 or keep 0.0.0.0? — OPERATOR-DECISION
BOB-104	Task		Medium	

ATM ID	Type	Status	Severity	Description
		In progress		§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescenc check helper for the unattributed auto commit path
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existe check for subagent task-brief source inputs
BOB-109	Task	Queued	Medium	BOB-074 followup: scaling-class test coverage absent from mandated test-ty matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test cover (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiting for boba’s 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation gold bad fixture for the rate-limit detector
BOB-120	Bug	Queued	Critical	3rd forced-logout incident 2026-08-18 23:45:49 — SIGKILL user@1000 + preventive-timer-inside-user-slice architectural gap
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logou architectural gap (task #85, incident #
BOB-129	Bug	Fixed (→ Fixed.md)	Medium	Potential production slowapi/starlette defect flagged by Task 105 subagent
BOB-131	Bug	Fixed (→ Fixed.md)	Medium	qbittorrent-proxy podman conmon cras pre-existing, surfaced during BOB-129 investigation
BOB-135	Bug	Queued	Low	Test isolation: test_list_hooks_after_cre fails in bulk suite (Permission denied / config)
BOB-136	Task	Ready for testing	High	Closure seam does not bind: 4 tracker rows found stale in one sweep, and workable-items diff is blind to body_mo drift
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while t same process still serves 7186 (GIL starvation by one spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile serv 7186/7187/7188 but its container bind only 7187 — doc contradicts the Docke
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolvable gitdir) pollute gate scan

ATM ID	Type	Status	Severity	Description
				scope and manufacture false BOB-126-class findings
BOB-144	Bug	Fixed (→ Fixed.md)	Low	/theme/stream calls the disconnect pro unguarded — fail-closed but via an uncaught traceback, inconsistent with two SSE generators
BOB-145	Bug	Fixed (→ Fixed.md)	High	Fix the 7187 wedge: offload and/or memoise Deduplicator.merge_results s O(N^2) regex work stops blocking the asyncio event loop
BOB-146	Bug	Fixed (→ Fixed.md)	High	Constitution §11.4.252 detector undercounts by 29% (30 vs 42 AST gro truth) — 4 distinct blind spots make its output a floor, not a census
BOB-148	Bug	Fixed (→ Fixed.md)	Medium	Standing red unit test nothing tracked test_no_credentials asserts has_sessior False, gets True — real defect or non-hermetic test, undecided
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/ across constitution, CLAUDE.md and t README badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labe read N/44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named ga with no implementation, and 24 of 36 scripts have no companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored thi party code: 82% of one gate's 38,291 findings come from submodules/
BOB-153	Bug	Fixed (→ Fixed.md)	Medium	Go profile cannot build: go.mod requir go 1.26.2 but the Dockerfile builder is golang:1.23-alpine
BOB-154	Bug	Ready for testing	Medium	Host venv and production container ru different starlette versions (1.4.1 vs 1.(
BOB-155	Bug	Fixed (→ Fixed.md)	High	workable-items diff reports 'DB and Markdown are in sync' having opened Markdown files when –issues/–fixed ar omitted
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard i load-sensitive and flaky: 8786ms under host load vs a 900-1500ms ceiling calibrated on a quiet host
BOB-157	Bug	Fixed (→ Fixed.md)	High	Our own BOB-137 stall watchdog can segfault the merge service: faulthandle

ATM ID	Type	Status	Severity	Description
				dump_traceback(all_threads=True) hit unpatched CPython 3.12 defect
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds asyncio.events._get_event_loop_policy, 3.13+ private API, while production is 3.12.13
BOB-159	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:01:17Z **Reported-By:** T041 independent revie (IMPORTANT-2), partially remediated **What (the report, verbatim):** The –recreate path was fixed this round: run_ownership_precondition -> stack_down -> run_ownership_repair - stack_up, so the repair walks a tree no container can write to. That ordering i now asserted behaviourally by three ne checks in tests/unit/ test_start_reload_recreate.sh, each kill by its paired 1.1 mutation (11/11 RED- The WARM path is NOT covered by tha and this item tracks the remainder. On warm ./start.sh the gate runs before T invocation brings anything up, but if th stack is ALREADY running, containers write while the repair walks. A contain can then create a new non-operator-ov file BEHIND the walk, after which the completion marker records complete o a tree that is not, and those stragglers never repaired without a manual –forc BOUNDED, not dismissed: after any successful pass the marker is valid for scope fingerprint and the repair short- circuits without walking (marker_is_va so the window requires a stale-or-abse marker AND a live stack together. Declaring a new scope entry re-arms th marker, which is exactly how that combination arises in practice. NOT D THIS ROUND, with the reason stated rather than hidden: the clean guard is cheap marker-only probe mode on scri ownership_repair.sh that answers has- scope-been-repaired without walking t tree. The existing –dry-run cannot serv deliberately SKIPS the marker fast-pat (FORCE=0 && DRY_RUN=0 guard) ar

ATM ID	Type	Status	Severity	Description
				always walks, so using it as a warm-sta probe would walk the tree twice on eve start. Adding that mode requires editin ownership_repair.sh, which another str was actively editing during this round; duplicating marker_is_valid into start.s instead would be a near-identical fork (11.4.251). Deferred to this item rathe than done unilaterally or silently. The misleading comment at the warm-start site (which claimed the gate runs befor any container writes) has been correct to state this boundary honestly. **Affect scope / file-scope manifest:** start.sh (warm-start dispatch, run_ownership_g call site); scripts/ownership_repair.sh (marker fast-path) **Reproduction / context:** 1. Ensure the stack is UP. 2. Change config/owned_paths.yaml so th scope fingerprint changes (this re-arm marker by design, data-model E2). 3. R a warm ./start.sh (no –recreate). 4. The ownership repair walks and chowns th declared tree while containers are still running and able to write into it. **Acceptance criteria:** A warm ./start.s cannot complete an ownership repair while a container that writes to a decla location is running. Either the repair is deferred with an actionable refusal nam –recreate, or the stack is quiesced for t walk. Asserted behaviourally in tests/u test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE_UP checks, each wit paired 1.1 mutation that kills it.
BOB-160	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:02:48 **Reported-By:** AI **What (the report, verbatim):** Independent §11.4.209 rev (IMPORTANT-3) found two of the featu strongest automated checks were neve executed by any runner: (1) tests/ ownership/ test_container_writes_owned_files.py – the §11.4.115 RED-turned-regression-guard for FR-011/FR-007, proven via a docker-compose.yml revert mutation –

ATM ID	Type	Status	Severity	Description
				was moved out of tests/integration/ (commit 58d340a, to escape an autous fixture hang) but no runner (ci.sh, test run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_b test_.sh, *the §1.1 paired-mutation meta tests for the scripts/pre_build/ check_cm_*.sh gate family, was execute NOTHING — self-reported honestly in commit 04742d7's own message ('STIL OPEN (reported, not fixed): tests/ pre_build/ is executed by NOTHING') l never tracked as a workable item (a §11.4.197 loss-of-requirements gap) ar never wired into scripts/ pre_build_verification.sh invariant 30, which globbed only tests/unit/test_*.sh.* ***Affected scope / file-scope manifest*** *ci.sh; scripts/pre_build_verification.sh invariant 30 (CM-BASH-UNIT-TESTS-EXECUTED)* ***Reproduction / context*** *grep -rn test_container_writes_owned_ ci.sh test.sh run-all-tests.sh scripts/ returned zero runner hits (only docs/sp mentioned the path); grep in scripts/ pre_build_verification.sh showed invar 30's for-loop globbing only tests/unit/ test_*.sh, never tests/pre_build/test_*.sh* ***Acceptance criteria:*** *ci.sh gains a runtime-gated pytest tests/ownership/ stage (skips honestly, rc=0, when no podman/docker present); scripts/ pre_build_verification.sh invariant 30's glob covers both tests/unit/test_*.sh anc tests/pre_build/test_*.sh, verified by ar extracted standalone run of the invaria exact block reporting a non-zero RAN count that includes files from both directories.
BOB-161	Task	Queued	High	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:31:12 **Reported-By:** BOB-092 remediation, residual finding **What (the report, verbatim):** §11.4.69 names CM-NO-FA OPEN-SKIP as one of three mandatory build gates: it audits sink-side probe helpers and FAILs if any code path converts an empty or unreachable response into a PASS-counting SKIP fo feature class with a sink-side probe. Th

ATM ID	Type	Status	Severity	Description
				project does not have it. Why this matt now rather than in the abstract: the BOB-092 remediation just removed TW fail-opens of exactly this class from tes e2e/test_live_stack_evidence.py — the nnmclub SKIP-on-404 (already remove 7baef2b) and a surviving sibling at test_iptorrents_is_authenticated_in_sea that skipped on 'not authenticated and status != success' while asserting 'trea as transient outage'. That assertion wa false for rejected credentials (upstream_http_403) and for a broken container (plugin_env_missing), both o which are definitive product failures th product already classifies via error_typ Two fail-opens of one class, found one time, is the §11.4.146 extend-to-all-cas signal and the §11.4.238 coverage-esca signal together: the regime did not sur either, an agent reading code did. The remediation's own guards are AST-structural and lethal under three discriminating mutations, but they LIV THE FILE THEY GUARD, so deleting th file evades them entirely. A pre-build g is the durable home because it audits t corpus rather than one file. Honest boundary (§11.4.6): this item does NOT claim the two remediated fail-opens ar unguarded — they are guarded, with runtime evidence (13 passed in 97.36s against the live stack). It claims the CI has no corpus-wide detector, so the ne instance in a different file is invisible again. **Affected scope / file-scope manifest:** scripts/pre_build/ (gate abs tests/e2e/test_live_stack_evidence.py (guards currently live inside the file th guard) **Reproduction / context:** grep 'CM-NO-FAIL-OPEN-SKIP' scripts/ test returns exactly ONE hit, and it is tests test_live_stack_evidence.py — the file t guards live in, not a gate. Control need the same query for CM-OWNERSHIP-INVARIANTS (a gate that does exist) returns 3 files, so the instrument can s gate tokens and the single hit is a real absence, not a blind zero (§11.4.201(7) **Acceptance criteria:** A scripts/pre_bu

ATM ID	Type	Status	Severity	Description
				gate named CM-NO-FAIL-OPEN-SKIP exists, is wired into scripts/ pre_build_verification.sh, and audits si side probe helpers for code paths converting an empty/unreachable/erro response into a PASS-counting SKIP fo feature class that HAS a sink-side prob ships a golden-TRUE fixture (a real fai open -> gate FIRES) and a golden-FAL with-carrier (an honest topology/geo s and a comment merely MENTIONING phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §1. mutation makes the gate FAIL before t gate is trusted.
BOB-162	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:35:02 **Reported-By:** BOB-106/BOB-107 remediation, residual **What (the repo verbatim):** The TESTS for both guard now executed: pre-build invariant 30's glob was extended from tests/unit + te pre_build to also cover tests/hooks. MEASURED before/after: the glob expanded to 27 suites with 0 under tes hooks while 3 existed there; it now expands to 30 with 3 under tests/hook three pass. But the GUARDS themselv are still invoked by nothing. check-brie inputs.sh is honestly conductor-run rat than automatic: a brief's required-inpu list lives in free-form prose that the Ag tool's structured input does not expose an automatic prompt-scraping gate wo itself be an unproven pattern-match. It seam is the dispatch procedure, and th a documentation + habit change, not a hook. unattributed-commit-guard.sh is different and is why this item exists. R against real history it names 14 violati commits since the last tag (and 21 bar Auto-commit subjects exist across all r Wiring it into the pre-build or commit seam AS-IS would therefore refuse eve commit immediately, on pre-existing d nobody can fix in the moment. That is precisely the 11.4.224(E) brownfield- adoption shape, and 11.4.66/11.4.122 the choice with the operator, not with a agent inventing a ratchet. The options

ATM ID	Type	Status	Severity	Description
				stated so the decision is a decision and a default: (a) immediate hard floor – th seam refuses until all 14 are attributed one-time monotone-decrease ratchet (t 11.4.135 pattern) – snapshot 14 as the baseline, the count may fall and may n rise, so day one is green and the disea cannot spread; (c) changed-commits-o gate only commits created from now o with a scheduled deadline for the historical 14; (d) report-only for a state period, then escalate. Recommendatio with the reasoning rather than just the pick: (b). It is the pattern this constitut already uses for exactly this situation, makes the debt visible and bounded immediately, and it cannot be satisfied deleting the guard's name (11.4.227 cl that gaming channel). But it is the operator's call and this item is not clos until that answer is recorded. Honest boundary (11.4.6): this item does NOT claim the 14 commits are harmful in themselves – it claims their ATTRIBUT is absent and that nothing currently prevents the 15th. **Affected scope / fi scope manifest:** scripts/hooks/ unattributed-commit-guard.sh; scripts/ hooks/check-brief-inputs.sh; scripts/ pre_build_verification.sh; scripts/comn push-all.sh **Reproduction / context:** I guards run correctly by hand and are covered by passing tests (9/9 and 15/1 each proven non-bluff against a no-op stub). Neither is invoked by scripts/ pre_build_verification.sh or scripts/ commit-push-all.sh, so neither exerts standing detection pressure (11.4.226 guard with no execution seam is not coverage; registration is not coverage) **Acceptance criteria:** Each guard is invoked by a named seam, OR carries registered deferral pointing at this iten For unattributed-commit-guard.sh specifically, the operator has answered adoption question below and the answ recorded as consumer DATA – never ar invented ratchet.
BOB-163	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:41:08Z

ATM ID	Type	Status	Severity	Description
				Reported-By: BOB-114 remediation, p existing defect surfaced by BOB-111 landing a real limiter **What (the repo verbatim):** PRE-EXISTING, NOT INTRODUCED – and proven so rather asserted: git diff shows assertion (c) is BYTE-IDENTICAL to HEAD, and the fa reproduces against the unmodified HE script. What changed is not the challen but the SYSTEM: BOB-111 appears at least partially delivered, because :718 now returns real 429s where it previou returned none. :7185 and :7189 still produce zero 429s. This is a 11.4.248-c corrosion risk and that is why it is filed rather than left as a footnote: run_all_challenges.sh will now fail INTERMITTENTLY, and an intermitten red trains everyone to re-run until gre at which point a real regression is dismissed as 'probably the flaky rate-li one'. THE DECISION, which is an ope call under 11.4.66 and which I delibera did NOT invent: Does a 429 from a WORKING limiter count as the sibling endpoint being RESPONSIVE? (a) YES 429 proves the endpoint is alive and correctly protecting itself. Assertion (c accepts 2xx OR 429, and only a conne failure / 5xx / timeout counts as degra Risk: a genuinely wedged endpoint tha happens to answer 429 would pass. (b) – keep requiring 2xx, but make the challenge limiter-aware: drain or wait the window before the sibling probe (r after is served, so the budget is knowa rather than guessed). (c) Probe sibling a path the limiter exempts (a healthz-c route), so the isolation question is aske without spending the bucket. Recommendation with reasoning, not j a pick: (a) composed with a bounded liveness check – a 429 carrying a well- formed Retry-After IS evidence of a liv correctly-behaving service, and (b) ma the challenge slower and couples it to window value that will drift. But the semantics of 'responsive' here are a product judgement, so the answer is recorded, not assumed. RELATED, stat

ATM ID	Type	Status	Severity	Description
				as fact rather than folded in: the detec counts 429 only, not 503, though the sc header says '429 (or equivalent)'. Coun 503 would collide with the crash detec 5xx tally. Worth deciding when BOB-11 lands limiters on :7185 and :7189. **Affected scope / file-scope manifest** challenges/scripts/ ddos_resilience_challenge.sh assertion cross-endpoint isolation; run_all_challenges.sh which invokes it **Reproduction / context:** Assertion (c requires a sibling-endpoint probe to answer ^2 (a 2xx). :7187 now enforces real limiter (measured live: x-ratelimit- limit: 120, x-ratelimit-remaining: 86, re after: 45, server: uvicorn). A full challe run sends roughly 250 requests to :718 so sibling probes fired during the OTH endpoints' tiers land on an exhausted bucket and receive 429 – which asserti (c) scores as 'endpoint degraded'. Tim dependent: one run measured PASS=5 FAIL=2; the UNMODIFIED HEAD scri against the same live stack ~30s later measured PASS=7 FAIL=0 SKIP=2. **Acceptance criteria:** The operator ha answered the classification question bo the answer is recorded as consumer D and assertion (c) implements it. The challenge then produces the same verd across 3 consecutive runs against a ra limited stack (11.4.50 deterministic consistency), and the fix ships a paired mutation proving assertion (c) still cat a genuinely degraded sibling endpoint narrowing it must not blind it (11.4.201(1)).
BOB-164	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:56:53Z **Reported-By:** BOB-110 UX-class coverage, discovered by the new axe-c suite on its first live run **What (the report, verbatim):** This is a REAL use facing defect, not a test-tuning artifact and it was found by the automated reg rather than by a human squinting at th page – which is exactly the 11.4.238 posture the project is aiming for. The static-grep oracle would NOT have fou

ATM ID	Type	Status	Severity	Description
				it. Measured: the served root ships an empty pre-hydration, so any check rea the raw HTML audits a page nobody s The violation only exists in the hydrate DOM, which is why 11.4.170 requires rendered oracle and forbids value-equa assertions as the proof a UI is correct. Severity reasoning, stated rather than assumed: this is user-visible and affect the primary dashboard heading, but it degrades legibility rather than breakin function, and the surface is operator-facing rather than public. Medium, no High. The failing test was left FAILIN purpose (11.4.238) instead of silenced marked xfail. tests/ux/ currently repor failed, 16 passed; that 1 is this defect. Anyone reading a red UX suite should it as this item, not as flakiness – and w this is fixed the suite goes fully green, which is the signal that it is closed. HONEST SCOPE LIMIT (11.4.6): only dashboard landing view was scanned. The /jackett/* sub-routes and the ng-se hosted :4200 route set were NOT audi the commands to close both are recor in docs/testing/ux_accessibility.md. So item's 21 nodes are a floor, not a total. **Affected scope / file-scope manifest** the Angular dashboard served at http:/ localhost:7187/ (same compiled SPA as frontend/); production component CSS test files **Reproduction / context:** Ru tests/ux/ test_live_dashboard_accessibility.py against the running merge service. axe core v4.13.0, scanning a real Playwrig rendered JS-hydrated DOM, reports co contrast violations on 21 nodes. Measu pairs: .brand / h1 text #9d001e on background #3c3f41 = 1.43-1.62:1 (W AA large-text floor is 3:1); tagline #80 on #3c3f41 = 2.68:1 (body-text floor is 4.5:1). **Acceptance criteria:** axe-core reports ZERO color-contrast violations against the live rendered dashboard, w the fix made in production component rather than by relaxing the assertion o excluding the rule (11.4.120: reconcile the correct mechanism, never weaken

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	check). The existing tests/ux/ suite is t guard and already fails today, so the R is captured – closure requires it flippin GREEN against the live surface, which runtime-class evidence per 11.4.226. **Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T20:00:08Z **Reported-By:** surfaced while capturin closure evidence for BOB-092; diagnos rather than worked around **What (the report, verbatim):** This is a STALE-A AFTER-INTERPRETER-UPGRADE defe not a missing package. The package is installed; its compiled half is built for t previous interpreter. That distinction matters because 'pip install rpds-py' st advice can appear to succeed while lea the 313 artefact in place. WHY IT WAS NOT NOTICED EARLIER, stated as fac the paths that DO work all avoid syste python3. ci.sh selects an interpreter through _select_python; the long-runni suites in this session ran .venv/bin/pyt -m pytest and passed. So the automate regime is green while the DOCUMENT operator command is broken – an 11.4.238-shaped gap, since the escape found by an agent capturing evidence rather than by the regime. WORKAROU (measured, not theorised): PYTEST_DISABLE_PLUGIN_AUTOLOA with the needed plugins passed explici p pytest_timeout ...) avoids the schemathesis autoload that drags in jsonschema -> referencing -> rpds. Th a workaround, NOT the fix: it silences import path rather than repairing the install, and it will surprise the next pe who runs the documented command verbatim. RELATED, and worth decidi together rather than twice: BOB-154 wants .venv rebuilt on 3.12 to match th container (which runs CPython 3.12.13 while both system and venv run 3.14.6 There are therefore THREE interprete play. Whoever rebuilds the venv shoul confirm rpds resolves for the TARGET interpreter afterwards, or this same cl reappears one directory over. HONEST BOUNDARY (11.4.6): this item does N

ATM ID	Type	Status	Severity	Description
				claim any test is wrong or any product code is broken. The product and the ve run suites are unaffected. What is brok is the documented entry point. **Affecte scope / file-scope manifest:** host use site-packages (~/.local/lib/python3/site packages/rpds/); every CLAUDE.md-documented ‘python3 -m pytest …’ invocation; any tooling that uses syster python3 rather than .venv/bin/python **Reproduction / context:** python3 -m pytest tests/e2e/ test_live_stack_evidence.py -q –import-mode=importlib -> ModuleNotFoundError: No module nar ‘rpds.rpds’, raised during collection via schemathesis plugin autoload -> jsonschema -> referencing._core -> rp MEASURED root cause: system pythor 3.14.6, but ~/.local/lib/python3/site-packages/rpds/ ships rpds.cpython-313 x86_64-linux-gnu.so – an extension bui for 3.13. rpds/**init**.py does ’from .rpds import *‘, and a 313-tagged .so is not importable by 3.14, so the submodule genuinely does not exist for this interpreter. The venv is CORRECT and unaffected: .venv/lib64/python3/site-packages/rpds/ ships rpds.cpython-314 x86_64-linux-gnu.so and’.venv/bin/pyth -c import rpds’ succeeds. **Acceptance criteria:** python3 -m pytest tests/unit/ –import-mode=importlib – the commar CLAUDE.md documents – reaches collection without ModuleNotFoundEr OR CLAUDE.md is corrected to docum the supported runner explicitly. Either the documented command and the working command agree (11.4.99: a gu that misleads is the documentation-lay equivalent of a PASS-bluff).