Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Fixed (→ Fixed.md)	9
Bug	In progress	2
Bug	Operator-blocked	1
Bug	Queued	16
Bug	Ready for testing	1
Task	Completed (→ Fixed.md)	1
Task	In progress	9
Task	Operator-blocked	3
Task	Queued	21
Task	Ready for testing	1
TOTAL		**64**

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-ho SOCKS routing (containers pkg/egress
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download-proxy + qBitTorrent-go + Jackett + compose env-forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Aut commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-chann ledger — ongoing coverage-escape bac
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absen from the mandated test-type matrix

ATM ID	Type	Status	Severity	Description
BOB-077	Task	Operator-blocked	High	RD2-10: Identify second host running t Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-cor mechanism through §11.4.234 dedicate commit/push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code review of the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workabl items for the four Lava-porting finding (closes GA-05)
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/ merge-service v1.0.0 readiness ledger (closes GA-10)
BOB-087	Task	Completed (→ Fixed.md)	High	RD2-20: Wire docs_chain / commit-sea sync hook per §11.4.106(F) so DB writ cannot land without MD mirror
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Tracked-Items + Status Documents tal row-completeness (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exercising all three start.sh subcomma end-to-end against real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + veri rutracker ReDoS regex bounds deploye container (closes runtime half of GA-12
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py wit §11.4.85 fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos for Go-side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage f exposed download-proxy/merge endpo (canonical impl of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one commit (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is –profile go parity still release goal? (gates RW-10..13) — OPERATOR-DECISION
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — tunnel 127.0.0.1 or keep 0.0.0.0? — OPERATOR-DECISION
BOB-104	Task		Medium	

ATM ID	Type	Status	Severity	Description
		In progress		§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescenc check helper for the unattributed auto commit path
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existe check for subagent task-brief source inputs
BOB-109	Task	Queued	Medium	BOB-074 followup: scaling-class test coverage absent from mandated test-ty matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test cover (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiting for boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation gold bad fixture for the rate-limit detector
BOB-120	Bug	Queued	Critical	3rd forced-logout incident 2026-08-18 23:45:49 — SIGKILL user@1000 + preventive-timer-inside-user-slice architectural gap
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logou architectural gap (task #85, incident #
BOB-129	Bug	Fixed (→ Fixed.md)	Medium	Potential production slowapi/starlette defect flagged by Task 105 subagent
BOB-131	Bug	Fixed (→ Fixed.md)	Medium	qbittorrent-proxy podman conmon cras pre-existing, surfaced during BOB-129 investigation
BOB-135	Bug	Queued	Low	Test isolation: test_list_hooks_after_cre fails in bulk suite (Permission denied / config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while t same process still serves 7186 (GIL starvation by one spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile serv 7186/7187/7188 but its container bind only 7187 — doc contradicts the Docke
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolvable gitdir) pollute gate scan scope and manufacture false BOB-126- class findings
BOB-144	Bug	Fixed (→ Fixed.md)	Low	/theme/stream calls the disconnect pro unguarded — fail-closed but via an

ATM ID	Type	Status	Severity	Description
				uncaught traceback, inconsistent with two SSE generators
BOB-145	Bug	Fixed (→ Fixed.md)	High	Fix the 7187 wedge: offload and/or memoise Deduplicator.merge_results s O(N^2) regex work stops blocking the asyncio event loop
BOB-146	Bug	Fixed (→ Fixed.md)	High	Constitution §11.4.252 detector undercounts by 29% (30 vs 42 AST gro truth) — 4 distinct blind spots make its output a floor, not a census
BOB-148	Bug	Fixed (→ Fixed.md)	Medium	Standing red unit test nothing tracked test_no_credentials asserts has_sessior False, gets True — real defect or non-hermetic test, undecided
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/ across constitution, CLAUDE.md and t README badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labe read N/44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named ga with no implementation, and 24 of 36 scripts have no companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored thi party code: 82% of one gate's 38,291 findings come from submodules/
BOB-153	Bug	Fixed (→ Fixed.md)	Medium	Go profile cannot build: go.mod requir go 1.26.2 but the Dockerfile builder is golang:1.23-alpine
BOB-154	Bug	Ready for testing	Medium	Host venv and production container ru different starlette versions (1.4.1 vs 1.
BOB-155	Bug	Fixed (→ Fixed.md)	High	workable-items diff reports 'DB and Markdown are in sync' having opened Markdown files when –issues/–fixed ar omitted
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard load-sensitive and flaky: 8786ms under host load vs a 900-1500ms ceiling calibrated on a quiet host
BOB-157	Bug	Fixed (→ Fixed.md)	High	Our own BOB-137 stall watchdog can segfault the merge service: faulthandle dump_traceback(all_threads=True) hit unpatched CPython 3.12 defect
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds

ATM ID	Type	Status	Severity	Description
				asyncio.events._get_event_loop_policy, 3.13+ private API, while production is 3.12.13
BOB-159	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:01:17Z **Reported-By:** T041 independent revie (IMPORTANT-2), partially remediated **What (the report, verbatim):** The –recreate path was fixed this round: run_ownership_precondition -> stack_down -> run_ownership_repair - stack_up, so the repair walks a tree no container can write to. That ordering i now asserted behaviourally by three n checks in tests/unit/ test_start_reload_recreate.sh, each kill by its paired 1.1 mutation (11/11 RED- The WARM path is NOT covered by tha and this item tracks the remainder. On warm ./start.sh the gate runs before T invocation brings anything up, but if th stack is ALREADY running, containers write while the repair walks. A contain can then create a new non-operator-ov file BEHIND the walk, after which the completion marker records complete c a tree that is not, and those stragglers never repaired without a manual –forc BOUNDED, not dismissed: after any successful pass the marker is valid for scope fingerprint and the repair short- circuits without walking (marker_is_va so the window requires a stale-or-abse marker AND a live stack together. Declaring a new scope entry re-arms t marker, which is exactly how that combination arises in practice. NOT D THIS ROUND, with the reason stated rather than hidden: the clean guard is cheap marker-only probe mode on scri ownership_repair.sh that answers has- scope-been-repaired without walking t tree. The existing –dry-run cannot serv deliberately SKIPS the marker fast-pat (FORCE=0 && DRY_RUN=0 guard) a always walks, so using it as a warm-sta probe would walk the tree twice on eve start. Adding that mode requires editi ownership_repair.sh, which another st was actively editing during this round;

ATM ID	Type	Status	Severity	Description
				duplicating marker_is_valid into start.s instead would be a near-identical fork (11.4.251). Deferred to this item rathe than done unilaterally or silently. The misleading comment at the warm-start site (which claimed the gate runs befo any container writes) has been correct to state this boundary honestly. **Affect scope / file-scope manifest:** start.sh (warm-start dispatch, run_ownership_g call site); scripts/ownership_repair.sh (marker fast-path) **Reproduction / context:** 1. Ensure the stack is UP. 2. Change config/owned_paths.yaml so th scope fingerprint changes (this re-arm marker by design, data-model E2). 3. R a warm ./start.sh (no –recreate). 4. The ownership repair walks and chowns th declared tree while containers are still running and able to write into it. **Acceptance criteria:** A warm ./start.s cannot complete an ownership repair while a container that writes to a decla location is running. Either the repair is deferred with an actionable refusal na –recreate, or the stack is quiesced for walk. Asserted behaviourally in tests/u test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE_UP checks, each wit paired 1.1 mutation that kills it.
BOB-160	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:02:48 **Reported-By:** AI **What (the report, verbatim):** Independent §11.4.209 rev (IMPORTANT-3) found two of the featu strongest automated checks were neve executed by any runner: (1) tests/ ownership/ test_container_writes_owned_files.py – the §11.4.115 RED-turned-regression-guard for FR-011/FR-007, proven via a docker-compose.yml revert mutation – was moved out of tests/integration/ (commit 58d340a, to escape an autous fixture hang) but no runner (ci.sh, test run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_b

ATM ID	Type	Status	Severity	Description
				test_*.sh, the §1.1 paired-mutation meta tests for the scripts/pre_build/ check_cm_*.sh gate family, was execute NOTHING — self-reported honestly in commit 04742d7's own message ('STIL OPEN (reported, not fixed): tests/ pre_build/ is executed by NOTHING') b never tracked as a workable item (a §11.4.197 loss-of-requirements gap) an never wired into scripts/ pre_build_verification.sh invariant 30, which globbed only tests/unit/test_*.sh. ***Affected scope / file-scope manifest*** *ci.sh; scripts/pre_build_verification.sh invariant 30 (CM-BASH-UNIT-TESTS-EXECUTED)* ***Reproduction / context*** *grep -rn test_container_writes_owned_ ci.sh test.sh run-all-tests.sh scripts/ returned zero runner hits (only docs/sp mentioned the path); grep in scripts/ pre_build_verification.sh showed invar 30's for-loop globbing only tests/unit/ test_*.sh, never tests/pre_build/test_*.sh ***Acceptance criteria:*** *ci.sh gains a runtime-gated pytest tests/ownership/ stage (skips honestly, rc=0, when no podman/docker present); scripts/ pre_build_verification.sh invariant 30's glob covers both tests/unit/test_*.sh and tests/pre_build/test_*.sh, verified by an extracted standalone run of the invaria exact block reporting a non-zero RAN count that includes files from both directories.
BOB-161	Task	Queued	High	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:31:12 **Reported-By:** BOB-092 remediation, residual finding **What (the report, verbatim):** §11.4.69 names CM-NO-FA OPEN-SKIP as one of three mandatory build gates: it audits sink-side probe helpers and FAILs if any code path converts an empty or unreachable response into a PASS-counting SKIP fo feature class with a sink-side probe. Th project does not have it. Why this matt now rather than in the abstract: the BOB-092 remediation just removed TW fail-opens of exactly this class from tes e2e/test_live_stack_evidence.py — the

ATM ID	Type	Status	Severity	Description
				nnmclub SKIP-on-404 (already remove 7baef2b) and a surviving sibling at test_iptorrents_is_authenticated_in_sea that skipped on 'not authenticated and status != success' while asserting 'trea as transient outage'. That assertion wa false for rejected credentials (upstream_http_403) and for a broken container (plugin_env_missing), both o which are definitive product failures th product already classifies via error_typ Two fail-opens of one class, found one time, is the §11.4.146 extend-to-all-cas signal and the §11.4.238 coverage-esca signal together: the regime did not sur either, an agent reading code did. The remediation's own guards are AST-structural and lethal under three discriminating mutations, but they LIV THE FILE THEY GUARD, so deleting th file evades them entirely. A pre-build g is the durable home because it audits t corpus rather than one file. Honest boundary (§11.4.6): this item does NO claim the two remediated fail-opens ar unguarded — they are guarded, with runtime evidence (13 passed in 97.36s against the live stack). It claims the CI has no corpus-wide detector, so the ne instance in a different file is invisible again. **Affected scope / file-scope manifest:** scripts/pre_build/ (gate abse tests/e2e/test_live_stack_evidence.py (guards currently live inside the file th guard) **Reproduction / context:** grep 'CM-NO-FAIL-OPEN-SKIP' scripts/ test returns exactly ONE hit, and it is tests test_live_stack_evidence.py — the file t guards live in, not a gate. Control need the same query for CM-OWNERSHIP-INVARIANTS (a gate that does exist) returns 3 files, so the instrument can s gate tokens and the single hit is a real absence, not a blind zero (§11.4.201(7) **Acceptance criteria:** A scripts/pre_bu gate named CM-NO-FAIL-OPEN-SKIP exists, is wired into scripts/ pre_build_verification.sh, and audits si side probe helpers for code paths converting an empty/unreachable/erro

ATM ID	Type	Status	Severity	Description
				response into a PASS-counting SKIP fo feature class that HAS a sink-side prob ships a golden-TRUE fixture (a real fai open -> gate FIRES) and a golden-FAL with-carrier (an honest topology/geo s and a comment merely MENTIONING phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §1. mutation makes the gate FAIL before t gate is trusted.
BOB-162	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:35:02 **Reported-By:** BOB-106/BOB-107 remediation, residual **What (the repo verbatim):** The TESTS for both guard now executed: pre-build invariant 30's glob was extended from tests/unit + te pre_build to also cover tests/hooks. MEASURED before/after: the glob expanded to 27 suites with 0 under tes hooks while 3 existed there; it now expands to 30 with 3 under tests/hook three pass. But the GUARDS themselv are still invoked by nothing. check-brie inputs.sh is honestly conductor-run rat than automatic: a brief's required-inpu list lives in free-form prose that the Ag tool's structured input does not expose an automatic prompt-scraping gate wo itself be an unproven pattern-match. It seam is the dispatch procedure, and th a documentation + habit change, not a hook. unattributed-commit-guard.sh is different and is why this item exists. R against real history it names 14 violati commits since the last tag (and 21 bar Auto-commit subjects exist across all r Wiring it into the pre-build or commit seam AS-IS would therefore refuse eve commit immediately, on pre-existing d nobody can fix in the moment. That is precisely the 11.4.224(E) brownfield-adoption shape, and 11.4.66/11.4.122 the choice with the operator, not with agent inventing a ratchet. The options stated so the decision is a decision and a default: (a) immediate hard floor – th seam refuses until all 14 are attributed one-time monotone-decrease ratchet (t 11.4.135 pattern) – snapshot 14 as the

ATM ID	Type	Status	Severity	Description
				baseline, the count may fall and may n rise, so day one is green and the diseas cannot spread; (c) changed-commits-or gate only commits created from now o with a scheduled deadline for the historical 14; (d) report-only for a state period, then escalate. Recommendatio with the reasoning rather than just the pick: (b). It is the pattern this constitut already uses for exactly this situation, makes the debt visible and bounded immediately, and it cannot be satisfied deleting the guard's name (11.4.227 cl that gaming channel). But it is the operator's call and this item is not clos until that answer is recorded. Honest boundary (11.4.6): this item does NOT claim the 14 commits are harmful in themselves – it claims their ATTRIBUT is absent and that nothing currently prevents the 15th. **Affected scope / fi scope manifest:** scripts/hooks/ unattributed-commit-guard.sh; scripts/ hooks/check-brief-inputs.sh; scripts/ pre_build_verification.sh; scripts/comm push-all.sh **Reproduction / context:** I guards run correctly by hand and are covered by passing tests (9/9 and 15/1 each proven non-bluff against a no-op stub). Neither is invoked by scripts/ pre_build_verification.sh or scripts/ commit-push-all.sh, so neither exerts standing detection pressure (11.4.226: guard with no execution seam is not coverage; registration is not coverage) **Acceptance criteria:** Each guard is invoked by a named seam, OR carries a registered deferral pointing at this iten For unattributed-commit-guard.sh specifically, the operator has answered adoption question below and the answ recorded as consumer DATA – never ar invented ratchet.
BOB-163	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:41:08Z **Reported-By:** BOB-114 remediation, p existing defect surfaced by BOB-111 landing a real limiter **What (the repo verbatim):** PRE-EXISTING, NOT INTRODUCED – and proven so rather

ATM ID	Type	Status	Severity	Description
				asserted: git diff shows assertion (c) is BYTE-IDENTICAL to HEAD, and the fa reproduces against the unmodified HE script. What changed is not the challen but the SYSTEM: BOB-111 appears at least partially delivered, because :718 now returns real 429s where it previou returned none. :7185 and :7189 still produce zero 429s. This is a 11.4.248-c corrosion risk and that is why it is filed rather than left as a footnote: run_all_challenges.sh will now fail INTERMITTENTLY, and an intermitten red trains everyone to re-run until gree at which point a real regression is dismissed as ‘probably the flaky rate-li one’. THE DECISION, which is an oper call under 11.4.66 and which I delibera did NOT invent: Does a 429 from a WORKING limiter count as the sibling endpoint being RESPONSIVE? (a) YES 429 proves the endpoint is alive and correctly protecting itself. Assertion (c accepts 2xx OR 429, and only a connec failure / 5xx / timeout counts as degrad Risk: a genuinely wedged endpoint tha happens to answer 429 would pass. (b – keep requiring 2xx, but make the challenge limiter-aware: drain or wait the window before the sibling probe (r after is served, so the budget is knowa rather than guessed). (c) Probe sibling a path the limiter exempts (a healthz-c route), so the isolation question is aske without spending the bucket. Recommendation with reasoning, not j a pick: (a) composed with a bounded liveness check – a 429 carrying a well-formed Retry-After IS evidence of a liv correctly-behaving service, and (b) ma the challenge slower and couples it to window value that will drift. But the semantics of ‘responsive’ here are a product judgement, so the answer is recorded, not assumed. RELATED, stat as fact rather than folded in: the detec counts 429 only, not 503, though the se header says ‘429 (or equivalent)’. Cour 503 would collide with the crash detec 5xx tally. Worth deciding when BOB-11

ATM ID	Type	Status	Severity	Description
				lands limiters on :7185 and :7189. **Affected scope / file-scope manifest** challenges/scripts/ ddos_resilience_challenge.sh assertion cross-endpoint isolation; run_all_challenges.sh which invokes it **Reproduction / context:** Assertion (c requires a sibling-endpoint probe to answer ^2 (a 2xx). :7187 now enforces real limiter (measured live: x-ratelimit- limit: 120, x-ratelimit-remaining: 86, re after: 45, server: uvicorn). A full challe run sends roughly 250 requests to :718 so sibling probes fired during the OTH endpoints' tiers land on an exhausted bucket and receive 429 – which asserti (c) scores as 'endpoint degraded'. Tim dependent: one run measured PASS=5 FAIL=2; the UNMODIFIED HEAD scri against the same live stack ~30s later measured PASS=7 FAIL=0 SKIP=2. **Acceptance criteria:** The operator ha answered the classification question b the answer is recorded as consumer D and assertion (c) implements it. The challenge then produces the same verd across 3 consecutive runs against a ra limited stack (11.4.50 deterministic consistency), and the fix ships a paired mutation proving assertion (c) still cat a genuinely degraded sibling endpoint narrowing it must not blind it (11.4.201(1)).
BOB-164	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:56:53Z **Reported-By:** BOB-110 UX-class coverage, discovered by the new axe-c suite on its first live run **What (the report, verbatim):** This is a REAL use facing defect, not a test-tuning artifact and it was found by the automated reg rather than by a human squinting at th page – which is exactly the 11.4.238 posture the project is aiming for. The static-grep oracle would NOT have fou it. Measured: the served root ships an empty pre-hydration, so any check rea the raw HTML audits a page nobody se The violation only exists in the hydrate DOM, which is why 11.4.170 requires

ATM ID	Type	Status	Severity	Description
				rendered oracle and forbids value-equa assertions as the proof a UI is correct. Severity reasoning, stated rather than assumed: this is user-visible and affect the primary dashboard heading, but it degrades legibility rather than breakin function, and the surface is operator-facing rather than public. Medium, not High. The failing test was left FAILING purpose (11.4.238) instead of silenced marked xfail. tests/ux/ currently report failed, 16 passed; that 1 is this defect. Anyone reading a red UX suite should it as this item, not as flakiness – and w this is fixed the suite goes fully green, which is the signal that it is closed. HONEST SCOPE LIMIT (11.4.6): only dashboard landing view was scanned. The /jackett/* sub-routes and the ng-se hosted :4200 route set were NOT audi the commands to close both are record in docs/testing/ux_accessibility.md. So item's 21 nodes are a floor, not a total. **Affected scope / file-scope manifest** the Angular dashboard served at http:/ localhost:7187/ (same compiled SPA as frontend/); production component CSS test files **Reproduction / context:** Ru tests/ux/ test_live_dashboard_accessibility.py against the running merge service. axe core v4.13.0, scanning a real Playwrig rendered JS-hydrated DOM, reports co contrast violations on 21 nodes. Measu pairs: .brand / h1 text #9d001e on background #3c3f41 = 1.43-1.62:1 (W AA large-text floor is 3:1); tagline #80 on #3c3f41 = 2.68:1 (body-text floor is 4.5:1). **Acceptance criteria:** axe-core reports ZERO color-contrast violations against the live rendered dashboard, w the fix made in production component rather than by relaxing the assertion o excluding the rule (11.4.120: reconcile the correct mechanism, never weaken check). The existing tests/ux/ suite is t guard and already fails today, so the R is captured – closure requires it flippin GREEN against the live surface, which runtime-class evidence per 11.4.226.

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	**Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T20:00:08Z **Reported-By:** surfaced while capturin closure evidence for BOB-092; diagnos rather than worked around **What (the report, verbatim):** This is a STALE-AI AFTER-INTERPRETER-UPGRADE defe not a missing package. The package is installed; its compiled half is built for t previous interpreter. That distinction matters because ‘pip install rpds-py’ st advice can appear to succeed while lea the 313 artefact in place. WHY IT WAS NOT NOTICED EARLIER, stated as fac the paths that DO work all avoid syster python3. ci.sh selects an interpreter through _select_python; the long-runni suites in this session ran .venv/bin/pytl -m pytest and passed. So the automate regime is green while the DOCUMENT operator command is broken – an 11.4.238-shaped gap, since the escape found by an agent capturing evidence rather than by the regime. WORKAROU (measured, not theorised): PYTEST_DISABLE_PLUGIN_AUTOLOA with the needed plugins passed explici p pytest_timeout ...) avoids the schemathesis autoload that drags in jsonschema -> referencing -> rpds. Th a workaround, NOT the fix: it silences import path rather than repairing the install, and it will surprise the next per who runs the documented command verbatim. RELATED, and worth decidir together rather than twice: BOB-154 wants .venv rebuilt on 3.12 to match tl container (which runs CPython 3.12.13 while both system and venv run 3.14.6 There are therefore THREE interprete play. Whoever rebuilds the venv shoulc confirm rpds resolves for the TARGET interpreter afterwards, or this same cl reappears one directory over. HONEST BOUNDARY (11.4.6): this item does N claim any test is wrong or any product code is broken. The product and the ve run suites are unaffected. What is brol is the documented entry point. **Affecte scope / file-scope manifest:** host use

ATM ID	Type	Status	Severity	Description
				site-packages (~/.local/lib/python3/site-packages/rpds/); every CLAUDE.md-documented ‘python3 -m pytest …’ invocation; any tooling that uses system python3 rather than .venv/bin/python **Reproduction / context:** python3 -m pytest tests/e2e/ test_live_stack_evidence.py -q –import-mode=importlib -> ModuleNotFoundError: No module nam ‘rpds.rpds’, raised during collection via schemathesis plugin autoload -> jsonschema -> referencing._core -> rp MEASURED root cause: system python 3.14.6, but ~/.local/lib/python3/site-packages/rpds/ ships rpds.cpython-313 x86_64-linux-gnu.so – an extension buil for 3.13. rpds/**init**.py does ’from .rpds import *‘, and a 313-tagged .so is not importable by 3.14, so the submodule genuinely does not exist for this interpreter. The venv is CORRECT and unaffected: .venv/lib64/python3/site-packages/rpds/ ships rpds.cpython-314 x86_64-linux-gnu.so and’.venv/bin/pyth -c import rpds’ succeeds. **Acceptance criteria:** python3 -m pytest tests/unit/ –import-mode=importlib – the comman CLAUDE.md documents – reaches collection without ModuleNotFoundEr OR CLAUDE.md is corrected to docum the supported runner explicitly. Either the documented command and the working command agree (11.4.99: a gu that misleads is the documentation-lay equivalent of a PASS-bluff).
BOB-166	Bug	Queued	High	WHAT. §11.4.19 requires closure migra to be ATOMIC: a resolving item moves Fixed.md, DISAPPEARS from Issues_Summary (open-only) and APPEARS in Fixed_Summary (closed-o Ten rows violate all three properties ri now — they carry a terminal status yet current_location=‘Issues’, so they rend in docs/Issues.md and are listed in doc Issues_Summary.md with a status that literally reads ‘(→ Fixed.md)’, while appearing 0 times in Fixed_Summary.n Any reader of the canonical open track told these are open work. MANIFEST

ATM ID	Type	Status	Severity	Description
				(measured 2026-08-21): BOB-087 (Completed), BOB-129, BOB-131, BOB-144, BOB-145, BOB-146, BOB-14 BOB-153, BOB-155, BOB-157 (all 'Fixe Fixed.md)'). Confirmed rendering: BOB-155/087/157 each grep 1 in Issues.md, 0 in Fixed.md, present in Issues_Summary.md, 0 hits in Fixed_Summary.md. ROOT CAUSE (reproduced at runtime, not inferred). 'close' performs the atomic migration a REQUIRES –evidence. 'update –status' any §11.4.15 closed-set value with NO evidence and NO migration — and it knows the location, because it prints it a COPY of the real DB: $ workable-iten update –id BOB-065 –db repro.db –stat 'Fixed (→ Fixed.md)' update: BOB-065 updated in Issues (status=Fixed (→ Fixed.md), type=Task) $ sqlite3 repro. 'SELECT atm_id,status,current_locatio BOB-065|Fixed (→ Fixed.md)|Issues Sc seam that is supposed to be the ONLY closure path (close, evidence-gated pe §11.4.146(D3)) has a parallel unguarde path that reaches the same status whil skipping BOTH the evidence requirem and the migration. WHY NOTHING CAUGHT IT (§11.4.238 coverage-escap audit). Three standing checks stay GR on a DB holding the forbidden row, ver on the poisoned copy from the repo roo no cwd artifact is involved: (1) 'validat 'OK — 164 items, all invariants satisfie (it has no status↔location coherence invariant); (2) 'diff' → 'in sync' (DB and Markdown agree — on the WRONG sta agreement is not correctness); (3) CM- CLOSURE-SEAM-BINDS CHECK A → I (it flags NON-terminal rows whose id appears in a work commit; these rows terminal, so they are outside its predic by construction). This was found by reading a status tally, not by the automated regime — a §11.4.238 discovery-channel escape, which is itse defect of equal standing to the mis-loca rows. ACCEPTANCE. (a) 'update' REFUSES a terminal status and names 'close' as the correct path, with a paire

ATM ID	Type	Status	Severity	Description
				§1.1 mutation proving the refusal (removing the guard must make the mutation pass). (b) ‘validate’ grows a status↔location coherence invariant th FAILS on the forbidden state, with a golden-bad fixture and a negative cont (a legitimately terminal row in Fixed m NOT fire — §11.4.201(1)). (c) The 10 existing rows are drained to Fixed with class-matched evidence per row, or, wh a row’s evidence cannot be produced, honestly re-opened rather than migrat on a bare assertion. (d) Honest bounda this closes the update-path hole and th detection gap; it does not claim every historical status write was evidence-backed. NOT CLAIMED. No fix is implemented by this filing. The 10 row are untouched; draining them is acceptance (c) and each needs its own evidence, not a bulk UPDATE.
BOB-167	Bug	Queued	Medium	WHAT. download-proxy exposes two Server-Sent-Events routes. They are th same expensive class — long-lived connections that hold a worker and a generator for their lifetime — but only is rate-limit classed: routes.py:801 @router.get(‘/search/stream/{search_i @_rl(‘sse_stream’) <- classed routes.py:150 @router.get(‘/theme/stream’) async def stream_theme(...) < NO limiter decorator _rl(cls) resolves t limiter.limit(limit_for(cls)), so /search/stream is bound to the sse_stream clas while /theme/stream falls through to th application default. Measured live on t running stack: /api/v1/theme/stream reports x-ratelimit-limit 120. PROVENANCE + A CORRECTION WO KEEPING (§11.4.6). This was surfaced the BOB-109 scaling agent, whose rep framed it as: ‘sse_stream_limit_decora is defined and imported/applied nowhe ... SSE falls through to the 120/minute default: 24x less protected’. That fram is WRONG and was NOT filed as given The agent searched for one symbol NA the wiring uses a different mechanism (@_rl(‘sse_stream’)), and it IS applied to /search/stream. Verified by reading

ATM ID	Type	Status	Severity	Description
				routes.py:795-806 and the _rl helper a routes.py:46-51, with a control needle confirming other *_limit_decorator sym show real usage in api/**init**.py so the ze hit was not a blind search. The agent's MEASUREMENT (120 on /theme/strea was correct and is what makes this a r finding; its MECHANISM was not. Bot halves are recorded so the next reader does not re-derive the same wrong cau WHY IT MATTERS. An unclassed SSE endpoint is the cheapest way to pin se resources: each connection is held ope and the default class permits 120/min them. The declared sse_stream class e precisely because this route shape nee tighter bound than ordinary GETs. ACCEPTANCE. (a) A decision, recorde whether /theme/stream belongs in the sse_stream class or genuinely warrant default — this is a policy question, not automatically a bug to patch. (b) If it belongs in sse_stream, the decorator is applied and a test drives BOTH SSE ro and asserts each returns its INTENDE class limit from x-ratelimit-limit, so a future route added without a class is caught. (c) A guard that enumerates S shaped routes and fails on any that ca no explicit rate-limit class — the gene form, so the third SSE route does not repeat this. (d) Honest boundary: this not claim 120/min is exploitable in this deployment; with network_mode host a no reverse proxy every caller shares th 127.0.0.1 bucket, which BOB-111 measured and recorded separately. NO CLAIMED. No change made. The limit values were read from headers, never driven to exhaustion — the limiter is p and shared with concurrent agents on host (§11.4.119).
BOB-168	Task	Queued	Low	WHAT. scripts/run_all_challenges.sh:6 lists "scaling_horizontal_challenge.sh" its challenge set. challenges/scripts/ scaling_horizontal_challenge.sh does n exist: $ sed -n '66p' scripts/ run_all_challenges.sh "scaling_horizontal_challenge.sh" $ ls challenges/scripts/

ATM ID	Type	Status	Severity	Description
				scaling_horizontal_challenge.sh ls: can access …: No such file or directory VERIFIED independently, not taken on report — surfaced by the BOB-109 sca agent and re-checked here by direct invocation. WHY IT MATTERS. Whethe this is cosmetic or a §11.4.201 gate-honesty defect depends entirely on how the runner treats a missing entry, and is the first thing to determine: if it SKI silently, the challenge bank advertises coverage it does not have (a §11.4.266 claim-vs-reality row with no passing challenge behind it); if it FAILs, the ru is permanently red for a reason unrela to the system under test, which trains readers to ignore it. Neither outcome i acceptable; they need different fixes. ACCEPTANCE. (a) Determine and reco the runner's actual behaviour on the missing entry by invoking it, not by reading it. (b) EITHER author the challenge, OR remove the entry — with §11.4.124 discipline: check git history whether it once existed and was delete since a silently-dropped challenge is th more interesting defect. (c) If the runn silently skips missing entries, that is it own finding: a missing challenge must loud (§11.4.3 SKIP-with-reason at minimum), never absent-and-quiet. SEVERITY. Low as a defect, but it sits the challenge-coverage seam, so (c) ma deserve its own item.
BOB-169	Bug	Queued	Medium	WHAT. The §11.4.65 markdown-export mandate requires every in-scope doc to ship .html/.pdf twins. 286 of the 326 .h files under docs/ (excluding dist/ and node_modules/) are pandoc FRAGMEN — they begin at
BOB-170	Task	Queued	Medium	WHAT. The BOB-109 scaling suite gate two timing tests on loadavg_1m/nproc 0.75 — on this 8-cpu host, a 1-minute l average at or below 6.0 — and SKIPs loudly otherwise. That gate is correct a was adopted for a good reason (§11.4.201(8)): the author validated a 2 growth-exponent threshold, then obse it FALSE-FAIL correct code at load 11. (baseline span 2.136), and at load 15-2

ATM ID	Type	Status	Severity	Description
				measured baseline spans of 2.358 agai injected-cubic mutants at 2.278-2.331 the signal is smaller than the noise and two are not separable in either directio Gating was the honest response. THE PROBLEM. The gate has consequently NEVER RUN under its own preconditio Every recorded run today measured lo 9.15-23; the independent reviewer measured 9.66 during review; the conductor measured 6.45 at its lowest. GREEN polarity WAS observed — at lo 9.4-11.6 (span 1.882) and in three stab runs (1.759/1.832/1.848) — but never beneath the shipped ceiling. The autho states this honestly as an inference fro data rather than a run they can paste, the reviewer confirmed the inference i sound (quiescence is strictly more favourable). Sound inference is still no observed run. WHY THAT MATTERS (§11.4.226). A registered, topology-pre guard that never executes is indistinguishable from a guard that wo fail if it did. Nothing currently ensures ever runs: no freshness contract, no scheduled quiet-window invocation, no tracked item — which is precisely the unexecuted-standing-guard class §11.4 names, where registration gets treated coverage while execution is nobody's j The forensic precedent in that anchor two standing guards that, when finally executed, immediately emitted latent FAILs. ACCEPTANCE. (a) Execute both quiescence-gated timing tests on a hos whose 1-minute loadavg is genuinely < 0.75/cpu — no agent fleet, no concurre build — and capture the run as an arti showing the resolved load reading alongside the measured span exponent the precondition is provable from the artifact and not asserted. (b) Record th observed GREEN polarity in the item a in the test source, replacing the curren honest-inference note. (c) If the gate F under genuine quiescence, that is the finding this item exists to surface, and reopens the threshold-calibration ques rather than being worked around. (d)

ATM ID	Type	Status	Severity	Description
				Establish a freshness contract per §11.4.226: how stale a quiescent verdi may be before the gate counts as uncovered again, so this does not silen return to never-executing. HONEST BOUNDARY. This item does not claim t gate is wrong. The evidence points the other way — quiescence is strictly mor favourable than the loads where GREE was already observed. It claims only th the run has not happened, and that an unexecuted guard cannot be cited as coverage. PROVENANCE. Raised as IMPORTANT-2 by the independent Fab substrate reviewer of the BOB-109 wo and independently corroborated: the reviewer verified the SKIP is loud and prints its resolved numbers, and confir the shipped tests do SKIP at today's lo
BOB-171	Bug	Queued	Low	WHAT. Both rate limiters key their per client bucket on the LEFTMOST eleme of X-Forwarded-For when TRUST_FORWARDED_FOR is enabled: download-proxy/src/api/ rate_limit.py:117-123 (:7187) and the r stdlib limiter in plugins/download_prox (:7186), which was deliberately built to exact policy parity with it. That elemen CLIENT-SUPPLIED. An honest reverse proxy APPENDS the peer address to th RIGHT of whatever arrived, so the left entry is whatever the client sent — attacker-controlled by construction, no misconfiguration. Consequences with t flag on: (1) rotating a forged leftmost v mints an unlimited sequence of fresh buckets, which is total bypass of the lin the flag is supposed to make MORE accurate; (2) it also defeats the LRU/id reap cap, since each forged identity is distinct key — the bucket map is a memory-growth surface bounded only the cap, and the cap's eviction then discards the budgets of REAL clients to make room for forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, that is why this is Low, not High. TRUST_FORWARDED_FOR defaults OI at both sites, and the deployed stack ru network_mode host with no reverse pr

ATM ID	Type	Status	Severity	Description
				in front (verified across all five compos services), so peer addresses are real cl IPs and no NAT collapse occurs. The de is latent: it arms the moment someone puts a proxy in front and turns the flag to recover real client IPs — which is exactly the situation the flag exists for. failure mode is therefore ‘correct-looki configuration change silently disables limiter’, not ‘currently broken’. PROVENANCE. Raised as MINOR-3 by independent reviewer of the BOB-111 limiter work and independently confirm by the implementing agent, which note at BOTH source sites as a tracked follo up rather than fixing it in that change. Filed here because neither agent has w access to the tracker. WHY IT COVERS BOTH PORTS. :7186’s limiter was buil deliberate policy parity with :7187’s, including this parsing. Fixing one and the other would leave the two surfaces disagreeing on client identity — worse than the shared defect, because it becomes surface-dependent and untestable as a single invariant. ACCEPTANCE. (a) Replace leftmost-element parsing with a trust-aware resolution at BOTH sites: either rightm minus-N-trusted-hops, or a trusted-pro CIDR allowlist where only a peer inside the allowlist may assert XFF at all — th choice is a deployment-topology decisi and should be recorded, not assumed. A test that forges a rotating leftmost element and asserts the bucket key do NOT rotate, per site. (c) Paired §1.1 mutation: restore leftmost parsing and test must FAIL. (d) Negative control (§11.4.201(1)): with TRUST_FORWARDED_FOR OFF, the header must be ignored entirely and keying must fall back to the peer addre — a guard that starts honouring XFF w the flag is off would be a worse defect the one being fixed. (e) Honest bounda this closes header-forgery bypass; it do not make per-IP limiting fair behind NA where many real users legitimately sha one address. NOT CLAIMED. No chang

ATM ID	Type	Status	Severity	Description
				made. Both sites still parse leftmost; th flag still defaults off.
BOB-172	Bug	Queued	High	WHAT. The rutracker SEARCH endpoi refused by bot protection while the site itself is fully reachable. Measured unauthenticated 2026-08-21 (no cooki read, no credential value in any artifac §11.4.10): GET /forum/index.php -> 20 96283 B,