

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	8
Bug	Operator-blocked	1
Bug	Queued	19
Bug	Ready for testing	4
Task	In progress	10
Task	Operator-blocked	3
Task	Queued	25
Task	Ready for testing	2
TOTAL		72

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-hos routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in proxy + qBitTorrent-go + Jackett + com forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel — ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
BOB-077	Task	Operator-blocked	High	

ATM ID	Type	Status	Severity	Description
				RD2-10: Identify second host running the commit rsync/sync mechanism (OPERATOR DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code re the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable the four Lava-porting findings (closes G
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes C
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Trac + Status Documents table row-completo (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exerc three start.sh subcommands end-to-end real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify ReDoS regex bounds deployed to contain (closes runtime half of GA-12)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos. side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for download-proxy/merge endpoints (canonical of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is -profile go parity still a goal? (gates RW-10..13) — OPERATOR-I
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — b 127.0.0.1 or keep 0.0.0.0? — OPERATOR DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence helper for the unattributed auto-commit
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence for subagent task-brief source inputs
BOB-109	Task		Medium	

ATM ID	Type	Status	Severity	Description
		Ready for testing		BOB-074 followup: scaling-class test coverage absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiters on boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golden fixture for the rate-limit detector
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout architectural gap (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_creation bulk suite (Permission denied /config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while the process still serves 7186 (GIL starvation spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile server serves 7186/7187/7188 but its container binds to 7187 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolved gitdir) pollute gate scan scope and manifest false BOB-126-class findings
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/41 from constitution, CLAUDE.md and the REALITY badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labels 44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate implementation, and 24 of 36 scripts have companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored third-party code: 82% of one gate's 38,291 findings from submodules/
BOB-154	Bug	Ready for testing	Medium	Host venv and production container run different starlette versions (1.4.1 vs 1.6.0)
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard is too sensitive and flaky: 8786ms under host 900-1500ms ceiling calibrated on a quiet machine
BOB-159	Bug	Queued	Medium	Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:01:17Z Reported-By: independent review (IMPORTANT-2), partially remediated What (the report, verbatim): -recreate path was fixed this round: run_ownership_precondition -> stack_down, run_ownership_repair -> stack_up, so the

ATM ID	Type	Status	Severity	Description
				walks a tree no container can write to. The ordering is now asserted behaviourally with new checks in tests/unit/test_start_reload_recreate.sh, each killed and paired 1.1 mutation (11/11 RED-OK). The path is NOT covered by that fix and this tracks the remainder. On a warm ./start gate runs before THIS invocation brings up, but if the stack is ALREADY running containers write while the repair walks. A container can then create a new non-owned file BEHIND the walk, after which a completion marker records complete ownership that is not, and those stragglers are never repaired without a manual -force. BOUQUET dismissed: after any successful pass the marker is valid for the scope fingerprint and the repair short-circuits without walking (marker_is_valid so the window requires a stale-or-absent marker AND a live stack together. Declaring a new entry re-arms the marker, which is exact to that combination arises in practice. NOT THIS ROUND, with the reason stated rather hidden: the clean guard is a cheap marker probe mode on scripts/ownership_repair.sh answers has-this-scope-been-repaired while walking the tree. The existing -dry-run serve: it deliberately SKIPS the marker (FORCE=0 && DRY_RUN=0 guard) and walks, so using it as a warm-start probe walk the tree twice on every start. Adding probe mode requires editing ownership_repair.sh another stream was actively editing during the round; duplicating marker_is_valid into scripts/ownership_repair.sh instead would be a near-identical fork (Deferred to this item rather than done unilaterally or silently. The misleading comment at the warm-start call site (which claimed gate runs before any container writes) has been corrected to state this boundary honestly. Affected scope / file-scope manifest: (warm-start dispatch, run_ownership_guard site); scripts/ownership_repair.sh (marker path) Reproduction / context: 1. Ensure stack is UP. 2. Change config/owned_path so the scope fingerprint changes (this requires marker by design, data-model E2). 3. Run warm ./start.sh (no -recreate). 4. The ownership repair walks and chowns the declared tree. Containers are still running and able to

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	it. Acceptance criteria: A warm ./start complete an ownership repair while a c that writes to a declared location is run. Either the repair is deferred with an act refusal naming -recreate, or the stack i for the walk. Asserted behaviourally in test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN REPAIR_AFTER_DOWN / REPAIR_BEFO checks, each with a paired 1.1 mutation it. Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:02:48Z Reported-By: What (the report, verbatim): Independ §11.4.209 review (IMPORTANT-3) found the feature's strongest automated check never executed by any runner: (1) tests/ownership/test_container_writes_owned — the §11.4.115 RED-turned-regression FR-011/FR-007, proven via a docker-con revert mutation — was moved out of tes integration/ (commit 58d340a, to escap autouse fixture hang) but no runner (ci. test.sh, run-all-tests.sh) was ever extend cover tests/ownership/. (2) tests/pre_bu test_.sh, the §1.1 paired-mutation meta- the scripts/pre_build/check_cm_.sh gate was executed by NOTHING — self-repo. honestly in commit 04742d7's own mes ('STILL OPEN (reported, not fixed): test pre_build/ is executed by NOTHING') b tracked as a workable item (a §11.4.197 requirements gap) and never wired into pre_build_verification.sh invariant 30, w globbed only tests/unit/test_.sh. Affecte file-scope manifest: ci.sh; scripts/ pre_build_verification.sh invariant 30 (C UNIT-TESTS-EXECUTED) Reproductio context: grep -rn test_container_writes_owned_files ci.sh run-all-tests.sh scripts/ returned zero ru (only docs/specs mentioned the path); g scripts/pre_build_verification.sh showed 30's for-loop globbing only tests/unit/te never tests/pre_build/test_.sh. Accepta criteria: ci.sh gains a runtime-gated py ownership/ stage (skips honestly, rc=0, podman/docker present); scripts/ pre_build_verification.sh invariant 30's covers both tests/unit/test_.sh and tests

ATM ID	Type	Status	Severity	Description
BOB-161	Task	Queued	High	pre_build/test_*.sh, verified by an extra standalone run of the invariant's exact l reporting a non-zero RAN count that inc files from both directories. Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:31:12Z Reported-By BOB-092 remediation, residual finding report, verbatim): §11.4.69 names CM OPEN-SKIP as one of three mandatory p gates: it audits sink-side probe helpers . if any code path converts an empty or unreachable response into a PASS-count for a feature class with a sink-side prob project does not have it. Why this matte rather than in the abstract: the BOB-09 remediation just removed TWO fail-oper exactly this class from tests/e2e/ test_live_stack_evidence.py — the nnmc on-404 (already removed at 7baef2b) an surviving sibling at test_iporrents_is_authenticated_in_sea skipped on 'not authenticated and statu success' while asserting 'treating as tra outage'. That assertion was false for rej credentials (upstream_http_403) and fo container (plugin_env_missing), both of definitive product failures the product a classifies via error_type. Two fail-opens class, found one at a time, is the §11.4.1 extend-to-all-cases signal and the §11.4 coverage-escape signal together: the re not surface either, an agent reading cod remediation's own guards are AST-struct lethal under three discriminating mutat they LIVE IN THE FILE THEY GUARD, deleting that file evades them entirely. A gate is the durable home because it auc corpus rather than one file. Honest bou (§11.4.6): this item does NOT claim the remediated fail-opens are unguarded — guarded, with runtime evidence (13 pas 97.36s against the live stack). It claims CLASS has no corpus-wide detector, so instance in a different file is invisible ag Affected scope / file-scope manifest: pre_build/ (gate absent); tests/e2e/ test_live_stack_evidence.py (guards cur inside the file they guard) Reproductio context: grep -rl 'CM-NO-FAIL-OPEN-S scripts/ tests/ returns exactly ONE hit, a

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	tests/e2e/test_live_stack_evidence.py — the guards live in, not a gate. Control no same query for CM-OWNERSHIP-INVALID (gate that does exist) returns 3 files, so the instrument can see gate tokens and there is a real absence, not a blind zero (§11.4.107(10)). Acceptance criteria: A scripts/pre_build_verification gate named CM-NO-FAIL-OPEN-SKIP is wired into scripts/pre_build_verification audits sink-side probe helpers for code converting an empty/unreachable/error into a PASS-counting SKIP for a feature HAS a sink-side probe. It ships a golden fixture (a real fail-open -> gate FIRES) golden-FALSE-with-carrier (an honest geo skip, and a comment merely MENT the phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §1.1 makes the gate FAIL before the gate is Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:35:02Z Reported-By: BOB-106/BOB-107 remediation, residual (the report, verbatim): The TESTS for guards are now executed: pre-build invalid glob was extended from tests/unit + tests/pre_build to also cover tests/hooks. BEFORE/after: the glob expanded to 27 sub 0 under tests/hooks while 3 existed there expands to 30 with 3 under tests/hooks. pass. But the GUARDS themselves are not invoked by nothing. check-brief-inputs.sh honestly conductor-run rather than auto brief's required-input list lives in free-form that the Agent tool's structured input does expose, so an automatic prompt-scraping would itself be an unproven pattern-matching seam is the dispatch procedure, and the documentation + habit change, not a honest unattributed-commit-guard.sh is different why this item exists. Run against real names 14 violating commits since the last (and 21 bare Auto-commit subjects exist all refs). Wiring it into the pre-build or seam AS-IS would therefore refuse even immediately, on pre-existing debt nobody in the moment. That is precisely the 11.4.66 brownfield-adoption shape, and 11.4.66 put the choice with the operator, not with agent inventing a ratchet. The options, the decision is a decision and not a default

ATM ID	Type	Status	Severity	Description
BOB-163	Bug	Queued	Medium	immediate hard floor - the seam refuses 14 are attributed; (b) one-time monotonic decrease ratchet (the 11.4.135 pattern) snapshot 14 as the baseline, the count may and may never rise, so day one is green disease cannot spread; (c) changed-commit - gate only commits created from now on scheduled deadline for the historical 14 report-only for a stated period, then escape Recommendation, with the reasoning ratchet just the pick: (b). It is the pattern this commit already uses for exactly this situation, if the debt visible and bounded immediately cannot be satisfied by deleting the guard (11.4.227 closes that gaming channel). the operator's call and this item is not covered until that answer is recorded. Honest bob (11.4.6): this item does NOT claim the 14 commits are harmful in themselves - it is their ATTRIBUTION is absent and that is currently prevents the 15th. Affected scope file-scope manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-bad-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh Reproduction context: Both guards run correctly by hand are covered by passing tests (9/9 and 15 proven non-bluff against a no-op stub). 14 invoked by scripts/pre_build_verification.sh; scripts/commit-push-all.sh, so neither existing standing detection pressure (11.4.226: 14 with no execution seam is not coverage registration is not coverage). Acceptance criteria: Each guard is invoked by a new seam, OR carries a registered deferral point this item. For unattributed-commit-guard specifically, the operator has answered adoption question below and the answer recorded as consumer DATA - never an ratchet. Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:41:08Z Reported-By: BOB-114 remediation, pre-existing defect surfaced by BOB-111 landing a real limit. (the report, verbatim): PRE-EXISTING INTRODUCED - and proven so rather than asserted: git diff shows assertion (c) is IDENTICAL to HEAD, and the failure reverts against the unmodified HEAD script. What changed is not the challenge but the SY

ATM ID	Type	Status	Severity	Description
				BOB-111 appears at least partially deliv because :7187 now returns real 429s w previously returned none. :7185 and :71 produce zero 429s. This is a 11.4.248-cl corrosion risk and that is why it is filed than left as a footnote: run_all_challeng now fail INTERMITTENTLY, and an inte red trains everyone to re-run until gree which point a real regression is dismiss 'probably the flaky rate-limit one'. THE DECISION, which is an operator call un 11.4.66 and which I deliberately did NO Does a 429 from a WORKING limiter co sibling endpoint being RESPONSIVE? (429 proves the endpoint is alive and co protecting itself. Assertion (c) accepts 2 429, and only a connection failure / 5xx counts as degraded. Risk: a genuinely w endpoint that happens to answer 429 w (b) NO - keep requiring 2xx, but make t challenge limiter-aware: drain or wait o window before the sibling probe (retry-served, so the budget is knowable rathe guessed). (c) Probe siblings on a path th exempts (a healthz-class route), so the i question is asked without spending the Recommendation with reasoning, not ju (a) composed with a bounded liveness o 429 carrying a well-formed Retry-After evidence of a live, correctly-behaving se (b) makes the challenge slower and cou window value that will drift. But the ser 'responsive' here are a product judgem answer is recorded, not assumed. RELA stated as fact rather than folded in: the counts 429 only, not 503, though the scr header says '429 (or equivalent)'. Count would collide with the crash detector's Worth deciding when BOB-111 lands lin :7185 and :7189. Affected scope / file-manifest: challenges/scripts/ddos_resilience_challenge.sh assertion endpoint isolation; run_all_challenges.s invokes it Reproduction / context: As requires a sibling-endpoint probe to ans 2xx). :7187 now enforces a real limiter live: x-ratelimit-limit: 120, x-ratelimit-re 86, retry-after: 45, server: uvicorn). A fu challenge run sends roughly 250 request :7187, so sibling probes fired during the

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	In progress	Medium	endpoints' tiers land on an exhausted b receive 429 - which assertion (c) scores 'endpoint degraded'. Timing-dependent measured PASS=5 FAIL=2; the UNMO HEAD script against the same live stack later measured PASS=7 FAIL=0 SKIP=. Acceptance criteria: The operator has the classification question below, the an recorded as consumer DATA, and assert implements it. The challenge then prod same verdict across 3 consecutive runs rate-limited stack (11.4.50 deterministic consistency), and the fix ships a paired mutation proving assertion (c) still catc genuinely degraded sibling endpoint - n it must not blind it (11.4.201(1)). Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:56:53Z Reported-By: BOB-110 UX-class coverage, discovered new axe-core suite on its first live run report, verbatim): This is a REAL user defect, not a test-tuning artifact, and it by the automated regime rather than by squinting at the page - which is exactly 11.4.238 posture the project is aiming f static-grep oracle would NOT have foun Measured: the served root ships an emp hydration, so any check reading the raw audits a page nobody sees. The violation exists in the hydrated DOM, which is w 11.4.170 requires a rendered oracle and value-equality assertions as the proof a correct. Severity reasoning, stated rath assumed: this is user-visible and affects primary dashboard heading, but it degr legibility rather than breaking function, surface is operator-facing rather than p Medium, not High. The failing test was FAILING on purpose (11.4.238) instead silenced or marked xfail. tests/ux/ curre reports 1 failed, 16 passed; that 1 is thi. Anyone reading a red UX suite should r this item, not as flakiness - and when th the suite goes fully green, which is the s it is closed. HONEST SCOPE LIMIT (11 the dashboard landing view was scanne jackson/* sub-routes and the ng-serve-ho :4200 route set were NOT audited - the commands to close both are recorded in testing/ux_accessibility.md. So this item

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	nodes are a floor, not a total. Affected s file-scope manifest: the Angular dashb served at http://localhost:7187/ (same c SPA as frontend/); production componen not test files Reproduction / context: ux/test_live_dashboard_accessibility.py the running merge service. axe-core v4. scanning a real Playwright-rendered JS DOM, reports color-contrast violations o nodes. Measured pairs: .brand / h1 text on background #3c3f41 = 1.43-1.62:1 (l large-text floor is 3:1); tagline #808080 #3c3f41 = 2.68:1 (body-text floor is 4.5 Acceptance criteria: axe-core reports color-contrast violations against the live dashboard, with the fix made in product component CSS rather than by relaxing assertion or excluding the rule (11.4.12 reconcile to the correct mechanism, nev weaken the check). The existing tests/u the guard and already fails today, so the captured - closure requires it flipping G against the live surface, which is runtim evidence per 11.4.226. Reported-Via: §11.4.202 reporting dire on 2026-08-21T20:00:08Z Reported-By surfaced while capturing closure eviden BOB-092; diagnosed rather than worke What (the report, verbatim): This is a ABI-AFTER-INTERPRETER-UPGRADE o a missing package. The package is insta compiled half is built for the previous in That distinction matters because 'pip in py' style advice can appear to succeed v leaving the 313 artefact in place. WHY NOT NOTICED EARLIER, stated as fact paths that DO work all avoid system pyt ci.sh selects an interpreter through _select_python; the long-running suites session ran .venv/bin/python -m pytest a passed. So the automated regime is gre the DOCUMENTED operator command - an 11.4.238-shaped gap, since the esc found by an agent capturing evidence r by the regime. WORKAROUND (measur theorised): PYTEST_DISABLE_PLUGIN_AUTOLOAD the needed plugins passed explicitly (-p pytest_timeout ...) avoids the schemath autoload that drags in jsonschema -> re

ATM ID	Type	Status	Severity	Description
				-> rpds. That is a workaround, NOT the silences the import path rather than rep install, and it will surprise the next pers runs the documented command verbatim RELATED, and worth deciding together than twice: BOB-154 wants .venv rebuild to match the container (which runs CPython 3.12.13 while both system and venv run There are therefore THREE interpreters. Whoever rebuilds the venv should configure resolves for the TARGET interpreter after or this same class reappears one directly HONEST BOUNDARY (11.4.6): this item NOT claim any test is wrong or any product is broken. The product and the venv-run are unaffected. What is broken is the documented entry point. Affected scope / file-scope manifest: host user site-packages (~/.local/lib/python3/site-packages/rpds/); every CLAUDE documented 'python3 -m pytest ...' involves any tooling that uses system python3 rather than .venv/bin/python Reproduction / python3 -m pytest tests/e2e/test_live_stack_evidence.py -q -importmode=importlib -> ModuleNotFoundError: module named 'rpds.rpds', raised during collection via the schemathesis plugin and -> jsonschema -> referencing. core -> 1 MEASURED root cause: system python3 but ~/.local/lib/python3/site-packages/rpds.cpython-313-x86_64-linux-gnu.so - extension built for 3.13. rpds/init.py does 'from .rpds import *', and a 313-tagged importable by 3.14, so the submodule does not exist for this interpreter. The venv CORRECT and unaffected: .venv/lib64/python3.13/site-packages/rpds/ ships rpds.cpython-313-x86_64-linux-gnu.so and '.venv/bin/python3 -m pytest tests/unit/ -v -importmode=importlib - the command CLAUDE documents - reaches collection without ModuleNotFoundError, OR CLAUDE.mn corrected to document the supported run explicitly. Either way the documented command and the working command agree (11.4.6). The guide that misleads is the documentation (equivalent of a PASS-bluff). WORKAROUND EFFECTS MEASURED 2026-08-21 — the product is not free, and the item previously imported

ATM ID	Type	Status	Severity	Description
				PYTEST_DISABLE_PLUGIN_AUTOLOAD disables ALL plugin autoload, not just the one, so anything the suite silently relied on can be re-enabled by hand: - It BREAKS 5 plugin-dependent tests in tests/unit/test_plugin_rutracker.py — TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace need an autoloaded env plugin. PROVEN TO BE CAUSED BY ANY in-flight change: running the OWN copy of that file under identical flags produces the same 5 failures. - It makes -timeout=60 (set in pyproject.toml) an UNKNOWN ARGUMENT unless -p pytest is passed explicitly, so pytest.mark.timeout is silently inert under the naive recipe — a test believed to be time-bounded is not. - With autoload ON and only schemathesis disabled (no:schemathesis), that same file is 96/99 green. CONSEQUENCE FOR ANYONE USING PYTEST: -p no:schemathesis alone is strictly better than blanket autoload-disabled, where it suffices, because it removes one broken plugin. Where the blanket form is used, -p pytest timeout must be added or timeout is inert, and the 5 env-dependent failures are recognised as workaround artifacts rather than filed as defects. That last point is the reason the workaround manufactures failures that are exactly like product defects. This does not fix the root cause (the venv's rpds ships a cpython-313 ABI tag CPython 3.14 cannot handle) or the fix (rebuild the venv — BOB-154). The documents that the interim recipe has a small radius, so nobody reads a workaround as a regression. WHAT. download-proxy exposes two Server-Sent Events routes. They are the same expensive — long-lived connections that hold a weakref generator for their lifetime — but only one is limit classed: routes.py:801 @router.get('/sse/{search_id}') @_rl('sse_stream') @_rl('sse_stream') classed routes.py:150 @router.get('/the/{search_id}') async def stream_theme(...) <- @_rl('sse_stream') decorator _rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/stream is bound to the sse_stream class while /the/stream is bound to the stream class.
BOB-167	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				stream falls through to the application c Measured live on the running stack: /ap theme/stream reports x-ratelimit-limit 1 PROVENANCE + A CORRECTION WOF KEEPING (§11.4.6). This was surfaced b BOB-109 scaling agent, whose report fr as: 'sse_stream_limit_decorator is defin imported/applied nowhere ... SSE falls t the 120/minute default: 24x less protect framing is WRONG and was NOT filed a The agent searched for one symbol NAM wiring uses a different mechanism (@_rl('sse_stream')), and it IS applied – search/stream. Verified by reading routes.py:795-806 and the _rl helper at routes.py:46-51, with a control needle c other *_limit_decorator symbols show re in api/init.py so the zero-hit was not a b search. The agent's MEASUREMENT (1 theme/stream) was correct and is what this a real finding; its MECHANISM wa halves are recorded so the next reader c re-derive the same wrong cause. WHY I MATTERS. An unclassified SSE endpoint cheapest way to pin server resources: e connection is held open, and the default permits 120/min of them. The declared sse_stream class exists precisely becaus route shape needs a tighter bound than GETs. ACCEPTANCE. (a) A decision, rec whether /theme/stream belongs in the s class or genuinely warrants the default policy question, not automatically a bug (b) If it belongs in sse_stream, the deco applied and a test drives BOTH SSE rou asserts each returns its INTENDED clas from x-ratelimit-limit, so a future route without a class is caught. (c) A guard th enumerates SSE-shaped routes and fail that carries no explicit rate-limit class – general form, so the third SSE route do repeat this. (d) Honest boundary: this d claim 120/min is exploitable in this depl with network_mode host and no reverse every caller shares the 127.0.0.1 bucke BOB-111 measured and recorded separ CLAIMED. No change made. The limit v were read from headers, never driven t exhaustion — the limiter is per-IP and s with concurrent agents on this host (§1

ATM ID	Type	Status	Severity	Description
BOB-168	Task	In progress	Low	WHAT. scripts/run_all_challenges.sh:66 “scaling_horizontal_challenge.sh” in its set. challenges/scripts/scaling_horizontal_challenge.sh does not exist. I tried to run it with sed -n ‘66p’ scripts/run_all_challenges.sh. “scaling_horizontal_challenge.sh” \$ ls -l scripts/scaling_horizontal_challenge.sh. access ...: No such file or directory VERIFIED. This is a defect, but it is not independently, not taken on report — see the BOB-109 scaling agent and re-checked by direct invocation. WHY IT MATTERS. This is cosmetic or a §11.4.201 gate-honour defect depends entirely on how the runner finds a missing entry, and that is the first thing to determine: if it SKIPS silently, the challenge advertises coverage it does not have (a claim-vs-reality row with no passing challenge behind it); if it FAILS, the runner is permitted for a reason unrelated to the system test, which trains readers to ignore it. Neither outcome is acceptable; they need different ACCEPTANCE. (a) Determine and record the runner’s actual behaviour on the missing challenge by invoking it, not by reading it. (b) EITHER remove the challenge, OR remove the entry — v §11.4.124 discipline: check git history for whether it once existed and was deleted. A silently-dropped challenge is the more important defect. (c) If the runner silently skips missing entries, that is its own finding: a missing challenge must be loud (§11.4.3 SKIP-warnings at minimum), never absent-and-quiet. So Low as a defect, but it sits on the challenge coverage seam, so (c) may deserve its own CORRECTION 2026-08-22 (§11.4.6) — THE ITEM’S PREMISE WAS FALSE, and the runner found mine. Recorded openly rather than quietly. scaling_horizontal_challenge.sh EXISTS in submodules/challenges/challenges/scripts/scaling_horizontal_challenge.sh, executed 14 bytes, dated Aug 7. So do all 14 other entries in the meta-runner. It was never deleted: a search across both repositories and all deletions finds ZERO deletions. HOW I GOT IT. I ran run_all_challenges.sh reads submodules/challenges/challenges/scripts/. I checked challenges/scripts/ — a DIFFERENT directory belonging to a DIFFERENT, glob-based

ATM ID	Type	Status	Severity	Description
BOB-169	Bug		Medium	Two similarly-named directories; I measured the wrong one. That is a §11.4.201(9) field-identity error. Worse, I wrote that it was “verified invocation”: I had listed a directory and the aggregator, so the claim overstated evidence class as well as getting the answer wrong. THE LESSON, which generalises the item. In the SAME commit I correctly caught the sibling agent making this exact class of error. BOB-167 — a zero-hit produced by searching by symbol name — and applied a control net there. Then I omitted the needle one path later on my own measurement. The distribution that would have caught it: A CONTROLLED EXPERIMENT PROVES THE INSTRUMENT CAN SEE, BUT DOES NOT PROVE IT IS POINTED AT THE TARGET UNDER TEST. §11.4.201(7)(b) as commonly applied guards blindness; it does not guard against aiming. The needle for THIS query would have been to resolve the path the runner itself would not to confirm that some directory listing was WHAT IS ACTUALLY WRONG, and it is more thorough than what I filed. The missing-entry path challenges by definition, so it can be exposed safely. Run against a checkout where the submodule is absent — the ordinary standard clone made without -recursive — the R runner (byte-identity sha-asserted, not a reports: PASS: 0 FAIL: 0 SKIP: 16 TOTAL OBSERVED EXIT CODE: 0 An ENTIRELY RUN BANK REPORTS SUCCESS. Not a dangling entry tolerated — the whole suite silently passes without executing anything is §11.4.201(6) false-null verbatim, and observed-reachable rather than constructed (§11.4.115(G)). REVISED ACCEPTANCE ANSWERED, though not as filed: the runner reports a loud, counted SKIP and then runs blocks, because the exit expression reaches FAIL count. (b) The dangling-entry half nothing is dangling. (c) The real work is contract: a roster entry the runner cannot execute must not be indistinguishable from a healthy run. See the decision recorded in BOB-168/decision_missing_vs_skip.md. REPORTED NOT FIXED: scripts/pre_build_verification.sh:1792 carries the SKIP/MISSING conflation one level up, other runner. Filed separately.

ATM ID	Type	Status	Severity	Description
		In progress		WHAT. The §11.4.65 markdown-export m requires every in-scope doc to ship .html twins. 286 of the 326 .html files under c (excluding dist/ and node_modules/) are FRAGMENTS — they begin at WHAT. The BOB-109 scaling suite gates timing tests on loadavg_1m/nproc <= 0 this 8-cpu host, a 1-minute load average below 6.0 — and SKIPS loudly otherwise gate is correct and was adopted for a go (§11.4.201(8)): the author validated a 2 growth-exponent threshold, then observed FALSE-FAIL correct code at load 11.06 span 2.136), and at load 15-23 measure spans of 2.358 against injected-cubic m 2.278-2.331 — the signal is smaller than and the two are not separable in either Gating was the honest response. THE P The gate has consequently NEVER RUN own precondition. Every recorded run t measured load 9.15-23; the independent measured 9.66 during review; the cond measured 6.45 at its lowest. GREEN po observed — at loads 9.4-11.6 (span 1.88 three stability runs (1.759/1.832/1.848) never once beneath the shipped ceiling author states this honestly as an inference data rather than a run they can paste, a reviewer confirmed the inference is sou (quiescence is strictly more favourable) inference is still not an observed run. W MATTERS (§11.4.226). A registered, top present guard that never executes is indistinguishable from a guard that would. Nothing currently ensures it ever r freshness contract, no scheduled quiet invocation, no tracked item — which is the unexecuted-standing-guard class §1 names, where registration gets treated coverage while execution is nobody's jo forensic precedent in that anchor is two guards that, when finally executed, imm emitted latent FAILs. ACCEPTANCE. (a both quiescence-gated timing tests on a whose 1-minute loadavg is genuinely < — no agent fleet, no concurrent build — capture the run as an artifact showing t resolved load reading alongside the me span exponent, so the precondition is p from the artifact and not asserted. (b) F
BOB-170	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				observed GREEN polarity in the item and test source, replacing the current honest inference note. (c) If the gate FAILS under genuine quiescence, that is the finding exists to surface, and it reopens the threshold calibration question rather than being worked around. (d) Establish a freshness contract §11.4.226: how stale a quiescent verdict before the gate counts as uncovered against this does not silently return to never-existence HONEST BOUNDARY. This item does not mean the gate is wrong. The evidence points the other way — quiescence is strictly more favorable than the loads where GREEN was already observed. It claims only that the run has happened, and that an unexecuted guarantee can be cited as coverage. PROVENANCE. RATIONALLY IMPORTANT-2 by the independent Fable substrate reviewer of the BOB-109 world independently corroborated: the review of the SKIP is loud and prints its resolved value and confirmed the shipped tests do SKIP today's load. WHAT. Both rate limiters key their per-client bucket on the LEFTMOST element of X-Forwarded-For when TRUST_FORWARDER is enabled: download-proxy/src/api/rate_limit.py:117-123 (:7187) and the naive limiter in plugins/download_proxy.py (:7187) which was deliberately built to exact parity with it. That element is CLIENT-SUPPLIED honest reverse proxy APPENDS the peer to the RIGHT of whatever arrived, so the entry is whatever the client sent — attached controlled by construction, not by misconfiguration. Consequences with the (1) rotating a forged leftmost value minimize an unlimited sequence of fresh buckets, with total bypass of the limit the flag is supposed to make MORE accurate; (2) it also defeats the idle-reap cap, since each forged identity has a distinct key — the bucket map is a memory growth surface bounded only by the cap; the cap's eviction then discards the budgets of honest clients to make room for forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, and that's why this is Low, not High. TRUST_FORWARDER defaults OFF for all sites, and the deployed stack runs network host with no reverse proxy in front (very
BOB-171	Bug	Queued	Low	

ATM ID	Type	Status	Severity	Description
BOB-172	Bug	Ready for testing	High	across all five compose services), so peer addresses are real client IPs and no NAT occurs. The defect is latent: it arms the someone puts a proxy in front and turns on to recover real client IPs — which is the situation the flag exists for. The failure is therefore ‘correct-looking configuration silently disables the limiter’, not ‘current broken’. PROVENANCE. Raised as MIN the independent reviewer of the BOB-1 limiter work and independently confirmed implementing agent, which noted it at E source sites as a tracked follow-up rather fixing it in that change. Filed here because neither agent has write access to the track WHY IT COVERS BOTH PORTS. :7186’s was built to deliberate policy parity with including this parsing. Fixing one and not the other would leave the two surfaces disjoint on client identity — worse than the shared because it becomes surface-dependent and untestable as a single invariant. ACCEPT (a) Replace leftmost-element parsing with aware resolution at BOTH sites: either minus-N-trusted-hops, or a trusted-proxy allowlist where only a peer inside the allowlist may assert XFF at all — the choice is a deployment-topology decision and should be recorded, not assumed. (b) A test that rotates leftmost element and asserts that the key does NOT rotate, per site. (c) Paired mutation: restore leftmost parsing and the test must FAIL. (d) Negative control (§11.4.2) with TRUST_FORWARDED_FOR OFF, the test must be ignored entirely and keying must be back to the peer address — a guard that honours XFF when the flag is off would be a worse defect than the one being fixed. (e) boundary: this closes header-forgery by does not make per-IP limiting fair behind where many real users legitimately share an address. NOT CLAIMED. No change made at sites still parse leftmost; the flag still does WHAT. The rutracker SEARCH endpoint refused by bot protection while the site is fully reachable. Measured unauthenticated 2026-08-21 (no cookie file read, no credentials value in any artifact — §11.4.10): GET /index.php -> 200, 96283 B,

ATM ID	Type	Status	Severity	Description
BOB-173	Bug	Ready for testing	High	WHAT. download-proxy/src/api/hooks.py <pre>def _save_hooks(hooks: list[dict[str, Any]]): if hooks is None: try: os.makedirs(os.path.dirname(HOOKS_FILE), exist_ok=True) with open(HOOKS_FILE, 'w') as f: json.dump(hooks, f, indent=2) except Exception as e: logger.error(f'Failed to save hooks: {e}') # catches EVERY exception, logs, and returns return None. The signature returns None, so the caller has no channel to learn the write failed. The caller sites then report success unconditionally. _save_hooks(hooks) :174 _save_hooks(hooks) :175 logger.info('Created hook: ...') :175 logger.info('Deleted hook: ...') :156 return HookResponse(hook_id=..., ...) :176 return {'message': 'Hook deleted', ...} USER-VIEW CONSEQUENCE, which is why this is HIGH. This is not not a code-hygiene nit. A user POSTs a hook, and receives HTTP 200 and a hook_id, and the hook was never written — their automation succeeds, never fires, and the API told them it existed. Symmetrically, a user DELETES a hook, and the 'Hook deleted', the file still holds it, and the again after the next restart. In both directions, product reports the opposite of what happened, and the only trace is a log line nobody is looking watching. THIS IS THE §11.4.252 SHADOWN. The path combines two dangerous capabilities: the that anchor's taxonomy — MUTATION of a resource resource (a filesystem write) and EXTENSION. The SIDE EFFECT (hooks are outbound calls) and the system will or will not make) — so it is not to FAIL CLOSED: verify the precondition before when it cannot be satisfied, and surface a graceful refusal. Instead it fails open, and a bare Exception:' that only logs is the exact anomaly §11.4.252 enumerates. At the product level, there is also a §11.4.201(6) false-null: a successful call and a swallowed failure are indistinguishable to the caller. PROVENANCE. Surfaced as a high-risk scope observation by the independent review of the BOB-135 test-isolation work — this is what what turned that defect into an 'assertion failure' mystery, because the EACCES on /config was logged and discarded while the endpoint was returning 200. Verified here directly from the log before filing (the function body and both the log read above), not taken on report. Recorded in §11.4.238 discovery-channel escape: found the agent reading code during an unrelated </pre>

ATM ID	Type	Status	Severity	Description
BOB-174	Bug	In progress	High	investigation, not by the automated QA. The coverage gap is a defect of equal status to the defect itself. ACCEPTANCE. (a) <code>_save</code> propagates failure — raise, or return a success body; callers must consume. (b) Both endpoints must translate a persistence failure into an HTTP 500-class, never a success body; a create that did not persist must not return a hook_id. (c) test drives each endpoint with the hook_id; unwritable (read-only dir or a patched code) raising OSError) and asserts a non-2xx response. AND that a subsequent GET does not list the phantom hook — assert on the user-observed outcome, not on the log line. (d) Paired mutation: restore the swallow; the test fails. FAIL. (e) Audit the same file for sibling files — this is a pattern, and one instance is enough alone. (f) Honest boundary: this does not mean the write currently fails in production; it means that WHEN it fails the user is told the outcome, and the BOB-135 investigation shows it has happened in at least one real environment. NOT CLOSING. No change made. No assessment of how often the write fails in the operator's deployment. RECORDING A SHELL ERROR OF MY OWN. (S11.4.6): the first version of this description was written with the anti-pattern snippet inside backticks in a double-quoted shell argument; when the shell ran it as command substitution, the text was replaced by nothing — the stored description read 'and is the exact anti-pattern'. This is the SECOND time this session the backticks-inside-double-quotes has corrupted the content (the first mangled a commit message). Fixed here by editing through a Python script with no shell quoting in the path. Noted that this is a silently-truncated defect description in the kind of quiet corruption S11.4.201(7) is all about — the path is part of the instrumentation, and failed without erroring. WHAT. A corrupt hooks file is silently indistinguishable from "no hooks configured" and the next create then DESTROYS every existing hook while returning HTTP 200 OK. defects on one path, filed together because any one alone leaves the data loss reachable — download-proxy/src/api/hooks.py:104. <code>_load_hooks</code> wraps the read in except OSError; <code>logger.error(...)</code> and falls through to read truncated or malformed hooks.json is the

ATM ID	Type	Status	Severity	Description
				reported to every caller as an empty, he list. Confirmed at source. A2 — the cons and the reason this is High. create_hook appends, saves. Given a corrupt file tha turns into [], the save writes a one-elem over the top. Every previously-configured gone. Measured by the implementing ag against the ALREADY-FIXED tree, so th the BOB-173 write-failure fix: BEFORE ['prod-hook-0', 'prod-hook-1', 'prod-hoo reports : 200 {'hooks': [], 'count': 0} <- ZERO hooks configured DELETE report {'detail': 'Hook not found'} <- for a hoo in the file POST reports : 200 hook_id=3c8a5930-... AFTER on disk : ['3c8a5930-...'] Three prod hooks destro HTTP 200 throughout, nothing surfaced user. A5 — _save_hooks writes with a pl open(path, "w"). A crash, ENOSPC, or a write truncates the file in place. That is the corruption A1 then reads as “no hoo A2 overwrites. The same codebase alrea the correct pattern: theme_state.py:115 tmp-file + os.replace. WHY THE THREE ONE ITEM. A5 manufactures the corrupt misreads it as empty, A2 destroys the co Fixing only A1 leaves truncation reacha only A5 leaves an existing corrupt file a trigger; fixing only A2 leaves the API lyi what is configured. The chain is the def DELIBERATELY NOT FIXED UNDER BO and the reasoning is sound. The implem agent identified all three while auditing swallows as that item required, and dec fix them in the same change because: th change GET semantics on an endpoint t frontend consumes (200 -> 5xx); they n CORRUPT-file reproduction rather than UNWRITABLE-file one BOB-173 built; a need a design decision that is genuinely obvious — a MISSING hooks file legitim means “no hooks configured”, while a C one does not, and today’s code cannot t apart. Expanding BOB-173’s scope to co would have meant shipping that decisio unexamined. ACCEPTANCE. (a) Disting MISSING from CORRUPT: a missing file an empty list; a corrupt one is an error, silently []. (b) Decide and RECORD wha does on corruption — 5xx, or 200 with a

ATM ID	Type	Status	Severity	Description
				degraded marker the frontend can render the design decision, and it should be stated rather than inferred from whatever the happens to do. (c) create/delete MUST not overwrite a file they could not parse — not clobber (§11.4.252: mutation plus erase side effect must fail closed). (d) Make the atomic via tmp + os.replace, reusing theme_state.py's existing pattern rather than inventing it (§11.4.28). (e) Tests asserting USER-OBSERVABLE outcome, not log level. a corrupt file, assert GET does not claim hooks, assert a create refuses rather than destroying, and assert the file still holds original hooks afterwards. (f) Paired §11.4.201(1): a MISSING file must still be empty list and a working create; a VALID file must behave exactly as today. A fix that every load fail closed by failing always in A3, RECORDED SEPARATELY, NOT PART OF THIS CHAIN. VALID_EVENTS and HookEventType are two sources of truth. closed set. Verified identical today, unguarded against drift: if they diverge a hook register successfully and then silently never fire assertion pinning them would close it. NOT CLAIMED. No change made. The BOB-166 failure fix is real and orthogonal — it makes FAILED write honest; it does nothing about SUCCESSFUL write of wrong data derived from misread file. WHAT. The workable-items update seams the status-location invariant in one direction. After BOB-166, update --status <terminal> Issues-located row is refused. The mirror open: update --location Fixed --status <terminal> exits 0 and creates a row that its own validator immediately rejects. Verified empirically by the independent reviewer. BOB-166: update -location Fixed -status 'progress' -> exit 0 validate -> FAILS, no row (check (f), fixedLocationNonTerminating). So the command can still mint a state it validate call refuses. That is the §11.4.1 shape at the seam layer: the invariant is CONFIGURED in validate but not ENFORCED every write path that can violate it. WHY DELIBERATELY LEFT OPEN IN BOB-166 why that was right. Three reasons, all c
BOB-175	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				rather than asserted: (1) the real corpus ZERO instances of this direction against the direction BOB-166 closed — the asymmetry the data matched the asymmetry in the detective coverage ALREADY existed and demonstrably fires, so unlike BOB-166's this cannot rot silently; (3) and the load one — reopenCmd documents and SANCT transient Fixed-plus-non-terminal state location Fixed operator override. A naive preventive guard at the update seam would collide with that override's semantics. On this therefore requires a design decision on how the two interact, which is a different question from the one BOB-166 was scored answer, and expanding that item would meant shipping the decision unexamined ACCEPTANCE. (a) Decide and RECORD preventive guard on this direction coexist with reopenCmd's sanctioned override — this is actual work, and the answer should be voted down rather than inferred from whatever patch does. Candidates worth weighing the reopen path explicitly, require an override flag on update mirroring reopen's, or leave the direction detective-only with the rationale recorded so the asymmetry is a decision rather than an oversight. (b) If a guard lands, in terminalStatuses() — BOB-166 establish predicate source specifically so the two cannot drift. (c) Paired §1.1 mutation. (d) NEGATIVE CONTROL (§11.4.201(1)), and sharp one here: the sanctioned reopen path MUST still work. A guard that breaks reopenCmd's documented override is a positive refusal, and worse than the gap because it breaks a workflow operators to use. HONEST BOUNDARY. This is not a data defect. Zero corpus instances, and the detective gate catches it at the next value filed so a scoped, deliberate omission stands as a tracked decision instead of living on code comment where the next reader would find it for an oversight (§11.4.197: a started path reaches a documented terminal state, not abandonment). PROVENANCE. Raised by the implementing agent of BOB-166 as a known asymmetry it deliberately did not close, independently verified and endorsed as an acceptable scope boundary by that item.

ATM ID	Type	Status	Severity	Description
BOB-176	Bug	Queued	Medium	reviewer, on the condition that it become a tracked item rather than a comment. The item. WHAT. <code>_get_enabled_trackers()</code> in <code>down/proxy/src/merge_service/search.py</code> gate <code>rutracker</code> on <code>RUTRACKER_USERNAME</code> and <code>RUTRACKER_PASSWORD</code> only. It never checks <code>RUTRACKER_COOKIES</code>. So an operator configured with cookies alone — no username or password — never has <code>rutracker</code> enabled and the tracker is silently absent from the output rather than failing visibly. WHY THIS IS INCOHERENT WITH THE REST OF THE CODE. <code>_search_rutracker</code> treats <code>COOKIES</code> as the <code>PREFERRED</code> auth path, not a fallback. A <code>nnmclub</code> sibling DOES check its own <code>*_COOKIE</code> variable. So the same codebase holds the two positions at once: cookies preferred at the <code>search</code> site, cookies ignored at the enablement gate. COOKIES honoured for a sibling tracker. WHY THIS MATTERS MORE THAN IT LOOKS. <code>CLAUSE 11.4.201(6)</code> documents a standing operator mandate (2026-08-15) that per-tracker Netscape cookies files at <code>\${TRACKER_COOKIE_DIR}/cookies_<tracker>.txt</code> are auto-loaded into <code>as <TRACKER>_COOKIES</code> before every gate. <code>up</code>, <code>restart</code>, <code>install Stage 6</code> and <code>start.sh</code> and <code>rutracker</code> is named in that set. So the <code>SANCTIONED</code>, documented configuration for this tracker produces exactly the state the gate ignores. An operator who follows the documented instructions gets a tracker that never runs, with no error to explain it. IN MODE. Silent absence, not a visible failure. same §11.4.201(6) false-null shape as BOB-172 one layer earlier. BOB-172 fixes a tracker that RAN and was REFUSED being reported. This is a tracker that never ran at all, and the merged result simply has one fewer component with nothing to indicate it. PROVENANCE by the BOB-172 implementing agent who was tracing the enablement path to locate the handling defect. Recorded here rather than in that change: it is a separate defect on a separate seam with its own oracle, and it would have widened BOB-172 past its evidence. ACCEPTANCE. (a) The enablement gate honours <code>RUTRACKER_COOKIES</code> as an independently sufficient credential, matching what <code>_search_rutracker</code> already prefers.

ATM ID	Type	Status	Severity	Description
				the nnmclub sibling already does. (b) A asserting that a cookies-only configurat ENABLES rutracker, driving the real <code>_get_enabled_trackers</code> rather than a rep Paired §1.1 mutation restoring the user password-only gate — the test must go NEGATIVE CONTROL (§11.4.201(1)): a configuration with NO credentials at all leave rutracker DISABLED. A fix that en unauthenticated tracker is worse than t because it produces failing searches ins absent ones. (e) Audit the other tracker for the same asymmetry rather than fixi the one that was noticed — three positi codebase suggests nobody has checked together (§11.4.118). HONEST BOUND verified against a live cookies-only depl read from source by the BOB-172 agent recorded on its report. The reading is s checkable, but whoever takes this shou by invocation before relying on it. BC Bug Queued High WHAT: BOB-172 identical 4-line HTTP-refusal guard at fi private-tracker fetch sites in download- merge_service/search.py (rutracker coo rutracker credential :1468, kinozal :164 nnmclub :1761, iptorrents :1934). Only rutracker COOKIE path is exercised by EVIDENCE (reviewer-authored mutatio 11.4.194(6)(d), during the BOB-172 ind review): deleting ONLY the kinozal guar (search.py:1655-1658) while leaving the intact left the full merge_service suite a passed, ZERO failures. The suite cannot of the five sites. FAILING SCENARIO: a drops or subtly breaks the wiring at kin nnmclub, iptorrents, or rutracker-crede 403 at that site silently folds back to status=empty with error=None -- the ex BOB-172 signature -- and no test redder DIRECTION (reviewer preferred): colla five duplicated wirings into one shared e.g. <code>_check_search_response(tracker_na status, body) -> bool</code>, so there is ONE c and a sixth site cannot be added unguar (11.4.251 byte-identical-fork extraction) Alternative: parametrise the guard tests all five sites with per-site stub sessions. ACCEPTANCE: mutating the wiring at A five sites reddens at least one test. Prov

ATM ID	Type	Status	Severity	Description
				running the same R1 deletion at each si BOB-178 Bug Queued Medium V download-proxy/src/merge_service/ search.py:1638-1640 handles a failed ki LOGIN response with if login_resp.statu (200, 301, 302): logger.error(...); return sets NO diagnostic. FAILING SCENARIO Cloudflare returns 403 on takelogin.php kinozal chip reads status=empty, error= That is the BOB-172 signature exactly: reported to the user as an empty result CONTRAST establishing this is an overs a design choice: the rutracker and nnm failures DO set diagnostics (upstream_c auth_failure), and the iptorrents login fa through to the search fetch where BOB guard catches it. Kinozal is the one leg neither. FIX DIRECTION: stash _classify_upstream_http_status(login_re "") before the early return, or an auth_f for the status-in-(200,301,302)-but-no-c case. ACCEPTANCE: a stubbed 403 on t kinozal login endpoint produces a non-M on the kinozal chip, with a RED capture the pre-fix code first (11.4.115). BOB Task Queued Medium WHAT: the BO independent review demonstrated two f paths that still report a refusing or unre tracker as an empty result. Both are PR EXISTING and neither is a regression fr BOB-172, but the fix evidence log prese five-site coverage without stating the bo (11.4.194(5) requires un-analysed dime explicit gaps, never silently assumed sa -- exception before resp.status is read. I probe B, captured: a connection-refused rutracker yields status=empty, error=N http_status=None, metadata.errors=[], metadata.status=completed. Each _sea method swallows exceptions (except Ex logger.error; return results) before _sea error handling can see them, so a track DOWN is indistinguishable from one tha genuinely empty. GAP B -- soft refusal a 200. Reviewer probe A, captured: _classify_upstream_http_status(200, <c body>) returns None, which is CORREC design (a 2xx is a usable response; body triggering would risk the over-fire the n controls exist to prevent). But all five se

ATM ID	Type	Status	Severity	Description
				use aiohttp's default allow_redirects=True session-expiry 302 -> login-page-200 ch to zero rows and reports empty. FIX DIRE Gap A needs the exception to reach a di rather than being swallowed. Gap B need DISTINCT detector (final-URL check or form marker) with its own RED -- explic widening of the status-code trigger. ACCEPTANCE: both gaps closed with th REDs, or explicitly closed per 11.4.112 evidence. Immediate sub-task: append a gaps paragraph to docs/qa/BOB-172/ fix_evidence_20260822.log. BOB-180 Queued Low WHAT: download-proxy/ routes.py:727-745 sets status=captcha_ with a hardcoded message naming RuT 'RuTracker requires CAPTCHA. Use /ap rutracker/captcha'. FAILING SCENARIO search returns zero results and the cap flavoured diagnostic came from NNMCL Turnstile, not RuTracker. The user is to a RuTracker captcha endpoint for an NI problem. SEVERITY BOUNDED: the ful and tracker_stats travel in the same res payload, so the truth is present and a cl reads them is not misled -- only the hea wrong. The branch was revived by BOB error propagation (correctly: suppressin propagation would recreate the same fa one layer up, 11.4.247) and is already c the routes layer by tests/unit/api_layer/ test_routes_coverage.py:782. FIX DIRE derive the message from the tracker(s) actually erred rather than hardcoding o SCOPE NOTE: api/ was owned by a sibl stream during BOB-172; the author was not to touch it. ACCEPTANCE: an NNM captcha diagnostic on a zero-result sear produces a headline naming NNMClub. BOB-181 Bug Queued High WHAT generate_markdown_exports.sh accepts arguments. It unconditionally scans PROJECT_ROOT/*.md, docs/ recursively scripts/ recursively (scan loops at lines 105/111/116 as of 2026-08-23; they wer when this was filed — line numbers shif the BOB-169 fix, the substance is uncha was independently verified by execution the generator with an explicit path proc out-of-scope export and NOT the reques

ATM ID	Type	Status	Severity	Description
BOB-182	Task	Queued	Medium	and PROJECT_ROOT is derived from BASH_SOURCE (line 17). A caller who v 'bash scripts/generate_markdown_exp docs/qa/SOME/file.md' gets a WHOLE-T with the argument silently discarded. W MATTERS: the caller reasonably believe was scoped to one file. It was not. In a s checkout with parallel work streams thi regenerates any export whose .md is ne its derived files anywhere in the tree -- artifacts in other streams' scopes and, v BOB-068 auto-commit hazard, into anot stream's commit. CLASS: a false afforda script's interface implies a capability it have (11.4.201: an interface must mean claims). FIX DIRECTION: either accept path arguments and honour them (scop run to exactly those files), or REFUSE v clear message when arguments are pas Silently discarding them is the one opti must not remain. ACCEPTANCE: passin either scopes the run to it, or exits non-naming the unsupported argument. A R capturing today's silent-discard behavior (11.4.115). LIVE INSTANCE, MEASURE 2026-08-23 (this is no longer hypothetic During the BOB-169 fix the conductor in the generator with an explicit path to re ONE document. The argument was disc the whole tree was scanned. It generate qa/BOB-174/DESIGN_DECISION.{html, at 11:50 -- a SIBLING STREAM's docum that stream's author was still editing th (.md mtime 12:00). The exports were th stale on arrival, and additionally carried fix BOB-169 charset defect. The sibling independently reported them as stale in hand-off, having no idea another stream created them. SECOND-ORDER LESSO keeping with this item: the conductor's verify the blast radius ALSO failed, and silently. 'git status --porcelain \\\n grep -E pdf)\$"' returned only its own files and r clean - but docs/qa/BOB-174/ is an UNT DIRECTORY, which porcelain collapses line, so the grep was structurally incapa seeing the files inside it (11.4.201(7)(a) the wrong surface; the quiet zero read a absence). The instrument that DOES se

ATM ID	Type	Status	Severity	Description
BOB-184	Task	Queued	Medium	WHAT: CM-EXPORT-CHARSET-VALID (p invariant 50) adopts its 301 pre-existing violations via a monotone-decrease rate than a hard floor. Two things are owed. ADOPTION DECISION IS THE OPERATOR NOT THE SCRIPT'S. 11.4.224(E) requires a brownfield adoption question be SURFACED by the operator and the answer recorded as consumer DATA - never an invented rate. What exists today is a header comment BOBA_EXPORT_CHARSET_BASELINE of 0 that DOCUMENTS a decision the agent makes does not RECORD an answer the operator makes. The independent reviewer judged the change defensible and would not reverse it (a hard floor on 301 violations makes the build unreasonable which 11.4.234 forbids; 11.4.101 favours a reversible option) but correctly held that the operator is yet compliant. Options for the operator: immediate hard floor (BASELINE=0) / keep monotone ratchet / per-corpus phase-in / changed-code-only with a scheduled full-code deadline. (2) THE RATCHET DOES NOT ACTUALLY RATCHET. BASELINE is a hard floor constant. On improvement the gate PASSES PRINTS the value to lower it to, but not lower it, so the gate keeps permitting new code all the way back to 301 after the corpus change. The header now says tightening is MANUAL rather than claiming automatic behavior. Manual tightening does not exist (11.4.6), but the honest solution is a stopgap, not the fix. WHY IT WAS NOT THE FIX: WITH THE FIX: a gate that writes its own threshold during a pre-build run becomes a PRODUCER as well as a GATE (11.4.24 requires role separation), and that is a design change the operator should approve rather than record as a side effect. ACCEPTANCE: the operator's adoption answer recorded as consumer DATA and either a persisted baseline the gate never raises (with the role-separation question answered), or an explicit decision that manual tightening is acceptable. Three icon-glyph controls (.bridge-retry, .toggle, .caret) render text made of non-points. axe-core declines to decide their role, reporting them under the incomplete change with messageKey nonBmp and resolving to NEITHER fgColor NOR bgColor, so the recomputation in docs/qa/BOB-164/

ATM ID	Type	Status	Severity	Description
				axe_contrast_scan.py cannot decide either such nodes across 16 palette x mode combinations. BOB-164 round 2 stopped being scored as silent passes by naming GLYPH_UNMEASURED, printing them, run and fencing them, so any NEW unmeasured glyph node blocks. What remains genuine and unproven is their contrast itself: as non-contrast interface components they owe 3:1 under SC 1.4.11, and no oracle in this repo measures that today. Acceptance: a measurement that resolves the effective foreground and background for glyph controls (for example reading getComputedStyle colour against the computed backdrop, or replacing the glyphs with controls carrying explicit fill tokens), each of the controls shown to clear 3:1 in all 16 palette mode combinations, and the GLYPH_UNMEASURED fence shrunk by the nodes then proven. docs/qa/BOB-164/axe_contrast_scan.py static dist with no backend, so every node's existence depends on fetched data is absent from the page it measures: the results table is empty, so .type-badge and .quality-badge never appear, and the hooks list renders empty, so .status pills never appear. That is precisely why the BOB-164 round-1 regression was unseen in a rendered DOM even though scanned in all 16 themes, and it is why we had to close the class in the arithmetic model instead. Those nodes are now covered by frontend/src/app/models/style-contrast.py which extracts declared foreground-on-background from the stylesheets, but that is an ARITMATIC claim about declared colour pairs, not a rendered-pixel one: it cannot see opacity by an ancestor, a composited backdrop, or cascade this repo's static extractor does not model. Acceptance: the scan drives the dashboard with seeded fixture data (a seeded backend, a route fixture, or an injected component harness) so the badge and status nodes are present in the DOM, axe measures them directly, and a control needle proves the rendered oracle sees a seeded sub-floor before any clean result from it is believed.
BOB-185	Task	Queued	Medium	
BOB-186	Bug	Queued	—	workable-items diff returned exit 0 and verdict "DB and Markdown are in sync" PARTIAL read: invoked with -issues but

ATM ID	Type	Status	Severity	Description
				-fixed it compared 77 Markdown items 184 DB items and called that in sync, pr both mismatched counts in its own outp 107 DB items were never compared aga anything. Measured 2026-08-25: (A) no markdown paths -> correctly REFUSES (B) -db-only -> honest, "no Markdown compared"; (C) -issues alone -> exit 0, (compared 77 Markdown item(s) agains item(s))". The 11.4.201(6) false-null: a p blind instrument returned the same qui genuinely-synced corpus returns. IMPA CORRECTED 2026-08-25, my original fi OVERSTATED this. I wrote that CHECK 11.4.106(F) commit seam was exposed. NOT. scripts/hooks/docs-sync-commit- seam.sh:277 passes BOTH -issues and - a repo-wide sweep found NO executable that breaks: the two real callers (that se pre_build_verification.sh:549) both pass paths; the single-path hits are prose evi files plus a static-scan test fixture. So th was real and latent, never live in this pr corrected claim is: any FUTURE single- caller would have been silently mis-told PROVENANCE, twice corrected. A suba attributed the false-null to the no-paths which actually refuses correctly; the cor re-measured and relocated it to the par Then the fixing agent corrected the con impact. Both corrections are recorded s wrong version propagates. FIXED (unco constitution submodule, fetched to upst 7f16739 before edit per 11.4.26). Desig by default (11.4.201(4)); the narrow per comparison PRESERVED behind an exp -partial-scope rather than deleted (dele existing documented capability would b 11.4.122 silent removal), mirroring how preserved the no-Markdown shape. Key MEASURED DB content, never on flag p — a DB with zero Fixed rows IS fully ac for by -issues alone and still passes, so keying would have traded the false-null 11.4.201(1) false-positive refusal. Resul property: "in sync" is reachable only wh two printed counts account for each oth fail/5 pass -> GREEN 9/9; paired 1.1 mu the fixs own revert, killing exactly the 4 guarding tests while the 5 preservation

ATM ID	Type	Status	Severity	Description
BOB-187	Bug	In progress	—	green. Case C now refuses naming “107 located in Fixed (supply -fixed)”, the 10 confirmed against an independent sqlite (107 Fixed + 78 Issues = 185). Incident closed an 11.4.238 gap: case As refusal test anywhere; TestDiffCmd_NoPathsSt now guards it. OPEN: -partial-scope is t naming choice, not an operator decision renaming is cheap now and expensive a consumers adopt it. Other consumers o shared submodule were not swept. scripts/ownership_precondition.sh cons SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe it, but it is NOT fenced. T028 remediati ownership_path_fence() to scripts/lib/ ownership.sh and wired it into ownership_repair.sh, closing the path-e class there; the precondition was outsid agents file scope and still accepts whate config/owned_paths.yaml expands to. B is lower than a recursive chown (it writ files rather than mutating ownership of existing tree) but the INPUT is identical equally unreviewed, so the same QBITTORRENT_DATA_DIR value that w driven a filesystem-wide chown will driv file creation at an arbitrary location. No escape is empirically confirmed for the path, not merely reasoned: during T028 remediation a probe declaring the verba shipped entry with QBITTORRENT_DAT hung the suite, and ps showed the real executing “find / (! -uid 1000 -o ! -gid 1 -printf ...” before it was killed by verifie Acceptance: ownership_precondition.sh SAME ownership_path_fence() predicat scripts/lib/ownership.sh (never a second 11.4.251), refuses fail-closed on a reject and ships a paired 1.1 mutation proving refusal fires PLUS a negative control pr six live shipped entries still ACCEPT so not a 11.4.201(1) false-positive refusal. Discovered out-of-band during T028 ren and reported by the agent as needing tr it is also a 11.4.238 coverage escape. pre_build_verification.sh invariant 17 (C WORKABLE-ITEMS-VALIDATE) resolves binary through the candidate loop at :5 FIRST entry is constitution/scripts/work
BOB-188	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
				items/bin/workable-items. That file is G TRACKED (md5 17644a248363, identical to workable-items-linux) and executable, so resolution over the current untracked source constitution/scripts/workable-items/workable-items (md5 43376a6d0184). The tracked file is STALE: it does not contain the guards it source now has. MEASURED, needle-prime, 2026-08-25. String presence in the tracked binary: "refusing to set terminal status" "Issues-location item has TERMINAL status control needle from the SAME updateC. At least one mutable field flag is required" the instrument sees through that path and zeros are real absences, not a blind read. untracked sibling returns 1 / 1 / 2 for the three queries. A check whose message is absent from the binary cannot fire. RUN PROOF: the same injected violation (a read of terminal status while current_location=run through both binaries -> stale repository violation (only the body_md desync class reports 2 including the location-status class verbatim. This is exactly why the ten BOB rows (BOB-087/129/131/144/145/146/148/15 survived undetected: the gate that was to catch them was running a binary structure incapable of seeing them. The 11.4.108 >ARTIFACT gap: source correct, artifact every gate reading the artifact reports q This is the THIRD instance of that class one session, alongside BOB-169 (a character that never opened a file) and BOB-183 (a bundle five days older than its sources). COMPOUNDING: the comment at :519- directly above the loop asserts "The naïve workable-items path never existed in the checkout (bin/ is a gitignored local build dir nothing ever populated)". That state is now FALSE - bin/ holds two tracked binaries stale comment asserting the absence of file that now shadows resolution is how stayed invisible. Also an 11.4.30 question operator owns: bin/workable-items and workable-items-linux are VERSIONED BY ARTIFACTS in a constitution submodule inherited by reference by every consumer every consumer inherits whichever binary last committed. Flagged independently

ATM ID	Type	Status	Severity	Description
BOB-189	Bug	Queued	—	separate agents this session. ACCEPTANCE: resolution must never prefer a stale artifact over a current one - either the tracked binary is removed and resolution falls through to demand, or a freshness check refuses a binary older than its sources (see the BOB-183 fingerprint gate for a working pattern); (c) the false comment at :519-523 corrected; (d) the 11.4.201 mutation proving the new refusal fix is a negative control proving a genuinely fresh fix still passes so the fix is not an 11.4.201 positive refusal; (d) the 11.4.30 tracked decision recorded as operator DATA. Diagnostics out-of-band during BOB-166 remediation. 11.4.238 coverage escape in its own right. WHAT: the fail-open scanner's shape-(A) heuristic cannot distinguish 'return False' meaning PROCEED-AS-IF-FINE from 'return False' meaning REFUSE. It flags six text-based CLOSED guards as fail-open defects: <code>_is_safe_fetch_url</code> (x3), <code>_qbit_add_success</code> and the hooks path-boundary guard. INDEPENDENTLY VERIFIED 2026-08-20. The conductor rather than taken on the trial word: <code>download-proxy/src/api/routes.py</code> defines <code>_is_safe_fetch_url</code> returning <code>bool</code> only call site at :1476 reads 'if not <code>_is_safe_fetch_url(url): logger.warning(ISSRF-unsafe download URL (non-public skipping); continue</code>' - a False return REJECTS the URL. WHY THIS MATTERS: per §11.4.252 a false-positive refusal is a FAIL-bluff of severity to a false-negative pass, and the consequence is worse than noise - acting on the finding would mean making the SSRF gate returning False, i.e. deleting an SSRF guard to satisfy a gate. A gate that instructs you to remove a security control is actively dangerous, not merely imprecise. REPRO: run the fail-open scan; observe six hits; read each call site. ACCEPTANCE: the scanner distinguishes the shaped from proceed-shaped False return, is site-aware, or an audited waiver list with entry justification), the six drop out, AND the golden-FALSE fixture containing a real guard is added so the discrimination is not falsifiable per §1.1. WHAT: the §11.4.252 fail-open scan report for qBitTorrent-go, and that zero is NOT evidence. The triage agent ran control n
BOB-191	Bug	Queued	—	WHAT: the §11.4.252 fail-open scan report for qBitTorrent-go, and that zero is NOT evidence. The triage agent ran control n

ATM ID	Type	Status	Severity	Description
				through the scanner's own path per §11.4.252 (b): a Python 'except Exception: pass' needle SEEN, a TypeScript 'catch (e) {}' needle SEEN (so frontend/src's zero IS a real zero). Go empty-'if err != nil {}' needle was NOT SEEN. An instrument that cannot see the idiom of the same quiet zero as a clean tree, and of those is honest. DISTINCT FROM BOB-189 deliberately not merged with it per §11.4.252. BOB-189 is a false MATCH (fail-closed gate reported as fail-open); this is a false NUANCE (entire language unscanned). Same scanner in opposite failure directions, different fixings, merging them would hide one behind the other. IMPACT: Go is the language of qbittorrent-go and boba-jackett (port 7189, which contains encrypted tracker credentials), so the user interface is exactly where §11.4.252's creation plus-mutation combination is most likely to have been assessed; the dashboard says clearly that ACCEPTANCE: the scanner grows a Go gate covering the empty-err-block and swallowing idioms, its needle is SEEN through the qBitTorrent-go's result is re-derived, and the finding is triaged as this Python/TS pass. Until then qBitTorrent-go's fail-open position is UNKNOWN and must be reported as UNKNOWN never as 0. WHAT: the BOB-161 gate lands with 6 more open skips RATCHETED rather than fixed. Ratcheting is the constitution's named last default (§11.4.135/§11.4.224(E)) and this is its own precedent, so the choice is correct. remediation it defers is real work that must be owned somewhere. WHY THIS ITEM EXISTS: the gate's own header asserted the 6 were SEPARATELY (§11.4.197)' while no trace of them existed. The §11.4.209 independent review verified the absence and raised it as IMPORTANT-5, noting that without a record of findings are precisely the parked-unverified class §11.4.226(4) names - the population operator samples and finds broken. A population of being tracked is not tracking. WHAT IT NEEDS: a skip that fires on evidence that the ANSWERED must either classify the response and FAIL on it, or take a §11.4.69 reason that is honestly derivable from the environment rather than from the response - verified against the stack, not asserted. Two of the six sit un-
BOB-192	Task	Queued	—	

ATM ID	Type	Status	Severity	Description
BOB-193	Bug	Queued	—	allow-skip:’ markers at tests/unit/test_tracker_auth_live.py:105 and :108, reviewer confirmed genuinely are fail-o that marker must not be treated as abs ACCEPTANCE: all 6 remediated with RE evidence per §11.4.115, the gate’s BAS ratcheted to 0, and the ratchet’s monot decreasing property preserved through (§11.4.227(A)). NOTE the reviewer’s MI count-baseline absorbs a one-out-one-in remediation progress must be checked the finding SET, not only the count. WHAT: in download-proxy/src/api/stream the dedup set is written BEFORE the yi seen_hashes.add() / seen_hashes_local.a execute at :317 and :356 ahead of emitt format_event raises, the result is never its hash is already recorded, so it is per dropped. Independently verified by the reviewer at both sites: mid-search the h persists across subsequent polls so the never re-emits; at completion the stream immediately after, so there is no later cl either. Results later in the SAME batch yet added and do re-emit on the next po is why the observed symptom is partial rather than total. WHY IT IS FILED SEF the §11.4.252 fail-open pass made this l OBSERVABLE by adding a log line, whic right first move — but observability is n The underlying at-most-once semantics and per §11.4.226(5) a mitigation cannot the defect it mitigates. THE DECISION TRADE, not an oversight to correct blind moving to add-after-successful-yield giv least-once, but a DETERMINISTIC seria failure would then retry the same result poll forever — trading silent loss for a h Options are (a) keep at-most-once and t log line as the contract, (b) add-after-su yield with a bounded per-hash retry buc add-before-yield but remove the hash on failure so exactly one retry occurs. ACCEPTANCE: an explicit decision reco implemented, and covered by a test tha raising emit and asserts the chosen sem with a RED observed first per §11.4.115 as MINOR-1 by the independent review fail-open batch; filed rather than absorb

ATM ID	Type	Status	Severity	Description
BOB-194	Bug	In progress	—	residual defect is not retired along with logging that revealed it. WHAT: agent worktrees under .claude/v are ephemeral copies pinned at arbitrary - never built, never shipped, deleted when agent finishes. Pre-build gates that walk with find(1) from the repo root descend anyway. MEASURED INSTANCE (2026-CM-GO-TOOLCHAIN-MATCHES-BUILD FAILED the main build with exactly one and it came from .claude/worktrees/age afda1df906c537ab4 - a worktree 86 com behind still carrying 'FROM golang:1.23 against a go.mod that had since moved The main tree was CORRECT throughout (Dockerfile golang:1.26.2-alpine, go.mod 1.26.2). Per §11.4.201(1) a false-positive a FAIL-bluff of equal severity to a false p blocks real work and teaches people to the gate. FIXED FOR ONE GATE, ONE INSTANCE OF A CLASS: '.claude' added gate's PRUNE_DIRS, with a determinist pair - a stale mismatch inside .claude/w must NOT fail a healthy tree, the SAME outside it MUST still fail, so the prune c widen into blindness (§11.4.201(6)). Mu verified: reverting the prune makes 2 cl SURVEY, measured not inferred: 8 of 13 gates walk the tree; only the fixed one prunes .claude. Gates using 'git ls-files' structurally immune - worktrees are un tracked file under .claude/). Empirically so far: check_cm_closure_seam_binds, check_cm_killpg_pgid_guard, check_cm_no_production_mutation_resi rc=0 with ZERO worktree references, s clean in practice today. check_cm_plugi TIMED OUT at 90s (rc=124) - unexplain worth its own look since a whole-tree w multiplied by 5 worktrees is a candidat UNKNOWN: check_cm_test_mock_pid_e and check_cm_test_mock_pid_patched_whe were not reached before the probe budg expired; their exposure is UNMEASURE clean. ACCEPTANCE: every tree-walkin either prunes agent scratch or is proven by measurement, each with a paired mu plugin_count's runtime explained. NOT asymmetry that makes this worth closin

ATM ID	Type	Status	Severity	Description
				properly: a stale worktree can only ever findings to a security-relevant scan (kill mutation-residue), never remove them - risk here is false refusal and wasted tru missed defect.
				WHAT: the §11.4.252 fail-open scanner Try-handler shapes (A1)/(A2). contextlib is a With node, so a swallow written tha invisible. VERIFIED BY THE CONDUCT control needle through the same invoca (§11.4.201(7)(b)), because a first attempt returned 0 for BOTH forms and was itse the gate takes -root, not a positional, so path exits on an unknown arg (§11.4.20 the invocation path is part of the instrum Correct run: try/except/pass around os.unlink(user_supplied_path) -> 'FAIL swallowed exception ... :5', the needle s contextlib.suppress(Exception) around t IDENTICAL call -> 'PASS - no swallowe exception ... anti-patterns found'. Both untrusted input with an irreversible unl swallow everything; one is caught and c WHY IT IS WORSE THAN A PLAIN GAP pyproject.toml enables ruff's SIM rules SIM105 is 'use contextlib.suppress inste except-pass'. Our own linter therefore i authors to rewrite the form this gate CA into the form it CANNOT. The two tools pulling in opposite directions and the li runs more often. Every SIM105 autofix shrinks this gate's coverage while the c down, which reads as progress. CONSE FOR THE NUMBER: the scanner's findi is a FLOOR over try/except shapes, not of swallows. Any statement of the form open sites remain' must be read as 'N a except shapes'. ACCEPTANCE: the scan recognises contextlib.suppress (a With v items call contextlib.suppress with a br exception type) and applies the same ≥ capability test; a golden-TRUE fixture in form; a golden-FALSE fixture where sup wraps a single-capability diagnostic em must NOT fire; and an explicit decision SIM105 tension - either exempt the sha gate governs, or accept suppress and te gate to read it. Discovered by the §11.4 remediation agent when its own MINOR required using contextlib.suppress at th
BOB-195	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				diagnostic-emitter sites; those three use legitimate and justified in-source.