

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	10
Bug	Operator-blocked	1
Bug	Queued	17
Bug	Ready for testing	4
Task	In progress	12
Task	Queued	27
Task	Ready for testing	2
TOTAL		73

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	OPERATOR DECISION (2026-08-26, interactive clarification): COMPLETE CAPTCHA FLOW ONCE Session established, path is DECIDED: the operator drives the interactive CAPTCHA flow at /api/v1/autotracker/captcha then /login once, and stores the resulting bb_session. The cool path and the cookies-file autoload path offered and NOT chosen. THE BLOCK M BUT DOES NOT LIFT: the operator must physically drive the flow — no agent can CAPTCHA (§11.4.52 honest operator_at boundary). AGENT PREP OWED BEFORE HAND-OFF, so the operator’s single attempt succeeds rather than discovering a broken endpoint mid-flow: verify both endpoints reachable and correctly wired end-to-end the session is actually PERSISTED after successful post (not merely accepted), a

ATM ID	Type	Status	Severity	Description
				a stored session is actually USED by a subsequent search. Honest limitation to front: a stored bb_session expires, so the requires repeating — that cost was accepted the decision. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent in and supersedes any prior agent-chosen this question. Options not chosen are not above so a future reader does not re-litigate settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior follows — RuTracker automated login by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-host routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in proxy + qBitTorrent-go + Jackett + com forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel — ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
				OPERATOR DECISION (2026-08-26, interactive clarification): UNKNOWN INSTRUMENT THE NEXT OCCURRENCE operator does not currently know which produces the +0500 Auto-commit fast-f so the three candidate answers (second session / intentional rsync job / stale job remain open. DECIDED ACTION: stop h host this session cannot reach (§11.4.6 state is not knowable without access) and INSTRUMENT the mechanism so the next occurrence identifies itself. Forensic capture: committer identity, hostname, time offset, push timing, and the git remote u recorded at Auto-commit time into a tra forensic log. ACCEPTANCE: the next Au commit event yields a captured record n origin host — at which point this item re one of the three original branches with rather than a guess. This is the §11.4.10 reversible-safe move: it costs little, bloc nothing, and converts a recurring myste self-identifying event. This answer is re
BOB-077	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
				consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent in and supersedes any prior agent-chosen this question. Options not chosen are not above so a future reader does not re-litigate settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior follows — RD2-10: Identify second host the Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code re the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable the four Lava-porting findings (closes G
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes C
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Trac + Status Documents table row-completo (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exerc three start.sh subcommands end-to-end real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify ReDoS regex bounds deployed to contain (closes runtime half of GA-12)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos. side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for download-proxy/merge endpoints (canonical of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one (canonical impl of RD2-09)
				OPERATOR DECISION (2026-08-26, interactive clarification): DEFER PA
BOB-101	Task	Queued	High	Go --profile go parity remains a goal but LONGER a v1.0.0 release blocker. RW-1 move to a post-1.0 milestone and stop g tag. Nothing advertised is removed, so stays clean and no §11.4.90 Obsolete cl

ATM ID	Type	Status	Severity	Description
BOB-102	Task	In progress	Medium	warranted. The BOB-141 measured fact unchanged: the Go container's Dockerfile ONE binary binding only :7187; nothing :7186 or :7188 despite compose setting PROXY_PORT/BRIDGE_PORT. CLAUDE map already describes what the container ACTUALLY does and needs no amendment to this decision. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent inference and supersedes any prior agent-chosen answer to this question. Options not chosen are named above so a future reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior item text follows — GA-19/RW-09: Is -profile go p a release goal? (gates RW-10..13) — OPERATOR-DECISION OPERATOR DECISION (2026-08-26, interactive clarification): KEEP 0.0.0.0 MANDATORY AUTH GUARD The tunnel is LAN-reachable — no bind address change to 127.0.0.1. The operator explicitly chose to preserve access from other devices on the network. THE DECISION CARRIES A BOUNDARY OBLIGATION: a permanent §11.4.135 requirement that guard MUST land that FAILS THE BUILD if LAN-reachable route ever stops demanding authentication. The operator accepted (and THE CONDITION that guard exists — with this decision is unprotected and the item is closeable. Guard requirements: enumerated from the authoritative source (FastAPI and handlers, boba-jackett) never a hand-made list; resolve real auth coverage not a guess string 'auth' (§11.4.201 real-condition); deliberately-public routes exempt ONLY if checked-in list with per-entry justification enumerated as honest gaps never silent TRUE + golden-FALSE fixtures per §11.4.35. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent inference, and supersedes any prior agent-chosen default on this question. Options not chosen are named above so a future reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior item text follows — RW-05: LAN threat model — bind tunnel 127.0.0.1 or 0.0.0.0? — OPERATOR-DECISION

ATM ID	Type	Status	Severity	Description
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence helper for the unattributed auto-commit
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence for subagent task-brief source inputs
BOB-109	Task	Ready for testing	Medium	BOB-074 followup: scaling-class test coverage absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limit boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golden fixture for the rate-limit detector
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout architectural gap (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_creation bulk suite (Permission denied /config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while the process still serves 7186 (GIL starvation spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile server 7186/7187/7188 but its container binds 7187 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolved gitdir) pollute gate scan scope and manifest false BOB-126-class findings
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/41 constitution, CLAUDE.md and the REAL badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant label 44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate implementation, and 24 of 36 scripts have companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored third code: 82% of one gate's 38,291 findings from submodules/
BOB-154	Bug	Ready for testing	Medium	Host venv and production container run starlette versions (1.4.1 vs 1.6.0)
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard is sensitive and flaky: 8786ms under host 900-1500ms ceiling calibrated on a quiet

ATM ID	Type	Status	Severity	Description
BOB-159	Bug	Queued	Medium	Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:01:17Z Reported-By: independent review (IMPORTANT-2), partially remediated What (the report, verbatim) -recreate path was fixed this round: run_ownership_precondition -> stack_down run_ownership_repair -> stack_up, so the walks a tree no container can write to. The ordering is now asserted behaviourally with new checks in tests/unit/test_start_reload_recreate.sh, each killed paired 1.1 mutation (11/11 RED-OK). The path is NOT covered by that fix and this tracks the remainder. On a warm ./start gate runs before THIS invocation brings up, but if the stack is ALREADY running containers write while the repair walks. A container can then create a new non-owned file BEHIND the walk, after which completion marker records complete ownership that is not, and those stragglers are never repaired without a manual -force. BOU dismissed: after any successful pass the valid for the scope fingerprint and the marker short-circuits without walking (marker_is_valid so the window requires a stale-or-absent AND a live stack together. Declaring a new entry re-arms the marker, which is exact that combination arises in practice. NOT THIS ROUND, with the reason stated rather hidden: the clean guard is a cheap marker probe mode on scripts/ownership_repair_answers has-this-scope-been-repaired while walking the tree. The existing -dry-run serve: it deliberately SKIPS the marker (FORCE=0 && DRY_RUN=0 guard) and walks, so using it as a warm-start probe walk the tree twice on every start. Additional mode requires editing ownership_repair another stream was actively editing during round; duplicating marker_is_valid into instead would be a near-identical fork (Deferred to this item rather than done unilaterally or silently. The misleading call at the warm-start call site (which claimed gate runs before any container writes) has been corrected to state this boundary honestly. Affected scope / file-scope manifest: (warm-start dispatch, run_ownership_gate site); scripts/ownership_repair.sh (marker

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	path) Reproduction / context: 1. Ensure stack is UP. 2. Change config/owned_path so the scope fingerprint changes (this requires marker by design, data-model E2). 3. Run warm ./start.sh (no -recreate). 4. The ownership repair walks and chowns the declared test containers are still running and able to run it. Acceptance criteria: A warm ./start.sh complete an ownership repair while a container that writes to a declared location is running. Either the repair is deferred with an acceptance refusal naming -recreate, or the stack is up for the walk. Asserted behaviourally in test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE checks, each with a paired 1.1 mutation to it. Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:02:48Z Reported-By: What (the report, verbatim): Independent §11.4.209 review (IMPORTANT-3) found the feature's strongest automated checks never executed by any runner: (1) tests/ownership/test_container_writes_owned_files — the §11.4.115 RED-turned-regression FR-011/FR-007, proven via a docker-compose revert mutation — was moved out of test integration/ (commit 58d340a, to escape autouse fixture hang) but no runner (ci.test.sh, run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_build/test.sh, the §1.1 paired-mutation meta-script the scripts/pre_build/check_cm.sh gate was executed by NOTHING — self-reported honestly in commit 04742d7's own message ('STILL OPEN (reported, not fixed): test_pre_build/ is executed by NOTHING') but tracked as a workable item (a §11.4.197 requirements gap) and never wired into pre_build_verification.sh invariant 30, which globbed only tests/unit/test.sh. Affected file-scope manifest: ci.sh; scripts/pre_build_verification.sh invariant 30 (CONTAINER_UNIT-TESTS-EXECUTED) Reproduction context: grep -rn test_container_writes_owned_files ci.sh run-all-tests.sh scripts/ returned zero results (only docs/specs mentioned the path); grep scripts/pre_build_verification.sh showed

ATM ID	Type	Status	Severity	Description
				30's for-loop globbing only tests/unit/te never tests/pre_build/test_.sh. Accepta criteria: ci.sh gains a runtime-gated py ownership/ stage (skips honestly, rc=0, podman/docker present); scripts/ pre_build_verification.sh invariant 30's covers both tests/unit/test_.sh and tests pre_build/test_*.sh, verified by an extra standalone run of the invariant's exact l reporting a non-zero RAN count that inc files from both directories. Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:31:12Z Reported-B BOB-092 remediation, residual finding V report, verbatim): §11.4.69 names CM OPEN-SKIP as one of three mandatory p gates: it audits sink-side probe helpers if any code path converts an empty or unreachable response into a PASS-count for a feature class with a sink-side prob project does not have it. Why this matte rather than in the abstract: the BOB-09 remediation just removed TWO fail-open exactly this class from tests/e2e/ test_live_stack_evidence.py — the nnmc on-404 (already removed at 7baef2b) an surviving sibling at test_iporrents_is_authenticated_in_sea skipped on 'not authenticated and statu success' while asserting 'treating as tra outage'. That assertion was false for rej credentials (upstream_http_403) and fo container (plugin_env_missing), both of definitive product failures the product a classifies via error_type. Two fail-opens class, found one at a time, is the §11.4.1 extend-to-all-cases signal and the §11.4. coverage-escape signal together: the re not surface either, an agent reading cod remediation's own guards are AST-struc lethal under three discriminating mutat they LIVE IN THE FILE THEY GUARD, deleting that file evades them entirely. A gate is the durable home because it auc corpus rather than one file. Honest bou (§11.4.6): this item does NOT claim the remediated fail-opens are unguarded — guarded, with runtime evidence (13 pas 97.36s against the live stack). It claims CLASS has no corpus-wide detector, so
BOB-161	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	instance in a different file is invisible again. Affected scope / file-scope manifest: pre_build/ (gate absent); tests/e2e/test_live_stack_evidence.py (guards current inside the file they guard) Reproduction context: grep -rl 'CM-NO-FAIL-OPEN-SKIP' scripts/ tests/ returns exactly ONE hit, at tests/e2e/test_live_stack_evidence.py — the guards live in, not a gate. Control not the same query for CM-OWNERSHIP-INVALID (a gate that does exist) returns 3 files, so the instrument can see gate tokens and the gate is a real absence, not a blind zero (§11.4.201(b)). Acceptance criteria: A scripts/pre_build gate named CM-NO-FAIL-OPEN-SKIP is wired into scripts/pre_build_verification audits sink-side probe helpers for code converting an empty/unreachable/error into a PASS-counting SKIP for a feature HAS a sink-side probe. It ships a golden fixture (a real fail-open -> gate FIRES) and a golden-FALSE-with-carrier (an honest to god geo skip, and a comment merely MENTION the phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §11.4.202 makes the gate FAIL before the gate is reported. Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:35:02Z Reported-By: BOB-106/BOB-107 remediation, residual (the report, verbatim): The TESTS for guards are now executed: pre-build invalid glob was extended from tests/unit + tests/pre_build to also cover tests/hooks. BEFORE/after: the glob expanded to 27 subjects 0 under tests/hooks while 3 existed there expands to 30 with 3 under tests/hooks. pass. But the GUARDS themselves are still invoked by nothing. check-brief-inputs.sh honestly conductor-run rather than auto-brief's required-input list lives in free-form that the Agent tool's structured input does expose, so an automatic prompt-scraping would itself be an unproven pattern-matching seam is the dispatch procedure, and that documentation + habit change, not a honest unattributed-commit-guard.sh is different why this item exists. Run against real names 14 violating commits since the last (and 21 bare Auto-commit subjects exist all refs). Wiring it into the pre-build or c

ATM ID	Type	Status	Severity	Description
				seam AS-IS would therefore refuse even immediately, on pre-existing debt nobody in the moment. That is precisely the 11.4.66 brownfield-adoption shape, and 11.4.66 put the choice with the operator, not with agent inventing a ratchet. The options, the decision is a decision and not a default immediate hard floor – the seam refuses 14 are attributed; (b) one-time monotonic decrease ratchet (the 11.4.135 pattern) snapshot 14 as the baseline, the count may and may never rise, so day one is green disease cannot spread; (c) changed-commit – gate only commits created from now on scheduled deadline for the historical 14 report-only for a stated period, then escape Recommendation, with the reasoning rather just the pick: (b). It is the pattern this already uses for exactly this situation, if the debt visible and bounded immediately cannot be satisfied by deleting the guard (11.4.227 closes that gaming channel). the operator’s call and this item is not closed until that answer is recorded. Honest Bob (11.4.6): this item does NOT claim the 14 commits are harmful in themselves – it is their ATTRIBUTION is absent and that is currently prevents the 15th. Affected s file-scope manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-bad-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh Reproduction context: Both guards run correctly by hand are covered by passing tests (9/9 and 13 proven non-bluff against a no-op stub). 14 invoked by scripts/pre_build_verification.sh; scripts/commit-push-all.sh, so neither existing standing detection pressure (11.4.226: 14 with no execution seam is not coverage registration is not coverage). Acceptance criteria: Each guard is invoked by a new seam, OR carries a registered deferral point this item. For unattributed-commit-guard specifically, the operator has answered adoption question below and the answer recorded as consumer DATA – never an ratchet. OPERATOR DECISION (2026-08-26, interactive clarification): YES — 429 RESPONSIVE IF Retry-After IS WELL
BOB-163	Bug	In progress	Medium	

ATM ID	Type	Status	Severity	Description
				FORMED A 429 carrying a VALID Retry header counts as the sibling endpoint b RESPONSIVE. Assertion (c) of the DDoS challenge accepts 2xx OR a well-formed (positive integer seconds or valid HTTP PARSED, not merely present/non-empty DEGRADED remains: connection failure any 5xx, and a 429 with absent or malformed Retry-After. RESIDUAL RISK, accepted and to be stated in the script source per a genuinely wedged endpoint that happens answer 429-with-Retry-After would pass Retry-After parse narrows that window, not close it. Options (b) limiter-aware d (c) exempt healthz probe were both offered NOT chosen. STILL OPEN, not covered decision: the detector counts 429 only, not though the script header says '429 (or equivalent)'; counting 503 would collide crash detector's 5xx tally. Decide when lands limiters on :7185 and :7189. This recorded as consumer DATA per §11.4.3 the operator's stated choice, not an age inference, and supersedes any prior age default on this question. Options not chosen named above so a future reader does not litigate a settled call (§11.4.112(5) bound verdict discipline applied to decisions). item text follows — Reported-Via: §11. reporting directive bug on 2026-08-21T1 Reported-By: BOB-114 remediation, per defect surfaced by BOB-111 landing a r What (the report, verbatim): PRE-EX NOT INTRODUCED - and proven so rather asserted: git diff shows assertion (c) is IDENTICAL to HEAD, and the failure re against the unmodified HEAD script. What changed is not the challenge but the SY BOB-111 appears at least partially delivered because :7187 now returns real 429s where previously returned none. :7185 and :7187 produce zero 429s. This is a 11.4.248-cl corrosion risk and that is why it is filed rather than left as a footnote: run <u>all</u> challenges now fail INTERMITTENTLY, and an intermediate red trains everyone to re-run until green which point a real regression is dismissed 'probably the flaky rate-limit one'. THE DECISION, which is an operator call under 11.4.66 and which I deliberately did NOT

ATM ID	Type	Status	Severity	Description
				Does a 429 from a WORKING limiter count as sibling endpoint being RESPONSIVE? (a) YES - 429 proves the endpoint is alive and correct, protecting itself. Assertion (c) accepts 2xx, 429, and only a connection failure / 5xx counts as degraded. Risk: a genuinely working endpoint that happens to answer 429 with (b) NO - keep requiring 2xx, but make the challenge limiter-aware: drain or wait out the window before the sibling probe (retry-served, so the budget is knowable rather than guessed). (c) Probe siblings on a path that exempts (a healthz-class route), so the question is asked without spending the budget. Recommendation with reasoning, not just a verdict. (a) composed with a bounded liveness check. 429 carrying a well-formed Retry-After header is evidence of a live, correctly-behaving server. (b) makes the challenge slower and could cause window value that will drift. But the server is 'responsive' here are a product judgement. If an answer is recorded, not assumed. RELAY should be stated as fact rather than folded in: the challenge counts 429 only, not 503, though the script header says '429 (or equivalent)'. Counting 503 would collide with the crash detector's logic. Worth deciding when BOB-111 lands line :7185 and :7189. Affected scope / file-manifest: challenges/scripts/ddos_resilience_challenge.sh assertion (c) endpoint isolation; run_all_challenges.sh (c) invokes it Reproduction / context: Assertion (c) requires a sibling-endpoint probe to answer 2xx. :7187 now enforces a real limiter (x-ratelimit-limit: 120, x-ratelimit-retry-after: 86, retry-after: 45, server: uvicorn). A full challenge run sends roughly 250 requests to :7187, so sibling probes fired during the challenge endpoints' tiers land on an exhausted backend and receive 429 - which assertion (c) scores as 'endpoint degraded'. Timing-dependent test results: measured PASS=5 FAIL=2; the UNMOVED HEAD script against the same live stack later measured PASS=7 FAIL=0 SKIP=0. Acceptance criteria: The operator has a question below, the answer is recorded as consumer DATA, and assertion (c) implements it. The challenge then produces the same verdict across 3 consecutive runs on a rate-limited stack (11.4.50 deterministic).

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	In progress	Medium	consistency), and the fix ships a paired mutation proving assertion (c) still catches genuinely degraded sibling endpoint – not it must not blind it (11.4.201(1)). Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:56:53Z Reported-By: BOB-110 UX-class coverage, discovered new axe-core suite on its first live run Verdict, verbatim): This is a REAL user defect, not a test-tuning artifact, and it is by the automated regime rather than by squinting at the page – which is exactly 11.4.238 posture the project is aiming for. static-grep oracle would NOT have found it. Measured: the served root ships an empty hydration, so any check reading the raw audits a page nobody sees. The violation exists in the hydrated DOM, which is what 11.4.170 requires a rendered oracle and value-equality assertions as the proof a correct. Severity reasoning, stated rather assumed: this is user-visible and affects primary dashboard heading, but it degrades legibility rather than breaking function, surface is operator-facing rather than p Medium, not High. The failing test was FAILING on purpose (11.4.238) instead silenced or marked xfail. tests/ux/ current reports 1 failed, 16 passed; that 1 is this. Anyone reading a red UX suite should read this item, not as flakiness – and when the suite goes fully green, which is the state it is closed. HONEST SCOPE LIMIT (11.4.238) the dashboard landing view was scanned by jackett/* sub-routes and the ng-serve-hot :4200 route set were NOT audited – the commands to close both are recorded in testing/ux_accessibility.md. So this item is nodes are a floor, not a total. Affected scope file-scope manifest: the Angular dashboard served at http://localhost:7187/ (same code SPA as frontend/); production components not test files Reproduction / context: ux/test_live_dashboard_accessibility.py on the running merge service. axe-core v4.0.0 scanning a real Playwright-rendered JS DOM, reports color-contrast violations on nodes. Measured pairs: .brand / h1 text on background #3c3f41 = 1.43-1.62:1 (large-text floor is 3:1); tagline #808080

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	#3c3f41 = 2.68:1 (body-text floor is 4.5) Acceptance criteria: axe-core reports color-contrast violations against the live dashboard, with the fix made in product component CSS rather than by relaxing assertion or excluding the rule (11.4.12 reconcile to the correct mechanism, never weaken the check). The existing tests/u the guard and already fails today, so the captured - closure requires it flipping G against the live surface, which is runtime evidence per 11.4.226. Reported-Via: §11.4.202 reporting direct on 2026-08-21T20:00:08Z Reported-By: surfaced while capturing closure evidence BOB-092; diagnosed rather than worked What (the report, verbatim): This is a ABI-AFTER-INTERPRETER-UPGRADE of a missing package. The package is installed compiled half is built for the previous in That distinction matters because 'pip in py' style advice can appear to succeed v leaving the 313 artefact in place. WHY NOT NOTICED EARLIER, stated as fact paths that DO work all avoid system python ci.sh selects an interpreter through _select_python; the long-running suites session ran .venv/bin/python -m pytest and passed. So the automated regime is gre the DOCUMENTED operator command - an 11.4.238-shaped gap, since the escape found by an agent capturing evidence reported by the regime. WORKAROUND (measured theorised): PYTEST_DISABLE_PLUGIN_AUTOLOAD the needed plugins passed explicitly (-p pytest_timeout ...) avoids the schemath autoload that drags in jsonschema -> re -> rpds. That is a workaround, NOT the silences the import path rather than re install, and it will surprise the next person runs the documented command verbatim RELATED, and worth deciding together than twice: BOB-154 wants .venv rebuilt to match the container (which runs CPython 3.12.13 while both system and venv run There are therefore THREE interpreters Whoever rebuilds the venv should configure resolves for the TARGET interpreter after or this same class reappears one direct

ATM ID	Type	Status	Severity	Description
				HONEST BOUNDARY (11.4.6): this item does NOT claim any test is wrong or any product is broken. The product and the venv-run are unaffected. What is broken is the documented entry point. Affected scope / file-scope manifest: host user site-packages (~/.local/lib/python3/site-packages/rpds/); every CLAUDE documented 'python3 -m pytest ...' invocation of any tooling that uses system python3 rather than .venv/bin/python Reproduction / python3 -m pytest tests/e2e/test_live_stack_evidence.py -q -importmode=importlib -> ModuleNotFoundError: No module named 'rpds.rpds', raised during collection via the schemathesis plugin and -> jsonschema -> referencing._core -> 1. EFFECTS MEASURED root cause: system python3 but ~/.local/lib/python3/site-packages/rpds.cpython-313-x86_64-linux-gnu.so - extension built for 3.13. rpds/init.py does 'from .rpds import *', and a 313-tagged extension importable by 3.14, so the submodule does not exist for this interpreter. The venv is CORRECT and unaffected: .venv/lib64/site-packages/rpds/ ships rpds.cpython-313-x86_64-linux-gnu.so and '.venv/bin/python3 -m pytest' succeeds. Acceptance criteria: python3 -m pytest tests/unit/ -v -importmode=importlib - the command CLAUDE documents - reaches collection without ModuleNotFoundError, OR CLAUDE.modified to document the supported run explicitly. Either way the documented command and the working command agree (11.4.6). The guide that misleads is the documentation (the equivalent of a PASS-bluff). WORKAROUND: EFFECTS MEASURED 2026-08-21 — the item is not free, and the item previously implied PYTEST_DISABLE_PLUGIN_AUTOLOAD=1 disables ALL plugin autoloading, not just the one, so anything the suite silently relied on can be re-enabled by hand: - It BREAKS 5 previously env-dependent tests in tests/unit/test_plugin_rutacker.py — TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace need an autoloading env plugin. PROVE IT!

ATM ID	Type	Status	Severity	Description
				caused by any in-flight change: running OWN copy of that file under identical flags produces the same 5 failures. - It makes -timeout=60 (set in pyproject.toml) an UNKNOWN ARGUMENT unless -p pytest is passed explicitly, so pytest.mark.time silently inert under the naive recipe — a believed to be time-bounded is not. - With autoload ON and only schemathesis disabled (no:schemathesis), that same file is 96/99 CONSEQUENCE FOR ANYONE USING WORKAROUND: -p no:schemathesis alone is strictly better than blanket autoload-disabled where it suffices, because it removes or broken plugin. Where the blanket form -p pytest timeout must be added or time inert, and the 5 env-dependent failures recognised as workaround artifacts rather than filed as defects. That last point is the reason the workaround manufactures failures that exactly like product defects. This does not fix the root cause (the venv's rpds ships a cpython-313 ABI tag CPython 3.14 cannot use) or the fix (rebuild the venv — BOB-154) documents that the interim recipe has a radius, so nobody reads a workaround as a regression. WHAT. download-proxy exposes two Server Events routes. They are the same expensive — long-lived connections that hold a worker generator for their lifetime — but only one limit classed: routes.py:801 @router.get('/search/{search_id}') @_rl('sse_stream') classed routes.py:150 @router.get('/theme/stream') async def stream_theme(...) <- limiter decorator _rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/stream bound to the sse_stream class while /theme/stream falls through to the application class. Measured live on the running stack: /ap theme/stream reports x-ratelimit-limit 120 PROVENANCE + A CORRECTION WORKAROUND KEEPING (§11.4.6). This was surfaced by BOB-109 scaling agent, whose report from as: 'sse_stream_limit_decorator is defined but not imported/applied nowhere ... SSE falls to the 120/minute default: 24x less protection framing is WRONG and was NOT filed as a defect. The agent searched for one symbol NAME wiring uses a different mechanism
BOB-167	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-168	Task	In progress	Low	(@_rl('sse_stream')), and it IS applied — search/stream. Verified by reading routes.py:795-806 and the _rl helper at routes.py:46-51, with a control needle of other *_limit_decorator symbols show re in api/init.py so the zero-hit was not a b search. The agent's MEASUREMENT (1 theme/stream) was correct and is what this a real finding; its MECHANISM was halves are recorded so the next reader c re-derive the same wrong cause. WHY IT MATTERS. An unclassified SSE endpoint c cheapest way to pin server resources: e connection is held open, and the default permits 120/min of them. The declared sse_stream class exists precisely because route shape needs a tighter bound than GETs. ACCEPTANCE. (a) A decision, re whether /theme/stream belongs in the s class or genuinely warrants the default policy question, not automatically a bug (b) If it belongs in sse_stream, the deco applied and a test drives BOTH SSE rou asserts each returns its INTENDED clas from x-ratelimit-limit, so a future route without a class is caught. (c) A guard th enumerates SSE-shaped routes and fail that carries no explicit rate-limit class — general form, so the third SSE route do repeat this. (d) Honest boundary: this d claim 120/min is exploitable in this depl with network_mode host and no reverse every caller shares the 127.0.0.1 bucke BOB-111 measured and recorded separ CLAIMED. No change made. The limit v were read from headers, never driven to exhaustion — the limiter is per-IP and s with concurrent agents on this host (§11.4.201 WHAT. scripts/run_all_challenges.sh:66 “scaling_horizontal_challenge.sh” in its set. challenges/scripts/ scaling_horizontal_challenge.sh does no sed -n '66p' scripts/run_all_challenges.s “scaling_horizontal_challenge.sh” \$ ls c scripts/scaling_horizontal_challenge.sh access ...: No such file or directory VER independently, not taken on report — su the BOB-109 scaling agent and re-check by direct invocation. WHY IT MATTERS this is cosmetic or a §11.4.201 gate-hon

ATM ID	Type	Status	Severity	Description
				defect depends entirely on how the runner handles a missing entry, and that is the first thing to determine: if it SKIPS silently, the challenge advertises coverage it does not have (a claim-vs-reality row with no passing challenge behind it); if it FAILS, the runner is penalised for a reason unrelated to the system test, which trains readers to ignore it. Neither outcome is acceptable; they need different solutions. ACCEPTANCE. (a) Determine and record the runner's actual behaviour on the missing entry by invoking it, not by reading it. (b) EITHER keep the challenge, OR remove the entry — with a §11.4.124 discipline: check git history for whether it once existed and was deleted; a silently-dropped challenge is the more interesting defect. (c) If the runner silently skips missing entries, that is its own finding: a missing challenge must be loud (§11.4.3 SKIP-words at minimum), never absent-and-quiet. So Low as a defect, but it sits on the challenge coverage seam, so (c) may deserve its own CORRECTION 2026-08-22 (§11.4.6) — THE ITEM'S PREMISE WAS FALSE, and the runner's mine. Recorded openly rather than quietly in scaling_horizontal_challenge.sh EXISTS in submodules/challenges/challenges/scripts/ scaling_horizontal_challenge.sh, executed 100 bytes, dated Aug 7. So do all 14 other entries in the meta-runner. It was never deleted: a search across both repositories and all 14 finds ZERO deletions. HOW I GOT IT WRONG: run_all_challenges.sh reads submodules/challenges/challenges/scripts/. I checked challenges/scripts/ — a DIFFERENT directory belonging to a DIFFERENT, glob-based search. Two similarly-named directories; I measured the wrong one. That is a §11.4.201(9) field-name error. Worse, I wrote that it was "verification invocation": I had listed a directory and the aggregator, so the claim overstated the evidence class as well as getting the answer wrong. THE LESSON, which generalises to any item. In the SAME commit I correctly caught a sibling agent making this exact class of error. BOB-167 — a zero-hit produced by searching a symbol name — and applied a control net there. Then I omitted the needle one paragraph

ATM ID	Type	Status	Severity	Description
				later on my own measurement. The dist that would have caught it: A CONTROL PROVES THE INSTRUMENT CAN SEE, NOT PROVE IT IS POINTED AT THE TH UNDER TEST. §11.4.201(7)(b) as comm applied guards blindness; it does not gu aiming. The needle for THIS query wou been to resolve the path the runner itse not to confirm that some directory listing WHAT IS ACTUALLY WRONG, and it is than what I filed. The missing-entry pat challenges by definition, so it can be ex safely. Run against a checkout where th submodule is absent — the ordinary sta clone made without -recursive — the R runner (byte-identity sha-asserted, not a reports: PASS: 0 FAIL: 0 SKIP: 16 TOTA OBSERVED EXIT CODE: 0 An ENTIREL RUN BANK REPORTS SUCCESS. Not o dangling entry tolerated — the whole su silently passes without executing anyth is §11.4.201(6) false-null verbatim, and observed-reachable rather than constru (§11.4.115(G)). REVISED ACCEPTANCE ANSWERED, though not as filed: the ru reports a loud, counted SKIP and then r blocks, because the exit expression reac FAIL count. (b) The dangling-entry half nothing is dangling. (c) The real work is contract: a roster entry the runner cann execute must not be indistinguishable f healthy run. See the decision recorded BOB-168/decision_missing_vs_skip.md. REPORTED NOT FIXED: scripts/ pre_build_verification.sh:1792 carries t SKIP/MISSING conflation one level up, other runner. Filed separately. WHAT. The §11.4.65 markdown-export r requires every in-scope doc to ship .htm twins. 286 of the 326 .html files under c (excluding dist/ and node_modules/) are FRAGMENTS — they begin at WHAT. The BOB-109 scaling suite gates timing tests on loadavg_1m/nproc <= 0 this 8-cpu host, a 1-minute load average below 6.0 — and SKIPS loudly otherwise gate is correct and was adopted for a g (§11.4.201(8)): the author validated a 2. growth-exponent threshold, then observ FALSE-FAIL correct code at load 11.06
BOB-169	Bug	In progress	Medium	
BOB-170	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				span 2.136), and at load 15-23 measure spans of 2.358 against injected-cubic m 2.278-2.331 — the signal is smaller than and the two are not separable in either Gating was the honest response. THE P The gate has consequently NEVER RUN own precondition. Every recorded run t measured load 9.15-23; the independent measured 9.66 during review; the cond measured 6.45 at its lowest. GREEN po observed — at loads 9.4-11.6 (span 1.88 three stability runs (1.759/1.832/1.848) never once beneath the shipped ceiling author states this honestly as an inference data rather than a run they can paste, a reviewer confirmed the inference is sou (quiescence is strictly more favourable) inference is still not an observed run. W MATTERS (§11.4.226). A registered, top present guard that never executes is indistinguishable from a guard that wou did. Nothing currently ensures it ever r freshness contract, no scheduled quiet- invocation, no tracked item — which is the unexecuted-standing-guard class §1 names, where registration gets treated coverage while execution is nobody's jo forensic precedent in that anchor is two guards that, when finally executed, imm emitted latent FAILs. ACCEPTANCE. (a both quiescence-gated timing tests on a whose 1-minute loadavg is genuinely <3 — no agent fleet, no concurrent build — capture the run as an artifact showing t resolved load reading alongside the me span exponent, so the precondition is pr from the artifact and not asserted. (b) R observed GREEN polarity in the item ar test source, replacing the current hones inference note. (c) If the gate FAILS unc genuine quiescence, that is the finding exists to surface, and it reopens the thr calibration question rather than being v around. (d) Establish a freshness contra §11.4.226: how stale a quiescent verdict before the gate counts as uncovered ag this does not silently return to never-ex HONEST BOUNDARY. This item does n the gate is wrong. The evidence points t way — quiescence is strictly more favou

ATM ID	Type	Status	Severity	Description
				than the loads where GREEN was already observed. It claims only that the run has happened, and that an unexecuted guarantee can be cited as coverage. PROVENANCE. Raised as MINOR-IMPORTANT-2 by the independent Fable substrate reviewer of the BOB-109 work, which was independently corroborated: the review of the SKIP is loud and prints its resolved status, and confirmed the shipped tests do SKIP on today's load. WHAT. Both rate limiters key their per-client bucket on the LEFTMOST element of X-Forwarded-For when TRUST_FORWARDER is enabled: download-proxy/src/api/rate_limit.py:117-123 (:7187) and the limiter in plugins/download_proxy.py (:7187) which was deliberately built to exact precision with it. That element is CLIENT-SUPPLIED, an honest reverse proxy APPENDS the peer address to the RIGHT of whatever arrived, so the entry is whatever the client sent — attached, controlled by construction, not by misconfiguration. Consequences with the (1) rotating a forged leftmost value into an unlimited sequence of fresh buckets, which is a total bypass of the limit the flag is supposed to make MORE accurate; (2) it also defeats the idle-reap cap, since each forged identity has a distinct key — the bucket map is a memory growth surface bounded only by the cap, and the cap's eviction then discards the budgeted clients to make room for forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, and that's why this is Low, not High. TRUST_FORWARDER defaults OFF on all sites, and the deployed stack runs networked on host with no reverse proxy in front (verified across all five compose services), so peer addresses are real client IPs and no NAT occurs. The defect is latent: it arms the limiter if someone puts a proxy in front and turns it on to recover real client IPs — which is the situation the flag exists for. The failure is therefore 'correct-looking configuration silently disables the limiter', not 'current configuration is broken'. PROVENANCE. Raised as MINOR by the independent reviewer of the BOB-109 work, which limiter work and independently confirmed by the implementing agent, which noted it at Fable source sites as a tracked follow-up rather than a bug.
BOB-171	Bug	Queued	Low	TRUST_FORWARDER defaults OFF on all sites, and the deployed stack runs networked on host with no reverse proxy in front (verified across all five compose services), so peer addresses are real client IPs and no NAT occurs. The defect is latent: it arms the limiter if someone puts a proxy in front and turns it on to recover real client IPs — which is the situation the flag exists for. The failure is therefore 'correct-looking configuration silently disables the limiter', not 'current configuration is broken'. PROVENANCE. Raised as MINOR by the independent reviewer of the BOB-109 work, which limiter work and independently confirmed by the implementing agent, which noted it at Fable source sites as a tracked follow-up rather than a bug.

ATM ID	Type	Status	Severity	Description
				fixing it in that change. Filed here because neither agent has write access to the trust. WHY IT COVERS BOTH PORTS. :7186's was built to deliberate policy parity with including this parsing. Fixing one and the other would leave the two surfaces disjoint on client identity — worse than the shared because it becomes surface-dependent and untestable as a single invariant. ACCEPT (a) Replace leftmost-element parsing with aware resolution at BOTH sites: either minus-N-trusted-hops, or a trusted-proxy allowlist where only a peer inside the allowlist may assert XFF at all — the choice is a deployment-topology decision and should be recorded, not assumed. (b) A test that rotates leftmost element and asserts that the key does NOT rotate, per site. (c) Paired mutation: restore leftmost parsing and the test must FAIL. (d) Negative control (§11.4.2) with TRUST_FORWARDED_FOR OFF, the test must be ignored entirely and keying must be back to the peer address — a guard that honours XFF when the flag is off would be a worse defect than the one being fixed. (e) boundary: this closes header-forgery by ensuring it does not make per-IP limiting fair behind a proxy where many real users legitimately share an address. NOT CLAIMED. No change made. Both sites still parse leftmost; the flag still does WHAT. The rutracker SEARCH endpoint was refused by bot protection while the site was fully reachable. Measured unauthenticated 2026-08-21 (no cookie file read, no credentials) value in any artifact — §11.4.10): GET /index.php -> 200, 96283 B, WHAT. download-proxy/src/api/hooks.py <pre>def _save_hooks(hooks: list[dict[str, Any]]): if hooks is None: try: os.makedirs(os.path.dirname(HOOKS_FILE), exist_ok=True) with open(HOOKS_FILE, 'w') as f: json.dump(hooks, f, indent=2) except Exception as e: logger.error(f'Failed to save hooks: {e}') except Exception: # catches EVERY exception, logs, and returns None. The signature returns None, so the caller has no channel to learn the write failed. Both sites then report success unconditionally. _save_hooks(hooks) :174 _save_hooks(hooks) :175 logger.info('Created hook: ...') :175 logger.info('Deleted hook: ...') :156 return hooks</pre>
BOB-172	Bug	Ready for testing	High	
BOB-173	Bug	Ready for testing	High	

ATM ID	Type	Status	Severity	Description
				HookResponse(hook_id=..., ...) :176 ret {‘message’: ‘Hook deleted’, ...} USER-V CONSEQUENCE, which is why this is H not a code-hygiene nit. A user POSTs a receives HTTP 200 and a hook_id, and t was never written — their automation s never fires, and the API told them it exi Symmetrically, a user DELETES a hook, ‘Hook deleted’, the file still holds it, and again after the next restart. In both dire product reports the opposite of what ha and the only trace is a log line nobody is watching. THIS IS THE §11.4.252 SHA path combines two dangerous capabilit that anchor’s taxonomy — MUTATION o resource (a filesystem write) and EXTE SIDE EFFECT (hooks are outbound call system will or will not make) — so it is r to FAIL CLOSED: verify the precondition when it cannot be satisfied, and surface refusal. Instead it fails open, and a bare Exception:’ that only logs is the exact a §11.4.252 enumerates. At the product la also a §11.4.201(6) false-null: a success and a swallowed failure are indistinguish the caller. PROVENANCE. Surfaced as a scope observation by the independent r the BOB-135 test-isolation work — this what turned that defect into an ‘assert (mystery, because the EACCES on /confi logged and discarded while the endpoint returning 200. Verified here directly fro before filing (the function body and botl read above), not taken on report. Recor §11.4.238 discovery-channel escape: fou agent reading code during an unrelated investigation, not by the automated QA the coverage gap is a defect of equal sta the defect itself. ACCEPTANCE. (a) _sav propagates failure — raise, or return a s callers must consume. (b) Both endpoint translate a persistence failure into an H (500-class), never a success body; a cre did not persist must not return a hook i test drives each endpoint with the hook unwritable (read-only dir or a patched c raising OSError) and asserts a non-2xx AND that a subsequent GET does not lis phantom hook — assert on the user-obs outcome, not on the log line. (d) Paired

ATM ID	Type	Status	Severity	Description
				mutation: restore the swallow; the test FAIL. (e) Audit the same file for sibling — this is a pattern, and one instance is alone. (f) Honest boundary: this does not the write currently fails in production; in that WHEN it fails the user is told the o and the BOB-135 investigation shows it in at least one real environment. NOT C No change made. No assessment of how write fails in the operator's deployment RECORDING A SHELL ERROR OF MY C (§11.4.6): the first version of this description written with the anti-pattern snippet inside backticks in a double-quoted shell argument the shell ran it as command substitution text was replaced by nothing — the stored description read 'and is the exact anti-p This is the SECOND time this session the backticks-inside-double-quotes has corrupted content (the first mangled a commit message Fixed here by editing through a Python with no shell quoting in the path. Noted a silently-truncated defect description is the kind of quiet corruption §11.4.201(7) about — the path is part of the instrument failed without erroring. WHAT. A corrupt hooks file is silently indistinguishable from "no hooks configured" and the next create then DESTROYS even existing hook while returning HTTP 200 defects on one path, filed together because any one alone leaves the data loss reachable — download-proxy/src/api/hooks.py:104 <code>_load_hooks</code> wraps the read in except <code>Exception</code> <code>logger.error(...)</code> and falls through to read truncated or malformed hooks.json is then reported to every caller as an empty, he list. Confirmed at source. A2 — the console and the reason this is High. <code>create_hooks</code> appends, saves. Given a corrupt file that turns into <code>[]</code>, the save writes a one-element over the top. Every previously-configured gone. Measured by the implementing against the ALREADY-FIXED tree, so the BOB-173 write-failure fix: BEFORE <code>['prod-hook-0', 'prod-hook-1', 'prod-hook-2']</code> reports : 200 <code>{'hooks': [], 'count': 0}</code> <- ZERO hooks configured DELETE report <code>{'detail': 'Hook not found'}</code> <- for a hook in the file POST reports : 200
BOB-174	Bug	In progress	High	

ATM ID	Type	Status	Severity	Description
				hook_id=3c8a5930-... AFTER on disk : ['3c8a5930-...'] Three prod hooks destroyed HTTP 200 throughout, nothing surfaced user. A5 — _save_hooks writes with a pl open(path, "w"). A crash, ENOSPC, or a write truncates the file in place. That is the corruption A1 then reads as “no hook A2 overwrites. The same codebase already the correct pattern: theme_state.py:115 tmp-file + os.replace. WHY THE THREE ONE ITEM. A5 manufactures the corrupt misreads it as empty, A2 destroys the cor Fixing only A1 leaves truncation reacha only A5 leaves an existing corrupt file a trigger; fixing only A2 leaves the API lyi what is configured. The chain is the def DELIBERATELY NOT FIXED UNDER BOB-173 and the reasoning is sound. The implem agent identified all three while auditing swallows as that item required, and dec fix them in the same change because: th change GET semantics on an endpoint t frontend consumes (200 -> 5xx); they n CORRUPT-file reproduction rather than UNWRITABLE-file one BOB-173 built; a need a design decision that is genuinely obvious — a MISSING hooks file legitim means “no hooks configured”, while a C one does not, and today’s code cannot t apart. Expanding BOB-173’s scope to co would have meant shipping that decisio unexamined. ACCEPTANCE. (a) Disting MISSING from CORRUPT: a missing file an empty list; a corrupt one is an error, silently []. (b) Decide and RECORD wha does on corruption — 5xx, or 200 with a degraded marker the frontend can rend the design decision, and it should be sta rather than inferred from whatever the happens to do. (c) create/delete MUST overwrite a file they could not parse — not clobber (§11.4.252: mutation plus e side effect must fail closed). (d) Make th atomic via tmp + os.replace, reusing theme_state.py’s existing pattern rather inventing it (§11.4.28). (e) Tests assertin USER-OBSERVABLE outcome, not log l a corrupt file, assert GET does not claim hooks, assert a create refuses rather th destroying, and assert the file still holds

ATM ID	Type	Status	Severity	Description
				original hooks afterwards. (f) Paired §1. mutation per guard. (g) NEGATIVE COM (§11.4.201(1)): a MISSING file must still empty list and a working create; a VALID must behave exactly as today. A fix that every load fail closed by failing always i A3, RECORDED SEPARATELY, NOT PA THIS CHAIN. VALID_EVENTS and HookEventType are two sources of truth. closed set. Verified identical today, ungu against drift: if they diverge a hook regi successfully and then silently never fire assertion pinning them would close it. M CLAIMED. No change made. The BOB-1 failure fix is real and orthogonal — it m FAILED write honest; it does nothing ab SUCCESSFUL write of wrong data deri misread file. WHAT. The workable-items update seam the status-location invariant in one dire After BOB-166, update --status <termina Issues-located row is refused. The mirror open: update --location Fixed --status < terminal> exits 0 and creates a row that own validator immediately rejects. Verifi empirically by the independent reviewe BOB-166: update -location Fixed -statu progress' -> exit 0 validate -> FAILS, na row (check (f), fixedLocationNonTermin So the command can still mint a state it validate call refuses. That is the §11.4.1 shape at the seam layer: the invariant is CONFIGURED in validate but not ENFC every write path that can violate it. WH DELIBERATELY LEFT OPEN IN BOB-16 why that was right. Three reasons, all c rather than asserted: (1) the real corpus ZERO instances of this direction agains the direction BOB-166 closed — the asy the data matched the asymmetry in the detective coverage ALREADY existed an demonstrably fires, so unlike BOB-166's this cannot rot silently; (3) and the load one — reopenCmd documents and SANCT transient Fixed-plus-non-terminal state location Fixed operator override. A naiv preventive guard at the update seam w collide with that override's semantics. C this therefore requires a design decision how the two interact, which is a differen
BOB-175	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				question from the one BOB-166 was scored as a question, and the answer, and expanding that item would mean shipping the decision unexamined. ACCEPTANCE. (a) Decide and RECORD. (b) Preventive guard on this direction coexist with reopenCmd's sanctioned override — this is not actual work, and the answer should be visible down rather than inferred from whatever the patch does. Candidates worth weighing the reopen path explicitly, require an override flag on update mirroring reopen's, or leave the direction detective-only with the rationale recorded so the asymmetry is a decision, not an oversight. (b) If a guard lands, it is terminalStatuses() — BOB-166 establishes a predicate source specifically so the two cannot drift. (c) Paired §1.1 mutation. (d) NEGATIVE CONTROL (§11.4.201(1)), and a sharp one here: the sanctioned reopen path MUST still work. A guard that breaks reopenCmd's documented override is a positive refusal, and worse than the gap because it breaks a workflow operators to use. HONEST BOUNDARY. This is not a data defect. Zero corpus instances, and the detective gate catches it at the next value filed so a scoped, deliberate omission stands as a tracked decision instead of living on code comment where the next reader will find it for an oversight (§11.4.197: a started reaches a documented terminal state, not abandonment). PROVENANCE. Raised by implementing agent of BOB-166 as a known asymmetry it deliberately did not close, independently verified and endorsed as acceptable scope boundary by that item's reviewer, on the condition that it become a tracked item rather than a comment. This is an item. WHAT. _get_enabled_trackers() in down proxy/src/merge_service/search.py gate. rutracker on RUTRACKER_USERNAME and RUTRACKER_PASSWORD only. It never checks RUTRACKER_COOKIES. So an operator configured with cookies alone — no user password — never has rutracker enabled, and the tracker is silently absent from evidence out rather than failing visibly. WHY THIS IS INCOHERENT WITH THE REST OF THE CODE. _search_rutracker treats COOKIES as the
BOB-176	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				PREFERRED auth path, not a fallback. A nnmclub sibling DOES check its own * _variable. So the same codebase holds the positions at once: cookies preferred at the site, cookies ignored at the enablement, cookies honoured for a sibling tracker. V MATTERS MORE THAN IT LOOKS. CLA documents a standing operator mandate (2026-08-15) that per-tracker Netscape files at \${TRACKER_COOKIE_DIR}/cookies_<tracker>.txt are auto-loaded in as <TRACKER>_COOKIES before every up, restart, install Stage 6 and start.sh and rutracker is named in that set. So the SANCTIONED, documented configuration for this tracker produces exactly the state gate ignores. An operator who follows the documented instructions gets a tracker that never runs, with no error to explain it. In MODE. Silent absence, not a visible failure. same §11.4.201(6) false-null shape as BOB one layer earlier. BOB-172 fixes a tracker RAN and was REFUSED being reported as this is a tracker that never ran at all, and the merged result simply has one fewer component with nothing to indicate it. PROVENANCE by the BOB-172 implementing agent while tracing the enablement path to locate the handling defect. Recorded here rather than in that change: it is a separate defect on a separate seam with its own oracle, and it in would have widened BOB-172 past its evidence. ACCEPTANCE. (a) The enablement gate honours RUTRACKER_COOKIES as an independently sufficient credential, matching what _search_rutracker already prefers to the nnmclub sibling already does. (b) A asserting that a cookies-only configuration ENABLES rutracker, driving the real _get_enabled_trackers rather than a replacement Paired §1.1 mutation restoring the user's password-only gate — the test must go to NEGATIVE CONTROL (§11.4.201(1)): a configuration with NO credentials at all leave rutracker DISABLED. A fix that erases an unauthenticated tracker is worse than the one that produces failing searches instead of absent ones. (e) Audit the other tracker for the same asymmetry rather than fixing the one that was noticed — three positions

ATM ID	Type	Status	Severity	Description
				codebase suggests nobody has checked together (§11.4.118). HONEST BOUND verified against a live cookies-only depl read from source by the BOB-172 agent recorded on its report. The reading is sj checkable, but whoever takes this shou by invocation before relying on it. BO Bug Queued High WHAT: BOB-172 identical 4-line HTTP-refusal guard at fi private-tracker fetch sites in download- merge_service/search.py (rutracker coo rutracker credential :1468, kinozal :164 nnmclub :1761, iptorrents :1934). Only rutracker COOKIE path is exercised by EVIDENCE (reviewer-authored mutatio 11.4.194(6)(d), during the BOB-172 ind review): deleting ONLY the kinozal guar (search.py:1655-1658) while leaving the intact left the full merge_service suite a passed, ZERO failures. The suite cannot of the five sites. FAILING SCENARIO: a drops or subtly breaks the wiring at kin nnmclub, iptorrents, or rutracker-crede 403 at that site silently folds back to status=empty with error=None -- the ex BOB-172 signature -- and no test redder DIRECTION (reviewer preferred): colla five duplicated wirings into one shared e.g. _check_search_response(tracker na status, body) -> bool, so there is ONE c and a sixth site cannot be added unguar (11.4.251 byte-identical-fork extraction). Alternative: parametrise the guard tests all five sites with per-site stub sessions. ACCEPTANCE: mutating the wiring at A five sites reddens at least one test. Prov running the same R1 deletion at each si BOB-178 Bug Queued Medium V download-proxy/src/merge_service/ search.py:1638-1640 handles a failed ki LOGIN response with if login_resp.statu (200, 301, 302): logger.error(...); return sets NO diagnostic. FAILING SCENARIO Cloudflare returns 403 on takelogin.php kinozal chip reads status=empty, error= That is the BOB-172 signature exactly: reported to the user as an empty result CONTRAST establishing this is an overs a design choice: the rutracker and nnm failures DO set diagnostics (upstream_o

ATM ID	Type	Status	Severity	Description
				auth_failure), and the iptorrents login fa through to the search fetch where BOB guard catches it. Kinozal is the one leg neither. FIX DIRECTION: stash _classify_upstream_http_status(login_re ""') before the early return, or an auth_f for the status-in-(200,301,302)-but-no-c case. ACCEPTANCE: a stubbed 403 on t kinozal login endpoint produces a non-M on the kinozal chip, with a RED capture the pre-fix code first (11.4.115). BOB Task Queued Medium WHAT: the BO independent review demonstrated two f paths that still report a refusing or unre tracker as an empty result. Both are PR EXISTING and neither is a regression fr BOB-172, but the fix evidence log prese five-site coverage without stating the bo (11.4.194(5) requires un-analysed dime explicit gaps, never silently assumed sa -- exception before resp.status is read. I probe B, captured: a connection-refused rutracker yields status=empty, error=N http_status=None, metadata.errors=[], metadata.status=completed. Each _sea method swallows exceptions (except Ex logger.error; return results) before _sea error handling can see them, so a track DOWN is indistinguishable from one tha genuinely empty. GAP B -- soft refusal a 200. Reviewer probe A, captured: _classify_upstream_http_status(200, <c body>) returns None, which is CORREC design (a 2xx is a usable response; body triggering would risk the over-fire the n controls exist to prevent). But all five se use aiohttp's default allow_redirects=Tr session-expiry 302 -> login-page-200 ch to zero rows and reports empty. FIX DIH Gap A needs the exception to reach a di rather than being swallowed. Gap B nee DISTINCT detector (final-URL check or form marker) with its own RED -- explic widening of the status-code trigger. ACCEPTANCE: both gaps closed with th REDs, or explicitly closed per 11.4.112 evidence. Immediate sub-task: append a gaps paragraph to docs/qa/BOB-172/ fix_evidence_20260822.log. BOB-180 Queued Low WHAT: download-proxy/

ATM ID	Type	Status	Severity	Description
				routes.py:727-745 sets status=captcha_ with a hardcoded message naming RuT 'RuTracker requires CAPTCHA. Use /ap rutracker/captcha'. FAILING SCENARIO search returns zero results and the captcha flavoured diagnostic came from NNMClub Turnstile, not RuTracker. The user is told a RuTracker captcha endpoint for an NNM problem. SEVERITY BOUNDED: the full and tracker_stats travel in the same response payload, so the truth is present and a client reads them is not misled -- only the headline is wrong. The branch was revived by BOB-172 error propagation (correctly: suppressing error propagation would recreate the same failure one layer up, 11.4.247) and is already covered in the routes layer by tests/unit/api_layer/test_routes_coverage.py:782. FIX DIRECTION: derive the message from the tracker(s) that actually erred rather than hardcoding one. SCOPE NOTE: api/ was owned by a sibling stream during BOB-172; the author was not to touch it. ACCEPTANCE: an NNMClub captcha diagnostic on a zero-result search produces a headline naming NNMClub. BOB-181 Bug Queued High WHAT generate_markdown_exports.sh accepts arguments. It unconditionally scans PROJECT_ROOT/*.md, docs/ recursively scripts/ recursively (scan loops at lines 105/111/116 as of 2026-08-23; they were when this was filed — line numbers shifted in the BOB-169 fix, the substance is unchanged) was independently verified by execution of the generator with an explicit path producing out-of-scope export and NOT the request and PROJECT_ROOT is derived from BASH_SOURCE (line 17). A caller who v 'bash scripts/generate_markdown_exports docs/qa/SOME/file.md' gets a WHOLE-T with the argument silently discarded. WHAT MATTERS: the caller reasonably believed was scoped to one file. It was not. In a stream checkout with parallel work streams this regenerates any export whose .md is new in its derived files anywhere in the tree -- but artifacts in other streams' scopes and, via BOB-068 auto-commit hazard, into another stream's commit. CLASS: a false affordance script's interface implies a capability it

ATM ID	Type	Status	Severity	Description
BOB-182	Task	In progress	Medium	have (11.4.201: an interface must mean claims). FIX DIRECTION: either accept path arguments and honour them (scope run to exactly those files), or REFUSE v clear message when arguments are passed. Silently discarding them is the one option must not remain. ACCEPTANCE: passing either scopes the run to it, or exits non-naming the unsupported argument. A R capturing today's silent-discard behavior (11.4.115). LIVE INSTANCE, MEASURE 2026-08-23 (this is no longer hypothetical). During the BOB-169 fix the conductor in the generator with an explicit path to re ONE document. The argument was disc the whole tree was scanned. It generate qa/BOB-174/DESIGN_DECISION.{html, at 11:50 -- a SIBLING_STREAM's docum that stream's author was still editing th (.md mtime 12:00). The exports were th stale on arrival, and additionally carried fix BOB-169 charset defect. The sibling independently reported them as stale in hand-off, having no idea another stream created them. SECOND-ORDER LESSO keeping with this item: the conductor's verify the blast radius ALSO failed, and silently. 'git status --porcelain \\\ grep -E pdf)\$"" returned only its own files and r clean - but docs/qa/BOB-174/ is an UNT DIRECTORY, which porcelain collapses line, so the grep was structurally incap seeing the files inside it (11.4.201(7)(a) the wrong surface; the quiet zero read absence). The instrument that DOES se OPERATOR DECISION (2026-08-26, interactive clarification): KEEP THE MONOTONE RATCHET §11.4.224(E) adoption answer, recorded as consumer the adoption model for CM-EXPORT-CH VALID is the MONOTONE-DECREASE I The count may only go down, never up. Immediate hard floor, changed-code-onl deadline, and per-corpus phase-in were offered and NOT chosen. This closes ha BOB-182 — the half that was genuinely operator's. THE REMAINING HALF IS A DEFECT STILL OWED: BASELINE is a constant, so the ratchet does not actual — on improvement the gate PASSES an

ATM ID	Type	Status	Severity	Description
				the value to lower it to, but nothing lower. The gate keeps permitting regression back to the original count after the corpus heals. The gate does NOT collapse §11.4.249 role separation. The gate that writes its own threshold during a phase-in run becomes a PRODUCER as well as a CONSUMER. Acceptable shapes include a persisted baseline. The gate READS but never writes with the baseline via a separate explicitly-invoked command. The gate that FAILS loudly when the live count falls below baseline so drift becomes an action, not a refusal. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's choice, not an agent inference, and supersedes any prior agent-chosen default on this question. Options not chosen are named above so the independent reader does not re-litigate a settled call. The gate (§11.4.112(5) bounded-verdict discipline) is subject to decisions). — prior item text follows - CM-EXPORT-CHARSET-VALID (pre-built) gate. The gate 50) adopts its 301 pre-existing violation. The gate monotone-decrease ratchet rather than a hard floor. Two things are owed. (1) THE ADOPTION DECISION IS THE OPERATOR'S, NOT THE GATE'S. SCRIPT'S. 11.4.224(E) requires the broad question of adoption question be SURFACED to the operator and the answer recorded as consumer DATA. The gate never an invented ratchet. What exists is a baseline. The header comment plus a BOBA_EXPORT_CHARSET_BASELINE comment. The gate that DOCUMENTS a decision the agent chooses. The gate does not RECORD an answer the operator chooses. The independent reviewer judged the choice defensible and would not reverse it (a hard floor on 301 violations makes the build unreasonable which 11.4.234 forbids; 11.4.101 favours a reversible option) but correctly held that the gate is yet compliant. Options for the operator: immediate hard floor (BASELINE=0) / keep the monotone ratchet / per-corpus phase-in. The gate changed-code-only with a scheduled full phase-in deadline. (2) THE RATCHET DOES NOT RATCHET. ACTUALLY RATCHET. BASELINE is a hard floor, not a constant. On improvement the gate PASSES. The gate PRINTS the value to lower it to, but nothing lower. The gate lowers it, so the gate keeps permitting regression all the way back to 301 after the corpus heals. The header now says tightening is MANIFEST. The gate rather than claiming automatic behavior. The gate does not exist (11.4.6), but the honest s

ATM ID	Type	Status	Severity	Description
				is a stopgap, not the fix. WHY IT WAS NOT THE FIX: a gate that writes its own threshold during a pre-build run becomes PRODUCER as well as a GATE (11.4.24 separation), and that is a design change the operator should approve rather than record as a side effect. ACCEPTANCE: the operator's adoption answer recorded as consumer and either a persisted baseline the gate and never raises (with the role-separation question answered), or an explicit decision manual tightening is acceptable. Three icon-glyph controls (.bridge-retry, .toggle, .caret) render text made of non-points. axe-core declines to decide their reporting them under the incomplete change with messageKey nonBmp and resolving NEITHER fgColor NOR bgColor, so the recomputation in docs/qa/BOB-164/axe_contrast_scan.py cannot decide either such nodes across 16 palette x mode combinations. BOB-164 round 2 stopped being scored as silent passes by naming GLYPH_UNMEASURED, printing them, running and fencing them, so any NEW unmeasured glyph node blocks. What remains genuinely unproven is their contrast itself: as non-interface components they owe 3:1 under SC 1.4.11, and no oracle in this repo measures that today. Acceptance: a measurement resolves the effective foreground and background for glyph controls (for example reading getComputedStyle colour against the color backdrop, or replacing the glyphs with markers carrying explicit fill tokens), each of the controls shown to clear 3:1 in all 16 palette mode combinations, and the GLYPH_UNMEASURED fence shrunk by the nodes then proven. docs/qa/BOB-164/axe_contrast_scan.py static dist with no backend, so every non-existence depends on fetched data is about the page it measures: the results table remains empty, so .type-badge and .quality-badge appear, and the hooks list renders empty, so .status pills never appear. That is precisely why the BOB-164 round-1 regression was unseen in a rendered DOM even though scanned in all 16 themes, and it is why I had to close the class in the arithmetic
BOB-184	Task	Queued	Medium	
BOB-185	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-186	Bug	Queued	—	instead. Those nodes are now covered by frontend/src/app/models/style-contrast., which extracts declared foreground-on- from the stylesheets, but that is an ART claim about declared colour pairs, not a rendered-pixel one: it cannot see opacity by an ancestor, a composited backdrop, cascade this repo's static extractor does model. Acceptance: the scan drives the dashboard with seeded fixture data (a s backend, a route fixture, or an injected component harness) so the badge and s nodes are present in the DOM, axe mea them directly, and a control needle prov rendered oracle sees a seeded sub-floor before any clean result from it is believe workable-items diff returned exit 0 and verdict "DB and Markdown are in sync" PARTIAL read: invoked with -issues but -fixed it compared 77 Markdown items 184 DB items and called that in sync, pr both mismatched counts in its own outp 107 DB items were never compared aga anything. Measured 2026-08-25: (A) no markdown paths -> correctly REFUSES (B) -db-only -> honest, "no Markdown compared"; (C) -issues alone -> exit 0, (compared 77 Markdown item(s) against item(s))". The 11.4.201(6) false-null: a p blind instrument returned the same qui genuinely-synced corpus returns. IMPA CORRECTED 2026-08-25, my original fi OVERSTATED this. I wrote that CHECK 11.4.106(F) commit seam was exposed. NOT. scripts/hooks/docs-sync-commit-seam.sh:277 passes BOTH -issues and - a repo-wide sweep found NO executable that breaks: the two real callers (that se pre_build_verification.sh:549) both pass paths; the single-path hits are prose evi files plus a static-scan test fixture. So th was real and latent, never live in this pr corrected claim is: any FUTURE single-caller would have been silently mis-told PROVENANCE, twice corrected. A suba attributed the false-null to the no-paths which actually refuses correctly; the co re-measured and relocated it to the par Then the fixing agent corrected the con impact. Both corrections are recorded s

ATM ID	Type	Status	Severity	Description
BOB-187	Bug	In progress	—	wrong version propagates. FIXED (uncorrupted constitution submodule, fetched to upstream 7f16739 before edit per 11.4.26). Designed by default (11.4.201(4)); the narrow permission comparison PRESERVED behind an experimental -partial-scope rather than deleted (deleted existing documented capability would be 11.4.122 silent removal), mirroring how preserved the no-Markdown shape. Key MEASURED DB content, never on flag page — a DB with zero Fixed rows IS fully accessible for by -issues alone and still passes, so keying would have traded the false-null 11.4.201(1) false-positive refusal. Result property: “in sync” is reachable only when two printed counts account for each other fail/5 pass -> GREEN 9/9; paired 1.1 must the fix's own revert, killing exactly the 4 guarding tests while the 5 preservation green. Case C now refuses naming “107 located in Fixed (supply -fixed)”, the 10 confirmed against an independent sqlite (107 Fixed + 78 Issues = 185). Incident closed an 11.4.238 gap: case A's refusal test anywhere; TestDiffCmd_NoPathsStill now guards it. OPEN: -partial-scope is the naming choice, not an operator decision; renaming is cheap now and expensive as consumers adopt it. Other consumers of shared submodule were not swept. scripts/ownership_precondition.sh contains SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe it, but it is NOT fenced. T028 remediation ownership_path_fence() to scripts/lib/ownership.sh and wired it into ownership_repair.sh, closing the path-escape class there; the precondition was outside agents file scope and still accepts whatever config/owned_paths.yaml expands to. But is lower than a recursive chown (it writes files rather than mutating ownership of existing tree) but the INPUT is identical, equally unreviewed, so the same QBITTORRENT_DATA_DIR value that would drive a filesystem-wide chown will drive file creation at an arbitrary location. No escape is empirically confirmed for the path, not merely reasoned: during T028 remediation a probe declaring the verb

ATM ID	Type	Status	Severity	Description
BOB-188	Bug	In progress	—	shipped entry with QBITTORRENT_DATA hung the suite, and ps showed the real executing “find / (! -uid 1000 -o ! -gid 1 -printf ...” before it was killed by verification. Acceptance: ownership_precondition.sh SAME ownership_path_fence() predicate scripts/lib/ownership.sh (never a second 11.4.251), refuses fail-closed on a rejection and ships a paired 1.1 mutation proving refusal fires PLUS a negative control proving six live shipped entries still ACCEPT so not a 11.4.201(1) false-positive refusal. Discovered out-of-band during T028 review and reported by the agent as needing tracking; it is also a 11.4.238 coverage escape. OPERATOR DECISION (2026-08-26, interactive clarification): UNTRACKED BINARY — BUILD ON DEMAND §11.4.35 answer, recorded as consumer DATA: the compiled workable-items binary is NOT tracked. ‘Any build derivate which we can recreate by executing proper mechanisms generating MUST NOT be versioned’ applies without exception here. Keeping it tracked only the invariant-52 fingerprint gate, and keeping it tracked with a commit-time record were both offered and NOT chosen. REJECTED END STATE: the gate builds from source REFUSES HONESTLY when the Go toolchain is absent (§11.4.201(11) — probe the artifact through its real invocation path, never fail pass, never silently fall back to a stale build. Staleness becomes structurally impossible than merely detected. ACCEPTED COST: the clone needs a Go toolchain before invariant can run — already true for the qBitTorrent backend, so this adds no new host requirement. The invariant-52 fingerprint gate remains as defence-in-depth for as long as any binary exists on disk. This answer is recorded as consumer DATA per §11.4.35 — it is the operator’s stated choice, not an agent inference and supersedes any prior agent-chosen answer to this question. Options not chosen are noted above so a future reader does not re-litigate the settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior decision follows — pre_build_verification.sh invariant (CM-WORKABLE-ITEMS-VALIDATE) resolves the binary through the candidate loop at :5.

ATM ID	Type	Status	Severity	Description
				FIRST entry is constitution/scripts/workable-items/bin/workable-items. That file is GIT TRACKED (md5 17644a248363, identical to workable-items-linux) and executable, so resolution over the current untracked source constitution/scripts/workable-items/workable-items (md5 43376a6d0184). The tracker is STALE: it does not contain the guards it the source now has. MEASURED, needle-prime, 2026-08-25. String presence in the tracked binary: "refusing to set terminal status". "Issues-location item has TERMINAL status control needle from the SAME updateC. At least one mutable field flag is required" the instrument sees through that path and zeros are real absences, not a blind read. An untracked sibling returns 1 / 1 / 2 for the three queries. A check whose message is absent from the binary cannot fire. RUN PROOF: the same injected violation (a read of terminal status while current_location=run through both binaries -> stale repository violation (only the body_md desync class) reports 2 including the location-status class verbatim. This is exactly why the ten BOB rows (BOB-087/129/131/144/145/146/148/150) survived undetected: the gate that was supposed to catch them was running a binary structure incapable of seeing them. The 11.4.108 >ARTIFACT gap: source correct, artifact incorrect. Every gate reading the artifact reports correct. This is the THIRD instance of that class in one session, alongside BOB-169 (a character that never opened a file) and BOB-183 (a bundle five days older than its sources). COMPOUNDING: the comment at :519- directly above the loop asserts "The naïve workable-items path never existed in the checkout (bin/ is a gitignored local build dir nothing ever populated)". That state is now FALSE - bin/ holds two tracked binaries, a stale comment asserting the absence of a file that now shadows resolution is how it stayed invisible. Also an 11.4.30 question: who owns: bin/workable-items and workable-items-linux are VERSIONED BY ARTIFACTS in a constitution submodule inherited by reference by every consumer. Every consumer inherits whichever binary

ATM ID	Type	Status	Severity	Description
				last committed. Flagged independently separate agents this session. ACCEPTANCE resolution must never prefer a stale art a current one - either the tracked binary removed and resolution falls through to demand, or a freshness check refuses a older than its sources (see the BOB-183 fingerprint gate for a working pattern); false comment at :519-523 corrected; (c) 1.1 mutation proving the new refusal fir negative control proving a genuinely fre still passes so the fix is not an 11.4.201 positive refusal; (d) the 11.4.30 tracked decision recorded as operator DATA. Di out-of-band during BOB-166 remediation 11.4.238 coverage escape in its own rig WHAT: the fail-open scanner's shape-(A heuristic cannot distinguish 'return Fals meaning PROCEED-AS-IF-FINE from 'r False' meaning REFUSE. It flags six tex CLOSED guards as fail-open defects: _is_safe_fetch_url (x3), _qbit_add_succe and the hooks path-boundary guard. INDEPENDENTLY VERIFIED 2026-08-2 conductor rather than taken on the tria word: download-proxy/src/api/routes.py defines _is_safe_fetch_url returning boo only call site at :1476 reads 'if not _is_safe_fetch_url(url): logger.warning(I SSRF-unsafe download URL (non-public skipping); continue' - a False return RE the URL. WHY THIS MATTERS: per §11.4 a false-positive refusal is a FAIL-bluff of severity to a false-negative pass, and he consequence is worse than noise - actin finding would mean making the SSRF g returning False, i.e. deleting an SSRF p to satisfy a gate. A gate that instructs y remove a security control is actively dan not merely imprecise. REPRO: run the f scan; observe six hits; read each call sit ACCEPTANCE: the scanner distinguishes shaped from proceed-shaped False retu site-aware, or an audited waiver list wit entry justification), the six drop out, AN golden-FALSE fixture containing a real guard is added so the discrimination is falsifiable per §1.1.
BOB-189	Bug	Queued	—	
BOB-191	Bug	Queued	—	WHAT: the §11.4.252 fail-open scan rep for qBitTorrent-go, and that zero is NOT

ATM ID	Type	Status	Severity	Description
				evidence. The triage agent ran control n through the scanner's own path per §11.4.252 (b): a Python 'except Exception: pass' n SEEN, a TypeScript 'catch (e) {}' needl SEEN (so frontend/src's zero IS a real z Go empty-'if err != nil {}' needle was N An instrument that cannot see the idiom the same quiet zero as a clean tree, and of those is honest. DISTINCT FROM BOB-189 deliberately not merged with it per §11.4.252. BOB-189 is a false MATCH (fail-closed q reported as fail-open); this is a false NU entire language unscanned). Same scan opposite failure directions, different fixe merging them would hide one behind th IMPACT: Go is the language of qbittorre go and boba-jackett (port 7189, which c encrypted tracker credentials), so the u surface is exactly where §11.4.252's cre plus-mutation combination is most likely has assessed it; the dashboard says clea ACCEPTANCE: the scanner grows a Go covering the empty-err-block and swallo idioms, its needle is SEEN through the qBitTorrent-go's result is re-derived, an finding is triaged as this Python/TS pass Until then qBitTorrent-go's fail-open po UNKNOWN and must be reported as U never as 0. WHAT: the BOB-161 gate lands with 6 r open skips RATCHETED rather than fixe Ratcheting is the constitution's named l default (§11.4.135/§11.4.224(E)) and thi own precedent, so the choice is correct remediation it defers is real work that n owned somewhere. WHY THIS ITEM EX gate's own header asserted the 6 were SEPARATELY (§11.4.197)' while no trac existed. The §11.4.209 independent rev verified the absence and raised it as IMPORTANT-5, noting that without a ro findings are precisely the parked-unver class §11.4.226(4) names - the populatio operator samples and finds broken. A p of being tracked is not tracking. WHAT NEEDS: a skip that fires on evidence th ANSWERED must either classify the res and FAIL on it, or take a §11.4.69 reaso honestly derivable from the environmen than from the response - verified against
BOB-192	Task	Queued	—	

ATM ID	Type	Status	Severity	Description
				stack, not asserted. Two of the six sit unallow-skip:' markers at tests/unit/test_tracker_auth_live.py:105 and :108, reviewer confirmed genuinely are fail-o that marker must not be treated as abs ACCEPTANCE: all 6 remediated with R evidence per §11.4.115, the gate's BAS ratcheted to 0, and the ratchet's monot decreasing property preserved through (§11.4.227(A)). NOTE the reviewer's MI count-baseline absorbs a one-out-one-in remediation progress must be checked the finding SET, not only the count. WHAT: in download-proxy/src/api/stream the dedup set is written BEFORE the yi seen_hashes.add() / seen_hashes_local.a execute at :317 and :356 ahead of emitt format_event raises, the result is never its hash is already recorded, so it is per dropped. Independently verified by the reviewer at both sites: mid-search the h persists across subsequent polls so the never re-emits; at completion the stream immediately after, so there is no later cl either. Results later in the SAME batch yet added and do re-emit on the next po is why the observed symptom is partial rather than total. WHY IT IS FILED SEF the §11.4.252 fail-open pass made this l OBSERVABLE by adding a log line, whic right first move — but observability is n The underlying at-most-once semantics and per §11.4.226(5) a mitigation cannot the defect it mitigates. THE DECISION TRADE, not an oversight to correct blind moving to add-after-successful-yield giv least-once, but a DETERMINISTIC seria failure would then retry the same result poll forever — trading silent loss for a h Options are (a) keep at-most-once and t log line as the contract, (b) add-after-su yield with a bounded per-hash retry buc add-before-yield but remove the hash on failure so exactly one retry occurs. ACCEPTANCE: an explicit decision reco implemented, and covered by a test tha raising emit and asserts the chosen sem with a RED observed first per §11.4.115 as MINOR-1 by the independent review fail-open batch; filed rather than absorb
BOB-193	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
BOB-194	Bug	In progress	—	residual defect is not retired along with logging that revealed it. WHAT: agent worktrees under .claude/v are ephemeral copies pinned at arbitrary - never built, never shipped, deleted when agent finishes. Pre-build gates that walk with find(1) from the repo root descend anyway. MEASURED INSTANCE (2026-CM-GO-TOOLCHAIN-MATCHES-BUILD FAILED the main build with exactly one from .claude/worktrees/agent-afda1df90 - a worktree 86 commits behind still can 'FROM golang:1.23-alpine' against a go since at 1.26.2. The main tree was COR throughout. Per §11.4.201(1) a false-positive refusal is a FAIL-bluff of equal severity to pass: it blocks real work and teaches people to bypass the gate. FIXED FOR ONE GATE added to that gate's PRUNE_DIRS with deterministic §1.1 pair - a stale mismatch inside .claude/worktrees must NOT fail tree, the SAME mismatch outside it MUST fail, so the prune cannot widen into blind (§11.4.201(6)). Mutation verified: revert prune fails 2 checks. SURVEY, measure inferred: 8 of 13 pre_build gates walk through only the fixed one prunes .claude. Gates 'ls-files' are structurally immune - worktrees untracked. Empirically probed: check_cm_closure_seam_binds, check_cm_killpg_pgid_guard and check_cm_no_production_mutation_resistant rc=0 with ZERO worktree references - common practice today. UNMEASURED, not clear check_cm_test_mock_pid_explicit int and check_cm_test_mock_pid_patched_when were not reached before the probe budget expired. HYPOTHESIS REFUTED, record nobody re-walks it: an earlier revision of speculated that check_cm_plugin_count timeout was caused by walking worktrees NOT. That gate's find is scoped to \$PLU (check_cm_plugin_count.sh:224), so worktrees are outside its walk entirely. Its slowness still running past 54s on a re-measure - different, still-unidentified cause and belongs to its own item, not this one. ACCEPTANCE tree-walking gate either prunes agent's is proven immune by measurement, each paired mutation. NOTE the asymmetry

ATM ID	Type	Status	Severity	Description
				bounds the risk: a stale worktree can on findings to a security-relevant scan (kill mutation-residue), never remove them - exposure is false refusal and eroded true missed defect. OPERATOR DECISION (2026-08-26, interactive clarification): TEACH THE SCANNER With NODES — KEEP SIM §11.4.252 fail-open scanner is FIXED to contextlib.suppress; ruff's SIM105 rule ENABLED in pyproject.toml. Disabling § and the belt-and-braces both-at-once op offered and NOT chosen. Rationale carries the decision: fixing the gate makes SIM harmless, whereas disabling SIM105 also leave the gate blind to any suppress wr hand, inherited from a dependency's style already present in the tree. REQUIRED BEHAVIOUR: with contextlib.suppress(with suppress(...)) (the from contextlib suppress binding) wrapping a dangerous combination call are DETECTED with the severity and message shape as the try/except pass form; a narrow suppress(SpecificError) around a NON-dangerous call must NOT since a false-positive refusal is a §11.4.252 FAIL-bluff of equal severity to a missed CONSEQUENCE FOR THE NUMBER: c landed, the scanner's finding count stops floor over try/except shapes and becomes genuine census — any prior count citing completeness was scoped to try/except answer is recorded as consumer DATA p §11.4.35 — it is the operator's stated choice an agent inference, and supersedes any agent-chosen default on this question. C not chosen are named above so a future does not re-litigate a settled call (§11.4. bounded-verdict discipline applied to de — prior item text follows — WHAT: the fail-open scanner matches Try-handler s (A1)/(A2). contextlib.suppress is a With swallow written that way is invisible. V BY THE CONDUCTOR with a control ne through the same invocation path (§11.4. (b)), because a first attempt returned 0 forms and was itself blind - the gate takes not a positional, so a bare path exits on unknown arg (§11.4.201(7)(c): the invoc is part of the instrument). Correct run:
BOB-195	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
				pass around <code>os.unlink(user_supplied_pa</code> 'FAIL - swallowed exception ... :5', the m sees. <code>contextlib.suppress(Exception)</code> are IDENTICAL call -> 'PASS - no swallowed exception ... anti-patterns found'. Both untrusted input with an irreversible unl swallow everything; one is caught and c WHY IT IS WORSE THAN A PLAIN GAP pyproject.toml enables ruff's SIM rules SIM105 is 'use <code>contextlib.suppress</code> inste except-pass'. Our own linter therefore i authors to rewrite the form this gate CA into the form it CANNOT. The two tools pulling in opposite directions and the li runs more often. Every SIM105 autofix shrinks this gate's coverage while the c down, which reads as progress. CONSE FOR THE NUMBER: the scanner's findi is a FLOOR over try/except shapes, not of swallows. Any statement of the form open sites remain' must be read as 'N a except shapes'. ACCEPTANCE: the scar recognises <code>contextlib.suppress</code> (a With v items call <code>contextlib.suppress</code> with a br exception type) and applies the same $\geq$ capability test; a golden-TRUE fixture in form; a golden-FALSE fixture where sup wraps a single-capability diagnostic em must NOT fire; and an explicit decision SIM105 tension - either exempt the sha gate governs, or accept suppress and te gate to read it. Discovered by the §11.4 remediation agent when its own MINOR required using <code>contextlib.suppress</code> at th diagnostic-emitter sites; those three use legitimate and justified in-source. WHAT: the CM-PLUGIN-COUNT pre-bu takes 189 SECONDS to verify 8 docume counts. Measured 2026-08-25: exceeded probe budget (<code>rc=124</code>), then a bounded completed at '<code>rc=0 wall=189s</code>' - it PASS simply slow. CAUSE, MEASURED NOT GUESSED: the marker loop at <code>check_cm_plugin_count.sh:273-290</code> iter EVERY LINE of every governed docume spawns a pipeline per line. Per iteration <code>grep -oE wc -l tr -d'</code> (4 processes) plu sed -n' (2) = ~6. Governed docs are AG 689 + CLAUDE.md 494 + docs/features Status.md 670 = 1853 lines, so ~11,118
BOB-196	Task	Queued	—	

ATM ID	Type	Status	Severity	Description
				spawns, measuring ~102 ms/line. Class line-subprocess shell trap. THE IRONY PRESERVING, because it explains why simplified it: the in-code comment shows -l' was chosen DELIBERATELY over 'gre measuring the §11.4.201(12) footgun - c host (ugrep 7.8.4) 'grep -coE' returned level but 1 inside a 'set -euo pipefail' su context-dependent reading. The correct right and must be kept; it is the PER-LI application of it that costs the three mir HYPOTHESIS REFUTED, recorded so n walks it: this is NOT worktree traversal gate's find is scoped to \$PLUGINS_DIR the five agent worktrees are outside its That guess was raised on BOB-194 and LIKELY FIX (unverified, stated as a can a conclusion): pre-filter with ONE grep containing 'CM-PLUGIN-COUNT:' before the loop, reducing ~1853 iterations to t lines that carry a marker, while leaving line parsing logic - and its footgun-avoid completely unchanged. WHY IT MATTE beyond impatience: this gate is on the p critical path, and CLAUDE.md leans on mechanical authority for the three disti rosters (43 curated / 43 engines / 12 bo whose conflation was BOB-149. A gate s enough to tempt anyone into skipping it nothing. ACCEPTANCE: runtime reduce counts and the footgun-avoiding parse p byte-for-byte in behaviour, proven by th paired mutation still biting plus a before wall-clock recorded. §11.4.6: nothing he questions the gate's VERDICTS - only it