

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	4
Bug	Operator-blocked	1
Bug	Queued	8
Bug	Ready for testing	2
Task	In progress	10
Task	Queued	24
Task	Ready for testing	2
TOTAL		51

Items

ATM ID	Type	Status	Severity	Description
				OPERATOR DECISION (2026-08-26, §11.4.66 interactive clarification): COMPLETE THE CAPTCHA FLOW ONCE Session establishment DECIDED: the operator drives the interactive CAPTCHA flow at /api/v1/auth/rutracker/cap then /login once, and the proxy stores the re bb_session. The cookie-paste path and the co autoload path were both offered and NOT ch THE BLOCK NARROWS BUT DOES NOT LIFT: operator must physically drive the flow — no can solve a CAPTCHA (§11.4.52 honest operator_attended boundary). AGENT PREP BEFORE THE HAND-OFF, so the operator’s sir attempt succeeds rather than discovering a b
BOB-008	Bug	Operator-blocked	—	

ATM ID	Type	Status	Severity	Description
				endpoint mid-flow: verify both endpoints are reachable and correctly wired end-to-end, verify session is actually PERSISTED after a successful (not merely accepted), and verify a stored session is actually USED by a subsequent search. Honesight limitation to state up front: a stored bb_session expires, so this path requires repeating — that was accepted with the decision. This answer is recorded as consumer DATA per §11.4.35 — in operator’s stated choice, not an agent inference supersedes any prior agent-chosen default on question. Options not chosen are named above. future reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict discipline applies to decisions). — prior item text follows — RuTra automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-host SOCKS routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download proxy + qBitTorrent-go + Jackett + compose e forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto-commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel ledger ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent from mandated test-type matrix
				OPERATOR DECISION (2026-08-26, §11.4.66 interactive clarification): UNKNOWN — INSTRUMENT THE NEXT OCCURRENCE The operator does not currently know which host produces the +0500 Auto-commit fast-forward three candidate answers (second Claude session intentional rsync job / stale job) all remain open. DECIDED ACTION: stop hunting a host this session cannot reach (§11.4.6 — remote state is not known without access) and instead INSTRUMENT the mechanism so the next occurrence identifies Forensic capture to add: committer identity, hostname, timezone offset, push timing, and remote used, recorded at Auto-commit time in
BOB-077	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
				tracked forensic log. ACCEPTANCE: the next Auto-commit event yields a captured record naming the origin host — at which point this item resolves one of the three original branches with evidence more than a guess. This is the §11.4.101 reversible-move: it costs little, blocks nothing, and converts a recurring mystery into a self-identifying event. The answer is recorded as consumer DATA per §11.4.102 it is the operator's stated choice, not an agent's inference, and supersedes any prior agent-choice default on this question. Options not chosen are named above so a future reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict dispositive applied to decisions). — prior item text follows RD2-10: Identify second host running the Auto-commit rsync/sync mechanism (OPERATOR-DECISION) RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated commit script OR retire it RD2-13: Retroactive Fable-xhigh code review of substantive Auto-commit diffs RD2-15: Create BOB-064..067 workable items from four Lava-porting findings (closes GA-05) RD2-18: Create top-level Boba proxy/merge-script v1.0.0 readiness ledger (closes GA-10) RD2-21: Complete/verify README Tracked-Item Status Documents table row-completeness (GA-11 remainder) RD2-25: HelixQA Challenge entry exercising a start.sh subcommands end-to-end against real compose stack RD2-28: Live compose bring-up + verify rutra ReDoS regex bounds deployed to container (closes runtime half of GA-12) RD2-29: Author tests/stress/test_tracker_fetch_stress_chaos.py with §11.4.101 injection
BOB-078	Task	Queued	High	
BOB-080	Task	Queued	High	
BOB-082	Task	Queued	High	
BOB-085	Task	Queued	Medium	
BOB-088	Task	Queued	Medium	
BOB-090	Task	In progress	Medium	
BOB-093	Task	In progress	High	
BOB-094	Task	In progress	Medium	
BOB-095	Task	In progress	Medium	

ATM ID	Type	Status	Severity	Description
				RD2-30: Author tests/stress/test_scheduler_hooks_sse_stress_chaos.py for triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for export/download-proxy/merge endpoints (canonical impl of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one commit (canonical impl of RD2-09)
				OPERATOR DECISION (2026-08-26, §11.4.66 interactive clarification): DEFER PAST v1.0.0 profile go parity remains a goal but is NO LONGER a v1.0.0 release blocker. RW-10..13 move to a post-v1 milestone and stop gating the tag. Nothing added is removed, so §11.4.122 stays clean and no §11.4.122 Obsolete closure is warranted. The BOB-141 requirement fact stands unchanged: the Go container's Dockerfile runs ONE binary binding only :7187; nothing is added :7186 or :7188 despite compose setting PROXY=BRIDGE_PORT. CLAUDE.md's port map already describes what the container ACTUALLY does and needs no amendment under this decision. The answer is recorded as consumer DATA per §11.4.112(5) it is the operator's stated choice, not an agent-chosen inference, and supersedes any prior agent-chosen default on this question. Options not chosen are named above so a future reader does not re-litigate the settled call (§11.4.112(5) bounded-verdict discipline applied to decisions). — prior item text follows GA-19/RW-09: Is -profile go parity still a release blocker (gates RW-10..13) — OPERATOR-DECISION
BOB-101	Task	Queued	High	
				OPERATOR DECISION (2026-08-26, §11.4.66 interactive clarification): KEEP 0.0.0.0 + MANDATORY AUTH GUARD The tunnel STA is reachable — no bind address changes to 127.0.0.1 operator explicitly chose to preserve access for other devices on the network. THE DECISION IS A BINDING OBLIGATION: a permanent §11.4.122 regression guard MUST land that FAILS THE test if any LAN-reachable route ever stops demanding authentication. The operator accepted 0.0.0.0 as the CONDITION that guard exists — without it the decision is unprotected and the item is NOT covered
BOB-102	Task	In progress	Medium	

ATM ID	Type	Status	Severity	Description
				Guard requirements: enumerate routes from authoritative source (FastAPI app, Go handle jackett) never a hand-maintained list; resolve auth coverage not a grep for the string ‘auth’ (§11.4.201 real-condition); deliberately-publicly-exempt ONLY via a checked-in list with per-e justification, enumerated as honest gaps never golden-TRUE + golden-FALSE fixtures per §11.4.107(10). This answer is recorded as con DATA per §11.4.35 — it is the operator’s state not an agent inference, and supersedes any p agent-chosen default on this question. Option chosen are named above so a future reader c re-litigate a settled call (§11.4.112(5) bounded discipline applied to decisions). — prior item follows — RW-05: LAN-exposure threat mode tunnel 127.0.0.1 or keep 0.0.0.0? — OPERATOR DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-, regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence-check for the unattributed auto-commit path
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence check subagent task-brief source inputs
BOB-109	Task	Ready for testing	Medium	BOB-074 followup: scaling-class test coverage from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limitin boba’s 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golden-bad for the rate-limit detector
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout arch gap (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_create fail suite (Permission denied /config)
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolvable) pollute gate scan scope and manufacture fals BOB-126-class findings

ATM ID	Type	Status	Severity	Description
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate with implementation, and 24 of 36 scripts have no companion doc OPERATOR DECISION (2026-08-26, §11.4.66), REPAIR, THEN RE-VERIFY UNTIL STABLE The warm ./start.sh path keeps running the owner repair against a live stack, but walks, re-verify repeats until a pass finds nothing new. Option chosen: refuse the repair while the stack is up direct the operator to –recreate; stop the stack warm path too (correct but ends ‘warm’); add document the window. WHY THIS MATTERS MORE THAN A NORMAL ITEM: this is the exact definition of the whole 002-user-owned-downloads feature exposed. The –recreate path was already fixed to precondition -> down -> repair -> up so the walk a tree no container can write. The warm path fails so a container can create a new non-owned file AFTER the repair has passed that directory — ownership problem leaking back in through the door. REQUIRED BY THE CHOSEN OPTION, and the hard part: the loop MUST be BOUNDED, and giving up it MUST report honestly that ownership UNPROVEN rather than silently exiting 0 (§11.4.66) — a loop that stops finding new files because out of iterations returns the same quiet zero (genuinely clean tree). An active download creates files faster than the walk completes is the real non-converging case and must be named in the up message. Recorded as consumer DATA per the operator’s stated choice, not an agent item. Options not chosen are named so a future repair not re-litigate a settled call. — prior item text — Reported-Via: §11.4.202 reporting directive 2026-08-21T19:01:17Z Reported-By: T041 independent review (IMPORTANT-2), partially remediated (the report, verbatim): The –recreate path works this round: run_ownership_precondition -> stack_down -> run_ownership_repair -> stack_up the repair walks a tree no container can write ordering is now asserted behaviourally by the checks in tests/unit/test_start_reload_recreate
BOB-159	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				killed by its paired 1.1 mutation (11/11 RED-C WARM path is NOT covered by that fix and the walk tracks the remainder. On a warm ./start.sh the repair runs before THIS invocation brings anything to the stack is ALREADY running, containers write the repair walks. A container can then create a non-operator-owned file BEHIND the walk, at which the completion marker records completion of a tree that is not, and those stragglers are never repaired without a manual -force. BOUNDED is dismissed: after any successful pass the marker is valid for the scope fingerprint and the repair walks circuits without walking (marker_is_valid), so the window requires a stale-or-absent marker ALREADY on the stack together. Declaring a new scope entry requires the marker, which is exactly how that combination arises in practice. NOT DONE THIS ROUND, with the reason stated rather than hidden: the clean gate is a cheap marker-only probe mode on scripts/ownership_repair.sh that answers has-this-scope-been-repaired without walking the tree. The -dry-run cannot serve: it deliberately SKIPS the marker fast-path (FORCE=0 && DRY_RUN=0 gate) and always walks, so using it as a warm-start would walk the tree twice on every start. Add a probe mode requires editing ownership_repair.sh, which was another stream was actively editing during the current round; duplicating marker_is_valid into start.sh instead would be a near-identical fork (11.4.2). Deferred to this item rather than done unilaterally. The misleading comment at the warm-start call site (which claimed the gate runs before the container writes) has been corrected to state the boundary honestly. Affected scope / file-scope manifest: start.sh (warm-start dispatch, run_ownership_gate call site); scripts/ownership_repair.sh (marker fast-path) Reproduction / context: 1. Ensure the stack is clean. Change config/owned_paths.yaml so the scope fingerprint changes (this re-arms the marker repair design, data-model E2). 3. Run a warm ./start.sh -recreate). 4. The ownership repair walks and the declared tree while containers are still running

ATM ID	Type	Status	Severity	Description
				and able to write into it. Acceptance criteria warm ./start.sh cannot complete an ownership while a container that writes to a declared lo running. Either the repair is deferred with an actionable refusal naming –recreate, or the s quiesced for the walk. Asserted behaviourall unit/test_start_reload_recreate.sh alongside t existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE_UP each with a paired 1.1 mutation that kills it. Reported-Via: §11.4.202 reporting directive + 2026-08-21T19:35:02Z Reported-By: BOB-106 remediation, residual What (the report, ver The TESTS for both guards are now executed build invariant 30's glob was extended from + tests/pre_build to also cover tests/hooks. MI before/after: the glob expanded to 27 suites v under tests/hooks while 3 existed there; it no expands to 30 with 3 under tests/hooks. All th But the GUARDS themselves are still invoked nothing. check-brief-inputs.sh is honestly con run rather than automatic: a brief's required list lives in free-form prose that the Agent too structured input does not expose, so an autom prompt-scraping gate would itself be an unpr pattern-match. Its seam is the dispatch proce and that is a documentation + habit change, n hook. unattributed-commit-guard.sh is differ is why this item exists. Run against real histo names 14 violating commits since the last tag bare Auto-commit subjects exist across all re Wiring it into the pre-build or commit seam would therefore refuse every commit immedi pre-existing debt nobody can fix in the mome is precisely the 11.4.224(E) brownfield-adopti shape, and 11.4.66/11.4.122 put the choice wi operator; not with an agent inventing a ratch options, stated so the decision is a decision an default: (a) immediate hard floor – the seam until all 14 are attributed; (b) one-time mono decrease ratchet (the 11.4.135 pattern) – snap as the baseline, the count may fall and may n
BOB-162	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				so day one is green and the disease cannot spread. changed-commits-only – gate only commits changed from now on, with a scheduled deadline for the historical 14; (d) report-only for a stated period to escalate. Recommendation, with the reasoning, is more than just the pick: (b). It is the pattern this constitution already uses for exactly this situation: makes the debt visible and bounded immediately if it cannot be satisfied by deleting the guard's name (11.4.227 closes that gaming channel). But it is the operator's call and this item is not closed until an answer is recorded. Honest boundary (11.4.6). This item does NOT claim the 14 commits are harmful to themselves – it claims their ATTRIBUTION is . . . and that nothing currently prevents the 15th commit. Affected scope / file-scope manifest: scripts/unattributed-commit-guard.sh; scripts/hooks/brief-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh Reproduction / coverage: Both guards run correctly by hand and are covered by passing tests (9/9 and 15/15, each proven non-interfering against a no-op stub). Neither is invoked by scripts/pre_build_verification.sh or scripts/commit-push-all.sh, so neither exerts standing detection pressure (11.4.226: a guard with no execution seam is not covered; registration is not coverage). Acceptance criteria: Each guard is invoked by a named script, carries a registered deferral pointing at this item, unattributed-commit-guard.sh specifically, the operator has answered the adoption question, and the answer is recorded as consumer DATA, never an invented ratchet. WHAT. The BOB-109 scaling suite gates its two tests on loadavg_1m/nproc <= 0.75 — on this host, a 1-minute load average at or below 6.0 SKIPS loudly otherwise. That gate is correct and adopted for a good reason (§11.4.201(8)): the validated a 2.10 growth-exponent threshold, observed it FALSE-FAIL correct code at load 1 (baseline span 2.136), and at load 15-23 measured baseline spans of 2.358 against injected-cubic at 2.278-2.331 — the signal is smaller than the
BOB-170	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				and the two are not separable in either direction. Gating was the honest response. THE PROBLEM is the gate has consequently NEVER RUN under its precondition. Every recorded run today measured load 9.15-23; the independent reviewer measured 6.45 during review; the conductor measured 6.45 lowest. GREEN polarity WAS observed — at load 9.4-11.6 (span 1.882) and in three stability runs (1.759/1.832/1.848) — but never once beneath the shipped ceiling. The author states this honest inference from data rather than a run they could not and the reviewer confirmed the inference is correct (quiescence is strictly more favourable). Sound inference is still not an observed run. WHY THIS MATTERS (§11.4.226). A registered, topology-aware guard that never executes is indistinguishable from a guard that would fail if it did. Nothing currently ensures it ever runs: no freshness contract, no scheduled quiet-window invocation, no tracking — which is precisely the unexecuted-standing class §11.4.226 names, where registration gets no coverage while execution is nobody's job. This is forensic precedent in that anchor is two standard guards that, when finally executed, immediately emitted latent FAILs. ACCEPTANCE. (a) Execute quiescence-gated timing tests on a host whose minute loadavg is genuinely $\leq 0.75/\text{cpu}$ — no fleet, no concurrent build — and capture the artifact showing the resolved load reading alongside the measured span exponent, so the precondition is provable from the artifact and asserted. (b) Record the observed GREEN polarity the item and in the test source, replacing the honest-inference note. (c) If the gate FAILS under genuine quiescence, that is the finding this item to surface, and it reopens the threshold-calibration question rather than being worked around. (d) Establish a freshness contract per §11.4.226: if a quiescent verdict may be before the gate could be uncovered again, so this does not silently return never-executing. HONEST BOUNDARY. This item does not claim the gate is wrong. The evidence points the other way — quiescence is strictly more favourable.

ATM ID	Type	Status	Severity	Description
				than the loads where GREEN was already observed. The bug claims only that the run has not happened, and that an unexecuted guard cannot be cited as covering the PROVENANCE. Raised as IMPORTANT-2 by the independent Fable-substrate reviewer of the work, and independently corroborated: the reviewer verified the SKIP is loud and prints its resolution numbers, and confirmed the shipped tests do not fail on today's load.
BOB-172	Bug	Ready for testing	High	WHAT. The rutracker SEARCH endpoint is returning a 404 bot protection while the site itself is fully reachable. Measured unauthenticated 2026-08-21 (no cookies, no read, no credential value in any artifact — \$1 GET /forum/index.php -> 200, 96283 B, WHAT. The workable-items update seam guard is not status-location invariant in one direction only. BOB-166, update --status <terminal> on an item located row is refused. The mirror is still open. update --location Fixed --status <non-terminal> exits 0 and creates a row that update's own validate immediately rejects. Verified empirically by the independent reviewer of BOB-166: update --location Fixed --status 'In progress' -> exit 0 validate --location Fixed naming the row (check (f), fixedLocationNonTerminalStatus) So the command can still mint a state its own validate call refuses. This is the §11.4.196(F) shape at the seam layer: the invariant is CONFIGURED in validate but not ENFORCED at every write path that can violate it. WHY IT WAS DELIBERATELY LEFT OPEN IN THE FIRST PLACE and why that was right. Three reasons, all checked rather than asserted: (1) the real corpus has 2 instances of this direction against TEN of the BOB-166 closed — the asymmetry in the data is the asymmetry in the code; (2) detective coverage ALREADY existed and demonstrably fires, so BOB-166's direction this cannot rot silently; (3) the load-bearing one — reopenCmd documents SANCTIONS a transient Fixed-plus-non-terminal via its --location Fixed operator override. A preventive guard at the update seam would conflict with that override's semantics. Closing this the
BOB-175	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				requires a design decision about how the two which is a different question from the one BOB-166 was scoped to answer, and expanding that item would have meant shipping the decision unexamined. BOB-166 ACCEPTANCE. (a) Decide and RECORD how a preventive guard on this direction coexists with reopenCmd's sanctioned override — this is the work, and the answer should be written down rather than inferred from whatever the patch does. Candidates worth weighing: exempt the reopenCmd explicitly, require an override flag on updateCmd mirroring reopen's, or leave the direction decision only with the rationale recorded so the asymmetry is a decision rather than an oversight. (b) If a guard lands, it reuses terminalStatuses() — BOB-166 established one predicate source specifically for the two directions cannot drift. (c) Paired §1.1 mirror (d) NEGATIVE CONTROL (§11.4.201(1)), and it is a sharp one here: the sanctioned reopen path MIGHT work. A guard that breaks reopenCmd's documented override is a false-positive refusal, and worse than the gap it closes, because it breaks a workflow open to be told to use. HONEST BOUNDARY. This is not a data defect. Zero corpus instances, and the defect gate catches it at the next validate. It is filed as a scoped, deliberate omission stays visible as a decision instead of living only in a code comment where the next reader will mistake it for an oversight (§11.4.197: a started thread reaches a documented terminal state, never quiet abandonment). BOB-166 PROVENANCE. Raised by the implementing a BOB-166 as a known asymmetry it deliberately does not close, independently verified and endorses an acceptable scope boundary by that item's review the condition that it become a tracked item rather than a comment. This is that item. WHAT: the BOB-172 independent review demonstrated two further paths that still represent refusing or unreachable tracker as an empty item. Both are PRE-EXISTING and neither is a regression from BOB-172, but the fix evidence log presents five-site coverage without stating the boundary.
BOB-179	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				(11.4.194(5) requires un-analysed dimensions explicit gaps, never silently assumed safe). G exception before resp.status is read. Reviewer B, captured: a connection-refused rutracker y status=empty, error=None, http_status=None, metadata.errors=[], metadata.status=complete <i>search*</i> method swallows exceptions (except Exception: logger.error; return results) before _search_one's error handling can see them, s tracker that is DOWN is indistinguishable from that is genuinely empty. GAP B – soft refusal 200. Reviewer probe A, captured: _classify_upstream_http_status(200,) returns which is CORRECT by design (a 2xx is a usable response; body-marker triggering would risk over-fire the negative controls exist to prevent all five search GETs use aiohttp's default allow_redirects=True, so a session-expiry 302 page-200 chain parses to zero rows and report empty. FIX DIRECTION: Gap A needs the exception reach a diagnostic rather than being swallowed. B needs a DISTINCT detector (final-URL check login-form marker) with its own RED – explicit a widening of the status-code trigger. ACCEPT both gaps closed with their own REDs, or explicit closed per 11.4.112 with evidence. Immediate task: append a stated-gaps paragraph to docs BOB-172/fix_evidence_20260822.log. Three icon-glyph controls (.bridge-retry, .then toggle, .caret) render text made of non-BMP points. axe-core declines to decide their contrast reporting them under the incomplete channel messageKey nonBmp and resolving NEITHER NOR bgColor, so the Python recomputation in BOB-164/axe_contrast_scan.py cannot decide 40 such nodes across 16 palette x mode combinations BOB-164 round 2 stopped them being scored passes by naming them GLYPH_UNMEASURED printing them on every run and fencing them NEW unmeasured glyph node blocks. What remains genuinely unproven is their contrast itself: as user-interface components they owe 3:1 under
BOB-184	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				SC 1.4.11, and no oracle in this repo measures today. Acceptance: a measurement path that the effective foreground and backdrop for glyph controls (for example reading getComputedStyle colour against the composited backdrop, or rendering the glyphs with inline SVG carrying explicit font tokens), each of the three controls shown to occur in all 16 palette x mode combinations, and the GLYPH_UNMEASURED fence shrunk by exact nodes then proven. docs/qa/BOB-164/axe_contrast_scan.py scans dist with no backend, so every node whose existence depends on fetched data is absent from the page measures: the results table renders empty, so .quality-badge and .quality-badge never appear, and the list renders empty, so .status pills never appear. This is precisely why the BOB-164 round-1 regression test is unseen in a rendered DOM even though it was scanned in all 16 themes, and it is why round-2 tests close the class in the arithmetic oracle instead. Nodes are now covered by frontend/src/app/mocks/style-contrast.spec.ts, which extracts declared foreground-on-fill pairs from the stylesheets, and is an ARITHMETIC claim about declared color, not a rendered-pixel one: it cannot see opacity by an ancestor, a composited backdrop, or a color. This repo's static extractor does not model. Acceptance: the scan drives the dashboard with seeded fixture data (a stub backend, a route for an injected component harness) so the badge and status nodes are present in the DOM, axe measures them directly, and a control needle proves the rendered oracle sees a seeded sub-floor badge. Any clean result from it is believed. scripts/ownership_precondition.sh consumes the SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe files in but it is NOT fenced. T028 remediation added ownership_path_fence() to scripts/lib/ownership and wired it into ownership_repair.sh, closing the path-escape class there; the precondition was that agents file scope and still accepts whatever
BOB-185	Task	Queued	Medium	
BOB-187	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
BOB-188	Bug	In progress	—	config/owned_paths.yaml expands to. Blast ra lower than a recursive chown (it writes prob rather than mutating ownership of an existin but the INPUT is identical and equally unrevi the same QBITTORRENT_DATA_DIR value tha have driven a filesystem-wide chown will dri probe-file creation at an arbitrary location. N escape is empirically confirmed for the repai not merely reasoned: during T028 remediation probe declaring the verbatim shipped entry v QBITTORRENT_DATA_DIR=/ hung the suite, a showed the real artifact executing “find / (! -o ! -gid 1000) -printf ...” before it was killed I verified pid. Acceptance: ownership_precond calls the SAME ownership_path_fence() predi from scripts/lib/ownership.sh (never a second 11.4.251), refuses fail-closed on a rejected ent ships a paired 1.1 mutation proving the refus PLUS a negative control proving the six live s entries still ACCEPT so the fix is not a 11.4.20 positive refusal. Discovered out-of-band duri remediation and reported by the agent as ne tracking, so it is also a 11.4.238 coverage esca OPERATOR DECISION (2026-08-26, §11.4.66 interactive clarification): UNTRACK THE B — BUILD ON DEMAND §11.4.30 answer, reco consumer DATA: the compiled workable-item is NOT to be git-tracked. ‘Any build derivate v can recreate by executing proper mechanism generating MUST NOT be versioned’ applies v exception here. Keeping it tracked with only invariant-52 fingerprint gate, and keeping it t with a commit-time rebuild, were both offere NOT chosen. REQUIRED END STATE: the gate from source, or REFUSES HONESTLY when th toolchain is absent (§11.4.201(11) — probe th through its real invocation path, never fake a never silently fall back to a stale binary). Stal becomes structurally impossible rather than detected. ACCEPTED COST: a fresh clone need toolchain before invariant 17 can run — alre for the qBitTorrent-go backend, so this adds n

ATM ID	Type	Status	Severity	Description
				host requirement. The invariant-52 fingerprint remains useful as defence-in-depth for as long as the binary exists on disk. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent inference, and supersedes any prior agent-chosen default on this question. Options not chosen are named above so a future reader does not re-litigate a settled call (§11.4.35). A bounded-verdict discipline applied to decision-making: prior item text follows — pre_build_verification: invariant 17 (CM-WORKABLE-ITEMS-VALIDATION) resolves its binary through the candidate loop, whose FIRST entry is constitution/scripts/workable-items/bin/workable-items. That file is GIT-TRACED (md5 17644a248363, identical to bin/workable-items/linux) and executable, so it WINS resolution over current untracked sibling constitution/scripts/workable-items/workable-items (md5 43376a248363). The tracked binary is STALE: it does not contain guards its own source now has. MEASURED, 2026-08-25. proven, 2026-08-25. String presence in the tracked binary: “refusing to set terminal status” -> 0; “location item has TERMINAL status” -> 0; “cannot find needle from the SAME updateCmd, “at least one mutable field flag is required” -> 2 (so the inspector sees through that path and the zeros are real absences, not a blind read). The untracked sibling returns 1 / 1 / 2 for the same three queries. A violation whose message string is absent from the binary cannot fire. RUNTIME PROOF: the same injection violation (a row set to terminal status while current_location=Issues) run through both binary stale reports 1 violation (only the body_md mismatch class), current reports 2 including the location mismatch class verbatim. This is exactly why the ten BOB rows (BOB-087/129/131/144/145/146/148/153/154/155) survived undetected: the gate that was supposed to catch them was running a binary structurally incapable of seeing them. The 11.4.108 SOURCE >ARTIFACT gap: source correct, artifact stale, gate reading the artifact reports green. This is the THIRD instance of that class found in one session alongside BOB-169 (a charset gate that never

ATM ID	Type	Status	Severity	Description
				file) and BOB-183 (a served bundle five days older than its sources). COMPOUNDING: the comment at :519-523 directly above the loop asserts “The bin/workable-items path never existed in this checkout (bin/ is a gitignored local build-output directory, nothing ever populated)”. That statement is not FALSE - bin/ holds two tracked binaries. A stale comment asserting the absence of the very files now shadows resolution is how this stayed in. Also an 11.4.30 question the operator owns: bin/workable-items and bin/workable-items-linux VERSIONED BUILD ARTIFACTS in a constituent submodule that is inherited by reference by a consumer, so every consumer inherits which binary was last committed. Flagged independent two separate agents this session. ACCEPTANCE resolution must never prefer a stale artifact over current one - either the tracked binaries are new and resolution falls through to build-on-demand, freshness check refuses a binary older than its sources (see the BOB-183 fingerprint gate for working pattern); (b) the false comment at :519 corrected; (c) a paired 1.1 mutation proving the refusal fires, plus a negative control proving a genuinely fresh binary still passes so the fix is 11.4.201(1) false-positive refusal; (d) the 11.4.201 tracked-binary decision recorded as operator error. Discovered out-of-band during BOB-166 removal of a 11.4.238 coverage escape in its own right. WHAT: the §11.4.252 fail-open scan reports 0 qBitTorrent-go, and that zero is NOT evidence of a triage agent ran control needles through the wrong own path per §11.4.201(7)(b): a Python ‘except Exception: pass’ needle was SEEN, a TypeScript (e) {} needle was SEEN (so frontend/src’s zero is real zero), but a Go empty-‘if err != nil {}’ needle was NOT SEEN. An instrument that cannot see the tree returns the same quiet zero as a clean tree, and one of those is honest. DISTINCT FROM BOB-189 deliberately not merged with it per §11.4.214. BOB-189 is a false MATCH (fail-closed guards as fail-open); this is a false NULL (an entire la
BOB-191	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				unscanned). Same scanner, opposite failure directions, different fixes - merging them would be one behind the other. IMPACT: Go is the language qbittorrent-proxy-go and boba-jackett (port 7711) which owns encrypted tracker credentials), so the unscanned surface is exactly where §11.4.252. credential-plus-mutation combination is most likely. Nobody has assessed it; the dashboard says c. ACCEPTANCE: the scanner grows a Go arm compared to the empty-err-block and swallowed-error idioms. The needle is SEEN through the real path, qBitTorrent result is re-derived, and every finding is triangulated in this Python/TS pass was. Until then qBitTorrent fail-open posture is UNKNOWN and must be as UNKNOWN, never as 0. === ROUND-2 INVESTIGATION, MEASURED 2026-08-26 — THE FILED TITLE WAS TOO NARROW (§11.4.6) == scanner is constitution/scripts/gates/cm_dangerous_combination_fail_closed.sh, d per-root by invariant 39 at scripts/pre_build_verification.sh:1481-1544. (scripts/pre_build/check_cm_no_fail_open_skip.sh has existed on any branch — it exists only in an uncommitted worktree. Any citation of that path citing a file that is not there.) IT IS A WEAK-M DEFECT, NOT A SCOPE HOLE. Line :243 sets e go rs c cc cpp h hpp java cs js ts jsx tsx php rb IS enumerated. But line :537 gates the Python analyser on case "\$f" in *.py), so every non file falls through to just two text matchers (:7 catch (...) {, :794 credential = x \\ \\ "lite has no catch keyword, so neither matcher CA Go by construction. NEEDLE TEST, BOTH ARM (planted / rc / result): Python 2 / 1 / SEEN · Typ 1 / 1 / SEEN · Java 1 / 1 / SEEN · Go 6 / 0 / NOT PASS · Rust 2 / 0 / NOT SEEN · Ruby 1 / 0 / NOT 2 / 0 / NOT SEEN. Go is ONE OF FIVE BLIND EXTENSIONS, not a special case. The Go need code: go build rc=0, go vet rc=0, gofmt -e clear DECISIVE DISCRIMINATOR — same bytes, ex changed: .zig (NOT enumerated) → “SKIP — topology_unsupported”; .go (enumerated) → ‘ The gate ALREADY OWNS an honest-blindness

ATM ID	Type	Status	Severity	Description
				mechanism and it works correctly; Go routes it precisely BY BEING ENUMERATED. A scope announces itself; this matcher hole prints gre makes it the worse defect and is why the corn title leads with it. QUANTIFIED SURFACE (ex stated, §11.4.224(E)): using the scanner's OW list and nothing added — 106 .go files / 20,389 production files / 6,858 LOC). WITH and WITH exclusions the count is 106 == 106; no vendor exists, so no unstated exclusion could manuf clean number. qBitTorrent-go holds 106 blind 0 analysable files — 100% blind — yet invari counts it among “5 first-party source root(s)” clean. DEFECTS THE BLINDNESS HID → filed BOB-204 (credential DELETE returns uncond 204 while the .env plaintext delete error is di plus the env_write_failed_db_rolled_back coo asserting an unconfirmed rollback). SEPARAT HOLE → filed as BOB-205 (cmd/boba-ctl, 947 I orchestrator, absent from DANGER_ROOTS er Distinct-but-similar per §11.4.214, deliberately merged. GOLDEN-FALSE FIXTURE FOR ANY I GO ARM: internal/jackettapi/ auth_middleware.go:45-66 is correctly fail-CL (constant-time compare, 401+return) — a Go MUST NOT flag it. LOAD-BEARING §11.4.250 PREDICTION: bolting a naive Go regex arm o IMMEDIATELY manufacture BOB-189-class fa positives in Go — if err != nil { return fa a Go validator is fail-CLOSED. Fixing this as l reproduces BOB-189 in a new language. THE the remediation order is fixed: (1) analyser R with an UNANALYSED verdict FIRST — extension→analyser map, unmapped extensi reports UNANALYSED never silent PASS, whi converts the false null into an honest gap and rust/ruby/C too; (2) only then a Go arm, call-s from day one (go/ast, or adopt errcheck/statio rather than hand-rolled regex), shipped with auth_middleware golden-FALSE. ONE PRIMIT TWO SYMPTOMS: the scanner classifies by lo syntactic shape and never consults semantic BOB-189 ignores the CALLER's context (false

ATM ID	Type	Status	Severity	Description
				BOB-191 searches for a shape from a DIFFER language family (false null). The gate's own h already concedes the primitive ("requires rea flow analysis this gate CANNOT honestly clai what is missing is propagating that concessio PER-LANGUAGE honesty. Keep two tracker it INSTRUMENT BUG RECORDED, NOT HIDDEN (§11.4.201(12)): the investigation's first verdic extraction used <code>grep -oE 'PASS\ FAIL\ SKIP'</code> matched FAIL inside the gate's own name (... CLOSED) and mislabelled three PASSes. Re-ru on exit code; the table above is the corrected EVIDENCE: docs/qa/BOB-191/ scanner_blindness_investigation_20260826.m lines, §11.4.44 header). No source, gate, or co modified by the investigation. WHAT: agent worktrees under .claude/worktr ephemeral copies pinned at arbitrary commi built, never shipped, deleted when the agent Pre-build gates that walk the tree with find(1 the repo root descend into them anyway. ME INSTANCE (2026-08-25): CM-GO-TOOLCHAIN- MATCHES-BUILDER FAILED the main build w exactly one finding, from .claude/worktrees/a afda1df906c537ab4 - a worktree 86 commits still carrying 'FROM golang:1.23-alpine' again go.mod long since at 1.26.2. The main tree wa CORRECT throughout. Per §11.4.201(1) a false refusal is a FAIL-bluff of equal severity to a fa it blocks real work and teaches people to hyp gate. FIXED FOR ONE GATE: '.claude' added t gate's PRUNE_DIRS with a deterministic §1.1 stale mismatch inside .claude/worktrees mus fail a healthy tree, the SAME mismatch outsid MUST still fail, so the prune cannot widen int blindness (§11.4.201(6)). Mutation verified: re the prune fails 2 checks. SURVEY, measured r inferred: 8 of 13 pre_build gates walk the tre the fixed one prunes .claude. Gates using 'git are structurally immune - worktrees are untr Empirically probed: check_cm_closure_seam check_cm_killpg_pgid_guard and
BOB-194	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
BOB-196	Task	Queued	—	check_cm_no_production_mutation_residue . with ZERO worktree references - clean in pra today. UNMEASURED, not clean: check_cm_test_mock_pid_explicit_int and check_cm_test_mock_pid_patched_when_rea were not reached before the probe budget ex HYPOTHESIS REFUTED, recorded so nobody it: an earlier revision of this item speculated check_cm_plugin_count's 90s timeout was ca walking worktrees. It is NOT. That gate's find to \$PLUGINS_DIR (check_cm_plugin_count.sh worktrees are outside its walk entirely. Its slo real - still running past 54s on a re-measure - different, still-unidentified cause and belong own item, not this one. ACCEPTANCE: every t walking gate either prunes agent scratch or i immune by measurement, each with a paire mutation. NOTE the asymmetry that bounds a stale worktree can only ADD findings to a s relevant scan (killpg, mutation-residue), neve remove them - so the exposure is false refusa eroded trust, not a missed defect. WHAT: the CM-PLUGIN-COUNT pre-build gat 189 SECONDS to verify 8 documented counts. Measured 2026-08-25: exceeded a 90s probe k (rc=124), then a bounded re-run completed a wall=189s' - it PASSES, it is simply slow. CAUS MEASURED NOT GUESSED: the marker loop a check_cm_plugin_count.sh:273-290 iterates E LINE of every governed document and spaw pipeline per line. Per iteration: 'printf grep -l tr -d' (4 processes) plus 'printf sed -n' (2 Governed docs are AGENTS.md 689 + CLAUD + docs/features/Status.md 670 = 1853 lines, so process spawns, measuring ~102 ms/line. Cla line-subprocess shell trap. THE IRONY WORT PRESERVING, because it explains why nobod simplified it: the in-code comment shows 'oE was chosen DELIBERATELY over 'grep -c' afte measuring the §11.4.201(12) footgun - on this (ugrep 7.8.4) 'grep -coE' returned 3 at top leve inside a 'set -euo pipefail' subshell, a context-

ATM ID	Type	Status	Severity	Description
				dependent reading. The correctness fix is right; it must be kept; it is the PER-LINE application of the fix that costs the three minutes. HYPOTHESIS REFUTED: the fix was recorded so nobody re-walks it: this is NOT what we want. The gate's find is scoped to \$PLUGIN_ROOT (224), so the five agent worktrees are outside the gate. That guess was raised on BOB-194 and is dead. The fix (unverified, stated as a candidate not a conclusion): pre-filter with ONE grep for lines containing 'CM-PLUGIN-COUNT:' before entering the loop, reducing ~1853 iterations to the ~8 lines that carry a marker, while leaving the per-line parsing logic - and its footgun-avoiding form - completely unchanged. WHY IT MATTERS beyond impact: this gate is on the pre-build critical path, and CLAUDE.md leans on it as the mechanical authority for the three distinct plugin rosters (43 curated engines / 12 bootstrap) whose conflation was the problem. A gate slow enough to tempt anyone into skipping it protects nothing. ACCEPTANCE: runtime reduction of the counts and the footgun-avoiding parse proved byte-for-byte in behaviour, proven by the existing test paired mutation still biting plus a before/after clock recorded. §11.4.6: nothing here questions the gate's VERDICTS - only its runtime. STATUS: LATENT on this host (measured: all system declared locations sit on ONE btrfs mount, full ownership-expressive) and OPERATOR-GATE verification. Retained from the dismissed BOB-210 rather than discarded with it, because the scenario is well-founded elsewhere: the declared mount is a udisks auto-mounted REMOVABLE NVMe, so the operator's portable drive being exFAT or NTFS is entirely ordinary. TWO RESIDUALS: (1) REPAIR DIAGNOSIS. On a uid-flattening mount whose flattened uid is NOT the operator, probe_location correctly REFUSES (verified: uid=0 and uid=1 both refuse). But it then hands ownership_repair_chown that will return EPERM, and nothing in the repair path is fstype-aware — so the operator's permission failure whose real remedy is a REPAIR_OPTION, which no message says. Reasoned, N
BOB-211	Bug	Queued	minor	

ATM ID	Type	Status	Severity	Description
				proven: the shim used to verify BOB-206 mode stat READ, not the chown WRITE, and cannot whether chown actually fails. (2) MODE CON vfat flattens MODE via fmask/dmask, so the preserve_mode: true / 600 contract on .env and boba.db would be silently unenforceable then §11.4.10-adjacent: a credential file believed to that is world-readable by mount option is a live mode check cannot see. WHY OPERATOR-GATED genuine end-to-end RED needs a real uid-flattening mount. The privilege-free route is bindfs -u (if fusermount IS present and user_allow_other but bindfs is ABSENT on this host, as are mkfs and losetup. So verification needs either a bindfs install or a privileged loopback mount. Neither attempted (§11.4.21 — high-blast-radius host is operator-gated). ALSO RECORDED: NO test anywhere exercises a uid-flattening filesystem exfat/ntfs/flatten/loopback/mkfs/losetup/bind fusermount all measure 0 across the three over test files, needle-proven (probe_location = 3 h through the same instrument). ACCEPTANCE operator decides whether to install bindfs or a privileged loopback so the RED can be built verified, the repair path gains fstype-aware naming the remount remedy; (3) the mode checker either detects mode-flattening mounts and records documents honestly that it cannot be enforced DISCOVERY CHANNEL (§11.4.238): retained from the BOB-206 verification stream. WHAT, conductor-verified: docs/guides/tracker-credentials.md is TRACKED. Its §11.4.65-manual export twins docs/guides/tracker-credentials. and .pdf both EXIST ON DISK, are NOT tracked ARE ignored — matched by the .gitignore credentials deny-all at :31. The .md is rescued by an explicit allowlist entry at :38; nobody added entries for twins. WHY IT IS A LIVE VIOLATION AND NOT HYPOTHETICAL: §11.4.65 requires every in-scope Markdown document to carry synchronized .html/.pdf siblings, and §11.4.2 requires every §11.4.65-scope doc to be reachable
BOB-219	Bug	Queued	major	

ATM ID	Type	Status	Severity	Description
				from the README. Two artifacts that exist lo can never be committed are, from any fresh ABSENT — so every clone of this repository h tracker-credentials guide with no exports, wh authoring host looks complete. The divergen invisible on the machine that would notice it NOBODY SAW IT: this is a §11.4.201(6) FALSE the BOB-212 class. git status never listed the t no gate and no author had a signal. The blind instrument and the clean tree return the ider quiet zero. ANOTHER INSTANCE OF THE SAM MECHANISM, filed here rather than separate because it is one fix: docs/qa/BOB-124/eviden loginctl_user_state.txt is swallowed by *_user at .gitignore:~55 and is therefore uncommitt §11.4.83 QA evidence. Its SIBLINGS IN THE SA DIRECTORY — user1000_grepped.log, user_scope_events_head.log, user_service_lif — ARE tracked, saved only by the robust glob negation !docs/qa/**/*log. That one file is lost because it is .txt rather than .log. The pattern robust glob negations hold; per-file and per-e rescues leak. ACCEPTANCE: (1) both twins be trackable and are committed alongside their the BOB-124 .txt evidence likewise; (3) which BOB-212 fix direction is chosen, it MUST cove — a fix that closes the future-swallow while l these three artifacts permanently uncommitt addressed the mechanism and not the damag check that a tracked .md in §11.4.65 scope ha COMMITTABLE twins, so this class cannot re silently. DISCOVERY CHANNEL (§11.4.238): fo the BOB-212 blast-radius sweep — and only c SECOND pass. The first pass filtered by a sour extension set that omitted .pdf, which hid thi entirely; the agent re-ran with no extension f all 420 non-artifact paths and recorded the b rather than shipping the first number. Worth an extension allowlist is itself a false-null ger which is the same shape as the defect being investigated.
BOB-223	Bug	Queued	major	

ATM ID	Type	Status	Severity	Description
				WHAT, both conductor-verified with control 1. CM-SCRIPT-DOCS-SYNC occurrences in script pre_build_verification.sh = 0 (control needle MARKDOWN-EXPORT-SYNC = 6, so the instrument sees). §11.4.44 revision-header enforcement 1 hits (control needle 3). Neither §11.4.18's companion doc mandate nor §11.4.44's revision header is enforced by any invariant. THE PERVERSE CONSEQUENCE, MEASURED: a new script with companion doc passes the gate. Writing the companion doc that §11.4.18 MANDATES creates scripts/.md, which invariant 16 (CM-MARKDOWN-EXPORT-SYNC) then requires to have .html and twins — and their absence is a BLOCKING failure. This session hit it exactly: docs/scripts/test_gitignore_swallow_is_loud.{html,pdf} are and are two of invariant 16's six violations. S gate's incentive gradient points AWAY from the constitution: complying with §11.4.18 is punishing, ignoring it is free. That is worse than an ungoverned — it is a rule the machinery actively discourages. §11.4.227 IS THE GOVERNING ANCHOR: a name that exists only in prose is gate debt, and the unimplemented CM-* names must be monotonically decreasing. CM-SCRIPT-DOCS-SYNC is named §11.4.18's own text and implemented nowhere. concrete row of exactly the 58%-unimplemented population §11.4.227 was minted to shrink. ACCEPTANCE: (1) either CM-SCRIPT-DOCS-SYNC implemented, or it is registered as a deferral at a tracked §11.4.197 item — silent absence is §11.4.227 forbids; (2) same for §11.4.44's header check; (3) whichever is chosen, the incentive inversion is closed: it must never be cheaper to mandate doc than to write one — if invariant demand twins, the doc-creation path must produce them, or invariant 16 must scope out docs/scripts it can; (4) a paired §1.1 mutation per gate implemented. HONEST NOTE ON SCOPE: this does not argue §11.4.18 should be enforced immediately — that is an operator call about priority (§11.4.66). It argues the CURRENT state is incoherent: an unenforced mandate whose

ATM ID	Type	Status	Severity	Description
BOB-226	Task	Queued	major	observance triggers a different gate's failure. enforce both ends or neither. DISCOVERY CH (§11.4.238): found by the T042 readiness pref while explaining why a new file broke invari Not by the automated QA regime. WHAT: ownership_repair walks a declared r repairs items whose uid is not the operator. T suite seeds a foreign uid only onto FILES and SYMLINKS; interior DIRECTORIES stay opera owned (case 22 seeds only a foreign declared on the failure path). Production first-start rep exactly the untested shape. MANIFEST: tests/ test_ownership_repair.sh seed_tree/seed_wro scripts/ownership_repair.sh walk at :994. REE a tree whose interior directories carry uid 10 podman unshare, run the repair, observe no automated assertion covers the outcome. WH UNTESTED: the unprivileged harness cannot symlinks inside a directory it no longer owns cases 8/9/18 require. The DECLARED GAPS no references tests/ownership/ test_container_writes_owned_files.py, but tha the CREATION side (FR-002), not the repair-si ACCEPTANCE: an integration-layer test (when unprivileged-harness constraint does not bin seeds foreign-owned interior directories, run repair, and asserts post-state ownership plus preservation. Surfaced by the BOB-207 indep review 2026-08-27.
				BOB-227 measured 3 of 4 artifacts of the CM- ROUTES-AUTHENTICATED gate untracked in criterion 1 (commit them) was already resolv unrelated prior commit before this session, c 2026-09-23. Criterion 2 was NOT addressed: 'a asserting that every executable a pre-build in invokes is itself tracked' – a general mechanis preventing this whole CLASS of defect (a pre- gate's own implementation files shipping unt invisible to a fresh clone, no committed base round-over-round diffs) from recurring for A future gate, not merely this one. ACCEPTANC enumerate every file path scripts/
BOB-231	Task	Queued	Major	

ATM ID	Type	Status	Severity	Description
				pre_build_verification.sh invokes (via bash/ti python3 calls to scripts under scripts/pre_bui every tests/pre_build/.sh and tests/hooks/.sh i likely via a static grep/parse of pre_build_verification.sh itself, or a runtime assert every one of those paths is git-tracked files –error-unmatch); (3) wire this as a new p invariant so a future untracked gate impleme is caught immediately, not discovered indepe weeks later; (4) a RED test creating an untrac gate-invocation target and asserting the new fails on it, GREEN once the mechanism exists fake target is either removed or tracked. Found 2026-09-23 while remediating BOB-192 the fail-open skips were fixed, the live run of integration/test_merge_api.py gave 10 x HTTP (hooks x5, magnet x3, download x2). The run merge service has BOBA_API_TOKEN set (BOB mandatory auth guard), so POST /api/v1/mag answer 401 ‘valid API token required’; the tes requests without the Authorization header. E failing assert precedes any line BOB-192 char HEAD fails identically: these tests were previ masked or never run against an armed servi the integration fixtures must read BOBA_API (env or .env, never printed, §11.4.10) and sen when no token is available and the service do one, FAIL loudly rather than skip. RED-first: l http.server fixture answering 401 without th Found 2026-09-23 running tests/integration/ test_tracker_auth_live.py against the live stac BOB-192: test_private_tracker_credentials_authenticat fails with authenticated=False status=‘empty’ UNKNOWN whether this is expired cookies/ credentials, tracker-side bot protection, or a p parse regression (compare BOB-176/BOB-178 touched the kinozal path and BOB-172 for ru Needs systematic-debugging: capture the raw response classification, distinguish credential from empty parse, and only then decide fix v operator action (credential refresh is an oper
BOB-234	Bug	Queued	—	
BOB-235	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				§11.4.10). Not to be re-skipped: the test verdict is genuine.