Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Operator-blocked	1
Bug	Queued	8
Task	In progress	2
Task	Operator-blocked	3
Task	Queued	32
TOTAL		**46**

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-host routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download-proxy + qBitTorrent-go + Jack compose env-forward
BOB-068	Bug	Queued	High	RD2-00: unattributed, unreviewed Auto-c mechanism pushing to main
BOB-069	Task	Queued	High	RD2-40: §11.4.238 QA-discovery-channel — ongoing coverage-escape backfill
BOB-070	Bug	Queued	High	RD2-41: pre-build mutation-marker scan false-positive silently defeats entire pre-b gate
BOB-071	Bug	Queued	Medium	RD2-01: guard-forbidden-commands.sh h live reproducible substring carrier false-
BOB-074	Task	Queued	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
BOB-076	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				RD2-09: submodules/jackett fork 1 comm behind upstream (informational)
BOB-077	Task	Operator-blocked	High	RD2-10: Identify second host running the commit rsync/sync mechanism (OPERAT DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-comm mechanism through §11.4.234 dedicated commit/push script OR retire it
BOB-079	Task	Queued	Medium	RD2-12: Retroactive attributed history n GA-18/21/22/25/26/27 changes (never re published history)
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code re the substantive Auto-commit diffs
BOB-081	Task	Queued	High	RD2-14: Author CONTINUATION.md Ses entry (currently 53 days / 24+ commits b HEAD)
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable for the four Lava-porting findings (closes
BOB-083	Task	Queued	Medium	RD2-16: Regenerate browser_extension/ codegraph Status.md + Summary/HTML siblings
BOB-084	Task	Queued	High	RD2-17: Reconcile BOB-008 DB/MD body via the workable-items tool
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes G
BOB-086	Task	Queued	Medium	RD2-19: Fix BOB-009/BOB-010 evidence backfill item_history for 56 silent closure
BOB-087	Task	Queued	High	RD2-20: Wire docs_chain / commit-seam hook per §11.4.106(F) so DB writes cann without MD mirror
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Track Items + Status Documents table row-completeness (GA-07 remainder)
BOB-089	Task	Queued	High	RD2-24: RED-first tests for start.sh reload_python/reload_plugins/recreate_s (closes test-half of GA-27)
BOB-090	Task	Queued	Medium	RD2-25: HelixQA Challenge entry exerci three start.sh subcommands end-to-end real compose stack
BOB-091	Bug	Queued	High	RD2-26: Relocate mocked SearchOrches tests to unit/ + author real-service repla (closes GA-14/15/16)
BOB-092	Bug	Queued	Medium	RD2-27: Remove test_live_stack_evidenc nnmclub SKIP-on-404 fallback + verify li (closes GA-13)

ATM ID	Type	Status	Severity	Description
BOB-093	Task	Queued	High	RD2-28: Live compose bring-up + verify rutracker ReDoS regex bounds deployed container (closes runtime half of GA-12)
BOB-094	Task	Queued	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with § fault injection
BOB-095	Task	Queued	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos.p Go-side triangle
BOB-096	Task	Queued	Medium	RD2-31: Extend qBitTorrent-go jackett_db_test.go with real process-kill/ resource-exhaustion fault injection
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for exposed download-proxy/merge endpoint (canonical impl of RD2-07)
BOB-098	Task	Queued	Medium	RD2-34: Parametrize 20 hardcoded /Volu paths in helixqa banks with PROJECT_RO (closes GA-23)
BOB-099	Bug	Queued	Medium	RD2-36: Fix guard-forbidden-commands. substring-match false-positive class + ad const033-poweroff-signal-triage carrier t EXCLUDE_PATHS
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one c (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is –profile go parity still a goal? (gates RW-10..13) — OPERATOR-DECISION
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — bi tunnel 127.0.0.1 or keep 0.0.0.0? — OPE DECISION
BOB-104	Task	In progress	Medium	Coverage-escape followup (docs/ QA_DISCOVERY_LEDGER.md, BOB-075 code-reading finding, commit e6162f7): CodeGraph 1.5.0 (up from documented 0 walked into nested-.gitignore-excluded fr node_modules and extension/node_modu (32,260 files / 514,456 nodes vs the 2026 baseline of 509 files / 8,906 nodes) instea honoring frontend/.gitignore + extension/.gitignore per git check-ignore Author a challenge (challenges/scripts/ codegraph_gitignore_honor_challenge.sh equivalent, e.g. wrapping 'codegraph doc –sniff-gitignore-honor' if that subcomman exists, else a real re-index + file-count as with §11.4.115 RED_MODE polarity: RED_MODE=1 reproduces the blowup ac

ATM ID	Type	Status	Severity	Description
				the live nested-.gitignore tree, RED_MOI asserts the resync stays within the docur baseline order of magnitude. Full eviden codegraph/Status.md lines 91-118. UPST FILED 2026-08-18: real upstream repo is github.com/colbymchenry/codegraph (np package @colbymchenry/codegraph, con via package.json + gh repo view; NOT va digital/codegraph, correcting this ledger original assumption). Issue: https://githu colbymchenry/codegraph/issues/1567 (fu reproduction evidence: git check-ignore verified first-hand; two bounded syntheti attempts up to 63 nested .gitignore files files against the actual installed v1.5.0 bi NOT reproduce the blowup in isolation – negative result, root cause not pinned to specific line in either scanning implemen Draft PR (diagnostic only, not a behavior https://github.com/colbymchenry/codegr pull/1568 – adds a logDebug() call so a fu report can confirm which of the two inde gitignore-respecting code paths (git-ls-fil based vs filesystem-walk fallback) actual Both PRs/issue filed via SSH (Hard Stop from a fork at github.com/milos85vasic/ codegraph. Status: In-progress pending upstream maintainer triage/merge – boba closure of this item still needs the RED/G challenge script (§11.4.115) authored pe original scope, independent of upstream' response.
BOB-105	Task	Queued	Medium	Coverage-escape followup (docs/ QA_DISCOVERY_LEDGER.md, meta-defe during this session's §11.4.140/141 collis resolution, boba commit 136a22c + cons commits 5ed8c80..e5f2891): §11.4.227(B anchor-block-integrity (exactly-once head anchor per governance file, byte-identica lockstep across mirrors, no anchor-numb collision) is fully specified in constitution but its own named gate CM-ANCHOR-BL INTEGRITY is explicitly 'gate-code = sep work item, NOT claimed shipped' — the §11.4.140/§11.4.141 double-mandate coll this session's Phase 0 resolved was verifi entirely BY HAND (md5 of moved anchor manual grep counts pasted into a commi message), not by a runnable script. Prop author a challenge (in this project's chall

ATM ID	Type	Status	Severity	Description
				scripts/, since this project cannot edit the constitution submodule's own gate-code this ticket) that greps every governance (CLAUDE.md/AGENTS.md/QWEN.md/ GEMINI.md/Constitution.md, wherever t project's constitution submodule checkou exposes them) for '### §11.4.NNN' hea occurrences and FAILs on >1-per-file-pe or on two DIFFERENT anchor bodies sha one NNN, per §11.4.227(B). Do NOT tou constitution-owned files while implement — the check reads them read-only.
BOB-106	Task	Queued	Medium	Coverage-escape followup (docs/ QA_DISCOVERY_LEDGER.md RD2-00/BO entry, docs/ GOVERNANCE_AUDIT_2026-08-08_ROU RD2-00): 20 bare 'Auto-commit' commits this repo's git history (confirmed via git l –oneline –all –grep, e.g. 54e313f/ 9c8f684/743097a/de9270b/ 1c36777/41179c2/7c529ca) with no ATM reference and no TDD trail, landing via o git pull fast-forward from a second sessio with push access to the same remotes (mismatched commit timezone vs the investigating host, per RD2-00 Update). § working-tree-quiescence has no mechani guard on this path — no gate flags a com reaching main with a bare/templated me and no ticket citation. Author a §11.4.84 quiescence-check helper (e.g. challenges no_unattributed_autocommit_challenge.s scans the commit range since the last kn good release tag and FAILs on any comm message matching a closed bare/templat pattern (e.g. ^Auto-commit$, ^sync:) wi ATM-NNN or task/PR reference, wired in scripts/pre_build_verification.sh or the §1 commit-push-all.sh entrypoint. BOB-068 (RD2-00) remains the tracking item for identifying/stopping the source; this item specifically the new automated CHECK.
BOB-107	Task	Queued	Medium	Coverage-escape followup (docs/ QA_DISCOVERY_LEDGER.md SCRATCH LOSS-2026-08-18 entry, evidence: .super sdd/task-phase1a-report.md line 21): the 1a subagent (a1cc331d) discovered at ta that 5 source files its brief named as req reads (curriculum_amendment_plan_v1.r ai_curriculum_modules_27_35_extracted

ATM ID	Type	Status	Severity	Description
				3 curriculum_analysis_modules_*.md gap analyses) were absent from the session scratchpad — root cause per that subage own investigation: the prior producer sub (ae59171f) hit its session rate limit (a §11.4.147(e) API-quota crash) before wri them. curriculum_amendment_plan_v1.m remains absent from the live scratchpad this session's re-check. No mechanical cl verifies a task brief's declared input files and are non-empty before the downstrea consumer subagent is dispatched. Autho dispatch precondition helper (project-sid orchestration tooling, e.g. a small script conductor runs before Task/Agent dispat a brief names required input paths) that closed with an actionable 'missing input, respawn the producer' message rather th silently letting a downstream agent proc absent evidence and fabricate content unsupported by its named sources.
BOB-109	Task	Queued	Medium	docs/testing/test_type_matrix.md's §11.4 type audit found zero scaling-class cover anywhere in the tree (no scaling-tagged directory, test file, or HelixQA bank distir growing-dataset/tracker-count/concurren scale-out from stress-under-burst). Scope least one scaling dimension, e.g. tracker- scale-out in merge search against challen helixqa-banks/boba-services.yaml's track or the qbittorrent-proxy-go –profile go sw with a real measured baseline.
BOB-110	Task	Queued	Medium	docs/testing/test_type_matrix.md's §11.4 type audit found UI functional coverage (Playwright) but nothing framed around u accessibility/UX outcomes specifically. Sc axe-core or equivalent accessibility pass Angular frontend, covering WCAG check keyboard-nav coverage, and screen-read labeling.
BOB-111	Task	Queued	High	Source inspection across the whole stack (2026-08-18) verified no rate-limit mecha exists for :7185 (qBittorrent WebUI prox (merge search service), or :7189 (boba-ja no slowapi/limiter/throttle import in dow proxy/src/, no rate-limit middleware in qBitTorrent-go/internal/middleware/ (onl cors.go + logging.go) or internal/jacketta (only auth/cors middleware tests, no rate

ATM ID	Type	Status	Severity	Description
				middleware), and no nginx/reverse-proxy in docker-compose.yml. Candidate remed documented in docs/testing/ddos_resilien nginx-in-container reverse proxy with limit_req_zone (most portable, adds a con a FastAPI slowapi dependency for the me search service; a Gin rate-limit middlewa boba-jackett following the existing intern middleware/ pattern. qBittorrent's own W bandwidth-shaping settings were NOT ve cover request-rate (only bandwidth) – do assume they close this gap without check
BOB-113	Task	Queued	Low	wrk is not installed on the development h ApacheBench (ab) is, and was used as th primary DDoS-resilience-challenge load t the task brief's 'wrk or ab' fallback path, curl-parallel-loop as the actual per-status assertion mechanism (more precise than aggregate summary) and ab's own outpu as supplementary real-tool evidence. Inst (e.g. via the project's setup.sh or a docum apt/build-from-source recipe) so future lo DDoS challenges have a modern HTTP benchmarking tool with latency-percenti histograms available out of the box, matc the brief's stated primary preference.
BOB-114	Task	Queued	Medium	challenges/scripts/ddos_resilience_challe –self-validate mode currently only ships a golden-bad fixture for the crash-resistanc detector (per §11.4.107(10)/§11.4.201). T limiting detector has no matching golden fixture proving it would actually FAIL a s no-rate-limit-enforced artifact, so an unv rate-limit detector could silently pass a b absent rate-limit deployment. Add a synt fixture (e.g. a stub server that never retu 429/503 under burst) and assert the dete correctly reports the absence, closing the validation gap for this detector class.
BOB-117	Bug	Queued	High	search.py:1329 (_search_rutracker passw login path) still classifies a no-session-co login result as error_type=auth_failure w message 'rutracker login returned no ses cookie — likely CAPTCHA wall or creden failure'. The adjacent code comment (line 1244-1248, same function) already states FACT that 'rutracker gates login.php beh CAPTCHA when logins spike… the CAPT the real blocker' — yet the error classific

ATM ID	Type	Status	Severity	Description
				still hedges with forbidden §11.4.6 vocab (likely) and reports the WRONG error_ty (auth_failure instead of upstream_captch unlike the exact same class of bug alread for nnmclub per docs/qa/nnmclub-login-diagnosis-20260616.md (which replaced credential failure' with a FACT-based upstream_captcha classification + explar comment explicitly citing 'This is NOT lik credential failure'). The rutracker code p could disambiguate the same way nnmcl does (or by reading the login response be captcha marker, mirroring the existing p used ~20 lines later in the SAME functio the search response) but currently discar opportunity and ships a guess. Found du §11.4.6 bluff audit (docs/qa/task-bluff-au direction: apply the same upstream_capt reclassification + FACT-based message u nnmclub, or add a captcha-marker check login response before falling back to a he message.
BOB-118	Bug	Queued	High	README.md line 28 ships a green/succe colored shields.io badge 'python tests-58 passing', and line 333 repeats '585+ test real, non-destructive verification this ses (timeout 90 nice -n 19 ionice -c 3 .venv/b python -m pytest tests/ –collect-only -q) c 5235 tests, not 585 — roughly 9x higher, badge is either wildly stale or was never machine-derived (violates §11.4.259 CM-MACHINE-DERIVED-SOURCE + §11.4.6 guessing). The README's own §11.4.259 provenance footnote (lines 40-57) explici discloses that the tests/vitest/plugins/me scan/license badges were 'carried forwar unverified this session (out of this task's and directs readers to docs/TESTING.md current authoritative per-suite counts' — docs/TESTING.md contains NO occurren '585' or '182' anywhere, so the cited authoritative source cannot actually corr the badge. The frontend vitest badge (18 passing) is also suspect: a source-level g it()/test() calls in frontend/src/**/*.spec.ts ~371, roughly 2x the claimed 182 (lower confidence proxy than the pytest collect-count, but consistent with the same stale pattern). Found during a §11.4.6 bluff au (docs/qa/task-bluff-audit/). Fix direction:

ATM ID	Type	Status	Severity	Description
				real badge-computer (README already f this as an owed §11.4.197 gate-code item regenerates the badge from a real collec count, and add the actual per-suite count docs/TESTING.md so the cited authoritat source is real.