

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Fixed (→ Fixed.md)	9
Bug	In progress	2
Bug	Operator-blocked	1
Bug	Queued	10
Bug	Ready for testing	1
Task	Completed (→ Fixed.md)	1
Task	In progress	9
Task	Operator-blocked	3
Task	Queued	19
Task	Ready for testing	2
TOTAL		57

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-ho SOCKS routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download-proxy + qBitTorrent-go + Jackett + compose env-forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Aut commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel ledger — ongoing coverage-escape backfill
BOB-074	Task		Medium	

ATM ID	Type	Status	Severity	Description
		In progress		RD2-07: DDoS-class testing fully absent from the mandated test-type matrix
BOB-077	Task	Operator-blocked	High	RD2-10: Identify second host running the Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated commit/push script OR retire
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code review of the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable items for the four Lava-porting findings (closes GA-05)
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/merge-service v1.0.0 readiness ledger (closes GA-10)
BOB-087	Task	Completed (→ Fixed.md)	High	RD2-20: Wire docs_chain / commit-sea sync hook per §11.4.106(F) so DB writes cannot land without MD mirror
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Tracked-Items + Status Documents table row-completeness (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exercising all three start.sh subcommands end-to-end against real compose stack
BOB-092	Bug	Queued	Medium	RD2-27: Remove test_live_stack_evidence.py:265 nnmcl SKIP-on-404 fallback + verify live 200 (closes GA-13)
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify rustracker ReDoS regex bounds deployment to container (closes runtime half of GA-13)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/test_tracker_fetch_stress_chaos.py with §11.4.85 fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/test_scheduler_hooks_sse_stress_chaos for Go-side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for exposed download-proxy/merge endpoints (canonical impl of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one commit (canonical impl of RD2-09)
BOB-101	Task		High	

ATM ID	Type	Status	Severity	Description
		Operator-blocked		GA-19/RW-09: Is -profile go parity still release goal? (gates RW-10..13) — OPERATOR-DECISION
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — bind tunnel 127.0.0.1 or keep 0.0.0.0? OPERATOR-DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence check helper for the unattributed auto-commit path
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence check for subagent task-brief source inputs
BOB-109	Task	Queued	Medium	BOB-074 followup: scaling-class test coverage absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiting for boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation gold bad fixture for the rate-limit detector
BOB-120	Bug	Queued	Critical	3rd forced-logout incident 2026-08-18 23:45:49 — SIGKILL user@1000 + preventive-timer-inside-user-slice architectural gap
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout architectural gap (task #85, incident #)
BOB-129	Bug	Fixed (→ Fixed.md)	Medium	Potential production slowapi/starlette defect flagged by Task 105 subagent
BOB-131	Bug	Fixed (→ Fixed.md)	Medium	qbittorrent-proxy podman common crash — pre-existing, surfaced during BOB-1 investigation
BOB-135	Bug	Queued	Low	Test isolation: test_list_hooks_after_create fails in bulk suite (Permission denied / config)
BOB-136	Task	Ready for testing	High	Closure seam does not bind: 4 tracker rows found stale in one sweep, and workable-items diff is blind to body_m drift
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while the same process still serves 7186 (GIL starvation by one spinning thread)

ATM ID	Type	Status	Severity	Description
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile serv 7186/7187/7188 but its container bind only 7187 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolvable gitdir) pollute gate scan scope and manufacture false BOB-126 class findings
BOB-144	Bug	Fixed (→ Fixed.md)	Low	/theme/stream calls the disconnect proc unguarded — fail-closed but via an uncaught traceback, inconsistent with two SSE generators
BOB-145	Bug	Fixed (→ Fixed.md)	High	Fix the 7187 wedge: offload and/or memoise Deduplicator.merge_results s O(N^2) regex work stops blocking the asyncio event loop
BOB-146	Bug	Fixed (→ Fixed.md)	High	Constitution §11.4.252 detector undercounts by 29% (30 vs 42 AST ground truth) — 4 distinct blind spots make its output a floor, not a census
BOB-148	Bug	Fixed (→ Fixed.md)	Medium	Standing red unit test nothing tracked test_no_credentials asserts has_session False, gets True — real defect or non-hermetic test, undecided
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/ across constitution, CLAUDE.md and t README badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant label read N/44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named g with no implementation, and 24 of 36 scripts have no companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored thi party code: 82% of one gate's 38,291 findings come from submodules/
BOB-153	Bug	Fixed (→ Fixed.md)	Medium	Go profile cannot build: go.mod requir go 1.26.2 but the Dockerfile builder is golang:1.23-alpine
BOB-154	Bug	Ready for testing	Medium	Host venv and production container ru different starlette versions (1.4.1 vs 1.
BOB-155	Bug	Fixed (→ Fixed.md)	High	workable-items diff reports 'DB and Markdown are in sync' having opened zero Markdown files when -issues/-fix are omitted

ATM ID	Type	Status	Severity	Description
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard load-sensitive and flaky: 8786ms under host load vs a 900-1500ms ceiling calibrated on a quiet host
BOB-157	Bug	Fixed (→ Fixed.md)	High	Our own BOB-137 stall watchdog can segfault the merge service: faulthandle.dump_traceback(all_threads=True) hit an unpatched CPython 3.12 defect
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds asyncio.events.get_event_loop_policy, 3.13+ private API, while production is 3.12.13 Reported-Via: §11.4.202 reporting directive bug on 2026-08-21T19:01:17Z Reported-By: T041 independent review (IMPORTANT-2), partially remediated What (the report, verbatim): The -recreate path was fixed this round: run_ownership_precondition -> stack_down -> run_ownership_repair -> stack_up, so the repair walks a tree no container can write to. That ordering is now asserted behaviourally by three new checks in tests/unit/test_start_reload_recreate.sh, each killed by its paired 1.1 mutation (11/11 RED-OK). The WARM path is NOT covered by that fix and this item tracks the remainder. On a warm ./start.sh the garbage runs before THIS invocation brings anything up, but if the stack is ALREADY running, containers write while the repair walks. A container can then create a new non-operator-owned file BEHIND the repair walk, after which the completion marker records complete over a tree that is not done and those stragglers are never repaired without a manual -force. BOUNDED, not dismissed: after any successful pass the marker is valid for the scope fingerprint and the repair short-circuits without walking (marker_is_valid), so the window requires a stale-or-absent marker AND a live stack together. Declaring a new scope entry re-arms the marker, which is exactly how that combination arises in practice. NOT DONE THIS ROUND, with the reason stated rather than hidden: the
BOB-159	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	clean guard is a cheap marker-only pro mode on scripts/ownership_repair.sh th answers has-this-scope-been-repaired without walking the tree. The existing -dry-run cannot serve: it deliberately SKIPS the marker fast-path (FORCE=0 && DRY_RUN=0 guard) and always walks, so using it as a warm-start prob would walk the tree twice on every sta Adding that mode requires editing ownership_repair.sh, which another stream was actively editing during this round; duplicating marker_is_valid into start.sh instead would be a near-identi fork (11.4.251). Deferred to this item rather than done unilaterally or silently. The misleading comment at the warm-start call site (which claimed the gate runs before any container writes) has been corrected to state this boundary honestly. Affected scope / file-scope manifest: start.sh (warm-start dispatcher_ownership_gate call site); scripts/ownership_repair.sh (marker fast-path). Reproduction / context: 1. Ensure the stack is UP. 2. Change config/owned_paths.yaml so the scope fingerprint changes (this re-arms the marker by design, data-model E2). 3. Run a warm ./start.sh (no -recreate). 4. The ownership repair walks and chowns the declared tree while containers are still running and able to write into it. Acceptance criteria: A warm ./start.sh cannot complete an ownership repair while a container that writes to a declared location is running. Either the repair is deferred with an actionable refusal naming -recreate, or the stack is quiescent for the walk. Asserted behaviourally in tests/unit/test_start_reload_recreate.sh alongside the existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE_UP checks, each with a paired 1.1 mutation that kills it. Reported-Via: §11.4.202 reporting directive task on 2026-08-21T19:02:48 Reported-By: AI What (the report, verbatim): Independent §11.4.209 rev

ATM ID	Type	Status	Severity	Description
				(IMPORTANT-3) found two of the features that the strongest automated checks were never executed by any runner: (1) tests/ownership/test_container_writes_owned_files.py – the §11.4.115 RED-turned-regression-guard for FR-011/FR-007, proven via a docker-compose.yml revert mutation – was moved out of tests/integration/ (commit 58d340a, to escape an autous fixture hang) but no runner (ci.sh, test-run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_build/test_.sh, the §1.1 paired-mutation meta-tests for the scripts/pre_build/check_cm_.sh gate family, was executed by NOTHING — self-reported honestly in commit 04742d7’s own message (‘STILL OPEN (reported, not fixed): tests/pre_build/ is executed by NOTHING’) but never tracked as a workable item (a §11.4.197 loss-of-requirements gap) and never wired into scripts/pre_build_verification.sh invariant 30, which globbed only tests/unit/test_ Affected scope / file-scope manifest ci.sh; scripts/pre_build_verification.sh invariant 30 (CM-BASH-UNIT-TESTS-EXECUTED) Reproduction / context grep -rn test_container_writes_owned_files ci.sh test.sh run-all-tests.sh scripts/ returned zero runner hits (only docs/specs mentioned the path); grep in scripts/pre_build_verification.sh showed invariant 30’s for-loop globbing only tests/unit/test_.sh, never tests/pre_build/test_.sh Acceptance criteria: ci.sh gains a runtime-gated pytest tests/ownership/ stage (skips honestly, rc=0, when no podman/docker present); scripts/pre_build_verification.sh invariant 30’s glob covers both tests/unit/test_.sh and tests/pre_build/test_*.sh, verified by an extracted standalone run of the invariant’s exact block reporting a non-zero RAN count that includes files from both directories. Reported-Via: §11.4.202 reporting directive task on 2026-08-21T19:31:12
BOB-161	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
				Reported-By: BOB-092 remediation, residual finding What (the report, verbatim): §11.4.69 names CM-NO-FA OPEN-SKIP as one of three mandatory pre-build gates: it audits sink-side probe helpers and FAILs if any code path converts an empty or unreachable response into a PASS-counting SKIP for feature class with a sink-side probe. The project does not have it. Why this matter now rather than in the abstract: the BOB-092 remediation just removed TW fail-opens of exactly this class from test_e2e/test_live_stack_evidence.py — the nnmclub SKIP-on-404 (already removed 7baef2b) and a surviving sibling at test_iporrents_is_authenticated_in_session that skipped on ‘not authenticated and status != success’ while asserting ‘treating as transient outage’. That assertion was false for rejected credentials (upstream_http_403) and for broken container (plugin_env_missing) both of which are definitive product failures the product already classifies with error_type. Two fail-opens of one class found one at a time, is the §11.4.146 extend-to-all-cases signal and the §11.4.238 coverage-escape signal together: the regime did not surface either, an agent reading code did. The remediation’s own guards are AST-structural and lethal under three discriminating mutations, but they LIVE IN THE FILE THEY GUARD, so deleting that file evades them entirely. A pre-build gate is the durable home because it audits the corpus rather than one file. Honest boundary (§11.4.6): this item does NOT claim the two remediated fail-opens are unguarded — they are guarded, with runtime evidence (13 passed in 97.36s against the live stack). It claims the CLASS has no corpus-wide detector, so the next instance in a different file is invisible again. Affected scope / file-scope manifest: scripts/pre_build/ (guards absent); tests/e2e/test_live_stack_evidence.py (guards currently live inside the file they guard)

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	Reproduction / context: <code>grep -rl 'CM-NO-FAIL-OPEN-SKIP' scripts/ tests/</code> returns exactly ONE hit, and it is <code>tests/e2e/test_live_stack_evidence.py</code> — the guards live in, not a gate. Control needle: the same query for <code>CM-OWNERSHIP-INVARIANTS</code> (a gate that does exist) returns 3 files, so the instrument can see gate tokens and the single hit is a real absence, not a blind zero (§11.4.201(7)(b)). Acceptance criteria: A <code>scripts/pre_build/gate_name_CM-NO-FAIL-OPEN-SKIP</code> exists, is wired into <code>scripts/pre_build_verification.sh</code>, audits sink-side probe helpers for code paths converting an empty/unreachable error response into a PASS-counting SUCCESS for a feature class that HAS a sink-side probe. It ships a golden-TRUE fixture (real fail-open -> gate FIRES) and a golden-FALSE-with-carrier (an honest topology/geo skip, and a comment mentioning the phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201. Paired §1.1 mutation makes the gate FIRE before the gate is trusted. Reported-Via: §11.4.202 reporting directive task on 2026-08-21T19:35:02Z Reported-By: BOB-106/BOB-107 remediation, residual What (the report verbatim): The TESTS for both guards are now executed: pre-build invariant .glob was extended from <code>tests/unit + tests/pre_build</code> to also cover <code>tests/hooks</code>. MEASURED before/after: the glob expanded to 27 suites with 0 under tests/hooks while 3 existed there; it now expands to 30 with 3 under tests/hooks. All three pass. But the GUARDS themselves are still invoked by nothing: <code>check-brief-inputs.sh</code> is honestly conductor-run rather than automatic: a brief's required-input list lives in free-form prose that the Agent tool's structured input does not expose, so an automatic prompt-scraping gate would itself be an unproven pattern-match. It seams the dispatch procedure, and there is a documentation + habit change, no hook. <code>unattributed-commit-guard.sh</code> is

ATM ID	Type	Status	Severity	Description
				different and is why this item exists. R against real history it names 14 violati commits since the last tag (and 21 bar Auto-commit subjects exist across all refs). Wiring it into the pre-build or commit seam AS-IS would therefore refuse every commit immediately, on p existing debt nobody can fix in the moment. That is precisely the 11.4.224 brownfield-adoption shape, and 11.4.66/11.4.122 put the choice with t operator, not with an agent inventing a ratchet. The options, stated so the decision is a decision and not a default (a) immediate hard floor – the seam refuses until all 14 are attributed; (b) c time monotone-decrease ratchet (the 11.4.135 pattern) – snapshot 14 as the baseline, the count may fall and may never rise, so day one is green and the disease cannot spread; (c) changed-commits-only – gate only commits crea from now on, with a scheduled deadline for the historical 14; (d) report-only for stated period, then escalate. Recommendation, with the reasoning rather than just the pick: (b). It is the pattern this constitution already uses f exactly this situation, it makes the deb visible and bounded immediately, and i cannot be satisfied by deleting the gua name (11.4.227 closes that gaming channel). But it is the operator’s call a this item is not closed until that answe recorded. Honest boundary (11.4.6): th item does NOT claim the 14 commits a harmful in themselves – it claims their ATTRIBUTION is absent and that noth currently prevents the 15th. Affected scope / file-scope manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-brief-inputs.sh; scripts/pre_build_verification.sh; script commit-push-all.sh Reproduction / context: Both guards run correctly by hand and are covered by passing tests (9/9 and 15/15, each proven non-bluff against a no-op stub). Neither is invoke by scripts/pre_build_verification.sh or scripts/commit-push-all.sh, so neither

ATM ID	Type	Status	Severity	Description
BOB-163	Bug	Queued	Medium	exerts standing detection pressure (11.4.226: a guard with no execution seam is not coverage; registration is not coverage). Acceptance criteria: Each guard is invoked by a named seam, OR carries a registered deferral pointing at this item. For unattributed-commit-guard.sh specifically, the operator has answered the adoption question below and the answer is recorded as consumed DATA – never an invented ratchet. Reported-Via: §11.4.202 reporting directive bug on 2026-08-21T19:41:08Z Reported-By: BOB-114 remediation, p existing defect surfaced by BOB-111 landing a real limiter What (the report verbatim): PRE-EXISTING, NOT INTRODUCED – and proven so rather than asserted: git diff shows assertion is BYTE-IDENTICAL to HEAD, and the failure reproduces against the unmodified HEAD script. What changed is not the challenge but the SYSTEM: BOB-111 appears at least partially delivered, because :7187 now returns real 429s where it previously returned none. :7188 and :7189 still produce zero 429s. This is a 11.4.248-class corrosion risk and that's why it is filed rather than left as a footnote: run_all_challenges.sh will not fail INTERMITTENTLY, and an intermittent red trains everyone to red until green, at which point a real regression is dismissed as ‘probably the flaky rate-limit one’. THE DECISION, which is an operator call under 11.4.60 and which I deliberately did NOT invert. Does a 429 from a WORKING limiter count as the sibling endpoint being RESPONSIVE? (a) YES – a 429 proves endpoint is alive and correctly protects itself. Assertion (c) accepts 2xx OR 429 and only a connection failure / 5xx / timeout counts as degraded. Risk: a genuinely wedged endpoint that happens to answer 429 would pass. (b) NO – keeping requiring 2xx, but make the challenge limiter-aware: drain or wait out the window before the sibling probe (retry) after is served, so the budget is known.

ATM ID	Type	Status	Severity	Description
				rather than guessed). (c) Probe sibling on a path the limiter exempts (a health class route), so the isolation question is asked without spending the bucket. Recommendation with reasoning, not just a pick: (a) composed with a bounded liveness check - a 429 carrying a well-formed Retry-After IS evidence of a live, correctly-behaving service, and (b) make the challenge slower and couples it to a window value that will drift. But the semantics of 'responsive' here are a product judgement, so the answer is recorded, not assumed. RELATED, state as fact rather than folded in: the detector counts 429 only, not 503, though the script header says '429 (or equivalent)'. Counting 503 would collide with the consumer detector's 5xx tally. Worth deciding whether BOB-111 lands limiters on :7185 and :7189. Affected scope / file-scope manifest: challenges/scripts/ddos_resilience_challenge.sh assertion cross-endpoint isolation; run_all_challenges.sh which invokes it. Reproduction / context: Assertion (c) requires a sibling-endpoint probe to answer ^2 (a 2xx). :7187 now enforces real limiter (measured live: x-ratelimit-limit: 120, x-ratelimit-remaining: 86, retry-after: 45, server: uvicorn). A full challenge run sends roughly 250 requests to :7187, so sibling probes fired during the OTHER endpoints' tiers land on an exhausted bucket and receive 429 - with assertion (c) scores as 'endpoint degraded'. Timing-dependent: one run measured PASS=5 FAIL=2; the UNMODIFIED HEAD script against the same live stack ~30s later measured PASS=7 FAIL=0 SKIP=2. Acceptance criteria: The operator has answered the classification question below, the answer is recorded as consumer DATA, and assertion (c) implements it. The challenge then produces the same verdict across consecutive runs against a rate-limited stack (11.4.50 deterministic consistency) and the fix ships a paired 1.1 mutation proving assertion (c) still catches a

ATM ID	Type	Status	Severity	Description
				genuinely degraded sibling endpoint - narrowing it must not blind it (11.4.201(1)).