Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Fixed (→ Fixed.md)	9
Bug	In progress	2
Bug	Operator-blocked	1
Bug	Queued	16
Bug	Ready for testing	3
Task	Completed (→ Fixed.md)	1
Task	In progress	9
Task	Operator-blocked	3
Task	Queued	21
Task	Ready for testing	2
TOTAL		**67**

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-ho SOCKS routing (containers pkg/egress
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download-proxy + qBitTorrent-go + Ja + compose env-forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Aut commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-chann ledger — ongoing coverage-escape bac
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully abser from the mandated test-type matrix

ATM ID	Type	Status	Severity	Description
BOB-077	Task	Operator-blocked	High	RD2-10: Identify second host running t Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-cor mechanism through §11.4.234 dedicat commit/push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code review of the substantive Auto-commit
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workabl items for the four Lava-porting finding (closes GA-05)
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/m service v1.0.0 readiness ledger (closes GA-10)
BOB-087	Task	Completed (→ Fixed.md)	High	RD2-20: Wire docs_chain / commit-sea sync hook per §11.4.106(F) so DB writ cannot land without MD mirror
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Tra Items + Status Documents table row-completeness (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exer all three start.sh subcommands end-to against real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + veri rutracker ReDoS regex bounds deploy container (closes runtime half of GA-12
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py wit §11.4.85 fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos for Go-side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage f exposed download-proxy/merge endpo (canonical impl of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one commit (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is –profile go parity still release goal? (gates RW-10..13) — OPERATOR-DECISION
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — tunnel 127.0.0.1 or keep 0.0.0.0? — OPERATOR-DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge

ATM ID	Type	Status	Severity	Description
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescen check helper for the unattributed auto commit path
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existe check for subagent task-brief source ir
BOB-109	Task	Ready for testing	Medium	BOB-074 followup: scaling-class test coverage absent from mandated test-ty matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test cover (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiting for boba's 3 public HTTP endp
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation gold bad fixture for the rate-limit detector
BOB-120	Bug	Queued	Critical	3rd forced-logout incident 2026-08-18 23:45:49 — SIGKILL user@1000 + preventive-timer-inside-user-slice architectural gap
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logou architectural gap (task #85, incident #
BOB-129	Bug	Fixed (→ Fixed.md)	Medium	Potential production slowapi/starlette flagged by Task 105 subagent
BOB-131	Bug	Fixed (→ Fixed.md)	Medium	qbittorrent-proxy podman conmon cras pre-existing, surfaced during BOB-129 investigation
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_cre fails in bulk suite (Permission denied / config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while t same process still serves 7186 (GIL starvation by one spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile serv 7186/7187/7188 but its container bind 7187 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolvable gitdir) pollute gate scan s and manufacture false BOB-126-class findings
BOB-144	Bug	Fixed (→ Fixed.md)	Low	/theme/stream calls the disconnect pro unguarded — fail-closed but via an unc traceback, inconsistent with the two S generators
BOB-145	Bug	Fixed (→ Fixed.md)	High	Fix the 7187 wedge: offload and/or me Deduplicator.merge_results so O(N^2)

ATM ID	Type	Status	Severity	Description
				regex work stops blocking the asyncio loop
BOB-146	Bug	Fixed (→ Fixed.md)	High	Constitution §11.4.252 detector undercounts by 29% (30 vs 42 AST gro truth) — 4 distinct blind spots make its output a floor, not a census
BOB-148	Bug	Fixed (→ Fixed.md)	Medium	Standing red unit test nothing tracked test_no_credentials asserts has_session False, gets True — real defect or non-hermetic test, undecided
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/ across constitution, CLAUDE.md and t README badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labe read N/44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named ga with no implementation, and 24 of 36 scripts have no companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored thi party code: 82% of one gate’s 38,291 findings come from submodules/
BOB-153	Bug	Fixed (→ Fixed.md)	Medium	Go profile cannot build: go.mod requir 1.26.2 but the Dockerfile builder is golang:1.23-alpine
BOB-154	Bug	Ready for testing	Medium	Host venv and production container ru different starlette versions (1.4.1 vs 1.(
BOB-155	Bug	Fixed (→ Fixed.md)	High	workable-items diff reports ‘DB and Markdown are in sync’ having opened Markdown files when –issues/–fixed ar omitted
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard i load-sensitive and flaky: 8786ms under load vs a 900-1500ms ceiling calibrate quiet host
BOB-157	Bug	Fixed (→ Fixed.md)	High	Our own BOB-137 stall watchdog can segfault the merge service: faulthandle dump_traceback(all_threads=True) hit unpatched CPython 3.12 defect
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds asyncio.events._get_event_loop_policy, 3.13+ private API, while production is 3.12.13
BOB-159	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				Reported-Via: §11.4.202 reporting directive bug on 2026-08-21T19:01:17Z **Reported-By:** T041 independent revie (IMPORTANT-2), partially remediated **(the report, verbatim):** The –recreat was fixed this round: run_ownership_precondition -> stack_ -> run_ownership_repair -> stack_up, repair walks a tree no container can w to. That ordering is now asserted behaviourally by three new checks in t unit/test_start_reload_recreate.sh, eac killed by its paired 1.1 mutation (11/11 OK). The WARM path is NOT covered b that fix and this item tracks the remain On a warm ./start.sh the gate runs befo THIS invocation brings anything up, bu the stack is ALREADY running, contain write while the repair walks. A contain can then create a new non-operator-ow file BEHIND the walk, after which the completion marker records complete c tree that is not, and those stragglers a never repaired without a manual –forc BOUNDED, not dismissed: after any successful pass the marker is valid for scope fingerprint and the repair short- circuits without walking (marker_is_va so the window requires a stale-or-abse marker AND a live stack together. Dec a new scope entry re-arms the marker, which is exactly how that combination arises in practice. NOT DONE THIS ROUND, with the reason stated rather hidden: the clean guard is a cheap mar only probe mode on scripts/ ownership_repair.sh that answers has- scope-been-repaired without walking t tree. The existing –dry-run cannot serv deliberately SKIPS the marker fast-pat (FORCE=0 && DRY_RUN=0 guard) an always walks, so using it as a warm-sta probe would walk the tree twice on eve start. Adding that mode requires editin ownership_repair.sh, which another str was actively editing during this round; duplicating marker_is_valid into start.s instead would be a near-identical fork (11.4.251). Deferred to this item rathe done unilaterally or silently. The mislea

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	comment at the warm-start call site (w claimed the gate runs before any conta writes) has been corrected to state this boundary honestly. **Affected scope / f scope manifest:** start.sh (warm-start dispatch, run_ownership_gate call site scripts/ownership_repair.sh (marker fa path) **Reproduction / context:** 1. Ens the stack is UP. 2. Change config/ owned_paths.yaml so the scope fingerp changes (this re-arms the marker by d data-model E2). 3. Run a warm ./start.s –recreate). 4. The ownership repair wa and chowns the declared tree while containers are still running and able to write into it. **Acceptance criteria:** A warm ./start.sh cannot complete an ownership repair while a container tha writes to a declared location is running Either the repair is deferred with an actionable refusal naming –recreate, o stack is quiesced for the walk. Asserte behaviourally in tests/unit/ test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DO REPAIR_AFTER_DOWN / REPAIR_BEFORE_UP checks, each wit paired 1.1 mutation that kills it. **Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:02:48 **Reported-By:** AI **What (the report, verbatim):** Independent §11.4.209 rev (IMPORTANT-3) found two of the featu strongest automated checks were neve executed by any runner: (1) tests/ ownership/ test_container_writes_owned_files.py – §11.4.115 RED-turned-regression-guar FR-011/FR-007, proven via a docker-compose.yml revert mutation — was m out of tests/integration/ (commit 58d34 to escape an autouse fixture hang) but runner (ci.sh, test.sh, run-all-tests.sh) ever extended to cover tests/ownershi tests/pre_build/test_*.sh, the §1.1 paired mutation meta-tests for the scripts/ pre_build/check_cm_*.sh gate family, wa executed by NOTHING — self-reported honestly in commit 04742d7's own me ('STILL OPEN (reported, not fixed): tes

ATM ID	Type	Status	Severity	Description
BOB-161	Task	Queued	High	pre_build/ is executed by NOTHING’) l never tracked as a workable item (a §11.4.197 loss-of-requirements gap) ar never wired into scripts/ pre_build_verification.sh invariant 30, globbed only tests/unit/test_*.sh. **Affect scope / file-scope manifest:** *ci.sh; sc pre_build_verification.sh invariant 30 (BASH-UNIT-TESTS-EXECUTED)* ***Repr ion / context:*** *grep -rn test_container_writes_owned_files ci.sh test.sh run-all-tests.sh scripts/ returne runner hits (only docs/specs mentioned path); grep in scripts/ pre_build_verification.sh showed invar 30’s for-loop globbing only tests/unit/ test_*.sh, never tests/pre_build/test_*.sh* ***Acceptance criteria:*** *ci.sh gains a ru gated pytest tests/ownership/ stage (sk honestly, rc=0, when no podman/docke present); scripts/pre_build_verification invariant 30’s glob covers both tests/u test_*.sh and tests/pre_build/test_*.sh,* verified by an extracted standalone ru the invariant’s exact block reporting a zero RAN count that includes files fro directories. **Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:31:12 **Reported-By:** BOB-092 remediation, residual finding **What (the report, verbatim):** §11.4.69 names CM-NO-FA OPEN-SKIP as one of three mandatory build gates: it audits sink-side probe h and FAILs if any code path converts an empty or unreachable response into a counting SKIP for a feature class with sink-side probe. This project does not h it. Why this matters now rather than in abstract: the BOB-092 remediation jus removed TWO fail-opens of exactly this from tests/e2e/test_live_stack_evidence — the nnmclub SKIP-on-404 (already removed at 7baef2b) and a surviving si at test_iptorrents_is_authenticated_in_sea that skipped on ‘not authenticated and status != success’ while asserting ‘trea as transient outage’. That assertion wa false for rejected credentials

ATM ID	Type	Status	Severity	Description
				(upstream_http_403) and for a broken container (plugin_env_missing), both o which are definitive product failures th product already classifies via error_typ Two fail-opens of one class, found one time, is the §11.4.146 extend-to-all-cas signal and the §11.4.238 coverage-esca signal together: the regime did not sur either, an agent reading code did. The remediation's own guards are AST-structural and lethal under three discriminating mutations, but they LIV THE FILE THEY GUARD, so deleting th file evades them entirely. A pre-build g the durable home because it audits the corpus rather than one file. Honest boundary (§11.4.6): this item does NOT claim the two remediated fail-opens ar unguarded — they are guarded, with runtime evidence (13 passed in 97.36s against the live stack). It claims the CI has no corpus-wide detector, so the ne instance in a different file is invisible a **Affected scope / file-scope manifest** scripts/pre_build/ (gate absent); tests/e test_live_stack_evidence.py (guards currently live inside the file they guard **Reproduction / context:** grep -rl 'CM FAIL-OPEN-SKIP' scripts/ tests/ return exactly ONE hit, and it is tests/e2e/ test_live_stack_evidence.py — the file t guards live in, not a gate. Control need the same query for CM-OWNERSHIP-INVARIANTS (a gate that does exist) r 3 files, so the instrument can see gate tokens and the single hit is a real abse not a blind zero (§11.4.201(7)(b)). **Acceptance criteria:** A scripts/pre_bu gate named CM-NO-FAIL-OPEN-SKIP is wired into scripts/ pre_build_verification.sh, and audits si side probe helpers for code paths conv an empty/unreachable/error response PASS-counting SKIP for a feature class HAS a sink-side probe. It ships a golde TRUE fixture (a real fail-open -> gate FIRES) and a golden-FALSE-with-carri honest topology/geo skip, and a comme merely MENTIONING the phrase -> ga MUST NOT fire), per §11.4.107(10)/

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	§11.4.201(1). Paired §1.1 mutation mal the gate FAIL before the gate is truste **Reported-Via:** §11.4.202 reporting directive `task` on 2026-08-21T19:35:02 **Reported-By:** BOB-106/BOB-107 remediation, residual **What (the repo verbatim):** The TESTS for both guards now executed: pre-build invariant 30's was extended from tests/unit + tests/ pre_build to also cover tests/hooks. MEASURED before/after: the glob exp to 27 suites with 0 under tests/hooks w existed there; it now expands to 30 wit under tests/hooks. All three pass. But t GUARDS themselves are still invoked l nothing. check-brief-inputs.sh is hones conductor-run rather than automatic: a brief's required-input list lives in free- prose that the Agent tool's structured does not expose, so an automatic prom scraping gate would itself be an unpro pattern-match. Its seam is the dispatch procedure, and that is a documentatio habit change, not a hook. unattributed commit-guard.sh is different and is wh item exists. Run against real history it names 14 violating commits since the l tag (and 21 bare Auto-commit subjects across all refs). Wiring it into the pre-b or commit seam AS-IS would therefore refuse every commit immediately, on p existing debt nobody can fix in the mo That is precisely the 11.4.224(E) brow adoption shape, and 11.4.66/11.4.122 the choice with the operator, not with agent inventing a ratchet. The options, stated so the decision is a decision and default: (a) immediate hard floor – the refuses until all 14 are attributed; (b) time monotone-decrease ratchet (the 11.4.135 pattern) – snapshot 14 as the baseline, the count may fall and may n rise, so day one is green and the disea cannot spread; (c) changed-commits-o gate only commits created from now o with a scheduled deadline for the histo 14; (d) report-only for a stated period, escalate. Recommendation, with the reasoning rather than just the pick: (b) the pattern this constitution already us

ATM ID	Type	Status	Severity	Description
BOB-163	Bug	Queued	Medium	exactly this situation, it makes the deb visible and bounded immediately, and i cannot be satisfied by deleting the gua name (11.4.227 closes that gaming channel). But it is the operator's call a this item is not closed until that answe recorded. Honest boundary (11.4.6): th item does NOT claim the 14 commits a harmful in themselves – it claims their ATTRIBUTION is absent and that noth currently prevents the 15th. **Affected scope / file-scope manifest:** scripts/h unattributed-commit-guard.sh; scripts/ hooks/check-brief-inputs.sh; scripts/ pre_build_verification.sh; scripts/comm push-all.sh **Reproduction / context:** I guards run correctly by hand and are covered by passing tests (9/9 and 15/1 each proven non-bluff against a no-op s Neither is invoked by scripts/ pre_build_verification.sh or scripts/com push-all.sh, so neither exerts standing detection pressure (11.4.226: a guard no execution seam is not coverage; registration is not coverage). **Accepta criteria:** Each guard is invoked by a n seam, OR carries a registered deferral pointing at this item. For unattributed- commit-guard.sh specifically, the opera has answered the adoption question be and the answer is recorded as consum DATA – never an invented ratchet. **Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:41:08Z **Reported-By:** BOB-114 remediation, p existing defect surfaced by BOB-111 la a real limiter **What (the report, verbatim):** PRE-EXISTING, NOT INTRODUCED – and proven so rather asserted: git diff shows assertion (c) is BYTE-IDENTICAL to HEAD, and the fa reproduces against the unmodified HE script. What changed is not the challen but the SYSTEM: BOB-111 appears at partially delivered, because :7187 now returns real 429s where it previously returned none. :7185 and :7189 still produce zero 429s. This is a 11.4.248-c corrosion risk and that is why it is filed rather than left as a footnote:

ATM ID	Type	Status	Severity	Description
				run_all_challenges.sh will now fail INTERMITTENTLY, and an intermitten trains everyone to re-run until green, a which point a real regression is dismis ‘probably the flaky rate-limit one’. THE DECISION, which is an operator call u 11.4.66 and which I deliberately did N invent: Does a 429 from a WORKING li count as the sibling endpoint being RESPONSIVE? (a) YES – a 429 proves endpoint is alive and correctly protecti itself. Assertion (c) accepts 2xx OR 429 only a connection failure / 5xx / timeou counts as degraded. Risk: a genuinely wedged endpoint that happens to answ 429 would pass. (b) NO – keep requirin 2xx, but make the challenge limiter-aw drain or wait out the window before th sibling probe (retry-after is served, so budget is knowable rather than guesse (c) Probe siblings on a path the limiter exempts (a healthz-class route), so the isolation question is asked without spe the bucket. Recommendation with reasoning, not just a pick: (a) compose with a bounded liveness check – a 429 carrying a well-formed Retry-After IS evidence of a live, correctly-behaving service, and (b) makes the challenge sl and couples it to a window value that w drift. But the semantics of ‘responsive’ are a product judgement, so the answe recorded, not assumed. RELATED, stat fact rather than folded in: the detector counts 429 only, not 503, though the sc header says ‘429 (or equivalent)’. Coun 503 would collide with the crash detec 5xx tally. Worth deciding when BOB-11 lands limiters on :7185 and :7189. **Affe scope / file-scope manifest:** challeng scripts/ddos_resilience_challenge.sh assertion (c) cross-endpoint isolation; run_all_challenges.sh which invokes it **Reproduction / context:** Assertion (c requires a sibling-endpoint probe to ar ^2 (a 2xx). :7187 now enforces a real l (measured live: x-ratelimit-limit: 120, x ratelimit-remaining: 86, retry-after: 45 server: uvicorn). A full challenge run s roughly 250 requests to :7187, so sibli

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	Queued	Medium	probes fired during the OTHER endpoi tiers land on an exhausted bucket and receive 429 – which assertion (c) score ‘endpoint degraded’. Timing-dependen run measured PASS=5 FAIL=2; the UNMODIFIED HEAD script against th same live stack ~30s later measured PASS=7 FAIL=0 SKIP=2. **Acceptance criteria:** The operator has answered t classification question below, the answ recorded as consumer DATA, and asse (c) implements it. The challenge then produces the same verdict across 3 consecutive runs against a rate-limited stack (11.4.50 deterministic consistenc and the fix ships a paired 1.1 mutation proving assertion (c) still catches a genuinely degraded sibling endpoint – narrowing it must not blind it (11.4.20 **Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T19:56:53Z **Reported-By:** BOB-110 UX-class cove discovered by the new axe-core suite o first live run **What (the report, verba** This is a REAL user-facing defect, not a tuning artifact, and it was found by the automated regime rather than by a hu squinting at the page – which is exactly 11.4.238 posture the project is aiming The static-grep oracle would NOT have found it. Measured: the served root sh empty pre-hydration, so any check rea the raw HTML audits a page nobody s The violation only exists in the hydrate DOM, which is why 11.4.170 requires rendered oracle and forbids value-equa assertions as the proof a UI is correct. Severity reasoning, stated rather than assumed: this is user-visible and affect primary dashboard heading, but it deg legibility rather than breaking function the surface is operator-facing rather th public. Medium, not High. The failing t was left FAILING on purpose (11.4.238 instead of silenced or marked xfail. tes currently reports 1 failed, 16 passed; t is this defect. Anyone reading a red UX should read it as this item, not as flaki and when this is fixed the suite goes fu green, which is the signal that it is clos

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	HONEST SCOPE LIMIT (11.4.6): only dashboard landing view was scanned. jackett/* sub-routes and the ng-serve-h :4200 route set were NOT audited – th commands to close both are recorded docs/testing/ux_accessibility.md. So th item's 21 nodes are a floor, not a total. **Affected scope / file-scope manifest** Angular dashboard served at http:// localhost:7187/ (same compiled SPA as frontend/); production component CSS test files **Reproduction / context:** Ru tests/ux/test_live_dashboard_accessibil against the running merge service. axe v4.13.0, scanning a real Playwright-rendered JS-hydrated DOM, reports co contrast violations on 21 nodes. Measu pairs: .brand / h1 text #9d001e on background #3c3f41 = 1.43-1.62:1 (W AA large-text floor is 3:1); tagline #808 on #3c3f41 = 2.68:1 (body-text floor is 4.5:1). **Acceptance criteria:** axe-core reports ZERO color-contrast violations against the live rendered dashboard, w the fix made in production component rather than by relaxing the assertion o excluding the rule (11.4.120: reconcile the correct mechanism, never weaken check). The existing tests/ux/ suite is t guard and already fails today, so the R captured – closure requires it flipping GREEN against the live surface, which runtime-class evidence per 11.4.226. **Reported-Via:** §11.4.202 reporting directive `bug` on 2026-08-21T20:00:08Z **Reported-By:** surfaced while capturin closure evidence for BOB-092; diagnos rather than worked around **What (the report, verbatim):** This is a STALE-AI AFTER-INTERPRETER-UPGRADE defe not a missing package. The package is installed; its compiled half is built for t previous interpreter. That distinction matters because 'pip install rpds-py' st advice can appear to succeed while lea the 313 artefact in place. WHY IT WAS NOTICED EARLIER, stated as fact: the paths that DO work all avoid system python3. ci.sh selects an interpreter through _select_python; the long-runni

ATM ID	Type	Status	Severity	Description
				suites in this session ran .venv/bin/pyt -m pytest and passed. So the automate regime is green while the DOCUMENT operator command is broken – an 11.4 shaped gap, since the escape was foun an agent capturing evidence rather tha the regime. WORKAROUND (measured theorised): PYTEST_DISABLE_PLUGIN_AUTOLOA with the needed plugins passed explici pytest_timeout …) avoids the schemath autoload that drags in jsonschema -> referencing -> rpds. That is a workaro NOT the fix: it silences the import path rather than repairing the install, and it surprise the next person who runs the documented command verbatim. RELA and worth deciding together rather tha twice: BOB-154 wants .venv rebuilt on to match the container (which runs CP 3.12.13 while both system and venv ru 3.14.6). There are therefore THREE interpreters in play. Whoever rebuilds venv should confirm rpds resolves for t TARGET interpreter afterwards, or this same class reappears one directory ov HONEST BOUNDARY (11.4.6): this ite does NOT claim any test is wrong or a product code is broken. The product a venv-run suites are unaffected. What is broken is the documented entry point. **Affected scope / file-scope manifest** user site-packages (~/.local/lib/python packages/rpds/); every CLAUDE.md-documented 'python3 -m pytest …' invocation; any tooling that uses syste python3 rather than .venv/bin/python **Reproduction / context:** python3 -m tests/e2e/test_live_stack_evidence.py - –import-mode=importlib -> ModuleNotFoundError: No module na 'rpds.rpds', raised during collection via schemathesis plugin autoload -> jsonschema -> referencing._core -> rp MEASURED root cause: system python 3.14.6, but ~/.local/lib/python3/site-packages/rpds/ ships rpds.cpython-313 x86_64-linux-gnu.so – an extension bui 3.13. rpds/**init**.py does 'from .rpds imp and a 313-tagged .so is not importable

ATM ID	Type	Status	Severity	Description
				3.14, so the submodule genuinely does exist for this interpreter. The venv is CORRECT and unaffected: .venv/lib64/ python3/site-packages/rpds/ ships rpds.cpython-314-x86_64-linux-gnu.so and'.venv/bin/python -c import rpds' succeeds. **Acceptance criteria:** pytho pytest tests/unit/ -v –import-mode=imp – the command CLAUDE.md document reaches collection without ModuleNotFoundError, OR CLAUDE.m corrected to document the supported runner explicitly. Either way the documented command and the workin command agree (11.4.99: a guide that misleads is the documentation-layer equivalent of a PASS-bluff). WORKARO SIDE EFFECTS MEASURED 2026-08-2 the recipe is not free, and the item previously implied it was. PYTEST_DISABLE_PLUGIN_AUTOLOA disables ALL plugin autoload, not just broken one, so anything the suite silen relied on must be re-enabled by hand: BREAKS 5 pre-existing env-dependent in tests/unit/test_plugin_rutracker.py – TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace which need an autoloaded env plugin. PROVEN not caused by any in-flight ch running HEAD's OWN copy of that file under identical flags produces the sam failures. - It makes –timeout=60 (set in pyproject.toml) an UNKNOWN ARGUM unless -p pytest_timeout is passed expl so pytest.mark.timeout is silently inert under the naive recipe — a test believe be time-bounded is not. - With autoloa and only schemathesis disabled (-p no:schemathesis), that same file is 96/9 green. CONSEQUENCE FOR ANYONE USING THE WORKAROUND: `-p no:schemathesis` alone is strictly better blanket autoload-disabling where it su because it removes only the broken plu Where the blanket form IS used, -p pytest_timeout must be added or timed

ATM ID	Type	Status	Severity	Description
				are inert, and the 5 env-dependent fail must be recognised as workaround art rather than filed as defects. That last p is the real risk: the workaround manufactures failures that look exactly product defects. This does not change root cause (the venv’s rpds ships a cpython-313 ABI tag CPython 3.14 can import) or the fix (rebuild the venv — BOB-154). It documents that the interi recipe has a blast radius, so nobody re workaround artifact as a regression.
BOB-166	Bug	Queued	High	WHAT. §11.4.19 requires closure migra to be ATOMIC: a resolving item moves Fixed.md, DISAPPEARS from Issues_Summary (open-only) and APPE in Fixed_Summary (closed-only). Ten r violate all three properties right now – carry a terminal status yet current_location=‘Issues’, so they rend docs/Issues.md and are listed in docs/ Issues_Summary.md with a status that literally reads ‘(→ Fixed.md)’, while appearing 0 times in Fixed_Summary.m Any reader of the canonical open track told these are open work. MANIFEST (measured 2026-08-21): BOB-087 (Completed), BOB-129, BOB-131, BOB BOB-145, BOB-146, BOB-148, BOB-15 BOB-155, BOB-157 (all ‘Fixed (→ Fixed.md)’). Confirmed rendering: BOB-155/087/157 each grep 1 in Issue 0 in Fixed.md, present in Issues_Summary.md, 0 hits in Fixed_Summary.md. ROOT CAUSE (reproduced at runtime, not inferred). ‘close’ performs the atomic migration a REQUIRES –evidence. ‘update –status’ any §11.4.15 closed-set value with NO evidence and NO migration — and it k the location, because it prints it. On a of the real DB: $ workable-items updat BOB-065 –db repro.db –status ‘Fixed (→ Fixed.md)’ update: BOB-065 updated i Issues (status=Fixed (→ Fixed.md), type=Task) $ sqlite3 repro.db ‘SELECT atm_id,status,current_location …’ BOB Fixed (→ Fixed.md)|Issues So the seam is supposed to be the ONLY closure pa (close, evidence-gated per §11.4.146(D

ATM ID	Type	Status	Severity	Description
				has a parallel unguarded path that rea the same status while skipping BOTH t evidence requirement and the migratio WHY NOTHING CAUGHT IT (§11.4.23 coverage-escape audit). Three standing checks stay GREEN on a DB holding th forbidden row, verified on the poisoned from the repo root so no cwd artifact is involved: (1) ‘validate’ → ‘OK — 164 ite all invariants satisfied’ (it has no status↔location coherence invariant); ‘diff’ → ‘in sync’ (DB and Markdown ag — on the WRONG state; agreement is correctness); (3) CM-CLOSURE-SEAM BINDS CHECK A → PASS (it flags NON terminal rows whose id appears in a w commit; these rows are terminal, so th are outside its predicate by constructio This was found by reading a status tall by the automated regime — a §11.4.23 discovery-channel escape, which is itse defect of equal standing to the mis-loca rows. ACCEPTANCE. (a) ‘update’ REFU a terminal status and names ‘close’ as correct path, with a paired §1.1 mutati proving the refusal (removing the guar must make the mutation pass). (b) ‘val grows a status↔location coherence invariant that FAILS on the forbidden s with a golden-bad fixture and a negativ control (a legitimately terminal row in must NOT fire — §11.4.201(1)). (c) The existing rows are drained to Fixed with class-matched evidence per row, or, wh row’s evidence cannot be produced, honestly re-opened rather than migrat a bare assertion. (d) Honest boundary: closes the update-path hole and the detection gap; it does not claim every historical status write was evidence-ba NOT CLAIMED. No fix is implemented this filing. The 10 rows are untouched; draining them is acceptance (c) and ea needs its own evidence, not a bulk UPI
BOB-167	Bug	Queued	Medium	WHAT. download-proxy exposes two Se Sent-Events routes. They are the same expensive class — long-lived connectio that hold a worker and a generator for lifetime — but only one is rate-limit cla routes.py:801 @router.get(‘/search/str

ATM ID	Type	Status	Severity	Description
				{search_id}’) @_rl(‘sse_stream’) <- cla routes.py:150 @router.get(‘/theme/stre async def stream_theme(…) <- NO limi decorator _rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/s is bound to the sse_stream class while theme/stream falls through to the application default. Measured live on t running stack: /api/v1/theme/stream re x-ratelimit-limit 120. PROVENANCE + CORRECTION WORTH KEEPING (§11 This was surfaced by the BOB-109 scal agent, whose report framed it as: ‘sse_stream_limit_decorator is defined imported/applied nowhere … SSE falls through to the 120/minute default: 24x protected’. That framing is WRONG an NOT filed as given. The agent searched one symbol NAME; the wiring uses a different mechanism (@_rl(‘sse_stream and it IS applied — to /search/stream. Verified by reading routes.py:795-806 the _rl helper at routes.py:46-51, with control needle confirming other *_limit_decorator symbols show real us in api/**init**.py so the zero-hit was not a search. The agent’s MEASUREMENT (on /theme/stream) was correct and is v makes this a real finding; its MECHAN was not. Both halves are recorded so t next reader does not re-derive the sam wrong cause. WHY IT MATTERS. An unclassed SSE endpoint is the cheapes to pin server resources: each connecti held open, and the default class permit 120/min of them. The declared sse_str class exists precisely because this rout shape needs a tighter bound than ordi GETs. ACCEPTANCE. (a) A decision, recorded, on whether /theme/stream belongs in the sse_stream class or gen warrants the default — this is a policy question, not automatically a bug to pa (b) If it belongs in sse_stream, the dec is applied and a test drives BOTH SSE routes and asserts each returns its INTENDED class limit from x-ratelimit so a future route added without a class caught. (c) A guard that enumerates S shaped routes and fails on any that ca

ATM ID	Type	Status	Severity	Description
				no explicit rate-limit class — the gener form, so the third SSE route does not r this. (d) Honest boundary: this does no claim 120/min is exploitable in this deployment; with network_mode host a no reverse proxy every caller shares th 127.0.0.1 bucket, which BOB-111 mea and recorded separately. NOT CLAIME change made. The limit values were re from headers, never driven to exhausti the limiter is per-IP and shared with concurrent agents on this host (§11.4.1
BOB-168	Task	Queued	Low	WHAT. scripts/run_all_challenges.sh:66 “scaling_horizontal_challenge.sh” in its challenge set. challenges/scripts/ scaling_horizontal_challenge.sh does n exist: $ sed -n ‘66p’ scripts/ run_all_challenges.sh “scaling_horizontal_challenge.sh” $ ls challenges/scripts/ scaling_horizontal_challenge.sh ls: can access …: No such file or directory VERIFIED independently, not taken on report — surfaced by the BOB-109 sca agent and re-checked here by direct invocation. WHY IT MATTERS. Whethe is cosmetic or a §11.4.201 gate-honest defect depends entirely on how the run treats a missing entry, and that is the f thing to determine: if it SKIPs silently, challenge bank advertises coverage it not have (a §11.4.266 claim-vs-reality r with no passing challenge behind it); if FAILs, the runner is permanently red f reason unrelated to the system under t which trains readers to ignore it. Neith outcome is acceptable; they need diffe fixes. ACCEPTANCE. (a) Determine an record the runner’s actual behaviour o missing entry by invoking it, not by rea it. (b) EITHER author the challenge, O remove the entry — with §11.4.124 discipline: check git history for whethe once existed and was deleted, since a silently-dropped challenge is the more interesting defect. (c) If the runner sile skips missing entries, that is its own fi a missing challenge must be loud (§11. SKIP-with-reason at minimum), never absent-and-quiet. SEVERITY. Low as a

ATM ID	Type	Status	Severity	Description
				defect, but it sits on the challenge-cove seam, so (c) may deserve its own item.
BOB-169	Bug	Queued	Medium	WHAT. The §11.4.65 markdown-export mandate requires every in-scope doc t ship .html/.pdf twins. 286 of the 326 .h files under docs/ (excluding dist/ and node_modules/) are pandoc FRAGMEN they begin at
BOB-170	Task	Queued	Medium	WHAT. The BOB-109 scaling suite gate two timing tests on loadavg_1m/nproc 0.75 — on this 8-cpu host, a 1-minute l average at or below 6.0 — and SKIPs l otherwise. That gate is correct and wa adopted for a good reason (§11.4.201(8 the author validated a 2.10 growth-exp threshold, then observed it FALSE-FAI correct code at load 11.06 (baseline sp 2.136), and at load 15-23 measured ba spans of 2.358 against injected-cubic mutants at 2.278-2.331 — the signal is smaller than the noise and the two are separable in either direction. Gating w honest response. THE PROBLEM. The has consequently NEVER RUN under i own precondition. Every recorded run measured load 9.15-23; the independe reviewer measured 9.66 during review conductor measured 6.45 at its lowest GREEN polarity WAS observed — at lo 9.4-11.6 (span 1.882) and in three stab runs (1.759/1.832/1.848) — but never beneath the shipped ceiling. The autho states this honestly as an inference fro data rather than a run they can paste, the reviewer confirmed the inference i sound (quiescence is strictly more favourable). Sound inference is still no observed run. WHY THAT MATTERS (§11.4.226). A registered, topology-pre guard that never executes is indistinguishable from a guard that wo fail if it did. Nothing currently ensures ever runs: no freshness contract, no scheduled quiet-window invocation, no tracked item — which is precisely the unexecuted-standing-guard class §11.4 names, where registration gets treated coverage while execution is nobody’s j The forensic precedent in that anchor standing guards that, when finally exe

ATM ID	Type	Status	Severity	Description
				immediately emitted latent FAILs. ACCEPTANCE. (a) Execute both quies gated timing tests on a host whose 1-m loadavg is genuinely <= 0.75/cpu — no agent fleet, no concurrent build — and capture the run as an artifact showing resolved load reading alongside the measured span exponent, so the precondition is provable from the artif and not asserted. (b) Record the obser GREEN polarity in the item and in the source, replacing the current honest- inference note. (c) If the gate FAILS ur genuine quiescence, that is the finding item exists to surface, and it reopens t threshold-calibration question rather t being worked around. (d) Establish a freshness contract per §11.4.226: how a quiescent verdict may be before the counts as uncovered again, so this doe silently return to never-executing. HOI BOUNDARY. This item does not claim t gate is wrong. The evidence points the way — quiescence is strictly more favourable than the loads where GREE was already observed. It claims only th run has not happened, and that an unexecuted guard cannot be cited as coverage. PROVENANCE. Raised as IMPORTANT-2 by the independent Fab substrate reviewer of the BOB-109 wo and independently corroborated: the reviewer verified the SKIP is loud and its resolved numbers, and confirmed th shipped tests do SKIP at today’s load.
BOB-171	Bug	Queued	Low	WHAT. Both rate limiters key their per bucket on the LEFTMOST element of X Forwarded-For when TRUST_FORWARDED_FOR is enabled: download-proxy/src/api/ rate_limit.py:117-123 (:7187) and the r stdlib limiter in plugins/download_prox (:7186), which was deliberately built to exact policy parity with it. That elemen CLIENT-SUPPLIED. An honest reverse APPENDS the peer address to the RIG whatever arrived, so the leftmost entry whatever the client sent — attacker- controlled by construction, not by misconfiguration. Consequences with t

ATM ID	Type	Status	Severity	Description
				flag on: (1) rotating a forged leftmost v mints an unlimited sequence of fresh buckets, which is total bypass of the li the flag is supposed to make MORE accurate; (2) it also defeats the LRU/id reap cap, since each forged identity is distinct key — the bucket map is a mer growth surface bounded only by the ca and the cap’s eviction then discards th budgets of REAL clients to make room forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, and that is why thi Low, not High. TRUST_FORWARDED_ defaults OFF at both sites, and the dep stack runs network_mode host with no reverse proxy in front (verified across five compose services), so peer addres are real client IPs and no NAT collapse occurs. The defect is latent: it arms the moment someone puts a proxy in front turns the flag on to recover real client which is exactly the situation the flag e for. The failure mode is therefore ‘corr looking configuration change silently disables the limiter’, not ‘currently bro PROVENANCE. Raised as MINOR-3 by independent reviewer of the BOB-111 limiter work and independently confirm by the implementing agent, which note at BOTH source sites as a tracked follo rather than fixing it in that change. Fil here because neither agent has write a to the tracker. WHY IT COVERS BOTH PORTS. :7186’s limiter was built to deliberate policy parity with :7187’s, including this parsing. Fixing one and the other would leave the two surfaces disagreeing on client identity — worse the shared defect, because it becomes surface-dependent and untestable as a single invariant. ACCEPTANCE. (a) Re leftmost-element parsing with a trust-a resolution at BOTH sites: either rightm minus-N-trusted-hops, or a trusted-pro CIDR allowlist where only a peer inside allowlist may assert XFF at all — the c is a deployment-topology decision and should be recorded, not assumed. (b) A that forges a rotating leftmost element asserts the bucket key does NOT rotat

ATM ID	Type	Status	Severity	Description
				site. (c) Paired §1.1 mutation: restore leftmost parsing and the test must FAI Negative control (§11.4.201(1)): with TRUST_FORWARDED_FOR OFF, the h must be ignored entirely and keying m fall back to the peer address — a guar starts honouring XFF when the flag is would be a worse defect than the one l fixed. (e) Honest boundary: this closes header-forgery bypass; it does not mak per-IP limiting fair behind NAT, where real users legitimately share one addre NOT CLAIMED. No change made. Both still parse leftmost; the flag still defaul
BOB-172	Bug	Queued	High	WHAT. The rutracker SEARCH endpoi refused by bot protection while the sit itself is fully reachable. Measured unauthenticated 2026-08-21 (no cookie read, no credential value in any artifac §11.4.10): GET /forum/index.php -> 20 96283 B,
BOB-173	Bug	Ready for testing	High	WHAT. download-proxy/src/api/ hooks.py:96-102: def _save_hooks(hook list[dict[str, Any]]) -> None: try: os.makedirs(os.path.dirname(HOOKS_ exist_ok=True) with open(HOOKS_FIL ‘w’) as f: json.dump(hooks, f, indent=2 except Exception as e: logger.error(f’F to save hooks: {e}’) It catches EVERY exception, logs, and returns None. The signature returns None, so the caller h channel to learn the write failed. Both sites then report success unconditiona :152 _save_hooks(hooks) :174 _save_hooks(hooks) :154 logger.info(‘Created hook: …’) :175 logger.info(‘Deleted hook: …’) :156 ret HookResponse(hook_id=…, …) :176 re {‘message’: ‘Hook deleted’, …} USER- VISIBLE CONSEQUENCE, which is wh is High and not a code-hygiene nit. A u POSTs a webhook, receives HTTP 200 hook_id, and the hook was never writt their automation silently never fires, a API told them it exists. Symmetrically, user DELETEs a hook, is told ‘Hook deleted’, the file still holds it, and it fir again after the next restart. In both directions the product reports the opp of what happened, and the only trace i

ATM ID	Type	Status	Severity	Description
				log line nobody is watching. THIS IS T §11.4.252 SHAPE. The path combines dangerous capabilities from that ancho taxonomy — MUTATION of a shared resource (a filesystem write) and EXTERNAL SIDE EFFECT (hooks are outbound calls the system will or will n make) — so it is required to FAIL CLOS verify the precondition, refuse when it cannot be satisfied, and surface the re Instead it fails open, and a bare ‘excep Exception:’ that only logs is the exact a pattern §11.4.252 enumerates. At the product layer it is also a §11.4.201(6) f null: a successful write and a swallowe failure are indistinguishable to the call PROVENANCE. Surfaced as an out-of-s observation by the independent review the BOB-135 test-isolation work — this swallow is what turned that defect into ‘assert 0 == 1’ mystery, because the EACCES on /config was logged and discarded while the endpoint kept retu 200. Verified here directly from source before filing (the function body and bo sites read above), not taken on report. Recorded as a §11.4.238 discovery-cha escape: found by an agent reading cod during an unrelated investigation, not automated QA regime — the coverage a defect of equal standing to the defec itself. ACCEPTANCE. (a) _save_hooks propagates failure — raise, or return a status the callers must consume. (b) Bo endpoints translate a persistence failu into an HTTP error (500-class), never a success body; a create that did not per must not return a hook_id. (c) A test dr each endpoint with the hooks file unwr (read-only dir or a patched open raisin OSError) and asserts a non-2xx status that a subsequent GET does not list th phantom hook — assert on the user-observable outcome, not on the log lin Paired §1.1 mutation: restore the swall the test must FAIL. (e) Audit the same for sibling swallows — this is a pattern one instance is rarely alone. (f) Honest boundary: this does not claim the write currently fails in production; it claims

ATM ID	Type	Status	Severity	Description
BOB-174	Bug	Queued	High	WHEN it fails the user is told the oppo and the BOB-135 investigation shows i fail in at least one real environment. N CLAIMED. No change made. No assess of how often the write fails in the oper deployment. RECORDING A SHELL EI OF MY OWN (§11.4.6): the first versio this description was written with the a pattern snippet inside backticks in a d quoted shell argument, so the shell ra command substitution and the text wa replaced by nothing — the stored description read ‘and is the exact anti- pattern’. This is the SECOND time this session that backticks-inside-double-qu has corrupted content (the first mangl commit message). Fixed here by editin through a Python client with no shell quoting in the path. Noted because a silently-truncated defect description is exactly the kind of quiet corruption §11.4.201(7)(c) warns about — the pat part of the instrument, and it failed wi erroring. WHAT. A corrupt hooks file is silently indistinguishable from “no hooks configured”, and the next create then DESTROYS every existing hook while returning HTTP 200. Three defects on path, filed together because fixing any alone leaves the data loss reachable. A download-proxy/src/api/hooks.py:104-1 _load_hooks wraps the read in `except Exception: logger.error(...)` and falls through to `return []`. A truncated or malformed hooks.json is therefore repo to every caller as an empty, healthy ho list. Confirmed at source. A2 — the consequence, and the reason this is Hi create_hook loads, appends, saves. Giv corrupt file that load turns into [], the writes a one-element list over the top. previously-configured hook is gone. Measured by the implementing agent against the ALREADY-FIXED tree, so t survives the BOB-173 write-failure fix: BEFORE on disk : [‘prod-hook-0’, ‘prod hook-1’, ‘prod-hook-2’] GET reports : 2 {‘hooks’: [], ‘count’: 0} <- claims ZERO hooks configured DELETE reports : 40

ATM ID	Type	Status	Severity	Description
				{‘detail’: ‘Hook not found’} <- for a ho that IS in the file POST reports : 200 hook_id=3c8a5930-… AFTER on disk : [‘3c8a5930-…’] Three prod hooks dest HTTP 200 throughout, nothing surface the user. A5 — _save_hooks writes with plain `open(path, "w")`. A crash, ENOSP kill mid-write truncates the file in plac That is precisely the corruption A1 the reads as “no hooks” and A2 overwrites same codebase already has the correct pattern: theme_state.py:115-126 uses file + os.replace. WHY THE THREE AR ONE ITEM. A5 manufactures the corru file, A1 misreads it as empty, A2 destro the contents. Fixing only A1 leaves truncation reachable; fixing only A5 le an existing corrupt file a data-loss trig fixing only A2 leaves the API lying abo what is configured. The chain is the de DELIBERATELY NOT FIXED UNDER BOB-173, and the reasoning is sound. implementing agent identified all three while auditing sibling swallows as that required, and declined to fix them in th same change because: they change GE semantics on an endpoint the frontend consumes (200 -> 5xx); they need a CORRUPT-file reproduction rather tha UNWRITABLE-file one BOB-173 built; they need a design decision that is genuinely not obvious — a MISSING h file legitimately means “no hooks configured”, while a CORRUPT one do not, and today’s code cannot tell those apart. Expanding BOB-173’s scope to d them would have meant shipping that decision unexamined. ACCEPTANCE. (Distinguish MISSING from CORRUPT: missing file remains an empty list; a co one is an error, never silently []. (b) De and RECORD what GET does on corru — 5xx, or 200 with an explicit degrade marker the frontend can render. This i design decision, and it should be state rather than inferred from whatever the patch happens to do. (c) create/delete MUST NOT overwrite a file they could parse — refuse, do not clobber (§11.4.2 mutation plus external side effect mus

ATM ID	Type	Status	Severity	Description
				closed). (d) Make the write atomic via os.replace, reusing theme_state.py's existing pattern rather than re-inventi (§11.4.28). (e) Tests asserting the USE OBSERVABLE outcome, not log lines: s corrupt file, assert GET does not claim hooks, assert a create refuses rather t destroying, and assert the file still hold original hooks afterwards. (f) Paired §1 mutation per guard. (g) NEGATIVE CONTROLS (§11.4.201(1)): a MISSING must still yield an empty list and a wor create; a VALID file must behave exact today. A fix that makes every load fail by failing always is not a fix. A3, RECORDED SEPARATELY, NOT PART THIS CHAIN. VALID_EVENTS and HookEventType are two sources of tru one closed set. Verified identical today unguarded against drift: if they diverg hook registers successfully and then si never fires. One assertion pinning ther would close it. NOT CLAIMED. No cha made. The BOB-173 write-failure fix is and orthogonal — it makes a FAILED v honest; it does nothing about a SUCCESSFUL write of wrong data der from a misread file.
BOB-175	Task	Queued	Low	WHAT. The workable-items `update` sear guards the status-location invariant in direction only. After BOB-166, `update - status <terminal>` on an Issues-located refused. The mirror is still open: `update location Fixed --status <non-terminal>` 0 and creates a row that `update`'s own validator immediately rejects. Verified empirically by the independent review BOB-166: update –location Fixed –stat progress' -> exit 0 validate -> FAILS, naming the row (check (f), fixedLocationNonTerminalStatus) So t command can still mint a state its own validate call refuses. That is the §11.4.196(F) shape at the seam layer: invariant is CONFIGURED in validate not ENFORCED at every write path th violate it. WHY IT WAS DELIBERATEL LEFT OPEN IN BOB-166, and why tha right. Three reasons, all checked rathe than asserted: (1) the real corpus has

ATM ID	Type	Status	Severity	Description
				instances of this direction against TEN the direction BOB-166 closed — the asymmetry in the data matched the asymmetry in the code; (2) detective coverage ALREADY existed and demonstrably fires, so unlike BOB-166 direction this cannot rot silently; (3) ar load-bearing one — `reopenCmd` documen and SANCTIONS a transient Fixed-plu terminal state via its `--location Fixed` operator override. A naive preventive g at the update seam would collide with override's semantics. Closing this ther requires a design decision about how t two interact, which is a different quest from the one BOB-166 was scoped to answer, and expanding that item would meant shipping the decision unexamin ACCEPTANCE. (a) Decide and RECOR a preventive guard on this direction co with reopenCmd's sanctioned override this is the actual work, and the answer should be written down rather than inf from whatever the patch does. Candida worth weighing: exempt the reopen pa explicitly, require an override flag on u mirroring reopen's, or leave the direct detective-only with the rationale recor so the asymmetry is a decision rather t an oversight. (b) If a guard lands, it re `terminalStatuses()` — BOB-166 establis one predicate source specifically so th directions cannot drift. (c) Paired §1.1 mutation. (d) NEGATIVE CONTROL (§11.4.201(1)), and it is the sharp one the sanctioned reopen path MUST still work. A guard that breaks reopenCmd documented override is a false-positive refusal, and worse than the gap it clos because it breaks a workflow operator told to use. HONEST BOUNDARY. This not a live data defect. Zero corpus instances, and the detective gate catch at the next validate. It is filed so a scop deliberate omission stays visible as a tracked decision instead of living only code comment where the next reader w mistake it for an oversight (§11.4.197: started thread reaches a documented terminal state, never quiet abandonme

ATM ID	Type	Status	Severity	Description
				PROVENANCE. Raised by the impleme agent of BOB-166 as a known asymme deliberately did not close, independent verified and endorsed as an acceptable scope boundary by that item’s reviewe the condition that it become a tracked rather than a comment. This is that ite