

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	5
Bug	Operator-blocked	2
Bug	Queued	7
Bug	Ready for testing	2
Task	In progress	10
Task	Queued	17
Task	Ready for testing	2
TOTAL		45

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	OPERATOR DECISION (2026-08-26, §11.4.66 clarification): COMPLETE THE CAPTCHA FLOW Session establishment path is DECIDED: the operator drives the interactive CAPTCHA flow at /api/rutracker/captcha then /login once, and the operator receives the resulting bb_session. The cookie-paste path and the cookies-file autoloading path were both offered but the operator chose the first. THE BLOCK NARROWS BUT DOES NOT PREVENT THE operator must physically drive the flow — not the system. The operator must solve a CAPTCHA (§11.4.52 honest operator_auditing boundary). AGENT PREP OWED BEFORE THE CAPTCHA so the operator’s single attempt succeeds rather than failing on discovering a broken endpoint mid-flow: verifying that all endpoints are reachable and correctly wired.

ATM ID	Type	Status	Severity	Description
				verify the session is actually PERSISTED after post (not merely accepted), and verify a stored actually USED by a subsequent search. Hones to state up front: a stored bb_session expires, requires repeating — that cost was accepted decision. This answer is recorded as consumed §11.4.35 — it is the operator’s stated choice, not inference, and supersedes any prior agent-choice on this question. Options not chosen are named for future reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict discipline applies to decisions). — prior item text follows — RuTra automated login blocked by CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-host SOCKS (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in download qBitTorrent-go + Jackett + compose env-forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto-commit mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel ledger coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent from test-type matrix
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit through §11.4.234 dedicated commit/push script
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code review substantive Auto-commit diffs
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/merge-script readiness ledger (closes GA-10)
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Tracked-Items Documents table row-completeness (GA-07 re)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exercising a start.sh subcommands end-to-end against real stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify rutra regex bounds deployed to container (closes re of GA-12)
BOB-094	Task		Medium	

ATM ID	Type	Status	Severity	Description
		In progress		RD2-29: Author tests/stress/test_tracker_fetch_stress_chaos.py with §11.4 injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/test_scheduler_hooks_sse_stress_chaos.py for triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for export/download-proxy/merge endpoints (canonical RD2-07)
				OPERATOR DECISION (2026-08-26, §11.4.66 clarification): DEFER PAST v1.0.0 Go --profile remains a goal but is NO LONGER a v1.0.0 release goal. RW-10..13 move to a post-1.0 milestone and stay on tag. Nothing advertised is removed, so §11.4.90 is clean and no §11.4.90 Obsolete closure is warranted. BOB-141 measured fact stands unchanged: the container's Dockerfile runs ONE binary binding to :7187; nothing binds :7186 or :7188 despite container setting PROXY_PORT/BRIDGE_PORT. CLAUDE's answer map already describes what the container ACCEPTS and needs no amendment under this decision. CLAUDE's answer is recorded as consumer DATA per §11.4.66 the operator's stated choice, not an agent inference. This supersedes any prior agent-chosen default on the question. Options not chosen are named above. The reader does not re-litigate a settled call (§11.4.66 bounded-verdict discipline applied to decision). The item text follows — GA-19/RW-09: Is --profile a release goal? (gates RW-10..13) — OPERATOR DECISION
BOB-101	Task	Queued	High	OPERATOR DECISION (2026-08-26, §11.4.66 clarification): KEEP 0.0.0.0 + MANDATORY GUARD The tunnel STAYS LAN-reachable — IP address changes to 127.0.0.1. The operator expects to preserve access from other devices on the network. DECISION CARRIES A BINDING OBLIGATION. §11.4.135 regression guard MUST land that FastAPI BUILD if any LAN-reachable route ever stops responding to authentication. The operator accepted 0.0.0.0 as a CONDITION that guard exists — without it the tunnel is unprotected and the item is NOT closeable. Guard requirements: enumerate routes from the authoritative source (FastAPI app, Go handlers, boba-jackpot)
BOB-102	Task	In progress	Medium	

ATM ID	Type	Status	Severity	Description
				hand-maintained list; resolve real auth coverage grep for the string ‘auth’ (§11.4.201 real-conditions) deliberately-public routes exempt ONLY via a list with per-entry justification, enumerated and gaps never silent; golden-TRUE + golden-FALSE per §11.4.107(10). This answer is recorded as DATA per §11.4.35 — it is the operator’s statement an agent inference, and supersedes any prior chosen default on this question. Options not chosen named above so a future reader does not re-litigate settled call (§11.4.112(5) bounded-verdict decisions applied to decisions). — prior item text follows LAN-exposure threat model — bind tunnel 12 keep 0.0.0.0? — OPERATOR-DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested- regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence-check the unattributed auto-commit path
BOB-109	Task	Ready for testing	Medium	BOB-074 followup: scaling-class test coverage mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (access usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiting public HTTP endpoints
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout arch (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_create fail suite (Permission denied /config)
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate with implementation, and 24 of 36 scripts have no doc
				OPERATOR DECISION (2026-08-26, §11.4.66) THEN RE-VERIFY UNTIL STABLE The warm path keeps running the ownership repair again stack, but walks, re-verifies, and repeats until nothing new. Options not chosen: refuse the the stack is up and direct the operator to –re- the stack on the warm path too (correct but e accept and document the window. WHY THIS MORE THAN A NORMAL ITEM: this is the exa
BOB-159	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				whole 002-user-owned-downloads feature ex The –recreate path was already fixed to orde precondition -> down -> repair -> up so the w tree no container can write. The warm path v container can create a new non-owned file A repair has passed that directory — the owner leaking back in through the other door. REQU CHOSEN OPTION, and it is the hard part: the BOUNDED, and on giving up it MUST report h ownership is UNPROVEN rather than silently (§11.4.201(6) — a loop that stops finding new it ran out of iterations returns the same quiet genuinely clean tree). An active download cr faster than the walk completes is the realistic converging case and must be named in the gi message. Recorded as consumer DATA per §1 operator’s stated choice, not an agent inferen not chosen are named so a future reader doe litigate a settled call. — prior item text follow Reported-Via: §11.4.202 reporting directive 1 2026-08-21T19:01:17Z Reported-By: T041 ind review (IMPORTANT-2), partially remediated report, verbatim): The –recreate path was fi round: run_ownership_precondition -> stack run_ownership_repair -> stack_up, so the rep tree no container can write to. That ordering asserted behaviourally by three new checks i test_start_reload_recreate.sh, each killed by i mutation (11/11 RED-OK). The WARM path is by that fix and this item tracks the remainder warm ./start.sh the gate runs before THIS inv brings anything up, but if the stack is ALREA containers write while the repair walks. A co then create a new non-operator-owned file B walk, after which the completion marker rec complete over a tree that is not, and those str never repaired without a manual –force. BOU dismissed: after any successful pass the mark the scope fingerprint and the repair short-cir walking (marker_is_valid), so the window rec or-absent marker AND a live stack together. I new scope entry re-arms the marker, which i that combination arises in practice. NOT DON

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	ROUND, with the reason stated rather than h clean guard is a cheap marker-only probe mo scripts/ownership_repair.sh that answers has been-repaired without walking the tree. The run cannot serve: it deliberately SKIPS the m path (FORCE=0 && DRY_RUN=0 guard) and al so using it as a warm-start probe would walk twice on every start. Adding that mode requi ownership_repair.sh, which another stream v editing during this round; duplicating marke into start.sh instead would be a near-identica (11.4.251). Deferred to this item rather than c unilaterally or silently. The misleading comm warm-start call site (which claimed the gate n any container writes) has been corrected to s boundary honestly. Affected scope / file-scope start.sh (warm-start dispatch, run_ownership site); scripts/ownership_repair.sh (marker fas Reproduction / context: 1. Ensure the stack Change config/owned_paths.yaml so the scop changes (this re-arms the marker by design, c E2). 3. Run a warm ./start.sh (no --recreate). 4 ownership repair walks and chowns the decl while containers are still running and able to Acceptance criteria: A warm ./start.sh cannot an ownership repair while a container that v declared location is running. Either the repai with an actionable refusal naming --recreate, is quiesced for the walk. Asserted behaviour unit/test_start_reload_recreate.sh alongside t PRECONDITION_BEFORE_DOWN / REPAIR_A REPAIR_BEFORE_UP checks, each with a pair mutation that kills it. Reported-Via: §11.4.202 reporting directive + 2026-08-21T19:35:02Z Reported-By: BOB-106 remediation, residual What (the report, ver TESTS for both guards are now executed: pre invariant 30's glob was extended from tests/u pre_build to also cover tests/hooks. MEASURE after: the glob expanded to 27 suites with 0 u hooks while 3 existed there; it now expands t under tests/hooks. All three pass. But the GUA

ATM ID	Type	Status	Severity	Description
				themselves are still invoked by nothing. check-brief-inputs.sh is honestly conductor-run rather than a brief's required-input list lives in free-form. The Agent tool's structured input does not expect an automatic prompt-scraping gate would itself be an unproven pattern-match. Its seam is the discrepancy in procedure, and that is a documentation + hook, not a hook. unattributed-commit-guard.sh is why this item exists. Run against real history. 14 violating commits since the last tag (and 20 commit subjects exist across all refs). Wiring pre-build or commit seam AS-IS would therefore trigger every commit immediately, on pre-existing data. We can fix in the moment. That is precisely the 11.4.66/11.4.67 brownfield-adoption shape, and 11.4.66/11.4.67 choice with the operator, not with an agent in the ratchet. The options, stated so the decision is made and not a default: (a) immediate hard floor – refuses until all 14 are attributed; (b) one-time decrease ratchet (the 11.4.135 pattern) – snapshot baseline, the count may fall and may never rise; (c) one is green and the disease cannot spread; (d) commits-only – gate only commits created from now with a scheduled deadline for the historical 14 commits only for a stated period, then escalate. Recommendation: with the reasoning rather than just the pick: 11.4.135 pattern this constitution already uses for example. In this situation, it makes the debt visible and bounded immediately, and it cannot be satisfied by deletion. unattributed-guard's name (11.4.227 closes that gaming check). This is the operator's call and this item is not closed. The answer is recorded. Honest boundary (11.4.66/11.4.67) does NOT claim the 14 commits are harmful. It claims their CONTRIBUTION is absent and therefore currently prevents the 15th. Affected scope / manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-brief-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh. Reproduction / context: Both guards run concurrently on hand and are covered by passing tests (9/9 are proven non-bluff against a no-op stub). Neither is covered by scripts/pre_build_verification.sh or scripts/commit-push-all.sh, so neither exerts standing detection.

ATM ID	Type	Status	Severity	Description
BOB-170	Task	Queued	Medium	(11.4.226: a guard with no execution seam is registration is not coverage). Acceptance cri guard is invoked by a named seam, OR carrier deferral pointing at this item. For unattributed guard.sh specifically, the operator has answered adoption question below and the answer is r consumer DATA – never an invented ratchet. WHAT. The BOB-109 scaling suite gates its two on loadavg_1m/nproc <= 0.75 — on this 8-cpu minute load average at or below 6.0 — and S otherwise. That gate is correct and was adopted reason (§11.4.201(8)): the author validated a exponent threshold, then observed it FALSE- code at load 11.06 (baseline span 2.136), and measured baseline spans of 2.358 against injected mutants at 2.278-2.331 — the signal is smaller noise and the two are not separable in either Gating was the honest response. THE PROBLEM has consequently NEVER RUN under its own Every recorded run today measured load 9.1 independent reviewer measured 9.66 during conductor measured 6.45 at its lowest. GREEN WAS observed — at loads 9.4-11.6 (span 1.882 three stability runs (1.759/1.832/1.848) — but beneath the shipped ceiling. The author stated honestly as an inference from data rather than can paste, and the reviewer confirmed the inference is sound (quiescence is strictly more favourable inference is still not an observed run. WHY THIS MATTERS (§11.4.226). A registered, topology- that never executes is indistinguishable from would fail if it did. Nothing currently ensures no freshness contract, no scheduled quiet-wi invocation, no tracked item — which is precisely unexecuted-standing-guard class §11.4.226 n registration gets treated as coverage while ex nobody’s job. The forensic precedent in that a standing guards that, when finally executed, emitted latent FAILs. ACCEPTANCE. (a) Execution quiescence-gated timing tests on a host whose loadavg is genuinely <= 0.75/cpu — no agent concurrent build — and capture the run as a

ATM ID	Type	Status	Severity	Description
				showing the resolved load reading alongside measured span exponent, so the precondition from the artifact and not asserted. (b) Record GREEN polarity in the item and in the test so replacing the current honest-inference note. FAILS under genuine quiescence, that is the item exists to surface, and it reopens the three calibration question rather than being worked. Establish a freshness contract per §11.4.226: 1 quiescent verdict may be before the gate could uncovered again, so this does not silently return executing. HONEST BOUNDARY. This item does the gate is wrong. The evidence points the other quiescence is strictly more favourable than the where GREEN was already observed. It claims the run has not happened, and that an unexecuted cannot be cited as coverage. PROVENANCE. R IMPORTANT-2 by the independent Fable-subreviewer of the BOB-109 work, and independent corroborated: the reviewer verified the SKIP prints its resolved numbers, and confirmed the tests do SKIP at today's load. WHAT. The rutracker SEARCH endpoint is reachable protection while the site itself is fully reachable unauthenticated 2026-08-21 (no cookie file reached credential value in any artifact — §11.4.10): C index.php -> 200, 96283 B, WHAT. The workable-items update seam guard location invariant in one direction only. After update --status <terminal> on an Issues-location refused. The mirror is still open: update --location --status <non-terminal> exits 0 and creates a update's own validator immediately rejects. V empirically by the independent reviewer of H update -location Fixed -status 'In progress' -> validate -> FAILS, naming the row (check (f), fixedLocationNonTerminalStatus) So the commit still mint a state its own validate call refuses. §11.4.196(F) shape at the seam layer: the invariant CONFIGURED in validate but not ENFORCED write path that can violate it. WHY IT WAS D. LEFT OPEN IN BOB-166, and why that was right
BOB-172	Bug	Ready for testing	High	
BOB-175	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				reasons, all checked rather than asserted: (1) corpus has ZERO instances of this direction a the direction BOB-166 closed — the asymmet matched the asymmetry in the code; (2) detect ALREADY existed and demonstrably fires, so BOB-166's direction this cannot rot silently; (3) load-bearing one — reopenCmd documents and a transient Fixed-plus-non-terminal state via location Fixed operator override. A naive pr guard at the update seam would collide with override's semantics. Closing this therefore r design decision about how the two interact, v different question from the one BOB-166 was answer, and expanding that item would have shipping the decision unexamined. ACCEPTA Decide and RECORD how a preventive guard direction coexists with reopenCmd's sanction — this is the actual work, and the answer sho written down rather than inferred from wha patch does. Candidates worth weighing: exer reopen path explicitly, require an override fl mirroring reopen's, or leave the direction det with the rationale recorded so the asymmetr rather than an oversight. (b) If a guard lands, terminalStatuses() — BOB-166 established o source specifically so the two directions cann Paired §1.1 mutation. (d) NEGATIVE CONTRO (§11.4.201(1)), and it is the sharp one here: th reopen path MUST still work. A guard that br reopenCmd's documented override is a false- refusal, and worse than the gap it closes, bec a workflow operators are told to use. HONES This is not a live data defect. Zero corpus inst the detective gate catches it at the next valida so a scoped, deliberate omission stays visible decision instead of living only in a code comm the next reader will mistake it for an oversigh a started thread reaches a documented termi never quiet abandonment). PROVENANCE. R implementing agent of BOB-166 as a known a deliberately did not close, independently ver endorsed as an acceptable scope boundary b

ATM ID	Type	Status	Severity	Description
BOB-184	Task	Queued	Medium	reviewer, on the condition that it become a tr rather than a comment. This is that item.
				Three icon-glyph controls (.bridge-retry, .then toggle, .caret) render text made of non-BMP c axe-core declines to decide their contrast, rep under the incomplete channel with message and resolving NEITHER fgColor NOR bgColor, Python recomputation in docs/qa/BOB-164/ axe_contrast_scan.py cannot decide either: 4 across 16 palette x mode combinations. BOB- stopped them being scored as silent passes by them GLYPH_UNMEASURED, printing them o and fencing them, so any NEW unmeasured g blocks. What remains genuinely unproven is contrast itself: as non-text user-interface com owe 3:1 under WCAG SC 1.4.11, and no oracle measures that today. Acceptance: a measure that resolves the effective foreground and ba glyph controls (for example reading getComp colour against the composited backdrop, or r glyphs with inline SVG carrying explicit fill to of the three controls shown to clear 3:1 in all mode combinations, and the GLYPH_UNMEA shrunk by exactly the nodes then proven.
BOB-185	Task	Queued	Medium	docs/qa/BOB-164/axe_contrast_scan.py scans with no backend, so every node whose existe on fetched data is absent from the page it me results table renders empty, so .type-badge a badge never appear, and the hooks list rende so .status pills never appear. That is precisely BOB-164 round-1 regression went unseen in a DOM even though it was scanned in all 16 th why round 2 had to close the class in the arit instead. Those nodes are now covered by fro app/models/style-contrast.spec.ts, which extr foreground-on-fill pairs from the stylesheets, ARITHMETIC claim about declared colour pa rendered-pixel one: it cannot see opacity app ancestor, a composited backdrop, or a cascad static extractor does not model. Acceptance: t drives the dashboard with seeded fixture dat backend, a route fixture, or an injected comp

ATM ID	Type	Status	Severity	Description
BOB-187	Bug	In progress	—	harness) so the badge and status nodes are p DOM, axe measures them directly, and a cont proves the rendered oracle sees a seeded sub before any clean result from it is believed. scripts/ownership_precondition.sh consumes unreviewed .env-driven scope as ownership_ creates probe files inside it, but it is NOT fenc remediation added ownership_path_fence() t ownership.sh and wired it into ownership_re closing the path-escape class there; the preco outside that agents file scope and still accepts config/owned_paths.yaml expands to. Blast ra than a recursive chown (it writes probe files mutating ownership of an existing tree) but t identical and equally unreviewed, so the sam QBITTORRENT_DATA_DIR value that would h filesystem-wide chown will drive probe-file c arbitrary location. Note the escape is empiric confirmed for the repair path, not merely rea during T028 remediation a probe declaring th shipped entry with QBITTORRENT_DATA_DIR suite, and ps showed the real artifact executi -uid 1000 -o ! -gid 1000) -printf ...” before it w verified pid. Acceptance: ownership_precond the SAME ownership_path_fence() predicate lib/ownership.sh (never a second dialect, 11.4 fail-closed on a rejected entry, and ships a pa mutation proving the refusal fires PLUS a neg proving the six live shipped entries still ACC is not a 11.4.201(1) false-positive refusal. Disc of-band during T028 remediation and report agent as needing tracking, so it is also a 11.4.2 escape. OPERATOR DECISION (2026-08-26, §11.4.66 clarification): UNTRACK THE BINARY — BU DEMAND §11.4.30 answer, recorded as consu the compiled workable-items binary is NOT t tracked. ‘Any build derivate which we can re executing proper mechanism for generating versioned’ applies without exception here. K tracked with only the invariant-52 fingerprin keeping it tracked with a commit-time rebuild
BOB-188	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
				offered and NOT chosen. REQUIRED END STA builds from source, or REFUSES HONESTLY v toolchain is absent (§11.4.201(11) — probe th through its real invocation path, never fake a silently fall back to a stale binary). Staleness structurally impossible rather than merely d ACCEPTED COST: a fresh clone needs a Go too before invariant 17 can run — already true f qBitTorrent-go backend, so this adds no new requirement. The invariant-52 fingerprint ga useful as defence-in-depth for as long as any on disk. This answer is recorded as consumer §11.4.35 — it is the operator’s stated choice, r inference, and supersedes any prior agent-ch on this question. Options not chosen are nam future reader does not re-litigate a settled cal (§11.4.112(5) bounded-verdict discipline appl decisions). — prior item text follows — pre_build_verification.sh invariant 17 (CM-W ITEMS-VALIDATE) resolves its binary through candidate loop at :534, whose FIRST entry is c scripts/workable-items/bin/workable-items. T GIT-TRACKED (md5 17644a248363, identical workable-items-linux) and executable, so it V resolution over the current untracked sibling scripts/workable-items/workable-items (md5 43376a6d0184). The tracked binary is STALE: contain the guards its own source now has. M needle-proven, 2026-08-25. String presence in binary: “refusing to set terminal status” -> 0; location item has TERMINAL status” -> 0; con from the SAME updateCmd, “at least one mut is required” -> 2 (so the instrument sees thro and the zeros are real absences, not a blind r untracked sibling returns 1 / 1 / 2 for the sam queries. A check whose message string is abs binary cannot fire. RUNTIME PROOF: the sam violation (a row set to terminal status while current_location=Issues) run through both bi stale reports 1 violation (only the body_md d current reports 2 including the location-statu verbatim. This is exactly why the ten BOB-16 (BOB-087/129/131/144/145/146/148/153/155/15

ATM ID	Type	Status	Severity	Description
				undetected: the gate that was supposed to catch stale artifacts while running a binary structurally incapable of seeing its sources. The 11.4.108 SOURCE->ARTIFACT gap: source artifacts are stale, every gate reading the artifact must be aware of this. This is the THIRD instance of that class found in this session, alongside BOB-169 (a charset gate that was never opened a file) and BOB-183 (a served bundle gate that is older than its sources). COMPOUNDING: the comment at :519-523 directly above the loop asserts “The workable-items path never existed in this checkout, so we have a gitignored local build-output dir nothing even close to populated”. That statement is now FALSE - because the tracked binaries. A stale comment asserting that the very file that now shadows resolution is hidden, but has stayed invisible. Also an 11.4.30 question the gate asks: owns: bin/workable-items and bin/workable-items. The answers are VERSIONED BUILD ARTIFACTS in a constant directory, a submodule that is inherited by reference by every consumer, so every consumer inherits which version of the gate was last committed. Flagged independently by several separate agents this session. ACCEPTANCE: (a) the gate must never prefer a stale artifact over a current one; (b) either the tracked binaries are removed and the gate falls through to build-on-demand, or a fresh binary is refused a binary older than its sources (see the fingerprint gate for a working pattern); (c) the comment at :519-523 corrected; (d) a paired 11.4.201(1) proving the new refusal fires, plus a negative test proving a genuinely fresh binary still passes the gate; (e) not an 11.4.201(1) false-positive refusal; (f) the tracked-binary decision recorded as operator error. Discovered out-of-band during BOB-166 removal. 11.4.238 coverage escape in its own right. WHAT: the §11.4.252 fail-open scan reports 0 for qBitTorrent-go, and that zero is NOT evidence that the agent ran control needles through the scanner. The scanner per §11.4.201(7)(b): a Python ‘except Exception: pass’ needle was SEEN, a TypeScript ‘catch (e) {}’ needle was NOT SEEN (so frontend/src’s zero IS a real zero), but the empty-‘if err != nil {}’ needle was NOT SEEN. The instrument that cannot see the idiom returns a quiet zero as a clean tree, and only one of the
BOB-191	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				DISTINCT FROM BOB-189, deliberately not m per §11.4.214: BOB-189 is a false MATCH (fail guards reported as fail-open); this is a false N entire language unscanned). Same scanner, o failure directions, different fixes - merging th hide one behind the other. IMPACT: Go is the qbittorrent-proxy-go and boba-jackett (port 7 owns encrypted tracker credentials), so the u surface is exactly where §11.4.252's credentia mutation combination is most likely. Nobody it; the dashboard says clean. ACCEPTANCE: th grows a Go arm covering the empty-err-block swallowed-error idioms, its needle is SEEN th real path, qBitTorrent-go's result is re-derived finding is triaged as this Python/TS pass was. qBitTorrent-go's fail-open posture is UNKNOW be reported as UNKNOWN, never as 0. === R INVESTIGATION, MEASURED 2026-08-26 — T TITLE WAS TOO NARROW (§11.4.6) === The l constitution/scripts/gates/ cm_dangerous_combination_fail_closed.sh, d root by invariant 39 at scripts/ pre_build_verification.sh:1481-1544. (scripts/ check_cm_no_fail_open_skip.sh has NEVER e branch — it exists only in an uncommitted w citation of that path is citing a file that is not WEAK-MATCHER DEFECT, NOT A SCOPE HOI sets exts="py go rs c cc cpp h hpp java cs js ts rb" — .go IS enumerated. But line :537 gates t analyser on case "\$f" in *.py), so every non falls through to just two text matchers (:772 e (...) {, :794 credential = x \\ \\ "literal"). catch keyword, so neither matcher CAN fire o construction. NEEDLE TEST, BOTH ARMS (pla result): Python 2 / 1 / SEEN · TypeScript 1 / 1 / 1 / 1 / SEEN · Go 6 / 0 / NOT SEEN — PASS · Ru SEEN · Ruby 1 / 0 / NOT SEEN · C 2 / 0 / NOT SE ONE OF FIVE BLIND EXTENSIONS, not a spec Go needle is real code: go build rc=0, go vet r clean. THE DECISIVE DISCRIMINATOR — sam extension changed: .zig (NOT enumerated) → topology_unsupported"; .go (enumerated) → gate ALREADY OWNS an honest-blindness m

ATM ID	Type	Status	Severity	Description
				it works correctly; Go routes around it precisely. ENUMERATED. A scope hole announces itself. matcher hole prints green. That makes it the and is why the corrected title leads with it. Q SURFACE (exclusions stated, §11.4.224(E)): us scanner’s OWN prune list and nothing added files / 20,389 LOC (52 production files / 6,858 I and WITHOUT exclusions the count is 106 == vendored tree exists, so no unstated exclusion manufacture a clean number. qBitTorrent-go blind files and 0 analysable files — 100% blind invariant 39 counts it among “5 first-party so reported clean. DEFECTS THE BLINDNESS HI BOB-204 (credential DELETE returns unconditional while the .env plaintext delete error is discarded env_write_failed_db_rolled_back code asserts unconfirmed rollback). SEPARATE SCOPE HO BOB-205 (cmd/boba-ctl, 947 LOC orchestrator DANGER_ROOTS entirely). Distinct-but-similar §11.4.214, deliberately not merged. GOLDEN-FIXTURE FOR ANY FUTURE GO ARM: internal auth_middleware.go:45-66 is correctly fail-C (constant-time compare, 401+return) — a Go NOT flag it. LOAD-BEARING §11.4.250 PREDIC bolting a naive Go regex arm on will IMMEDIATE manufacture BOB-189-class false positives in err != nil { return false } in a Go validator CLOSED. Fixing this as layer N+1 reproduces new language. THEREFORE the remediation (1) analyser REGISTRY with an UNANALYSED FIRST — extension→analyser map, unmapped reports UNANALYSED never silent PASS, which the false null into an honest gap and covers not too; (2) only then a Go arm, call-site-aware from (go/ast, or adopt errcheck/staticcheck rather than rolled regex), shipped with the auth_middleware FALSE. ONE PRIMITIVE, TWO SYMPTOMS: the classifies by local syntactic shape and never considers semantic role — BOB-189 ignores the CALLED (false match), BOB-191 searches for a shape from DIFFERENT language family (false null). The header already concedes the primitive (“require data-flow analysis this gate CANNOT honestly

ATM ID	Type	Status	Severity	Description
				what is missing is propagating that concession to the LANGUAGE honesty. Keep two tracker items. INSTRUMENT BUG RECORDED, NOT HIDDEN (§11.4.201(12)): the investigation's first verdict used <code>grep -oE 'PASS\ FAIL\ SKIP'</code>, which made inside the gate's own name (...FAIL-CLOSED) mislabelled three PASSes. Re-run keyed on example table above is the corrected run. EVIDENCE: <code>BOB-191/scanner_blindness_investigation_2026-09-25</code> (255 lines, §11.4.44 header). No source, gate, or modification by the investigation. === ROUND-3 MEASURED 2026-09-25 — STEP (1) OF THE REMEDIATION ORDER LANDED; STEP (2) STILL PENDING ACCEPTANCE NOT YET MET === Per the round-3 remediation order, step (1) — the analyser RECORDED an honest UNANALYSED verdict — is now implemented and independently re-verified (conductor re-compiled mutation suite from a clean shell, 206/206 green) — a real semantic Go arm — is DELIBERATELY RECORDED so this item stays Queued: the ACCEPTANCE CHECKED above (“its needle is SEEN through the real picture”) is unmet for Go/Rust/Ruby/C. CORRECTION TO MAKE: OWN FIGURE (§11.4.6, re-verified live, not re-verified) round-2 said “5 blind extensions (go/rs/rb/c/h) live against all candidates this round: C++ (.cpp) empty catch() {} is SEEN — rc=1 FAIL, already caught. It was never actually blind; .h is content ambiguous (same header extension serves both) and is not a distinct 5th case. Exactly 4 languages structurally blind: Go, Rust, Ruby, C. This now widens, the round-2 figure. WHAT LANDED (scripts/gates/cm_dangerous_combination_fair) a constitution-submodule file, inherited by real landing it for real still needs the submodule's fetch/pull/push workflow, separate from bob's path): - New DANGEROUS_COMBO_UNANALYSED_EXTENSION (default go rs rb c). Verdict now has 3 branches → FAIL (unchanged); zero hits + UNANALYSED present → new honest Δ NOTE – N file(s) UNANALYSED PASS (PARTIAL COVERAGE) (still exit 0, advisory blocks a build on an unanalysed language); zero unanalysed → unchanged original clean scripts/pre_build_verification.sh invariants

ATM ID	Type	Status	Severity	Description
				owned): fixed to read the UNANALYSED count unconditionally, not only on nonzero exit — 100%-UNANALYSED root would still have silent “no fail-open anti-pattern across N clean root caller never seeing the NOTE. - Mutation test with 17 new assertions covering Go/Rust/Ruby UNANALYSED-branch RED/GREEN, a C++ neg (still caught, never marked UNANALYSED — must not swallow a working extension), a clean negative control, and a mixed-real-hit-not-ma Full suite 206/206 (189 pre-existing + 17 new) independently re-run by the conductor from LIVE FULL-REPO RESCAN RESULT (real DANC this round): Python backlog unchanged at 59 (download-proxy/src:25, plugins:32, scripts:2 existing, out of scope, advisory since 2026-08 honest signal: 121 files now reported UNANALYSED (qBitTorrent-go:117, cmd/boba-ctl:4 — BOB-2 hole now visibly folded into the same honest mechanism) instead of silently counting toward Zero new FAIL findings in Go — by design; a checker needs call-site-aware semantic analysis deferred. WHY STEP (2) STAYS DEFERRED, NOT REGRESSION: the round-2 §11.4.250 PREDICTION unrefuted — a naive Go regex arm would implement flag internal/jackettapi/auth_middleware.go golden-FALSE fixture) as a BOB-189-class false real arm needs go/ast or an adopted tool (error staticcheck), shipped WITH that golden-FALSE before it can fire. REMAINING WORK TO CLOSE ITEM: implement the Go semantic arm (go/ast errcheck/staticcheck-adopted) covering the error block and swallowed-error idioms, verified against auth_middleware.go golden-FALSE fixture, recheck qBitTorrent-go’s real result (currently UNKNOWN honestly reported as such), and triage every item as the Python/TS pass already is. Rust/Ruby/C first-party files in boba’s tree today (verified that their own arms are lower priority until such as the UNANALYSED registry already covers the error in the meantime. EVIDENCE: docs/qa/BOB-19 closure_evidence_20260925.md (full RED/GREEN transcripts, the C++ correction with live proof

ATM ID	Type	Status	Severity	Description
BOB-194	Bug	In progress	—	repo scan table, anti-bluff provenance). State uncommitted by design (constitution-submod the §11.4.26 workflow, not boba's commit-pu WHAT: agent worktrees under .claude/worktrees ephemeral copies pinned at arbitrary commit built, never shipped, deleted when the agent build gates that walk the tree with find(1) from root descend into them anyway. MEASURED (2026-08-25): CM-GO-TOOLCHAIN-MATCHES-FAILED the main build with exactly one finding from .claude/worktrees/agent-afda1df906c53 worktree 86 commits behind still carrying 'F1 golang:1.23-alpine' against a go.mod long since The main tree was CORRECT throughout. Perhaps a false-positive refusal is a FAIL-bluff of equal false pass: it blocks real work and teaches people bypass the gate. FIXED FOR ONE GATE: '.claude' that gate's PRUNE_DIRS with a deterministic stale mismatch inside .claude/worktrees must healthy tree, the SAME mismatch outside it M so the prune cannot widen into blindness (§1 Mutation verified: reverting the prune fails 2 SURVEY, measured not inferred: 8 of 13 pre_k walk the tree; only the fixed one prunes .claude using 'git ls-files' are structurally immune - w untracked. Empirically probed: check_cm_closure_seam_binds, check_cm_killpg_pgid_guard and check_cm_no_production_mutation_residue . ZERO worktree references - clean in practice UNMEASURED, not clean: check_cm_test_mock_pid_explicit_int and check_cm_test_mock_pid_patched_when_rea not reached before the probe budget expired HYPOTHESIS REFUTED, recorded so nobody an earlier revision of this item speculated that check_cm_plugin_count's 90s timeout was ca walking worktrees. It is NOT. That gate's find \$PLUGINS_DIR (check_cm_plugin_count.sh:2 worktrees are outside its walk entirely. Its slo - still running past 54s on a re-measure - but different, still-unidentified cause and belongs

ATM ID	Type	Status	Severity	Description
BOB-196	Task	Queued	—	item, not this one. ACCEPTANCE: every tree-v either prunes agent scratch or is proven imm measurement, each with a paired mutation. I asymmetry that bounds the risk: a stale work ADD findings to a security-relevant scan (kill residue), never remove them - so the exposur refusal and eroded trust, not a missed defect. WHAT: the CM-PLUGIN-COUNT pre-build gat SECONDS to verify 8 documented counts. Me 2026-08-25: exceeded a 90s probe budget (rc= bounded re-run completed at 'rc=0 wall=189s it is simply slow. CAUSE, MEASURED NOT GU marker loop at check_cm_plugin_count.sh:27 EVERY LINE of every governed document and pipeline per line. Per iteration: 'printf grep tr -d' (4 processes) plus 'printf sed -n' (2) = ~ docs are AGENTS.md 689 + CLAUDE.md 494 + features/Status.md 670 = 1853 lines, so ~11,11 spawns, measuring ~102 ms/line. Classic per- subprocess shell trap. THE IRONY WORTH PE because it explains why nobody simplified it: comment shows 'oE wc -l' was chosen DEL over 'grep -c' after measuring the \$11.4.201(1 on this host (ugrep 7.8.4) 'grep -coE' returned but 1 inside a 'set -euo pipefail' subshell, a co dependent reading. The correctness fix is right be kept; it is the PER-LINE application of it th three minutes. HYPOTHESIS REFUTED, recor nobody re-walks it: this is NOT worktree trav gate's find is scoped to \$PLUGINS_DIR (:224), agent worktrees are outside its walk. That gu raised on BOB-194 and is dead. LIKELY FIX (u stated as a candidate not a conclusion): pre-fi grep for lines containing 'CM-PLUGIN-COUNT entering the loop, reducing ~1853 iterations t that carry a marker, while leaving the per-lin logic - and its footgun-avoiding form - comple unchanged. WHY IT MATTERS beyond impac gate is on the pre-build critical path, and CLA leans on it as the mechanical authority for th distinct plugin rosters (43 curated / 43 engine bootstrap) whose conflation was BOB-149. A

ATM ID	Type	Status	Severity	Description
BOB-211	Bug	Queued	minor	enough to tempt anyone into skipping it prot ACCEPTANCE: runtime reduced with the coun footgun-avoiding parse preserved byte-for-by behaviour, proven by the existing paired mut biting plus a before/after wall-clock recorded nothing here questions the gate's VERDICTS - runtime. STATUS: LATENT on this host (measured: all s locations sit on ONE btrfs mount, fully owner expressive) and OPERATOR-GATED for verific Retained from the dismissed BOB-206 rather discarded with it, because the scenario is wel elsewhere: the declared mount is a udisks au REMOVABLE NVMe, so another operator's po being exFAT or NTFS is entirely ordinary. TW RESIDUALS: (1) REPAIR-LAYER DIAGNOSIS. C flattening mount whose flattened uid is NOT probe_location correctly REFUSES (verified: u uid=100999 both refuse). But it then hands ownership_repair.sh a chown that will return nothing in the repair path is fstype-aware — operator sees a permission failure whose rea REMOUNT OPTION, which no message says. I NOT proven: the shim used to verify BOB-206 stat READ, not the chown WRITE, and cannot whether chown actually fails. (2) MODE CON flattens MODE via fmask/dmask, so the prese true / 600 contract on .env and config/boba.d silently unenforceable there. §11.4.10-adjace credential file believed to be 600 that is worl mount option is a leak the mode check canno OPERATOR-GATED: a genuine end-to-end REL uid-flattening mount. The privilege-free rout (FUSE; fusermount IS present and user_allow — but bindfs is ABSENT on this host, as are m losetup. So verification needs either a bindfs privileged loopback mount. Neither was att (§11.4.21 — high-blast-radius host mutation i gated). ALSO RECORDED: NO test anywhere e uid-flattening filesystem — vfat/exfat/ntfs/fla loopback/mkfs/losetup/bindfs/fusermount all across the three ownership test files, needle-p

ATM ID	Type	Status	Severity	Description
BOB-219	Bug	In progress	major	(probe_location = 3 hits through the same ins ACCEPTANCE: (1) operator decides whether t bindfs or provide a privileged loopback so th built; (2) if verified, the repair path gains fsty diagnosis naming the remount remedy; (3) th contract either detects mode-flattening mour refuses, or documents honestly that it cannot there. DISCOVERY CHANNEL (§11.4.238): reta from the BOB-206 verification stream. WHAT, conductor-verified: docs/guides/tracker credentials.md is TRACKED. Its §11.4.65-man twins docs/guides/tracker-credentials.html an EXIST ON DISK, are NOT tracked, and ARE ign matched by the .gitignore <i>credentials</i> deny-al The .md is rescued by an explicit allowlist en nobody added entries for the twins. WHY IT VIOLATION AND NOT A HYPOTHETICAL: §11 every in-scope Markdown document to carry synchronized .html/.pdf siblings, and §11.4.2 every §11.4.65-scope doc to be reachable from README. Two artifacts that exist locally but committed are, from any fresh clone, ABSEN clone of this repository has a tracker-creden with no exports, while the authoring host loo The divergence is invisible on the machine th notice it. WHY NOBODY SAW IT: this is a §11. FALSE-NULL of the BOB-212 class. git status r the twins, so no gate and no author had a sig instrument and the clean tree return the iden zero. ANOTHER INSTANCE OF THE SAME ME filed here rather than separately because it is docs/qa/BOB-124/evidence/loginctl_user_state swallowed by *_user* at .gitignore:~55 and is uncommitted §11.4.83 QA evidence. Its SIBLI SAME DIRECTORY — user1000_grepped.log, user_scope_events_head.log, user_service_lif ARE tracked, saved only by the robust glob n qa/**/*log. That one file is lost purely becaus rather than .log. The pattern is exact: robust negations hold; per-file and per-extension re ACCEPTANCE: (1) both twins become trackab committed alongside their .md; (2) the BOB-1

ATM ID	Type	Status	Severity	Description
				evidence likewise; (3) whichever BOB-212 fix chosen, it MUST cover these — a fix that close swallow while leaving these three artifacts present; (4) an uncommittable has addressed the mechanism damage; (4) a check that a tracked .md in §11.4.238 has COMMITTABLE twins, so this class cannot be added silently. DISCOVERY CHANNEL (§11.4.238): for BOB-212 blast-radius sweep — and only on it pass. The first pass filtered by a source-extension filter omitted .pdf, which hid this finding entirely; the second ran with no extension filter over all 420 non-extensions and recorded the blind spot rather than shipping a file number. Worth keeping: an extension allowing for a false-null generator, which is the same shape as the one being investigated. === ROUND-2 PARTIAL CLOSURE MEASURED 2026-09-25 — criterion 1 already satisfied, criterion 2 unsatisfiable, criteria 3+4 done ONCE the .md-export-twin class === Independently verified via <code>git check-ignore -v</code> on all three named as CRITERION 1 (tracker-credentials.{html,pdf} uncommittable) — ALREADY DONE, pre-dating commit 9f6f517 already added <code>!docs/guides/tracker-credentials.{html,pdf}</code> rescue lines (<code>.gitignore</code> and <code>credentials.{html,pdf}</code> both are tracked and clean. Nothing to do. CRITERION 2 (BOB-124 <code>loginctl_user_state.txt</code> trackable) — UNSATISFIABLE. The file does not exist on disk. A trace in <code>git log --all --diff-filter=A --pretty=format:"%s" --</code> <code>"*loginctl_user_state"</code>, and is absent from the working tree. Its .log siblings in the same directory are tracked (rescued by the existing <code>!docs/qa/**/loginctl_user_state.log</code> negation), confirming the file was lost purely by a .txt extension, exactly as the original report claimed, but there is no recovery path and no legitimate way to “make it trackable” without fabricating QA evidence, which is itself an anti-bluff violation. This criterion is permanently closed, not deferred. CRITERION 3 (mechanism preventing recurrence) — PARTIALLY SATISFIED. A new guard (see criterion 4) closes the class of uncommittables exemplified: a TRACKED .md file whose §11.4.238 twin is gitignore-swallowed. It does NOT close the class of criterion 2 exemplified: a NON-.md QA-evidence file (with extension) swallowed by an unrelated gitignore. Extending the existing <code>!docs/qa/**/*.*.log</code> negation to

ATM ID	Type	Status	Severity	Description
				cover .txt was concretely measured this round. PROVEN UNSAFE for this project: .gitignore cookies_*.txt as this repo's own secret-file naming convention, and a scratch-tree experiment cookies_rutracker.txt / RUTRACKER_session.txt my_password.txt — real credential-shaped files relying on that exact pattern staying broad. The fix is now closed off with evidence rather than a future session to re-attempt and re-measure. REMAINING GAP: a general mechanism for “evidence file legitimately meant to be tracked” swallowed by a broad secret-shaped-filename pattern, distinguished from an actual secret file genuinely open — it needs a smarter signal than extension (e.g. content-sniffing, or an explicit allowlist-on-creation discipline) that this round attempt to design, since inventing one hastily exactly the credential-leak class the constraints prevent. CRITERION 4 (a check exists) — DONE the .md-export-twins class. New gate scripts/check_md_export_twins_committable.sh + self-tests/pre_build/test_check_md_export_twins_committable.sh GREEN, 2 golden-FALSE, fail-closed x3, real-world paired mutation — all independently re-run and confirmed passing by the conductor). Wired to pre_build_verification.sh as invariant 58/58 (EXPORT-TWINS-COMMITTABLE), mirroring to CM-GITIGNORE-SWALLOW-GUARD (BOB-212) exactly — BLOCKING, not advisory, since it is a scan needing no live stack. REMAINING WORK: CLOSE THIS ITEM: design and implement the residual — a mechanism distinguishing “a legitimate evidence file lost to a secret-shaped gitignore” from “an actual secret correctly caught by the gitignore” without narrowing the existing credential-pattern globs. Until then this item stays Queued; criterion genuinely done, criterion 3 is genuinely partially done. EVIDENCE: docs/qa/BOB-219/closure_evidence_20260925.md (full command before/after git check-ignore, the .txt-wide-re

ATM ID	Type	Status	Severity	Description
				experiment, 7-arm test output, standalone ga against the real repo). WHAT, both conductor-verified with control SCRIPT-DOCS-SYNC occurrences in scripts/pre_build_verification.sh = 0 (control needle MARKDOWN-EXPORT-SYNC = 6, so the instru §11.4.44 revision-header enforcement likewise (control needle 3). Neither §11.4.18’s companion mandate nor §11.4.44’s revision header is en invariant. THE PERVERSE CONSEQUENCE, M new script with NO companion doc passes th Writing the companion doc that §11.4.18 MA creates docs/scripts/.md, which invariant 16 (MARKDOWN-EXPORT-SYNC) then requires to and .pdf twins — and their absence is a BLOC This session hit it exactly: docs/scripts/test_gitignore_swallow_is_loud.{html,pdf} ar are two of invariant 16’s six violations. So the incentive gradient points AWAY from the con complying with §11.4.18 is punished, ignorin That is worse than an ungated rule — it is a r machinery actively discourages. §11.4.227 IS GOVERNING ANCHOR: a named gate that exi prose is gate debt, and the ledger of unimple names must be monotone-decreasing. CM-SC SYNC is named in §11.4.18’s own text and imp nowhere — one concrete row of exactly the 5 unimplemented population §11.4.227 was mi shrink. ACCEPTANCE: (1) either CM-SCRIPT-D implemented, or it is registered as a deferral tracked §11.4.197 item — silent absence is wh forbids; (2) same for §11.4.44’s header check; whichever is chosen, the incentive inversion must never be cheaper to skip a mandated d write one — if invariant 16 will demand twin creation path must produce them, or invaria scope out docs/scripts until it can; (4) a paire mutation per gate implemented. HONEST NO SCOPE: this item does not argue §11.4.18 shor enforced immediately — that is an operator o gate debt priority (§11.4.66). It argues the CUI is incoherent: an unenforced mandate whose
BOB-223	Bug	Queued	major	

ATM ID	Type	Status	Severity	Description
				triggers a different gate's failure. Either enforced or not. DISCOVERY CHANNEL (§11.4.238). T042 readiness preflight while explaining what broke invariant 16. Not by the automated QA. BOB-227 measured 3 of 4 artifacts of the CM-3 AUTHENTICATED gate untracked in git; criterion (commit them) was already resolved by an untracked commit before this session, closed 2026-09-23. was NOT addressed: 'a gate asserting that even if executable a pre-build invariant invokes is it a general mechanism preventing this whole (a) defect (a pre-build gate's own implementation shipping untracked, invisible to a fresh clone committed baseline for round-over-round diff) recurring for ANY future gate, not merely this. ACCEPTANCE: (1) enumerate every file path so pre_build_verification.sh invokes (via bash/ti python3 calls to scripts under scripts/pre_build every tests/pre_build/.sh and tests/hooks/.sh is likely via a static grep/parse of pre_build_verification itself, or a runtime trace; (2) assert every one paths is git-tracked (git ls-files --error-unmatch this as a new pre-build invariant so a future gate implementation is caught immediately, not discovered independently weeks later; (4) a fix creating an untracked fake gate-invocation target asserting the new check fails on it, GREEN on mechanism exists and the fake target is either tracked. Live kinozal credential test reports authentication status='empty', error=''' (found 2026-09-23). R (docs/qa/BOB-235/investigation_20260923.md) operator/environment: kinozal.tv publishes A 127.0.0.1 (host, container, DoH Cloudflare+Go agree; rutracker.org resolves normally as con KINOZAL_MIRRORS is empty so search.py fails kinozal.tv; the live mirror kinozal.guru answers non-browser client with a Cloudflare 403 'Just challenge (login POST, curl, aiohttp, operator two user-agents; no cf_clearance cookie; KINOZAL_COOKIES is loaded into the container code reads it). (B) code defect FIXED under T
BOB-231	Task	Queued	Major	
BOB-235	Bug	Operator-blocked	—	

ATM ID	Type	Status	Severity	Description
BOB-237	Bug	Queued	high	failed/1 passed, GREEN 19 passed): _search_k except branch only logged and stashed no dia the chip read empty with error=None; it now error_type + 'Kinozal request failed: '. Needs -reload-python to be live. Credential validity UNKNOWN (no request ever reached a login do not rotate credentials on this evidence. Op Details: WHAT: choose (1) reachable domain KINOZAL_MIRRORS=https://kinozal.guru + ./s -recreate, or change the roster primary in tra how to pass the Cloudflare challenge - FlareS deployed), cf_clearance cookie + code that ac KINOZAL_COOKIES, Jackett's kinozal indexer kinozal operator-blocked; (3) optionally re-ex kinozal-only cookie file. WHY: every agent-re was exhausted - DNS, mirrors, login POST, co propagation all measured. UNBLOCK CONDI test authenticates or the operator marks kind unsupported. WHO: repository operator. WHAT: CM-GATE-LEDGER-RATCHET (§11.4.234) the pre-build sweep during this session's bat attempt: unimplemented=414 exceeds the ch baseline=403 in constitution/scripts/gates/ gate_ledger_baseline.txt. This is a HARD, bloc (unlike the other gate findings surfaced in the sweep, which are all explicitly ADVISORY/nor per §11.4.234). INVESTIGATION (§11.4.102): th newest anchors this session's constitution pu (§11.4.268 tamper-evident evidence chain, §11 consensus advisory-only ban, §11.4.270 depe existence-verdict register, §11.4.271 waiver m were checked FIRST as the likely cause, since most recently landed (2026-08-26) and each r recommended mechanism gates in its own te OUT: every one of those gates' names (CM-CR CONSENSUS-ADVISORY-ONLY, CM-DEPENDEN EXISTENCE-VERDICT-REGISTER, CM-EVIDENC ANCHOR-CATCHES-RECOMPUTE-AND-TRUNC EVIDENCE-CHAIN-DELETE-REORDER-DETECT EVIDENCE-CHAIN-INCOMPLETE-VERIFICATION CM-WAIVER-ROSTERED-EXPIRY-TRACKED) is correctly registered DEFERRED against an OV

ATM ID	Type	Status	Severity	Description
				NNN tracked item, and each anchor's own §11.4.209 propagation gate is IMPLEMENTED. So this boba-side edit in this session (confirmed: git commit 34e42f0) shows the ledger run (docs/qa/BOB-237/unimplemented_gates_snapshot_20260925.txt) captured verbatim from bash constitution/scripts/gate_ledger_ratchet.sh) shows the state is 414 UNIMPLEMENTED + 121 IMPLEMENTED + 121 DEFERRED. The checked-in baseline (403) was legitimately bumped by commit 34e42f0 ("cit §11.4.209 gate rename, bump baseline 397->403" point after that, the unimplemented count grew more) without a matching baseline bump or registration for the newly-named gates, across more subsequent anchor-landing rounds in the constitution submodule. This is accumulated governance-corpus debt, not something introduced by a boba-side edit in this session (confirmed: git commit eba38e8..HEAD -- constitution/scripts/gate_ledger_baseline.txt inside the constitution submodule is empty — the baseline file itself was touched across the two merges this session per §11.4.209 WHY DEFERRED RATHER THAN FIXED HERE). Of the 414 unimplemented names are 11 that pushed past baseline, and then either implemented 11 gates' worth of real UNIVERSE code or registering 11 proper OWED-GATE-NNN with tracked-item citations, is itself real engineering scoped to the constitution submodule (universe project-specific per §11.4.17) — it is not a boba-side edit and doing it hastily mid-session, unrelated to actual mandate (Jackett incorporation + backdoor risks exactly the kind of rushed, unreviewed corpus edit §11.4.209's Fable-xhigh review requires to prevent. ACCEPTANCE: (1) a future session constitution's own maintainers) triages the 403 snapshot against the 403 baseline to identify unaccounted gate names; (2) each is either implemented or registered as a proper OWED-GATE-NNN definition (genuinely retired) cited in constitution/scripts/gate_ledger_removals.tsv; (3) the constitution baseline is re-bumped with a citation once accepted following the same pattern as commit 34e42f0

ATM ID	Type	Status	Severity	Description
				own pinned constitution pointer is then re-sy item closed with the passing CM-GATE-LEDG re-run as evidence. DISPOSITION FOR THIS S COMMIT: per §11.4.234(D) (the commit/push MUST always be able to complete; a failing ga skippable ONLY via an explicit recorded defe the skip recorded in the commit message so t tracked, never forgotten), this session's batch used the sanctioned skip (BOBA_SYNC_SKIP_ this tracked item BOB-237, rather than silentl or spending unbounded effort resolving an in universal-corpus debt mid-session. DISCOVER (§11.4.238): found by boba's own pre-build sw pre_build_verification.sh, invariant 38/57) du routine session commit — the automated gat discoverer, exactly as §11.4.238 requires. WHAT: while independently verifying this se BOB-077 subagent report, the conductor disc hooks/test_check_brief_inputs.sh — a track proving BOB-107 was ALREADY fully implem commit 7b45113c6b03524d9c799bdd68496ec202 (2026-08-21), weeks before this session dispat subagent that built a redundant duplicate (ca removed before commit — see BOB-107's own closure evidence). The SAME commit's own m independently confirmed it also closed BOB-1 testing/ddos_resilience.md:571 carries an exp “CLOSED 2026-08-21” note this session found for BOB-114 separately). Two confirmed stale cases from ONE commit raised a systemic qu 7b45113 close MORE currently-open items th this session happened to stumble onto? MEAS log -1 --format="%B" 7b45113 \ grep -oE " \ sort -u returns 14 distinct item reference BOB-088, BOB-092, BOB-106, BOB-110, BOB-12 BOB-154, BOB-158, BOB-159, BOB-160, BOB-16 BOB-163. Cross-checked against the live DB (s workable_items.db "SELECT atm_id, status F WHERE atm_id IN (...)": 9 of the 14 are ALR (BOB-079/092/120/136/154/158/160/161/163, al Fixed) — consistent with genuine closures al reconciled. 5 remain Queued: BOB-088, BOB-
BOB-238	Bug	Queued	medium	

ATM ID	Type	Status	Severity	Description
				BOB-110, BOB-159, BOB-162. HONEST BOUNDARY (§11.4.6): a bare mention of an item-ID in a late remediation-round commit message does NOT prove that commit closed it — the commit could have referenced an item as CONTEXT, PARTIAL, or addressed it, or introduced work that only later The 9-of-14-already-closed ratio is suggestive for the remaining 5. This item does NOT claim BOB-088/106/110/159/162 are stale — it claims UNVERIFIED against this specific commit and pattern so nobody re-derives the same discovery accident path this session took. ACCEPTANCE BOB-088, BOB-106, BOB-110, BOB-159, BOB-162 7b45113's full diff for the file(s)/section relevant item's own acceptance criteria (not just the commit MESSAGE's mention of the ID), and CONFIRMATION of these three outcomes actually applies — (a) already closes it — close it with a corrected citation to 7b45113, mirroring BOB-107/BOB-108-corrected closure pattern exactly; (b) the commit addresses it — update its description with the progress split, same pattern as this session's BOB-219; (c) the commit's mention is unrelated only — leave it open, note the mention was in and ruled out so nobody re-checks it. DISCOVERY CHANNEL (§11.4.238): found by the conductor's independent-verification discipline (re-running subagent's regression suite surfaced the stale not by any automated gate — itself worth not coverage-escape class: no mechanical check catches "a large remediation commit's own m references an item-id that never gets tracker- WHAT: §11.4.33's closed vocabulary mapping >Fixed, Feature->Implemented, Task->Completed comprehensive audit of the whole workable- (2026-09-25) found 30 Fixed-location items with Status pairing violates this mapping — all 30 (from earlier sessions, spanning at least back BOB-005/015/016/042..059), found AFTER this already self-corrected 4 of its OWN fresh instances identical mistake (BOB-077/100/179/226, all T mistakenly closed with status=Fixed instead of
BOB-239	Bug	Queued	low	

ATM ID	Type	Status	Severity	Description
				— caught by this same audit discipline and first session, before commit, and therefore NOT accounted here). REPRODUCTION (exact count independently re-verified twice, §11.4.6 — this item mis-stated the count as 25; the second recount below is the corrected, accurate figure) <pre>\$ sqlite3 docs/workable_items.db "SELECT a status FROM items WHERE current_location=' ((type='Task' AND status NOT LIKE 'Comple status NOT LIKE 'Obsolete%') OR (type='Bug NOT LIKE 'Fixed%' AND status NOT LIKE 'Obs (type='Feature' AND status NOT LIKE 'Imple status NOT LIKE 'Obsolete%')));" \ wc -l 3</pre> Type=Task): BOB-005, BOB-015, BOB-016, BOB-043, BOB-044, BOB-045, BOB-046, BOB-049, BOB-050, BOB-051, BOB-052, BOB-055, BOB-056, BOB-057, BOB-058, BOB-150, BOB-161, BOB-168, BOB-182, BOB-190, BOB-228, BOB-232 — 19 of the 30 (BOB-042..058) consecutive run, plus none others) are closed as “Implemented” (the Feature-only word); the 11 (BOB-005/015/016/059/141/150/161/168/182/190) — recount: that is 13, not 11; see honest note above closed as “Fixed” (the Bug-only word). Every item is Type=Task and should read “Completed”. FOLLOW-UP ON SUB-COUNTS (§11.4.6): the 19-Implemented (30-19)=11-Fixed split stated above does not hold against the full 30-item list (13 “Fixed” entries visible in the raw query output, not 11) — this does not re-verify the sub-split further; the AUTHORITATIVE figure is the machine-count of 30, independently reproducible via the command. Whoever picks up this item’s ACCEPTANCE criteria re-run the reproduction command fresh (never either sub-count stated in prose here) before the same discipline that caught this descriptive first-draft miscount. WHY NOT BULK-FIXED IN THIS SESSION: changing 30 historical closure Status is consequential, tracker-wide action — it rewrites every one of these items reads in Fixed.md/ Fixed_Summary.md and their exported HTML twins, touches items this session did zero independent investigation on, and deserves its own dedicated

ATM ID	Type	Status	Severity	Description
				regenerated docs and a focused commit, not the effect of an unrelated backlog-triage session carrying substantial other work. §11.4.33 its mandate WHEN a pre-existing violation must be corrected, only that new closures follow the CANDIDATE ROOT CAUSE (stated as a hypothesis proven — §11.4.6): the CLI's close subcommand status <fixed\ implemented\ completed\ observed> literal flag value the caller must choose correct item's own Type; nothing in the CLI itself chooses status word against the item's Type and mismatch. This session's own 4 fresh violations exactly this way — the closing agent (this correct simply typed --status fixed out of habit with Type first, on three separate close calls, before unrelated audit query caught the pattern. If 3 independent historical instances of the identical exist, the likely mechanism is the same human habit, uncaught because no mechanical gate Type<->Status mapping at write time — correct the fact workable-items validate ran clean through this session despite 30 (34 counting this session's transient instances) live violations existing simultaneously. ACCEPTANCE: (1) each of the existing items is individually corrected via workable-items update --status "Completed Fixed.md)" --location Fixed, docs regenerated, focused commit; (2) per the stated candidate extend workable-items close/update to VALID status word against the item's own Type and mismatch (closing the write-time gap that lets instances of this exact mistake land, 30 of the uncaught as of this filing) — a genuine candidate mechanical fix, not yet built; (3) extend workable-items validate's existing invariant set to catch this defense-in-depth (it currently does not, confirm session's own repeated clean validate: OK run database that simultaneously held these violations). DISCOVERY CHANNEL (§11.4.238): found by the conductor's own routine self-audit after catching instance of the same mistake on BOB-077 — no pre-build gate (workable-items validate ran

ATM ID	Type	Status	Severity	Description
				throughout this session despite 30+ live insta confirming this class is currently invisible to