Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	10
Bug	Operator-blocked	1
Bug	Queued	19
Bug	Ready for testing	4
Task	In progress	12
Task	Queued	28
Task	Ready for testing	2
TOTAL		**76**

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	**OPERATOR DECISION (2026-08-26, interactive clarification): COMPLET CAPTCHA FLOW ONCE** Session estab path is DECIDED: the operator drives tl interactive CAPTCHA flow at `/api/v1/au rutracker/captcha` then `/login` once, and stores the resulting `bb_session`. The cool path and the cookies-file autoload path offered and NOT chosen. THE BLOCK N BUT DOES NOT LIFT: the operator mus physically drive the flow — no agent can CAPTCHA (§11.4.52 honest operator_at boundary). AGENT PREP OWED BEFO HAND-OFF, so the operator's single att succeeds rather than discovering a brol endpoint mid-flow: verify both endpoint reachable and correctly wired end-to-er the session is actually PERSISTED after successful post (not merely accepted), a

ATM ID	Type	Status	Severity	Description
				a stored session is actually USED by a subsequent search. Honest limitation to front: a stored bb_session expires, so th requires repeating — that cost was acce the decision. This answer is recorded as consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent in and supersedes any prior agent-chosen this question. Options not chosen are na above so a future reader does not re-liti settled call (§11.4.112(5) bounded-verdi discipline applied to decisions). — prior follows — RuTracker automated login b CAPTCHA
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-hos routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in proxy + qBitTorrent-go + Jackett + com forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channe — ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
BOB-077	Task	Queued	High	**OPERATOR DECISION (2026-08-26, interactive clarification): UNKNOWI INSTRUMENT THE NEXT OCCURRI** operator does not currently know which produces the +0500 Auto-commit fast-fo so the three candidate answers (second session / intentional rsync job / stale job remain open. DECIDED ACTION: stop h host this session cannot reach (§11.4.6 state is not knowable without access) an INSTRUMENT the mechanism so the ne occurrence identifies itself. Forensic cap add: committer identity, hostname, time offset, push timing, and the git remote u recorded at Auto-commit time into a tra forensic log. ACCEPTANCE: the next Au commit event yields a captured record n origin host — at which point this item re one of the three original branches with rather than a guess. This is the §11.4.10 reversible-safe move: it costs little, bloc nothing, and converts a recurring myste self-identifying event. This answer is rec

ATM ID	Type	Status	Severity	Description
				consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent in and supersedes any prior agent-chosen this question. Options not chosen are na above so a future reader does not re-liti settled call (§11.4.112(5) bounded-verdi discipline applied to decisions). — prior follows — RD2-10: Identify second host the Auto-commit rsync/sync mechanism (OPERATOR-DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-com mechanism through §11.4.234 dedicate push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code r the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable the four Lava-porting findings (closes G
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes (
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Trac + Status Documents table row-complete (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exerc three start.sh subcommands end-to-end real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify ReDoS regex bounds deployed to contai (closes runtime half of GA-12)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos. side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage fo download-proxy/merge endpoints (cano of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one (canonical impl of RD2-09)
BOB-101	Task	Queued	High	**OPERATOR DECISION (2026-08-26, interactive clarification): DEFER PA** Go `--profile go` parity remains a goal bu LONGER a v1.0.0 release blocker. RW-1 move to a post-1.0 milestone and stop g tag. Nothing advertised is removed, so stays clean and no §11.4.90 Obsolete cl

ATM ID	Type	Status	Severity	Description
				warranted. The BOB-141 measured fact unchanged: the Go container's Dockerfi ONE binary binding only :7187; nothing :7186 or :7188 despite compose setting PROXY_PORT/BRIDGE_PORT. CLAUDE. map already describes what the contain ACTUALLY does and needs no amendm this decision. This answer is recorded a consumer DATA per §11.4.35 — it is the operator's stated choice, not an agent i and supersedes any prior agent-chosen this question. Options not chosen are n above so a future reader does not re-liti settled call (§11.4.112(5) bounded-verdi discipline applied to decisions). — prior follows — GA-19/RW-09: Is –profile go p a release goal? (gates RW-10..13) — OP DECISION
BOB-102	Task	In progress	Medium	**OPERATOR DECISION (2026-08-26, interactive clarification): KEEP 0.0.0 MANDATORY AUTH GUARD** The tun LAN-reachable — no bind address chan 127.0.0.1. The operator explicitly chose preserve access from other devices on t network. THE DECISION CARRIES A B OBLIGATION: a permanent §11.4.135 r guard MUST land that FAILS THE BUIL LAN-reachable route ever stops deman authentication. The operator accepted THE CONDITION that guard exists — w this decision is unprotected and the ite closeable. Guard requirements: enume from the authoritative source (FastAPI a handlers, boba-jackett) never a hand-m list; resolve real auth coverage not a gr string 'auth' (§11.4.201 real-condition); deliberately-public routes exempt ONL checked-in list with per-entry justificati enumerated as honest gaps never silent TRUE + golden-FALSE fixtures per §11. This answer is recorded as consumer D §11.4.35 — it is the operator's stated ch an agent inference, and supersedes any agent-chosen default on this question. (not chosen are named above so a future does not re-litigate a settled call (§11.4. bounded-verdict discipline applied to de — prior item text follows — RW-05: LA threat model — bind tunnel 127.0.0.1 o 0.0.0.0? — OPERATOR-DECISION

ATM ID	Type	Status	Severity	Description
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescenc helper for the unattributed auto-commit
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existe for subagent task-brief source inputs
BOB-109	Task	Ready for testing	Medium	BOB-074 followup: scaling-class test co absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test covera (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate l boba’s 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golde fixture for the rate-limit detector
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logou architectural gap (task #85, incident #
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_cre bulk suite (Permission denied /config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while th process still serves 7186 (GIL starvatio spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile serve 7186/7187/7188 but its container binds — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unres gitdir) pollute gate scan scope and man false BOB-126-class findings
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/4 constitution, CLAUDE.md and the REA badge; the badge is hand-maintained a unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labe 44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named ga implementation, and 24 of 36 scripts ha companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored thi code: 82% of one gate’s 38,291 finding from submodules/
BOB-154	Bug	Ready for testing	Medium	Host venv and production container ru starlette versions (1.4.1 vs 1.6.0)
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard i sensitive and flaky: 8786ms under host 900-1500ms ceiling calibrated on a qui

ATM ID	Type	Status	Severity	Description
BOB-159	Bug	Queued	Medium	**Reported-Via:** §11.4.202 reporting dire on 2026-08-21T19:01:17Z **Reported-By** independent review (IMPORTANT-2), pa remediated **What (the report, verbati** –recreate path was fixed this round: run_ownership_precondition -> stack_d run_ownership_repair -> stack_up, so th walks a tree no container can write to. T ordering is now asserted behaviourally new checks in tests/unit/ test_start_reload_recreate.sh, each kille paired 1.1 mutation (11/11 RED-OK). Th path is NOT covered by that fix and this tracks the remainder. On a warm ./start gate runs before THIS invocation brings up, but if the stack is ALREADY running containers write while the repair walks. container can then create a new non-op owned file BEHIND the walk, after whic completion marker records complete ov that is not, and those stragglers are nev repaired without a manual –force. BOU dismissed: after any successful pass the valid for the scope fingerprint and the r short-circuits without walking (marker_ so the window requires a stale-or-absen AND a live stack together. Declaring a n entry re-arms the marker, which is exac that combination arises in practice. NOT THIS ROUND, with the reason stated ra hidden: the clean guard is a cheap mark probe mode on scripts/ownership_repai answers has-this-scope-been-repaired w walking the tree. The existing –dry-run serve: it deliberately SKIPS the marker (FORCE=0 && DRY_RUN=0 guard) and walks, so using it as a warm-start probe walk the tree twice on every start. Addi mode requires editing ownership_repair another stream was actively editing dur round; duplicating marker_is_valid into instead would be a near-identical fork (Deferred to this item rather than done unilaterally or silently. The misleading c at the warm-start call site (which claime gate runs before any container writes) h corrected to state this boundary honest **Affected scope / file-scope manifest:** (warm-start dispatch, run_ownership_ga site); scripts/ownership_repair.sh (mark

ATM ID	Type	Status	Severity	Description
				path) **Reproduction / context:** 1. Ensu stack is UP. 2. Change config/owned_pa so the scope fingerprint changes (this re marker by design, data-model E2). 3. Ru warm ./start.sh (no –recreate). 4. The ov repair walks and chowns the declared t containers are still running and able to it. **Acceptance criteria:** A warm ./start complete an ownership repair while a c that writes to a declared location is run Either the repair is deferred with an act refusal naming –recreate, or the stack is for the walk. Asserted behaviourally in t test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOV REPAIR_AFTER_DOWN / REPAIR_BEFO checks, each with a paired 1.1 mutation it.
BOB-160	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting dire on 2026-08-21T19:02:48Z **Reported-By What (the report, verbatim):** Indeper §11.4.209 review (IMPORTANT-3) found the feature's strongest automated checl never executed by any runner: (1) tests/ ownership/test_container_writes_ownec — the §11.4.115 RED-turned-regression FR-011/FR-007, proven via a docker-cor revert mutation — was moved out of tes integration/ (commit 58d340a, to escap autouse fixture hang) but no runner (ci. test.sh, run-all-tests.sh) was ever extenc cover tests/ownership/. (2) tests/pre_bu test_*.sh, the §1.1 paired-mutation meta- the scripts/pre_build/check_cm_*.sh gate was executed by NOTHING — self-repor honestly in commit 04742d7's own mes ('STILL OPEN (reported, not fixed): test pre_build/ is executed by NOTHING') bu tracked as a workable item (a §11.4.197 requirements gap) and never wired into pre_build_verification.sh invariant 30, w globbed only tests/unit/test_*.sh. ***Affecte file-scope manifest:*** *ci.sh; scripts/ pre_build_verification.sh invariant 30 (C UNIT-TESTS-EXECUTED)* ***Reproductio context:*** *grep -rn test_container_writes_owned_files ci.sh run-all-tests.sh scripts/ returned zero ru (only docs/specs mentioned the path); g scripts/pre_build_verification.sh showec*

ATM ID	Type	Status	Severity	Description
				30's for-loop globbing only tests/unit/tes never tests/pre_build/test_.*sh.* ***Acceptan criteria:*** *ci.sh gains a runtime-gated py ownership/ stage (skips honestly, rc=0, podman/docker present); scripts/ pre_build_verification.sh invariant 30's covers both tests/unit/test*_.sh and tests pre_build/test_*.sh, verified by an extrac standalone run of the invariant's exact l reporting a non-zero RAN count that inc files from both directories.
BOB-161	Task	Queued	High	**Reported-Via:** §11.4.202 reporting dire on 2026-08-21T19:31:12Z **Reported-By** BOB-092 remediation, residual finding **report, verbatim):** §11.4.69 names CM OPEN-SKIP as one of three mandatory p gates: it audits sink-side probe helpers if any code path converts an empty or unreachable response into a PASS-coun for a feature class with a sink-side prob project does not have it. Why this matte rather than in the abstract: the BOB-09 remediation just removed TWO fail-open exactly this class from tests/e2e/ test_live_stack_evidence.py — the nnmo on-404 (already removed at 7baef2b) an surviving sibling at test_iptorrents_is_authenticated_in_sea skipped on 'not authenticated and statu success' while asserting 'treating as tra outage'. That assertion was false for rej credentials (upstream_http_403) and fo container (plugin_env_missing), both of definitive product failures the product a classifies via error_type. Two fail-opens class, found one at a time, is the §11.4.1 extend-to-all-cases signal and the §11.4 coverage-escape signal together: the re not surface either, an agent reading cod remediation's own guards are AST-struc lethal under three discriminating mutat they LIVE IN THE FILE THEY GUARD, deleting that file evades them entirely. A gate is the durable home because it aud corpus rather than one file. Honest bou (§11.4.6): this item does NOT claim the remediated fail-opens are unguarded — guarded, with runtime evidence (13 pas 97.36s against the live stack). It claims CLASS has no corpus-wide detector, so

ATM ID	Type	Status	Severity	Description
				instance in a different file is invisible ag **Affected scope / file-scope manifest:** pre_build/ (gate absent); tests/e2e/ test_live_stack_evidence.py (guards cur inside the file they guard) **Reproductio context:** grep -rl ‘CM-NO-FAIL-OPEN-S scripts/ tests/ returns exactly ONE hit, a tests/e2e/test_live_stack_evidence.py — the guards live in, not a gate. Control n same query for CM-OWNERSHIP-INVAL gate that does exist) returns 3 files, so t instrument can see gate tokens and the is a real absence, not a blind zero (§11.4 (b)). **Acceptance criteria:** A scripts/pre gate named CM-NO-FAIL-OPEN-SKIP e wired into scripts/pre_build_verification audits sink-side probe helpers for code converting an empty/unreachable/error into a PASS-counting SKIP for a feature HAS a sink-side probe. It ships a golden fixture (a real fail-open -> gate FIRES) golden-FALSE-with-carrier (an honest t geo skip, and a comment merely MENT the phrase -> gate MUST NOT fire), pe §11.4.107(10)/§11.4.201(1). Paired §1.1 makes the gate FAIL before the gate is
BOB-162	Task	Queued	Medium	**Reported-Via:** §11.4.202 reporting dir on 2026-08-21T19:35:02Z **Reported-By** BOB-106/BOB-107 remediation, residua **(the report, verbatim):** The TESTS fo guards are now executed: pre-build inv glob was extended from tests/unit + tes pre_build to also cover tests/hooks. ME before/after: the glob expanded to 27 s 0 under tests/hooks while 3 existed the expands to 30 with 3 under tests/hooks pass. But the GUARDS themselves are s invoked by nothing. check-brief-inputs.s honestly conductor-run rather than aut brief’s required-input list lives in free-f that the Agent tool’s structured input d expose, so an automatic prompt-scrapin would itself be an unproven pattern-ma seam is the dispatch procedure, and tha documentation + habit change, not a ho unattributed-commit-guard.sh is differe why this item exists. Run against real h names 14 violating commits since the la (and 21 bare Auto-commit subjects exis all refs). Wiring it into the pre-build or

ATM ID	Type	Status	Severity	Description
				seam AS-IS would therefore refuse ever immediately, on pre-existing debt nobod in the moment. That is precisely the 11. brownfield-adoption shape, and 11.4.66 put the choice with the operator, not wi agent inventing a ratchet. The options, the decision is a decision and not a defa immediate hard floor – the seam refuses 14 are attributed; (b) one-time monoton decrease ratchet (the 11.4.135 pattern) snapshot 14 as the baseline, the count r and may never rise, so day one is green disease cannot spread; (c) changed-com – gate only commits created from now o scheduled deadline for the historical 14 report-only for a stated period, then esc Recommendation, with the reasoning ra just the pick: (b). It is the pattern this co already uses for exactly this situation, it the debt visible and bounded immediate cannot be satisfied by deleting the guar (11.4.227 closes that gaming channel). the operator's call and this item is not c until that answer is recorded. Honest bo (11.4.6): this item does NOT claim the 1 commits are harmful in themselves – it their ATTRIBUTION is absent and that r currently prevents the 15th. **Affected s file-scope manifest:** scripts/hooks/una commit-guard.sh; scripts/hooks/check-b inputs.sh; scripts/pre_build_verification scripts/commit-push-all.sh **Reproducti context:** Both guards run correctly by h are covered by passing tests (9/9 and 15 proven non-bluff against a no-op stub). invoked by scripts/pre_build_verification scripts/commit-push-all.sh, so neither ex standing detection pressure (11.4.226: with no execution seam is not coverage registration is not coverage). **Acceptan criteria:** Each guard is invoked by a na seam, OR carries a registered deferral p this item. For unattributed-commit-guar specifically, the operator has answered adoption question below and the answe recorded as consumer DATA – never an ratchet.
BOB-163	Bug	In progress	Medium	**OPERATOR DECISION (2026-08-26, interactive clarification): YES — 429 RESPONSIVE IF Retry-After IS WEL**

ATM ID	Type	Status	Severity	Description
				FORMED A 429 carrying a VALID Retr header counts as the sibling endpoint b RESPONSIVE. Assertion (c) of the DDoS challenge accepts 2xx OR a well-formed (positive integer seconds or valid HTTP- PARSED, not merely present/non-empty DEGRADED remains: connection failure any 5xx, and a 429 with absent or malfo Retry-After. RESIDUAL RISK, accepted and to be stated in the script source per a genuinely wedged endpoint that happ answer 429-with-Retry-After would pass Retry-After parse narrows that window, not close it. Options (b) limiter-aware di (c) exempt healthz probe were both offe NOT chosen. STILL OPEN, not covered decision: the detector counts 429 only, r though the script header says ‘429 (or equivalent)’; counting 503 would collide crash detector’s 5xx tally. Decide when lands limiters on :7185 and :7189. This recorded as consumer DATA per §11.4.3 the operator’s stated choice, not an age inference, and supersedes any prior age default on this question. Options not ch named above so a future reader does no litigate a settled call (§11.4.112(5) boun verdict discipline applied to decisions). item text follows — **Reported-Via:** §11. reporting directive `bug` on 2026-08-21T1 **Reported-By:** BOB-114 remediation, pr defect surfaced by BOB-111 landing a r **What (the report, verbatim):** PRE-EX NOT INTRODUCED – and proven so rat asserted: git diff shows assertion (c) is I IDENTICAL to HEAD, and the failure re against the unmodified HEAD script. W changed is not the challenge but the SY BOB-111 appears at least partially deliv because :7187 now returns real 429s wh previously returned none. :7185 and :71 produce zero 429s. This is a 11.4.248-cl corrosion risk and that is why it is filed than left as a footnote: run_all_challeng now fail INTERMITTENTLY, and an inte red trains everyone to re-run until gree which point a real regression is dismiss ‘probably the flaky rate-limit one’. THE DECISION, which is an operator call un 11.4.66 and which I deliberately did NO

ATM ID	Type	Status	Severity	Description
				Does a 429 from a WORKING limiter co sibling endpoint being RESPONSIVE? (429 proves the endpoint is alive and co protecting itself. Assertion (c) accepts 2 429, and only a connection failure / 5xx counts as degraded. Risk: a genuinely w endpoint that happens to answer 429 w (b) NO – keep requiring 2xx, but make t challenge limiter-aware: drain or wait o window before the sibling probe (retry- served, so the budget is knowable rathe guessed). (c) Probe siblings on a path th exempts (a healthz-class route), so the i question is asked without spending the Recommendation with reasoning, not ju (a) composed with a bounded liveness c 429 carrying a well-formed Retry-After evidence of a live, correctly-behaving se (b) makes the challenge slower and cou window value that will drift. But the se ‘responsive’ here are a product judgem answer is recorded, not assumed. RELA stated as fact rather than folded in: the counts 429 only, not 503, though the sc header says ‘429 (or equivalent)’. Coun would collide with the crash detector’s Worth deciding when BOB-111 lands li :7185 and :7189. **Affected scope / file- manifest:** challenges/scripts/ ddos_resilience_challenge.sh assertion endpoint isolation; run_all_challenges.s invokes it **Reproduction / context:** As requires a sibling-endpoint probe to an 2xx). :7187 now enforces a real limiter live: x-ratelimit-limit: 120, x-ratelimit-re 86, retry-after: 45, server: uvicorn). A f challenge run sends roughly 250 reques :7187, so sibling probes fired during the endpoints’ tiers land on an exhausted b receive 429 – which assertion (c) scores ‘endpoint degraded’. Timing-dependent measured PASS=5 FAIL=2; the UNMOD HEAD script against the same live stac later measured PASS=7 FAIL=0 SKIP= **Acceptance criteria:** The operator ha the classification question below, the an recorded as consumer DATA, and asser implements it. The challenge then prod same verdict across 3 consecutive runs rate-limited stack (11.4.50 deterministi

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	In progress	Medium	consistency), and the fix ships a paired mutation proving assertion (c) still catc genuinely degraded sibling endpoint – r it must not blind it (11.4.201(1)). **Reported-Via:** §11.4.202 reporting dire on 2026-08-21T19:56:53Z **Reported-By** BOB-110 UX-class coverage, discovered new axe-core suite on its first live run **V report, verbatim):** This is a REAL user defect, not a test-tuning artifact, and it by the automated regime rather than by squinting at the page – which is exactly 11.4.238 posture the project is aiming f static-grep oracle would NOT have foun Measured: the served root ships an emp hydration, so any check reading the raw audits a page nobody sees. The violation exists in the hydrated DOM, which is wh 11.4.170 requires a rendered oracle and value-equality assertions as the proof a correct. Severity reasoning, stated rath assumed: this is user-visible and affects primary dashboard heading, but it degr legibility rather than breaking function, surface is operator-facing rather than p Medium, not High. The failing test was FAILING on purpose (11.4.238) instead silenced or marked xfail. tests/ux/ curre reports 1 failed, 16 passed; that 1 is thi Anyone reading a red UX suite should r this item, not as flakiness – and when th the suite goes fully green, which is the s it is closed. HONEST SCOPE LIMIT (11 the dashboard landing view was scanne jackett/* sub-routes and the ng-serve-ho :4200 route set were NOT audited – the commands to close both are recorded in testing/ux_accessibility.md. So this item nodes are a floor, not a total. **Affected s file-scope manifest:** the Angular dash served at http://localhost:7187/ (same c SPA as frontend/); production componen not test files **Reproduction / context:** ux/test_live_dashboard_accessibility.py the running merge service. axe-core v4. scanning a real Playwright-rendered JS- DOM, reports color-contrast violations o nodes. Measured pairs: .brand / h1 text on background #3c3f41 = 1.43-1.62:1 (large-text floor is 3:1); tagline #808080

ATM ID	Type	Status	Severity	Description
				#3c3f41 = 2.68:1 (body-text floor is 4.5 **Acceptance criteria:** axe-core reports color-contrast violations against the live dashboard, with the fix made in product component CSS rather than by relaxing assertion or excluding the rule (11.4.12 reconcile to the correct mechanism, nev weaken the check). The existing tests/u the guard and already fails today, so the captured – closure requires it flipping G against the live surface, which is runtim evidence per 11.4.226.
BOB-165	Bug	Queued	High	**Reported-Via:** §11.4.202 reporting dire on 2026-08-21T20:00:08Z **Reported-By** surfaced while capturing closure eviden BOB-092; diagnosed rather than worke **What (the report, verbatim):** This is a ABI-AFTER-INTERPRETER-UPGRADE a missing package. The package is insta compiled half is built for the previous in That distinction matters because 'pip in py' style advice can appear to succeed v leaving the 313 artefact in place. WHY NOT NOTICED EARLIER, stated as fact paths that DO work all avoid system pyt ci.sh selects an interpreter through _select_python; the long-running suites session ran .venv/bin/python -m pytest a passed. So the automated regime is gre the DOCUMENTED operator command – an 11.4.238-shaped gap, since the esc found by an agent capturing evidence ra by the regime. WORKAROUND (measur theorised): PYTEST_DISABLE_PLUGIN_AUTOLOA the needed plugins passed explicitly (-p pytest_timeout ...) avoids the schemath autoload that drags in jsonschema -> re -> rpds. That is a workaround, NOT the silences the import path rather than rep install, and it will surprise the next pers runs the documented command verbati RELATED, and worth deciding together than twice: BOB-154 wants .venv rebuil to match the container (which runs CPy 3.12.13 while both system and venv run There are therefore THREE interpreters Whoever rebuilds the venv should confi resolves for the TARGET interpreter aft or this same class reappears one directo

ATM ID	Type	Status	Severity	Description
				HONEST BOUNDARY (11.4.6): this iten NOT claim any test is wrong or any proc is broken. The product and the venv-rur are unaffected. What is broken is the dc entry point. **Affected scope / file-scop manifest:** host user site-packages (~/.l python3/site-packages/rpds/); every CL documented 'python3 -m pytest …' invo any tooling that uses system python3 ra than .venv/bin/python **Reproduction /** python3 -m pytest tests/e2e/ test_live_stack_evidence.py -q –import-mode=importlib -> ModuleNotFoundEr module named 'rpds.rpds', raised durin collection via the schemathesis plugin a -> jsonschema -> referencing._core -> MEASURED root cause: system python3 but ~/.local/lib/python3/site-packages/r rpds.cpython-313-x86_64-linux-gnu.so – extension built for 3.13. rpds/**init**.py do 'from .rpds import *', and a 313-tagged importable by 3.14, so the submodule g does not exist for this interpreter. The v CORRECT and unaffected: .venv/lib64/p site-packages/rpds/ ships rpds.cpython-x86_64-linux-gnu.so and'.venv/bin/pytho import rpds' succeeds. **Acceptance cri** python3 -m pytest tests/unit/ -v –import-mode=importlib – the command CLAUD documents – reaches collection without ModuleNotFoundError, OR CLAUDE.mc corrected to document the supported ru explicitly. Either way the documented c and the working command agree (11.4. guide that misleads is the documentatio equivalent of a PASS-bluff). WORKAROU EFFECTS MEASURED 2026-08-21 — th is not free, and the item previously impl PYTEST_DISABLE_PLUGIN_AUTOLOAI disables ALL plugin autoload, not just tl one, so anything the suite silently relied be re-enabled by hand: - It BREAKS 5 p env-dependent tests in tests/unit/ test_plugin_rutracker.py — TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace need an autoloaded env plugin. PROVE

ATM ID	Type	Status	Severity	Description
				caused by any in-flight change: running OWN copy of that file under identical fla produces the same 5 failures. - It makes –timeout=60 (set in pyproject.toml) an UNKNOWN ARGUMENT unless -p pyte is passed explicitly, so pytest.mark.time silently inert under the naive recipe — a believed to be time-bounded is not. - Wi autoload ON and only schemathesis disa no:schemathesis), that same file is 96/9 CONSEQUENCE FOR ANYONE USING WORKAROUND: `-p no:schemathesis` alon strictly better than blanket autoload-dis where it suffices, because it removes on broken plugin. Where the blanket form -p pytest_timeout must be added or time inert, and the 5 env-dependent failures recognised as workaround artifacts rath filed as defects. That last point is the re the workaround manufactures failures t exactly like product defects. This does n the root cause (the venv's rpds ships a cpython-313 ABI tag CPython 3.14 cann or the fix (rebuild the venv — BOB-154) documents that the interim recipe has a radius, so nobody reads a workaround a a regression.
BOB-167	Bug	Queued	Medium	WHAT. download-proxy exposes two Ser Events routes. They are the same expen — long-lived connections that hold a wo generator for their lifetime — but only o limit classed: routes.py:801 @router.get stream/{search_id}') @_rl('sse_stream') classed routes.py:150 @router.get('/the stream') async def stream_theme(...) <- limiter decorator _rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/st bound to the sse_stream class while /the stream falls through to the application d Measured live on the running stack: /ap theme/stream reports x-ratelimit-limit 1 PROVENANCE + A CORRECTION WOR KEEPING (§11.4.6). This was surfaced b BOB-109 scaling agent, whose report fr as: 'sse_stream_limit_decorator is define imported/applied nowhere ... SSE falls t the 120/minute default: 24x less protect framing is WRONG and was NOT filed a The agent searched for one symbol NAM wiring uses a different mechanism

ATM ID	Type	Status	Severity	Description
				(@_rl('sse_stream')), and it IS applied — search/stream. Verified by reading routes.py:795-806 and the _rl helper at routes.py:46-51, with a control needle c other *_limit_decorator symbols show re in api/**init**.py so the zero-hit was not a b search. The agent's MEASUREMENT (1 theme/stream) was correct and is what this a real finding; its MECHANISM was halves are recorded so the next reader re-derive the same wrong cause. WHY I MATTERS. An unclassed SSE endpoint i cheapest way to pin server resources: e connection is held open, and the default permits 120/min of them. The declared sse_stream class exists precisely becaus route shape needs a tighter bound than GETs. ACCEPTANCE. (a) A decision, rec whether /theme/stream belongs in the s class or genuinely warrants the default policy question, not automatically a bug (b) If it belongs in sse_stream, the deco applied and a test drives BOTH SSE rou asserts each returns its INTENDED clas from x-ratelimit-limit, so a future route without a class is caught. (c) A guard th enumerates SSE-shaped routes and fail that carries no explicit rate-limit class — general form, so the third SSE route do repeat this. (d) Honest boundary: this d claim 120/min is exploitable in this depl with network_mode host and no reverse every caller shares the 127.0.0.1 bucket BOB-111 measured and recorded separ CLAIMED. No change made. The limit v were read from headers, never driven to exhaustion — the limiter is per-IP and s with concurrent agents on this host (§11
BOB-168	Task	In progress	Low	WHAT. scripts/run_all_challenges.sh:66 "scaling_horizontal_challenge.sh" in its set. challenges/scripts/ scaling_horizontal_challenge.sh does no sed -n '66p' scripts/run_all_challenges.s "scaling_horizontal_challenge.sh" $ ls c scripts/scaling_horizontal_challenge.sh access …: No such file or directory VER independently, not taken on report — su the BOB-109 scaling agent and re-check by direct invocation. WHY IT MATTERS this is cosmetic or a §11.4.201 gate-hon

ATM ID	Type	Status	Severity	Description
				defect depends entirely on how the runn a missing entry, and that is the first thin determine: if it SKIPs silently, the challe advertises coverage it does not have (a claim-vs-reality row with no passing cha behind it); if it FAILs, the runner is perr red for a reason unrelated to the system test, which trains readers to ignore it. N outcome is acceptable; they need differe ACCEPTANCE. (a) Determine and recor runner's actual behaviour on the missin invoking it, not by reading it. (b) EITHE the challenge, OR remove the entry — w §11.4.124 discipline: check git history fo whether it once existed and was deleted silently-dropped challenge is the more i defect. (c) If the runner silently skips m entries, that is its own finding: a missin challenge must be loud (§11.4.3 SKIP-w at minimum), never absent-and-quiet. S Low as a defect, but it sits on the challe coverage seam, so (c) may deserve its o CORRECTION 2026-08-22 (§11.4.6) — T ITEM'S PREMISE WAS FALSE, and the mine. Recorded openly rather than quie scaling_horizontal_challenge.sh EXISTS submodules/challenges/challenges/scrip scaling_horizontal_challenge.sh, execut bytes, dated Aug 7. So do all 14 other e the meta-runner. It was never deleted: a the roster at 1c959ba, the challenge its in the submodule at 873c0b1, and a pick search across both repositories and all l finds ZERO deletions. HOW I GOT IT W run_all_challenges.sh reads submodules challenges/challenges/scripts/. I checke challenges/scripts/ — a DIFFERENT dir belonging to a DIFFERENT, glob-based Two similarly-named directories; I meas wrong one. That is a §11.4.201(9) field-i error. Worse, I wrote that it was "verifie invocation": I had listed a directory and the aggregator, so the claim overstated evidence class as well as getting the an wrong. THE LESSON, which generalise item. In the SAME commit I correctly ca sibling agent making this exact class of BOB-167 — a zero-hit produced by sear symbol name — and applied a control n there. Then I omitted the needle one pa

ATM ID	Type	Status	Severity	Description
				later on my own measurement. The dist that would have caught it: A CONTROL PROVES THE INSTRUMENT CAN SEE; NOT PROVE IT IS POINTED AT THE TH UNDER TEST. §11.4.201(7)(b) as comm applied guards blindness; it does not gu aiming. The needle for THIS query wou been to resolve the path the runner itse not to confirm that some directory listin WHAT IS ACTUALLY WRONG, and it is than what I filed. The missing-entry pat challenges by definition, so it can be exe safely. Run against a checkout where th submodule is absent — the ordinary sta clone made without –recursive — the R runner (byte-identity sha-asserted, not a reports: PASS: 0 FAIL: 0 SKIP: 16 TOTA OBSERVED EXIT CODE: 0 An ENTIREL RUN BANK REPORTS SUCCESS. Not o dangling entry tolerated — the whole su silently passes without executing anythi is §11.4.201(6) false-null verbatim, and observed-reachable rather than constru (§11.4.115(G)). REVISED ACCEPTANCE ANSWERED, though not as filed: the ru reports a loud, counted SKIP and then r blocks, because the exit expression rea FAIL count. (b) The dangling-entry half nothing is dangling. (c) The real work is contract: a roster entry the runner cann execute must not be indistinguishable f healthy run. See the decision recorded BOB-168/decision_missing_vs_skip.md. REPORTED NOT FIXED: scripts/ pre_build_verification.sh:1792 carries th SKIP/MISSING conflation one level up, other runner. Filed separately.
BOB-169	Bug	In progress	Medium	WHAT. The §11.4.65 markdown-export r requires every in-scope doc to ship .htm twins. 286 of the 326 .html files under d (excluding dist/ and node_modules/) are FRAGMENTS — they begin at
BOB-170	Task	Queued	Medium	WHAT. The BOB-109 scaling suite gates timing tests on loadavg_1m/nproc <= 0 this 8-cpu host, a 1-minute load average below 6.0 — and SKIPs loudly otherwise gate is correct and was adopted for a go (§11.4.201(8)): the author validated a 2. growth-exponent threshold, then observ FALSE-FAIL correct code at load 11.06

ATM ID	Type	Status	Severity	Description
				span 2.136), and at load 15-23 measure spans of 2.358 against injected-cubic m 2.278-2.331 — the signal is smaller thar and the two are not separable in either Gating was the honest response. THE P The gate has consequently NEVER RUN own precondition. Every recorded run t measured load 9.15-23; the independen measured 9.66 during review; the cond measured 6.45 at its lowest. GREEN po observed — at loads 9.4-11.6 (span 1.88 three stability runs (1.759/1.832/1.848) never once beneath the shipped ceiling. author states this honestly as an inferer data rather than a run they can paste, a reviewer confirmed the inference is sou (quiescence is strictly more favourable) inference is still not an observed run. W MATTERS (§11.4.226). A registered, top present guard that never executes is indistinguishable from a guard that wou did. Nothing currently ensures it ever r freshness contract, no scheduled quiet- invocation, no tracked item — which is the unexecuted-standing-guard class §1 names, where registration gets treated coverage while execution is nobody's jo forensic precedent in that anchor is two guards that, when finally executed, imm emitted latent FAILs. ACCEPTANCE. (a) both quiescence-gated timing tests on a whose 1-minute loadavg is genuinely <= — no agent fleet, no concurrent build — capture the run as an artifact showing t resolved load reading alongside the mea span exponent, so the precondition is pr from the artifact and not asserted. (b) R observed GREEN polarity in the item ar test source, replacing the current hones inference note. (c) If the gate FAILS unc genuine quiescence, that is the finding exists to surface, and it reopens the thre calibration question rather than being w around. (d) Establish a freshness contra §11.4.226: how stale a quiescent verdic before the gate counts as uncovered ag this does not silently return to never-ex HONEST BOUNDARY. This item does n the gate is wrong. The evidence points t way — quiescence is strictly more favou

ATM ID	Type	Status	Severity	Description
				than the loads where GREEN was alrea observed. It claims only that the run ha happened, and that an unexecuted guar be cited as coverage. PROVENANCE. R IMPORTANT-2 by the independent Fabl substrate reviewer of the BOB-109 wor independently corroborated: the review the SKIP is loud and prints its resolved and confirmed the shipped tests do SKI today's load.
BOB-171	Bug	Queued	Low	WHAT. Both rate limiters key their per-c bucket on the LEFTMOST element of X- Forwarded-For when TRUST_FORWARI is enabled: download-proxy/src/api/ rate_limit.py:117-123 (:7187) and the n limiter in plugins/download_proxy.py (:7 which was deliberately built to exact po with it. That element is CLIENT-SUPPL honest reverse proxy APPENDS the pee to the RIGHT of whatever arrived, so th entry is whatever the client sent — atta controlled by construction, not by misconfiguration. Consequences with th (1) rotating a forged leftmost value min unlimited sequence of fresh buckets, wl total bypass of the limit the flag is supp make MORE accurate; (2) it also defeat idle-reap cap, since each forged identity distinct key — the bucket map is a mem growth surface bounded only by the cap cap's eviction then discards the budgets clients to make room for forged ones. N EXPLOITABLE AS DEPLOYED TODAY, a why this is Low, not High. TRUST_FORWARDED_FOR defaults OF sites, and the deployed stack runs netw host with no reverse proxy in front (veri across all five compose services), so pee addresses are real client IPs and no NA occurs. The defect is latent: it arms the someone puts a proxy in front and turns on to recover real client IPs — which is the situation the flag exists for. The failu is therefore 'correct-looking configurati silently disables the limiter', not 'curren broken'. PROVENANCE. Raised as MIN the independent reviewer of the BOB-1 limiter work and independently confirm implementing agent, which noted it at E source sites as a tracked follow-up rathe

ATM ID	Type	Status	Severity	Description
				fixing it in that change. Filed here beca neither agent has write access to the tr WHY IT COVERS BOTH PORTS. :7186's was built to deliberate policy parity wit including this parsing. Fixing one and n other would leave the two surfaces disa on client identity — worse than the sha because it becomes surface-dependent untestable as a single invariant. ACCEP (a) Replace leftmost-element parsing wi aware resolution at BOTH sites: either minus-N-trusted-hops, or a trusted-prox allowlist where only a peer inside the a may assert XFF at all — the choice is a deployment-topology decision and shou recorded, not assumed. (b) A test that f rotating leftmost element and asserts t key does NOT rotate, per site. (c) Paire mutation: restore leftmost parsing and must FAIL. (d) Negative control (§11.4. with TRUST_FORWARDED_FOR OFF, t must be ignored entirely and keying mu back to the peer address — a guard tha honouring XFF when the flag is off wou worse defect than the one being fixed. (boundary: this closes header-forgery by does not make per-IP limiting fair behin where many real users legitimately sha address. NOT CLAIMED. No change ma sites still parse leftmost; the flag still d
BOB-172	Bug	Ready for testing	High	WHAT. The rutracker SEARCH endpoin refused by bot protection while the site fully reachable. Measured unauthentica 2026-08-21 (no cookie file read, no cre value in any artifact — §11.4.10): GET / index.php -> 200, 96283 B,
BOB-173	Bug	Ready for testing	High	WHAT. download-proxy/src/api/hooks.p def _save_hooks(hooks: list[dict[str, An None: try: os.makedirs(os.path.dirname(HOOKS_ exist_ok=True) with open(HOOKS_FIL json.dump(hooks, f, indent=2) except E as e: logger.error(f'Failed to save hooks catches EVERY exception, logs, and ret None. The signature returns None, so t has no channel to learn the write failed sites then report success unconditional _save_hooks(hooks) :174 _save_hooks(h logger.info('Created hook: ...') :175 logger.info('Deleted hook: ...') :156 ret

ATM ID	Type	Status	Severity	Description
				HookResponse(hook_id=..., ...) :176 ret {'message': 'Hook deleted', ...} USER-V CONSEQUENCE, which is why this is H not a code-hygiene nit. A user POSTs a receives HTTP 200 and a hook_id, and t was never written — their automation s never fires, and the API told them it exi Symmetrically, a user DELETEs a hook, 'Hook deleted', the file still holds it, and again after the next restart. In both dire product reports the opposite of what ha and the only trace is a log line nobody i watching. THIS IS THE §11.4.252 SHA path combines two dangerous capabiliti that anchor's taxonomy — MUTATION o resource (a filesystem write) and EXTE SIDE EFFECT (hooks are outbound call system will or will not make) — so it is to FAIL CLOSED: verify the preconditio when it cannot be satisfied, and surface refusal. Instead it fails open, and a bare Exception:' that only logs is the exact a §11.4.252 enumerates. At the product l also a §11.4.201(6) false-null: a success and a swallowed failure are indistinguis the caller. PROVENANCE. Surfaced as a scope observation by the independent r the BOB-135 test-isolation work — this what turned that defect into an 'assert mystery, because the EACCES on /confi logged and discarded while the endpoin returning 200. Verified here directly fro before filing (the function body and bot read above), not taken on report. Recor §11.4.238 discovery-channel escape: fou agent reading code during an unrelated investigation, not by the automated QA the coverage gap is a defect of equal sta the defect itself. ACCEPTANCE. (a) _sav propagates failure — raise, or return a s callers must consume. (b) Both endpoin translate a persistence failure into an H (500-class), never a success body; a cre did not persist must not return a hook_i test drives each endpoint with the hook unwritable (read-only dir or a patched o raising OSError) and asserts a non-2xx AND that a subsequent GET does not lis phantom hook — assert on the user-obs outcome, not on the log line. (d) Paired

ATM ID	Type	Status	Severity	Description
				mutation: restore the swallow; the test FAIL. (e) Audit the same file for sibling — this is a pattern, and one instance is alone. (f) Honest boundary: this does no the write currently fails in production; i that WHEN it fails the user is told the o and the BOB-135 investigation shows it in at least one real environment. NOT C No change made. No assessment of how write fails in the operator's deployment RECORDING A SHELL ERROR OF MY (§11.4.6): the first version of this descri written with the anti-pattern snippet ins backticks in a double-quoted shell argu the shell ran it as command substitution text was replaced by nothing — the stor description read 'and is the exact anti-p This is the SECOND time this session th backticks-inside-double-quotes has corr content (the first mangled a commit me Fixed here by editing through a Python with no shell quoting in the path. Noted a silently-truncated defect description is the kind of quiet corruption §11.4.201(7 about — the path is part of the instrume failed without erroring.
BOB-174	Bug	In progress	High	WHAT. A corrupt hooks file is silently indistinguishable from "no hooks config and the next create then DESTROYS ev existing hook while returning HTTP 200 defects on one path, filed together beca any one alone leaves the data loss reach — download-proxy/src/api/hooks.py:104 _load_hooks wraps the read in `except Ex` `logger.error(...)` and falls through to `re` truncated or malformed hooks.json is th reported to every caller as an empty, he list. Confirmed at source. A2 — the cons and the reason this is High. create_hook appends, saves. Given a corrupt file tha turns into [], the save writes a one-elem over the top. Every previously-configure gone. Measured by the implementing ag against the ALREADY-FIXED tree, so th the BOB-173 write-failure fix: BEFORE ['prod-hook-0', 'prod-hook-1', 'prod-hook reports : 200 {'hooks': [], 'count': 0} <- ZERO hooks configured DELETE report {'detail': 'Hook not found'} <- for a hoo in the file POST reports : 200

ATM ID	Type	Status	Severity	Description
				hook_id=3c8a5930-… AFTER on disk : [‘3c8a5930-…’] Three prod hooks destro HTTP 200 throughout, nothing surfaced user. A5 — _save_hooks writes with a pl `open(path, "w")`. A crash, ENOSPC, or a write truncates the file in place. That is the corruption A1 then reads as “no hoo A2 overwrites. The same codebase alrea the correct pattern: theme_state.py:115 tmp-file + os.replace. WHY THE THREE ONE ITEM. A5 manufactures the corru misreads it as empty, A2 destroys the c Fixing only A1 leaves truncation reacha only A5 leaves an existing corrupt file a trigger; fixing only A2 leaves the API ly what is configured. The chain is the def DELIBERATELY NOT FIXED UNDER B and the reasoning is sound. The implem agent identified all three while auditing swallows as that item required, and dec fix them in the same change because: th change GET semantics on an endpoint t frontend consumes (200 -> 5xx); they n CORRUPT-file reproduction rather than UNWRITABLE-file one BOB-173 built; a need a design decision that is genuinely obvious — a MISSING hooks file legitim means “no hooks configured”, while a C one does not, and today’s code cannot t apart. Expanding BOB-173’s scope to co would have meant shipping that decisio unexamined. ACCEPTANCE. (a) Disting MISSING from CORRUPT: a missing file an empty list; a corrupt one is an error, silently []. (b) Decide and RECORD wha does on corruption — 5xx, or 200 with a degraded marker the frontend can rend the design decision, and it should be sta rather than inferred from whatever the happens to do. (c) create/delete MUST overwrite a file they could not parse — not clobber (§11.4.252: mutation plus e side effect must fail closed). (d) Make th atomic via tmp + os.replace, reusing theme_state.py’s existing pattern rather inventing it (§11.4.28). (e) Tests asserti USER-OBSERVABLE outcome, not log l a corrupt file, assert GET does not clai hooks, assert a create refuses rather th destroying, and assert the file still holds

ATM ID	Type	Status	Severity	Description
				original hooks afterwards. (f) Paired §1. mutation per guard. (g) NEGATIVE CON (§11.4.201(1)): a MISSING file must stil empty list and a working create; a VALI must behave exactly as today. A fix that every load fail closed by failing always i A3, RECORDED SEPARATELY, NOT PAR THIS CHAIN. VALID_EVENTS and HookEventType are two sources of trut closed set. Verified identical today, ungu against drift: if they diverge a hook regi successfully and then silently never fire assertion pinning them would close it. N CLAIMED. No change made. The BOB-1 failure fix is real and orthogonal — it ma FAILED write honest; it does nothing ab SUCCESSFUL write of wrong data deri misread file.
BOB-175	Task	Queued	Low	WHAT. The workable-items `update` seam the status-location invariant in one dire After BOB-166, `update --status <termina` Issues-located row is refused. The mirro open: `update --location Fixed --status` `terminal>` exits 0 and creates a row that own validator immediately rejects. Verif empirically by the independent reviewe BOB-166: update –location Fixed –statu progress’ -> exit 0 validate -> FAILS, na row (check (f), fixedLocationNonTermin So the command can still mint a state it validate call refuses. That is the §11.4.1 shape at the seam layer: the invariant is CONFIGURED in validate but not ENFO every write path that can violate it. WH DELIBERATELY LEFT OPEN IN BOB-16 why that was right. Three reasons, all c rather than asserted: (1) the real corpus ZERO instances of this direction agains the direction BOB-166 closed — the asy the data matched the asymmetry in the detective coverage ALREADY existed an demonstrably fires, so unlike BOB-166’s this cannot rot silently; (3) and the load one — `reopenCmd` documents and SANCT transient Fixed-plus-non-terminal state `location Fixed` operator override. A naiv preventive guard at the update seam w collide with that override’s semantics. C this therefore requires a design decisio how the two interact, which is a differe

ATM ID	Type	Status	Severity	Description
				question from the one BOB-166 was sco answer, and expanding that item would meant shipping the decision unexamine ACCEPTANCE. (a) Decide and RECORD preventive guard on this direction coexi reopenCmd's sanctioned override — thi actual work, and the answer should be down rather than inferred from whateve patch does. Candidates worth weighing the reopen path explicitly, require an ov flag on `update` mirroring reopen's, or lea direction detective-only with the rationa recorded so the asymmetry is a decision than an oversight. (b) If a guard lands, `terminalStatuses()` — BOB-166 establish predicate source specifically so the two cannot drift. (c) Paired §1.1 mutation. (NEGATIVE CONTROL (§11.4.201(1)), a sharp one here: the sanctioned reopen MUST still work. A guard that breaks reopenCmd's documented override is a positive refusal, and worse than the gap because it breaks a workflow operators to use. HONEST BOUNDARY. This is no data defect. Zero corpus instances, and detective gate catches it at the next val filed so a scoped, deliberate omission st as a tracked decision instead of living o code comment where the next reader w it for an oversight (§11.4.197: a started reaches a documented terminal state, n abandonment). PROVENANCE. Raised implementing agent of BOB-166 as a kn asymmetry it deliberately did not close, independently verified and endorsed as acceptable scope boundary by that item reviewer, on the condition that it becom tracked item rather than a comment. Th item.
BOB-176	Bug	Queued	Medium	WHAT. _get_enabled_trackers() in down proxy/src/merge_service/search.py gate rutracker on RUTRACKER_USERNAME RUTRACKER_PASSWORD only. It never RUTRACKER_COOKIES. So an operator configured with cookies alone — no use password — never has rutracker enable and the tracker is silently absent from e out rather than failing visibly. WHY THA INCOHERENT WITH THE REST OF TH _search_rutracker treats COOKIES as th

ATM ID	Type	Status	Severity	Description
				PREFERRED auth path, not a fallback. nnmclub sibling DOES check its own *_ variable. So the same codebase holds th positions at once: cookies preferred at t site, cookies ignored at the enablement cookies honoured for a sibling tracker. MATTERS MORE THAN IT LOOKS. CLA documents a standing operator mandate (2026-08-15) that per-tracker Netscape files at ${TRACKER_COOKIE_DIR}/ cookies_<tracker>.txt are auto-loaded i as <TRACKER>_COOKIES before every up, restart, install Stage 6 and start.sh and rutracker is named in that set. So th SANCTIONED, documented configurati for this tracker produces exactly the sta gate ignores. An operator who follows t documented instructions gets a tracker never runs, with no error to explain it. MODE. Silent absence, not a visible fail same §11.4.201(6) false-null shape as B one layer earlier. BOB-172 fixes a track RAN and was REFUSED being reported this is a tracker that never ran at all, an merged result simply has one fewer co with nothing to indicate it. PROVENAN by the BOB-172 implementing agent wh tracing the enablement path to locate th handling defect. Recorded here rather t in that change: it is a separate defect o separate seam with its own oracle, and in would have widened BOB-172 past it evidence. ACCEPTANCE. (a) The enable gate honours RUTRACKER_COOKIES a independently sufficient credential, ma what _search_rutracker already prefers the nnmclub sibling already does. (b) A asserting that a cookies-only configurat ENABLES rutracker, driving the real _get_enabled_trackers rather than a rep Paired §1.1 mutation restoring the user password-only gate — the test must go NEGATIVE CONTROL (§11.4.201(1)): a configuration with NO credentials at all leave rutracker DISABLED. A fix that en unauthenticated tracker is worse than t because it produces failing searches ins absent ones. (e) Audit the other tracker for the same asymmetry rather than fix the one that was noticed — three positi

ATM ID	Type	Status	Severity	Description
				codebase suggests nobody has checked together (§11.4.118). HONEST BOUND verified against a live cookies-only deplo read from source by the BOB-172 agent recorded on its report. The reading is sp checkable, but whoever takes this shou by invocation before relying on it. | | BO Bug | Queued | High | WHAT: BOB-172 identical 4-line HTTP-refusal guard at fi private-tracker fetch sites in download- merge_service/search.py (rutracker coo rutracker credential :1468, kinozal :164 nnmclub :1761, iptorrents :1934). Only rutracker COOKIE path is exercised by EVIDENCE (reviewer-authored mutatio 11.4.194(6)(d), during the BOB-172 ind review): deleting ONLY the kinozal guar (search.py:1655-1658) while leaving the intact left the full merge_service suite a passed, ZERO failures. The suite canno of the five sites. FAILING SCENARIO: a drops or subtly breaks the wiring at kin nnmclub, iptorrents, or rutracker-crede 403 at that site silently folds back to status=empty with error=None -- the ex BOB-172 signature -- and no test redde DIRECTION (reviewer preferred): colla five duplicated wirings into one shared e.g. _check_search_response(tracker_na status, body) -> bool, so there is ONE c and a sixth site cannot be added ungua (11.4.251 byte-identical-fork extraction) Alternative: parametrise the guard tests all five sites with per-site stub sessions. ACCEPTANCE: mutating the wiring at A five sites reddens at least one test. Prov running the same R1 deletion at each si | | BOB-178 | Bug | Queued | Medium | download-proxy/src/merge_service/ search.py:1638-1640 handles a failed ki LOGIN response with if login_resp.statu (200, 301, 302): logger.error(...); return sets NO diagnostic. FAILING SCENARIO Cloudflare returns 403 on takelogin.php kinozal chip reads status=empty, error= That is the BOB-172 signature exactly: reported to the user as an empty result CONTRAST establishing this is an overs a design choice: the rutracker and nnm failures DO set diagnostics (upstream_

<table>
<tr><th>ATM ID</th><th>Type</th><th>Status</th><th>Severity</th><th>Description</th></tr>
<tr><td></td><td></td><td></td><td></td><td>auth_failure), and the iptorrents login fa
through to the search fetch where BOB-
guard catches it. Kinozal is the one leg
neither. FIX DIRECTION: stash
_classify_upstream_http_status(login_re
"") before the early return, or an auth_fa
for the status-in-(200,301,302)-but-no-c
case. ACCEPTANCE: a stubbed 403 on t
kinozal login endpoint produces a non-N
on the kinozal chip, with a RED capture
the pre-fix code first (11.4.115). | | BOB
Task | Queued | Medium | WHAT: the BO
independent review demonstrated two f
paths that still report a refusing or unre
tracker as an empty result. Both are PR
EXISTING and neither is a regression fr
BOB-172, but the fix evidence log prese
five-site coverage without stating the bo
(11.4.194(5) requires un-analysed dime
explicit gaps, never silently assumed sa
-- exception before resp.status is read. I
probe B, captured: a connection-refused
rutracker yields status=empty, error=N
http_status=None, metadata.errors=[],
metadata.status=completed. Each _sear
method swallows exceptions (except Ex
logger.error; return results) before _sea
error handling can see them, so a track
DOWN is indistinguishable from one tha
genuinely empty. GAP B -- soft refusal a
200. Reviewer probe A, captured:
_classify_upstream_http_status(200, <c
body>) returns None, which is CORREC
design (a 2xx is a usable response; body
triggering would risk the over-fire the n
controls exist to prevent). But all five se
use aiohttp's default allow_redirects=Tr
session-expiry 302 -> login-page-200 ch
to zero rows and reports empty. FIX DIR
Gap A needs the exception to reach a di
rather than being swallowed. Gap B nee
DISTINCT detector (final-URL check or
form marker) with its own RED -- explic
widening of the status-code trigger.
ACCEPTANCE: both gaps closed with th
REDs, or explicitly closed per 11.4.112
evidence. Immediate sub-task: append a
gaps paragraph to docs/qa/BOB-172/
fix_evidence_20260822.log. | | BOB-180
Queued | Low | WHAT: download-proxy/</td></tr>
</table>

ATM ID	Type	Status	Severity	Description
				routes.py:727-745 sets status=captcha_ with a hardcoded message naming RuT 'RuTracker requires CAPTCHA. Use /ap rutracker/captcha'. FAILING SCENARIO search returns zero results and the capt flavoured diagnostic came from NNMCl Turnstile, not RuTracker. The user is tol a RuTracker captcha endpoint for an NI problem. SEVERITY BOUNDED: the ful and tracker_stats travel in the same res payload, so the truth is present and a cl reads them is not misled -- only the head wrong. The branch was revived by BOB error propagation (correctly: suppressin propagation would recreate the same fa one layer up, 11.4.247) and is already c the routes layer by tests/unit/api_layer/ test_routes_coverage.py:782. FIX DIREC derive the message from the tracker(s) actually erred rather than hardcoding o SCOPE NOTE: api/ was owned by a sibl stream during BOB-172; the author was not to touch it. ACCEPTANCE: an NNM captcha diagnostic on a zero-result sear produces a headline naming NNMClub. BOB-181 | Bug | Queued | High | WHAT generate_markdown_exports.sh accepts arguments. It unconditionally scans PROJECT_ROOT/*.md, docs/ recursively scripts/ recursively (scan loops at lines 105/111/116 as of 2026-08-23; they wer when this was filed — line numbers shif the BOB-169 fix, the substance is uncha was independently verified by execution the generator with an explicit path prod out-of-scope export and NOT the reques and PROJECT_ROOT is derived from BASH_SOURCE (line 17). A caller who v 'bash scripts/generate_markdown_expor docs/qa/SOME/file.md' gets a WHOLE-T with the argument silently discarded. W MATTERS: the caller reasonably believe was scoped to one file. It was not. In a s checkout with parallel work streams thi regenerates any export whose .md is ne its derived files anywhere in the tree -- l artifacts in other streams' scopes and, v BOB-068 auto-commit hazard, into anot stream's commit. CLASS: a false afforda script's interface implies a capability it

ATM ID	Type	Status	Severity	Description
				have (11.4.201: an interface must mean claims). FIX DIRECTION: either accept path arguments and honour them (scop run to exactly those files), or REFUSE w clear message when arguments are pas Silently discarding them is the one optic must not remain. ACCEPTANCE: passin either scopes the run to it, or exits non- naming the unsupported argument. A R capturing today's silent-discard behavio (11.4.115). LIVE INSTANCE, MEASURE 2026-08-23 (this is no longer hypothetic During the BOB-169 fix the conductor in the generator with an explicit path to re ONE document. The argument was disc the whole tree was scanned. It generate qa/BOB-174/DESIGN_DECISION.{html, at 11:50 -- a SIBLING STREAM's docum that stream's author was still editing th (.md mtime 12:00). The exports were th stale on arrival, and additionally carried fix BOB-169 charset defect. The sibling independently reported them as stale in hand-off, having no idea another stream created them. SECOND-ORDER LESSO keeping with this item: the conductor's verify the blast radius ALSO failed, and silently. 'git status --porcelain | grep -E pdf)$"' returned only its own files and r clean – but docs/qa/BOB-174/ is an UNT DIRECTORY, which porcelain collapses line, so the grep was structurally incapa seeing the files inside it (11.4.201(7)(a) the wrong surface; the quiet zero read a absence). The instrument that DOES se
BOB-182	Task	In progress	Medium	**OPERATOR DECISION (2026-08-26, interactive clarification): KEEP THE MONOTONE RATCHET** §11.4.224(E) adoption answer, recorded as consumer the adoption model for CM-EXPORT-CH VALID is the MONOTONE-DECREASE The count may only go down, never up. Immediate hard floor, changed-code-onl deadline, and per-corpus phase-in were offered and NOT chosen. This closes ha BOB-182 — the half that was genuinely operator's. THE REMAINING HALF IS DEFECT STILL OWED: BASELINE is a constant, so the ratchet does not actual — on improvement the gate PASSES an

ATM ID	Type	Status	Severity	Description
				the value to lower it to, but nothing low the gate keeps permitting regression ba original count after the corpus heals. Th NOT collapse §11.4.249 role separation that writes its own threshold during a p run becomes a PRODUCER as well as a Acceptable shapes include a persisted b the gate READS but never writes with l via a separate explicitly-invoked comma gate that FAILS loudly when the live co below baseline so drift becomes an acti refusal. This answer is recorded as cons DATA per §11.4.35 — it is the operator's choice, not an agent inference, and sup any prior agent-chosen default on this Options not chosen are named above so reader does not re-litigate a settled call (§11.4.112(5) bounded-verdict disciplin to decisions). — prior item text follows CM-EXPORT-CHARSET-VALID (pre-build 50) adopts its 301 pre-existing violation monotone-decrease ratchet rather than floor. Two things are owed. (1) THE AD DECISION IS THE OPERATOR'S, NOT SCRIPT'S. 11.4.224(E) requires the bro adoption question be SURFACED to the and the answer recorded as consumer never an invented ratchet. What exists header comment plus a BOBA_EXPORT_CHARSET_BASELINE that DOCUMENTS a decision the agent does not RECORD an answer the operat The independent reviewer judged the c defensible and would not reverse it (a h on 301 violations makes the build unrea which 11.4.234 forbids; 11.4.101 favour reversible option) but correctly held tha yet compliant. Options for the operator immediate hard floor (BASELINE=0) / k monotone ratchet / per-corpus phase-in changed-code-only with a scheduled ful deadline. (2) THE RATCHET DOES NO ACTUALLY RATCHET. BASELINE is a h constant. On improvement the gate PAS PRINTS the value to lower it to, but no lowers it, so the gate keeps permitting all the way back to 301 after the corpus The header now says tightening is MAN rather than claiming automatic behavio does not exist (11.4.6), but the honest s

ATM ID	Type	Status	Severity	Description
				is a stopgap, not the fix. WHY IT WAS N WITH THE FIX: a gate that writes its ov threshold during a pre-build run becom PRODUCER as well as a GATE (11.4.24 separation), and that is a design change operator should approve rather than rec side effect. ACCEPTANCE: the operator adoption answer recorded as consumer and either a persisted baseline the gate and never raises (with the role-separati question answered), or an explicit decis manual tightening is acceptable.
BOB-184	Task	Queued	Medium	Three icon-glyph controls (.bridge-retry toggle, .caret) render text made of non- points. axe-core declines to decide their reporting them under the incomplete ch with messageKey nonBmp and resolving NEITHER fgColor NOR bgColor, so the recomputation in docs/qa/BOB-164/ axe_contrast_scan.py cannot decide eith such nodes across 16 palette x mode combinations. BOB-164 round 2 stopped being scored as silent passes by naming GLYPH_UNMEASURED, printing them run and fencing them, so any NEW unm glyph node blocks. What remains genui unproven is their contrast itself: as non- interface components they owe 3:1 und SC 1.4.11, and no oracle in this repo me that today. Acceptance: a measurement resolves the effective foreground and ba for glyph controls (for example reading getComputedStyle colour against the co backdrop, or replacing the glyphs with carrying explicit fill tokens), each of the controls shown to clear 3:1 in all 16 pal mode combinations, and the GLYPH_UNMEASURED fence shrunk by the nodes then proven.
BOB-185	Task	Queued	Medium	docs/qa/BOB-164/axe_contrast_scan.py static dist with no backend, so every no existence depends on fetched data is ab the page it measures: the results table empty, so .type-badge and .quality-badg appear, and the hooks list renders empt so .status pills never appear. That is pre why the BOB-164 round-1 regression w unseen in a rendered DOM even though scanned in all 16 themes, and it is why had to close the class in the arithmetic

ATM ID	Type	Status	Severity	Description
				instead. Those nodes are now covered b frontend/src/app/models/style-contrast.s which extracts declared foreground-on- from the stylesheets, but that is an ARIT claim about declared colour pairs, not a rendered-pixel one: it cannot see opacit by an ancestor, a composited backdrop, cascade this repo's static extractor does model. Acceptance: the scan drives the dashboard with seeded fixture data (a s backend, a route fixture, or an injected component harness) so the badge and s nodes are present in the DOM, axe mea them directly, and a control needle prov rendered oracle sees a seeded sub-floor before any clean result from it is believe
BOB-186	Bug	Queued	—	workable-items diff returned exit 0 and verdict "DB and Markdown are in sync" PARTIAL read: invoked with –issues but –fixed it compared 77 Markdown items 184 DB items and called that in sync, pr both mismatched counts in its own outp 107 DB items were never compared aga anything. Measured 2026-08-25: (A) no markdown paths -> correctly REFUSES (B) –db-only -> honest, "no Markdown compared"; (C) –issues alone -> exit 0, (compared 77 Markdown item(s) agains item(s))". The 11.4.201(6) false-null: a p blind instrument returned the same qui genuinely-synced corpus returns. IMPAC CORRECTED 2026-08-25, my original fi OVERSTATED this. I wrote that CHECK 11.4.106(F) commit seam was exposed. NOT. scripts/hooks/docs-sync-commit-seam.sh:277 passes BOTH –issues and - a repo-wide sweep found NO executable that breaks: the two real callers (that se pre_build_verification.sh:549) both pass paths; the single-path hits are prose evi files plus a static-scan test fixture. So th was real and latent, never live in this pr corrected claim is: any FUTURE single- caller would have been silently mis-told PROVENANCE, twice corrected. A suba attributed the false-null to the no-paths which actually refuses correctly; the co re-measured and relocated it to the par Then the fixing agent corrected the con impact. Both corrections are recorded s

ATM ID	Type	Status	Severity	Description
				wrong version propagates. FIXED (unco constitution submodule, fetched to upst 7f16739 before edit per 11.4.26). Desig by default (11.4.201(4)); the narrow per comparison PRESERVED behind an exp –partial-scope rather than deleted (dele existing documented capability would b 11.4.122 silent removal), mirroring how preserved the no-Markdown shape. Key MEASURED DB content, never on flag p — a DB with zero Fixed rows IS fully ac for by –issues alone and still passes, so keying would have traded the false-null 11.4.201(1) false-positive refusal. Resul property: "in sync" is reachable only wh two printed counts account for each oth fail/5 pass -> GREEN 9/9; paired 1.1 mu the fixs own revert, killing exactly the 4 guarding tests while the 5 preservation green. Case C now refuses naming "107 located in Fixed (supply –fixed)", the 10 confirmed against an independent sqlite (107 Fixed + 78 Issues = 185). Incident closed an 11.4.238 gap: case As refusal test anywhere; TestDiffCmd_NoPathsSti now guards it. OPEN: –partial-scope is t naming choice, not an operator decision renaming is cheap now and expensive a consumers adopt it. Other consumers of shared submodule were not swept.
BOB-187	Bug	In progress	—	scripts/ownership_precondition.sh cons SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe it, but it is NOT fenced. T028 remediati ownership_path_fence() to scripts/lib/ ownership.sh and wired it into ownership_repair.sh, closing the path-es class there; the precondition was outsid agents file scope and still accepts whate config/owned_paths.yaml expands to. B is lower than a recursive chown (it write files rather than mutating ownership of existing tree) but the INPUT is identical equally unreviewed, so the same QBITTORRENT_DATA_DIR value that w driven a filesystem-wide chown will driv file creation at an arbitrary location. No escape is empirically confirmed for the path, not merely reasoned: during T028 remediation a probe declaring the verba

ATM ID	Type	Status	Severity	Description
BOB-188	Bug	In progress	—	shipped entry with QBITTORRENT_DAT hung the suite, and ps showed the real executing “find / (! -uid 1000 -o ! -gid 1 -printf …” before it was killed by verifie Acceptance: ownership_precondition.sh SAME ownership_path_fence() predicat scripts/lib/ownership.sh (never a secon 11.4.251), refuses fail-closed on a reject and ships a paired 1.1 mutation proving refusal fires PLUS a negative control pr six live shipped entries still ACCEPT so not a 11.4.201(1) false-positive refusal. Discovered out-of-band during T028 rer and reported by the agent as needing tr it is also a 11.4.238 coverage escape. **OPERATOR DECISION (2026-08-26, interactive clarification): UNTRACK BINARY — BUILD ON DEMAND** §11. answer, recorded as consumer DATA: th compiled `workable-items` binary is NOT t tracked. ‘Any build derivate which we ca recreate by executing proper mechanis generating MUST NOT be versioned’ ap without exception here. Keeping it track only the invariant-52 fingerprint gate, a keeping it tracked with a commit-time r were both offered and NOT chosen. RE END STATE: the gate builds from sourc REFUSES HONESTLY when the Go tool absent (§11.4.201(11) — probe the artif through its real invocation path, never f pass, never silently fall back to a stale b Staleness becomes structurally impossi than merely detected. ACCEPTED COS clone needs a Go toolchain before invar can run — already true for the qBitTorr backend, so this adds no new host requ The invariant-52 fingerprint gate remai as defence-in-depth for as long as any b exists on disk. This answer is recorded a consumer DATA per §11.4.35 — it is the operator’s stated choice, not an agent i and supersedes any prior agent-chosen this question. Options not chosen are na above so a future reader does not re-liti settled call (§11.4.112(5) bounded-verdi discipline applied to decisions). — prior follows — pre_build_verification.sh inva (CM-WORKABLE-ITEMS-VALIDATE) res binary through the candidate loop at :5

ATM ID	Type	Status	Severity	Description
				FIRST entry is constitution/scripts/work items/bin/workable-items. That file is GI TRACKED (md5 17644a248363, identic workable-items-linux) and executable, s resolution over the current untracked si constitution/scripts/workable-items/wor items (md5 43376a6d0184). The tracke STALE: it does not contain the guards it source now has. MEASURED, needle-pr 2026-08-25. String presence in the trac binary: “refusing to set terminal status” “Issues-location item has TERMINAL st control needle from the SAME updateC least one mutable field flag is required” the instrument sees through that path a zeros are real absences, not a blind rea untracked sibling returns 1 / 1 / 2 for th three queries. A check whose message s absent from the binary cannot fire. RUN PROOF: the same injected violation (a r terminal status while current_location= run through both binaries -> stale repo violation (only the body_md desync clas reports 2 including the location-status c verbatim. This is exactly why the ten BO rows (BOB-087/129/131/144/145/146/148/15 survived undetected: the gate that was to catch them was running a binary stru incapable of seeing them. The 11.4.108 >ARTIFACT gap: source correct, artifac every gate reading the artifact reports g This is the THIRD instance of that class one session, alongside BOB-169 (a char that never opened a file) and BOB-183 (bundle five days older than its sources). COMPOUNDING: the comment at :519- directly above the loop asserts “The nai workable-items path never existed in th checkout (bin/ is a gitignored local build dir nothing ever populated)”. That state now FALSE - bin/ holds two tracked bin stale comment asserting the absence of file that now shadows resolution is how stayed invisible. Also an 11.4.30 questio operator owns: bin/workable-items and workable-items-linux are VERSIONED B ARTIFACTS in a constitution submodule inherited by reference by every consum every consumer inherits whichever bina

ATM ID	Type	Status	Severity	Description
				last committed. Flagged independently separate agents this session. ACCEPTA resolution must never prefer a stale art a current one - either the tracked binar removed and resolution falls through to demand, or a freshness check refuses a older than its sources (see the BOB-183 fingerprint gate for a working pattern); false comment at :519-523 corrected; (c 1.1 mutation proving the new refusal fi negative control proving a genuinely fre still passes so the fix is not an 11.4.201 positive refusal; (d) the 11.4.30 tracked decision recorded as operator DATA. Di out-of-band during BOB-166 remediatio 11.4.238 coverage escape in its own rig
BOB-189	Bug	Queued	—	WHAT: the fail-open scanner's shape-(A heuristic cannot distinguish 'return Fals meaning PROCEED-AS-IF-FINE from 'r False' meaning REFUSE. It flags six tex CLOSED guards as fail-open defects: _is_safe_fetch_url (x3), _qbit_add_succe and the hooks path-boundary guard. INDEPENDENTLY VERIFIED 2026-08-2 conductor rather than taken on the tria word: download-proxy/src/api/routes.py defines _is_safe_fetch_url returning boo only call site at :1476 reads 'if not _is_safe_fetch_url(url): logger.warning(SSRF-unsafe download URL (non-public skipping); continue' - a False return RE the URL. WHY THIS MATTERS: per §11 a false-positive refusal is a FAIL-bluff of severity to a false-negative pass, and he consequence is worse than noise - actin finding would mean making the SSRF g returning False, i.e. deleting an SSRF p to satisfy a gate. A gate that instructs y remove a security control is actively da not merely imprecise. REPRO: run the f scan; observe six hits; read each call si ACCEPTANCE: the scanner distinguish shaped from proceed-shaped False retu site-aware, or an audited waiver list wit entry justification), the six drop out, AN golden-FALSE fixture containing a real guard is added so the discrimination is falsifiable per §1.1.
BOB-191	Bug	Queued	—	WHAT: the §11.4.252 fail-open scan rep for qBitTorrent-go, and that zero is NO

ATM ID	Type	Status	Severity	Description
				evidence. The triage agent ran control n through the scanner's own path per §11 (b): a Python 'except Exception: pass' n SEEN, a TypeScript 'catch (e) {}' needl SEEN (so frontend/src's zero IS a real z Go empty-'if err != nil {}' needle was N An instrument that cannot see the idiom the same quiet zero as a clean tree, and of those is honest. DISTINCT FROM BO deliberately not merged with it per §11. BOB-189 is a false MATCH (fail-closed g reported as fail-open); this is a false NU entire language unscanned). Same scan opposite failure directions, different fixe merging them would hide one behind th IMPACT: Go is the language of qbittorre go and boba-jackett (port 7189, which c encrypted tracker credentials), so the u surface is exactly where §11.4.252's cre plus-mutation combination is most likely has assessed it; the dashboard says clea ACCEPTANCE: the scanner grows a Go covering the empty-err-block and swallo idioms, its needle is SEEN through the r qBitTorrent-go's result is re-derived, an finding is triaged as this Python/TS pass Until then qBitTorrent-go's fail-open po UNKNOWN and must be reported as UI never as 0.
BOB-192	Task	Queued	—	WHAT: the BOB-161 gate lands with 6 r open skips RATCHETED rather than fixe Ratcheting is the constitution's named l default (§11.4.135/§11.4.224(E)) and thi own precedent, so the choice is correct remediation it defers is real work that n owned somewhere. WHY THIS ITEM EX gate's own header asserted the 6 were SEPARATELY (§11.4.197)' while no trac existed. The §11.4.209 independent revi verified the absence and raised it as IMPORTANT-5, noting that without a ro findings are precisely the parked-unver class §11.4.226(4) names - the populatio operator samples and finds broken. A pr of being tracked is not tracking. WHAT NEEDS: a skip that fires on evidence th ANSWERED must either classify the res and FAIL on it, or take a §11.4.69 reaso honestly derivable from the environmen than from the response - verified agains

ATM ID	Type	Status	Severity	Description
				stack, not asserted. Two of the six sit ur allow-skip:' markers at tests/unit/ test_tracker_auth_live.py:105 and :108, reviewer confirmed genuinely are fail-o that marker must not be treated as abso ACCEPTANCE: all 6 remediated with RI evidence per §11.4.115, the gate's BASI ratcheted to 0, and the ratchet's monoto decreasing property preserved through (§11.4.227(A)). NOTE the reviewer's MI count-baseline absorbs a one-out-one-in remediation progress must be checked the finding SET, not only the count.
BOB-193	Bug	Queued	—	WHAT: in download-proxy/src/api/strea the dedup set is written BEFORE the yi seen_hashes.add() / seen_hashes_local.a execute at :317 and :356 ahead of emitt format_event raises, the result is never its hash is already recorded, so it is per dropped. Independently verified by the reviewer at both sites: mid-search the h persists across subsequent polls so the never re-emits; at completion the strear immediately after, so there is no later cl either. Results later in the SAME batch yet added and do re-emit on the next po is why the observed symptom is partial rather than total. WHY IT IS FILED SEF the §11.4.252 fail-open pass made this l OBSERVABLE by adding a log line, whic right first move — but observability is n The underlying at-most-once semantics and per §11.4.226(5) a mitigation canno the defect it mitigates. THE DECISION TRADE, not an oversight to correct blin moving to add-after-successful-yield giv least-once, but a DETERMINISTIC seria failure would then retry the same result poll forever — trading silent loss for a h Options are (a) keep at-most-once and t log line as the contract, (b) add-after-su yield with a bounded per-hash retry buc add-before-yield but remove the hash on failure so exactly one retry occurs. ACCEPTANCE: an explicit decision reco implemented, and covered by a test tha raising emit and asserts the chosen sem with a RED observed first per §11.4.115 as MINOR-1 by the independent review fail-open batch; filed rather than absorb

ATM ID	Type	Status	Severity	Description
				residual defect is not retired along with logging that revealed it.
BOB-194	Bug	In progress	—	WHAT: agent worktrees under .claude/v are ephemeral copies pinned at arbitrar - never built, never shipped, deleted wh agent finishes. Pre-build gates that wall with find(1) from the repo root descend anyway. MEASURED INSTANCE (2026- CM-GO-TOOLCHAIN-MATCHES-BUILD FAILED the main build with exactly one from .claude/worktrees/agent-afda1df90 - a worktree 86 commits behind still car 'FROM golang:1.23-alpine' against a go since at 1.26.2. The main tree was COR throughout. Per §11.4.201(1) a false-pos refusal is a FAIL-bluff of equal severity t pass: it blocks real work and teaches pe bypass the gate. FIXED FOR ONE GATE added to that gate's PRUNE_DIRS with deterministic §1.1 pair - a stale mismatc inside .claude/worktrees must NOT fail tree, the SAME mismatch outside it MU fail, so the prune cannot widen into blin (§11.4.201(6)). Mutation verified: revert prune fails 2 checks. SURVEY, measured inferred: 8 of 13 pre_build gates walk th only the fixed one prunes .claude. Gates ls-files' are structurally immune - workt untracked. Empirically probed: check_cm_closure_seam_binds, check_cm_killpg_pgid_guard and check_cm_no_production_mutation_resi rc=0 with ZERO worktree references - practice today. UNMEASURED, not clea check_cm_test_mock_pid_explicit_int an check_cm_test_mock_pid_patched_when were not reached before the probe budg expired. HYPOTHESIS REFUTED, recor nobody re-walks it: an earlier revision o speculated that check_cm_plugin_count timeout was caused by walking worktre NOT. That gate's find is scoped to $PLU (check_cm_plugin_count.sh:224), so wo are outside its walk entirely. Its slownes still running past 54s on a re-measure - different, still-unidentified cause and be its own item, not this one. ACCEPTANC tree-walking gate either prunes agent s is proven immune by measurement, eac paired mutation. NOTE the asymmetry

ATM ID	Type	Status	Severity	Description
BOB-195	Bug	In progress	—	bounds the risk: a stale worktree can or findings to a security-relevant scan (kill mutation-residue), never remove them - exposure is false refusal and eroded tru missed defect. **OPERATOR DECISION (2026-08-26, interactive clarification): TEACH TH SCANNER With NODES — KEEP SIM** §11.4.252 fail-open scanner is FIXED to `contextlib.suppress`; ruff's SIM105 rule ENABLED in pyproject.toml. Disabling and the belt-and-braces both-at-once op offered and NOT chosen. Rationale carr the decision: fixing the gate makes SIM harmless, whereas disabling SIM105 al leave the gate blind to any suppress wr hand, inherited from a dependency's sty already present in the tree. REQUIRED BEHAVIOUR: `with contextlib.suppress(` `with suppress(...)` (the `from contextlib` `suppress` binding) wrapping a dangerous combination call are DETECTED with th severity and message shape as the try/e pass form; a narrow `suppress(SpecificEr` around a NON-dangerous call must NO since a false-positive refusal is a §11.4.2 FAIL-bluff of equal severity to a missed CONSEQUENCE FOR THE NUMBER: c landed, the scanner's finding count sto floor over try/except shapes and becom genuine census — any prior count citing completeness was scoped to try/except answer is recorded as consumer DATA p §11.4.35 — it is the operator's stated ch an agent inference, and supersedes any agent-chosen default on this question. C not chosen are named above so a future does not re-litigate a settled call (§11.4. bounded-verdict discipline applied to de — prior item text follows — WHAT: the fail-open scanner matches Try-handler s (A1)/(A2). contextlib.suppress is a With swallow written that way is invisible. V BY THE CONDUCTOR with a control ne through the same invocation path (§11. (b)), because a first attempt returned 0 forms and was itself blind - the gate tak not a positional, so a bare path exits on unknown arg (§11.4.201(7)(c): the invoc is part of the instrument). Correct run:

ATM ID	Type	Status	Severity	Description
				pass around os.unlink(user_supplied_pa ‘FAIL - swallowed exception … :5’, the r sees. contextlib.suppress(Exception) arc IDENTICAL call -> ‘PASS - no swallowe exception … anti-patterns found’. Both untrusted input with an irreversible unl swallow everything; one is caught and WHY IT IS WORSE THAN A PLAIN GAP pyproject.toml enables ruff’s SIM rules SIM105 is ‘use contextlib.suppress inste except-pass’. Our own linter therefore i authors to rewrite the form this gate CA into the form it CANNOT. The two tools pulling in opposite directions and the lin runs more often. Every SIM105 autofix shrinks this gate’s coverage while the c down, which reads as progress. CONSE FOR THE NUMBER: the scanner’s findi is a FLOOR over try/except shapes, not of swallows. Any statement of the form open sites remain’ must be read as ‘N a except shapes’. ACCEPTANCE: the scar recognises contextlib.suppress (a With items call contextlib.suppress with a bro exception type) and applies the same ≥ capability test; a golden-TRUE fixture in form; a golden-FALSE fixture where sup wraps a single-capability diagnostic emi must NOT fire; and an explicit decision SIM105 tension - either exempt the sha gate governs, or accept suppress and te gate to read it. Discovered by the §11.4 remediation agent when its own MINOR required using contextlib.suppress at th diagnostic-emitter sites; those three use legitimate and justified in-source.
BOB-196	Task	Queued	—	WHAT: the CM-PLUGIN-COUNT pre-bu takes 189 SECONDS to verify 8 docume counts. Measured 2026-08-25: exceeded probe budget (rc=124), then a bounded completed at ‘rc=0 wall=189s’ - it PASS simply slow. CAUSE, MEASURED NOT GUESSED: the marker loop at check_cm_plugin_count.sh:273-290 iter EVERY LINE of every governed docume spawns a pipeline per line. Per iteration grep -oE | wc -l | tr -d’ (4 processes) plu sed -n’ (2) = ~6. Governed docs are AG 689 + CLAUDE.md 494 + docs/features Status.md 670 = 1853 lines, so ~11,118

<table>
<tr><th>ATM ID</th><th>Type</th><th>Status</th><th>Severity</th><th>Description</th></tr>
<tr><td></td><td></td><td></td><td></td><td>spawns, measuring ~102 ms/line. Class
line-subprocess shell trap. THE IRONY
PRESERVING, because it explains why
simplified it: the in-code comment show
-l’ was chosen DELIBERATELY over ‘gre
measuring the §11.4.201(12) footgun - c
host (ugrep 7.8.4) ‘grep -coE’ returned
level but 1 inside a ‘set -euo pipefail’ su
context-dependent reading. The correct
right and must be kept; it is the PER-LI
application of it that costs the three mi
HYPOTHESIS REFUTED, recorded so n
walks it: this is NOT worktree traversal
gate’s find is scoped to $PLUGINS_DIR
the five agent worktrees are outside its
That guess was raised on BOB-194 and
LIKELY FIX (unverified, stated as a cand
a conclusion): pre-filter with ONE grep
containing 'CM-PLUGIN-COUNT:' befor
the loop, reducing ~1853 iterations to t
lines that carry a marker, while leaving
line parsing logic - and its footgun-avoic
completely unchanged. WHY IT MATTE
beyond impatience: this gate is on the p
critical path, and CLAUDE.md leans on
mechanical authority for the three disti
rosters (43 curated / 43 engines / 12 bo
whose conflation was BOB-149. A gate s
enough to tempt anyone into skipping it
nothing. ACCEPTANCE: runtime reduce
counts and the footgun-avoiding parse p
byte-for-byte in behaviour, proven by th
paired mutation still biting plus a befor
wall-clock recorded. §11.4.6: nothing he
questions the gate's VERDICTS - only it
| | BOB-197 | Bug | Queued | — | OPERA
DECISION (2026-08-26): arm BOBA_AP
and keep 0.0.0.0, backed by a BOOT-TIM
invariant that refuses to start LAN-bour
token unset. This item IS that invariant.
MEASURED FACT, not inference. downl
proxy/src/api/routes.py:83 require_api_t
reads BOBA_API_TOKEN at request tim
returns immediately when unset or emp
own docstring states it plainly:
"BOBA_API_TOKEN unset/empty -> ret
(OPEN). This is the DEFAULT and prese
current no-auth contract + dev workflo
docker-compose.yml:237 passes
BOBA_API_TOKEN=${BOBA_API_TOK</td></tr>
</table>

ATM ID	Type	Status	Severity	Description
				the container inherits empty unless the declares it. The operator .env does NOT it, needle-verified per 11.4.201(7)(b): th matcher hit QBITTORRENT_DATA_DIR same file, so the miss is a real absence a blind read. The qbittorrent-proxy contai RUNNING and LAN-bound at measuren Consequence: the ten routes the BOB-1 resolves as auth-wired are open in prac default deploy. WHY A STATIC GATE CA COVER THIS. The BOB-102 gate assert route WIRING - that a Depends(require_api_token) marker exi handler, decorator or router. It cannot a runtime ARMING, which depends on an read per request. Claiming the static ga it would be exactly the bluff the gate ex prevent (11.4.262: machine evidence m the layer it claims). This is a boot-time i class per 11.4.254. ACCEPTANCE: a bo check that, when the listener is LAN-bo loopback), refuses to start with a non-ze and a named cause if BOBA_API_TOKEN or empty, so the open state becomes UNREACHABLE rather than the default 1.1 mutation: remove the check, prove t starts LAN-bound-and-open. Golden-FA 11.4.201(1): a loopback-bound listener v token unset must NOT be refused, or th becomes a false-positive refusal blockin legitimate dev work. 11.4.238 COVERA ESCAPE NOTE: found by a subagent re source during BOB-102 guard construct by the automated QA regime - which is defect class 11.4.238 names. The boot-t IS the new automated check that would caught it.
BOB-198	Bug	Queued	—	MEASURED, conductor-verified indepe the reporting subagent. qBitTorrent-go/ qbittorrent-proxy/main.go registers exa middlewares, at lines 62, 63 and 68: middleware.CORS, middleware.Logger, middleware.GinRateLimit. There is NO authentication middleware and no per-r dependency. main.go:119-121 binds add which is ALL interfaces, so every route reachable when the profile runs. SCOPE routes, of which the mutating ones matt POST /api/v1/download, POST /api/v1/m DELETE /api/v1/hooks/:id, POST and DE

ATM ID	Type	Status	Severity	Description
				api/v1/schedules. A LAN peer can queue downloads and delete webhook configu with no credential. PARITY GAP, stated PUT /api/v1/theme is auth-gated in the service and open in Go. The same asym holds for the download, hooks and sche mutations. Two implementations of one advertised surface disagree about whet needs a credential. RELATIONSHIP TO (Go parity deferred past v1.0.0, operato 2026-08-26): that deferral covered POR the Go container binding :7186 and :71 as :7187. It did NOT cover an authentic and the operator was not asked about o because the hole had not been measure the question was posed. Treating this a by the deferral would be a 11.4.112(5) leak: citing a bounded decision outside its evidence supports. This item stands and its severity is NOT settled by BOB- MITIGATING FACT, so severity is not ov the Go profile is opt-in via –profile go a NOT running at measurement time (the list showed qbittorrent-proxy, the Pytho with no Go sibling). The exposure is lat live. ACCEPTANCE: either auth middle landed in the Go service with parity to route set, or an explicit operator decisi recording the Go profile as dev-only tha refuse to start LAN-bound - with a guar enforcing whichever is chosen.
BOB-199	Task	Queued	—	RESIDUAL ASYMMETRY surfaced by th BOB-195 fix, reported rather than silen After teaching the scanner With nodes, suppress (Exception / BaseException) is with the same severity as try/except/pa NARROW suppress (suppress(FileNotFoundError)) is not - semantically identical narrow try/excep FileNotFoundError: pass IS flagged. So SIM105 rewriting a narrow handler stil that site out of the gate's scope, a smal of the exact mechanism BOB-195 was f close. WHY IT WAS NOT CLOSED IN T with the measurement that decided it: t authoring agent reported 37 narrow su sites, all idiomatic. CORRECTION (202 independent review): that figure DOES REPRODUCE and this conductor propa into this item as fact without verifying -

ATM ID	Type	Status	Severity	Description
				is mine, not the reviewer's. Measured: t constitution repo has 0 narrow `with sup` sites; boba first-party has exactly 1 (a suppress(OSError)). The nearest corpus 35 narrow `except X: pass` handlers in bo party - a DIFFERENT construct - and th census (OSError x7, ImportError x7, BrokenPipeError x3, RuntimeError x3, ValueError x2, FileNotFoundError x2, I x2, SystemExit x2) contradicts 'every or idiomatic': a bare `SystemExit: pass` is no tolerance. The two populations must no conflated, and the distinction is the who substance of this item. The false-positiv argument therefore rests on the 35 narr except-handlers, not on 37 suppress site strength should be re-judged on that ba under 11.4.201(1) is a FAIL-bluff of equ to the gap it would close, and worse in because it trains readers to ignore the g subagent recorded the asymmetry in th header as a known gap rather than ship storm. That was the right call, and is wh separate tracked item rather than an ur one. THE DECISION IS THE OPERATOI mirroring BOB-195 itself: (a) accept the asymmetry permanently, with the gate l stating it so no reader mistakes the cou census; (b) flag narrow suppress ONLY combined with an irreversible capability truncate / kill), catching the shape that matters and leaving the 37 idiomatic sit (c) flag all narrow suppress and absorb a justified exemption list like the LAN-re uses. RELATED FACT worth carrying: th newly-visible production sites in downlo src/merge_service/search.py (1294, 130 1329, 1333) wrap proc.kill / os.killpg / p in the BOB-126 cleanup path. The cond verified the 11.4.263 pgid guard IS pres correct at both killpg sites - _pid and _p checked isinstance(int) and > 1 before t syscall, with the BOB-126 forensic reas inline - so those sites are true-by-the-sca definition but SAFE in fact, and want a exemption entry rather than a code cha