

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	7
Bug	Operator-blocked	1
Bug	Queued	20
Bug	Ready for testing	4
Task	In progress	10
Task	Operator-blocked	3
Task	Queued	24
Task	Ready for testing	2
TOTAL		71

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-hos routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in proxy + qBitTorrent-go + Jackett + com forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel — ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
BOB-077	Task	Operator-blocked	High	

ATM ID	Type	Status	Severity	Description
				RD2-10: Identify second host running the commit rsync/sync mechanism (OPERATOR DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code re the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable the four Lava-porting findings (closes G
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes C
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Trac + Status Documents table row-completo (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exerc three start.sh subcommands end-to-end real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify ReDoS regex bounds deployed to contain (closes runtime half of GA-12)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos. side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for download-proxy/merge endpoints (canonical of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is -profile go parity still a goal? (gates RW-10..13) — OPERATOR-I
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — b 127.0.0.1 or keep 0.0.0.0? — OPERATOR DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence helper for the unattributed auto-commit
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence for subagent task-brief source inputs
BOB-109	Task		Medium	

ATM ID	Type	Status	Severity	Description
		Ready for testing		BOB-074 followup: scaling-class test coverage absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limiters on boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golden fixture for the rate-limit detector
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout architectural gap (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_creation bulk suite (Permission denied /config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while the process still serves 7186 (GIL starvation, spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile server serves 7186/7187/7188 but its container binds to 7187 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolved gitdir) pollute gate scan scope and manifest false BOB-126-class findings
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/41 from constitution, CLAUDE.md and the REAL badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labels 44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate implementation, and 24 of 36 scripts have companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored third-party code: 82% of one gate's 38,291 findings from submodules/
BOB-154	Bug	Ready for testing	Medium	Host venv and production container run different starlette versions (1.4.1 vs 1.6.0)
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard is too sensitive and flaky: 8786ms under host vs 900-1500ms ceiling calibrated on a quiet host
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds <code>asyncio.events.get_event_loop_policy</code> , a private API, while production is 3.12.13
BOB-159	Bug	Queued	Medium	Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:01:17Z Reported-By: independent review (IMPORTANT-2), page 1

ATM ID	Type	Status	Severity	Description
				remediated What (the report, verbatim) -recreate path was fixed this round: run_ownership_precondition -> stack_d run_ownership_repair -> stack_up, so th walks a tree no container can write to. ordering is now asserted behaviourally new checks in tests/unit/ test_start_reload_recreate.sh, each kille paired 1.1 mutation (11/11 RED-OK). TH path is NOT covered by that fix and this tracks the remainder. On a warm ./start gate runs before THIS invocation brings up, but if the stack is ALREADY running containers write while the repair walks. container can then create a new non-op owned file BEHIND the walk, after whic completion marker records complete ov that is not, and those stragglers are nev repaired without a manual -force. BOU dismissed: after any successful pass the valid for the scope fingerprint and the r short-circuits without walking (marker_ so the window requires a stale-or-absen AND a live stack together. Declaring a r entry re-arms the marker, which is exac that combination arises in practice. NO THIS ROUND, with the reason stated ra hidden: the clean guard is a cheap mark probe mode on scripts/ownership_repair answers has-this-scope-been-repaired w walking the tree. The existing -dry-run serve: it deliberately SKIPS the marker (FORCE=0 && DRY_RUN=0 guard) and walks, so using it as a warm-start probe walk the tree twice on every start. Addi mode requires editing ownership_repair another stream was actively editing dur round; duplicating marker_is_valid into instead would be a near-identical fork (Deferred to this item rather than done unilaterally or silently. The misleading c at the warm-start call site (which claim gate runs before any container writes) l corrected to state this boundary honest Affected scope / file-scope manifest: (warm-start dispatch, run_ownership_g site); scripts/ownership_repair.sh (mark path) Reproduction / context: 1. Ens stack is UP. 2. Change config/owned_pa so the scope fingerprint changes (this r

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	marker by design, data-model E2). 3. Run warm ./start.sh (no -recreate). 4. The ownership repair walks and chowns the declared test containers are still running and able to run it. Acceptance criteria: A warm ./start.sh complete an ownership repair while a container that writes to a declared location is running. Either the repair is deferred with an acceptance refusal naming -recreate, or the stack is ready for the walk. Asserted behaviourally in test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE_DOWN checks, each with a paired 1.1 mutation to fail it. Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:02:48Z Reported-By: What (the report, verbatim): Independent §11.4.209 review (IMPORTANT-3) found the feature's strongest automated checks never executed by any runner: (1) tests/ownership/test_container_writes_owned_files — the §11.4.115 RED-turned-regression FR-011/FR-007, proven via a docker-compose revert mutation — was moved out of test integration/ (commit 58d340a, to escape autouse fixture hang) but no runner (ci.test.sh, run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_build/test.sh, the §1.1 paired-mutation meta-test the scripts/pre_build/check_cm.sh gate was executed by NOTHING — self-reported honestly in commit 04742d7's own message ('STILL OPEN (reported, not fixed): test_pre_build/ is executed by NOTHING') but tracked as a workable item (a §11.4.197 requirements gap) and never wired into pre_build_verification.sh invariant 30, which globbed only tests/unit/test.sh. Affected file-scope manifest: ci.sh; scripts/pre_build_verification.sh invariant 30 (C UNIT-TESTS-EXECUTED) Reproduction context: grep -rn test_container_writes_owned_files ci.sh run-all-tests.sh scripts/ returned zero results (only docs/specs mentioned the path); grep scripts/pre_build_verification.sh showed invariant 30's for-loop globbing only tests/unit/test.sh never tests/pre_build/test.sh. Acceptance criteria: ci.sh gains a runtime-gated py

ATM ID	Type	Status	Severity	Description
				ownership/ stage (skips honestly, rc=0, podman/docker present); scripts/ pre_build_verification.sh invariant 30's covers both tests/unit/test_.sh and tests pre_build/test_*.sh, verified by an extra standalone run of the invariant's exact l reporting a non-zero RAN count that inc files from both directories. Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:31:12Z Reported-BY: BOB-092 remediation, residual finding report, verbatim): §11.4.69 names CM OPEN-SKIP as one of three mandatory p gates: it audits sink-side probe helpers if any code path converts an empty or unreachable response into a PASS-count for a feature class with a sink-side prob project does not have it. Why this matte rather than in the abstract: the BOB-09 remediation just removed TWO fail-oper exactly this class from tests/e2e/ test_live_stack_evidence.py — the nnmc on-404 (already removed at 7baef2b) an surviving sibling at test_iporrents_is_authenticated_in_sea skipped on 'not authenticated and statu success' while asserting 'treating as tra outage'. That assertion was false for rej credentials (upstream_http_403) and fo container (plugin_env_missing), both of definitive product failures the product a classifies via error_type. Two fail-opens class, found one at a time, is the §11.4.1 extend-to-all-cases signal and the §11.4 coverage-escape signal together: the re not surface either, an agent reading cod remediation's own guards are AST-struct lethal under three discriminating mutat they LIVE IN THE FILE THEY GUARD, deleting that file evades them entirely. A gate is the durable home because it auc corpus rather than one file. Honest bou (§11.4.6): this item does NOT claim the remediated fail-opens are unguarded — guarded, with runtime evidence (13 pas 97.36s against the live stack). It claims CLASS has no corpus-wide detector, so instance in a different file is invisible ag Affected scope / file-scope manifest: pre_build/ (gate absent); tests/e2e/
BOB-161	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	test_live_stack_evidence.py (guards currently inside the file they guard) Reproduction context: <code>grep -rl 'CM-NO-FAIL-OPEN-S' scripts/ tests/</code> returns exactly ONE hit, at <code>tests/e2e/test_live_stack_evidence.py</code> — the guards live in, not a gate. Control net same query for CM-OWNERSHIP-INVALID (gate that does exist) returns 3 files, so the instrument can see gate tokens and there is a real absence, not a blind zero (§11.4.201(b)). Acceptance criteria: A script/probe named CM-NO-FAIL-OPEN-SKIP is wired into <code>scripts/pre_build_verification</code> audits sink-side probe helpers for code converting an empty/unreachable/error into a PASS-counting SKIP for a feature HAS a sink-side probe. It ships a golden fixture (a real fail-open -> gate FIRES) a golden-FALSE-with-carrier (an honest test geo skip, and a comment merely MENTION the phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §11.4.202 makes the gate FAIL before the gate is reported. Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:35:02Z Reported-By: BOB-106/BOB-107 remediation, residual (the report, verbatim): The TESTS for guards are now executed: pre-build invalid glob was extended from <code>tests/unit + tests/pre_build</code> to also cover <code>tests/hooks</code>. BEFORE/after: the glob expanded to 27 sub 0 under <code>tests/hooks</code> while 3 existed there expands to 30 with 3 under <code>tests/hooks</code> pass. But the GUARDS themselves are still invoked by nothing. <code>check-brief-inputs.sh</code> honestly conductor-run rather than auto-brief's required-input list lives in free-form that the Agent tool's structured input does expose, so an automatic prompt-scraping would itself be an unproven pattern-matching seam is the dispatch procedure, and that documentation + habit change, not a honest unattributed-commit-guard.sh is different why this item exists. Run against real names 14 violating commits since the last (and 21 bare Auto-commit subjects exist all refs). Wiring it into the pre-build or seam AS-IS would therefore refuse even immediately, on pre-existing debt nobody in the moment. That is precisely the 11.

ATM ID	Type	Status	Severity	Description
				brownfield-adoption shape, and 11.4.66 put the choice with the operator, not with agent inventing a ratchet. The options, the decision is a decision and not a default immediate hard floor – the seam refuses 14 are attributed; (b) one-time monotonic decrease ratchet (the 11.4.135 pattern) snapshot 14 as the baseline, the count may and may never rise, so day one is green disease cannot spread; (c) changed-commit – gate only commits created from now on scheduled deadline for the historical 14 report-only for a stated period, then escape Recommendation, with the reasoning rather just the pick: (b). It is the pattern this commit already uses for exactly this situation, if the debt visible and bounded immediately cannot be satisfied by deleting the guard (11.4.227 closes that gaming channel). the operator’s call and this item is not closed until that answer is recorded. Honest boy (11.4.6): this item does NOT claim the 14 commits are harmful in themselves – it is their ATTRIBUTION is absent and that is currently prevents the 15th. Affected scope file-scope manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-build-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh Reproduction context: Both guards run correctly by hand are covered by passing tests (9/9 and 13 proven non-bluff against a no-op stub). 14 invoked by scripts/pre_build_verification.sh; scripts/commit-push-all.sh, so neither escape standing detection pressure (11.4.226: with no execution seam is not coverage registration is not coverage). Acceptance criteria: Each guard is invoked by a name seam, OR carries a registered deferral point this item. For unattributed-commit-guard specifically, the operator has answered adoption question below and the answer recorded as consumer DATA – never an ratchet. Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:41:08Z Reported-By: BOB-114 remediation, pre-existing defect surfaced by BOB-111 landing a real limit. (the report, verbatim): PRE-EXISTING INTRODUCED – and proven so rather than
BOB-163	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				asserted: git diff shows assertion (c) is IDENTICAL to HEAD, and the failure re-against the unmodified HEAD script. What changed is not the challenge but the SYNOPSIS. BOB-111 appears at least partially delivered because :7187 now returns real 429s whereas previously returned none. :7185 and :7189 produce zero 429s. This is a 11.4.248-class corrosion risk and that is why it is filed rather than left as a footnote: run_all_challenges.sh now fail INTERMITTENTLY, and an intermediate trains everyone to re-run until green, which point a real regression is dismissed as 'probably the flaky rate-limit one'. THE DECISION, which is an operator call under 11.4.66 and which I deliberately did NOT make. Does a 429 from a WORKING limiter count as a sibling endpoint being RESPONSIVE? (a) YES - 429 proves the endpoint is alive and corner-protecting itself. Assertion (c) accepts 2xx/429, and only a connection failure / 5xx counts as degraded. Risk: a genuinely working endpoint that happens to answer 429 with (b) NO - keep requiring 2xx, but make the challenge limiter-aware: drain or wait out the window before the sibling probe (retry-served, so the budget is knowable rather than guessed). (c) Probe siblings on a path that exempts (a healthz-class route), so the initial question is asked without spending the budget. Recommendation with reasoning, not just (a) composed with a bounded liveness of 429 carrying a well-formed Retry-After header as evidence of a live, correctly-behaving sibling. (b) makes the challenge slower and could increase window value that will drift. But the server 'responsive' here are a product judgement: an answer is recorded, not assumed. RELATED: stated as fact rather than folded in: the script counts 429 only, not 503, though the script header says '429 (or equivalent)'. Counting 503 would collide with the crash detector's logic. Worth deciding when BOB-111 lands line :7185 and :7189. Affected scope / file-manifest: challenges/scripts/ddos_resilience_challenge.sh assertion endpoint isolation; run_all_challenges.sh invokes it Reproduction / context: Assertion requires a sibling-endpoint probe to answer 2xx). :7187 now enforces a real limiter (

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	In progress	Medium	live: x-ratelimit-limit: 120, x-ratelimit-re 86, retry-after: 45, server: uvicorn). A full challenge run sends roughly 250 requests :7187, so sibling probes fired during the endpoints' tiers land on an exhausted b receive 429 - which assertion (c) scores 'endpoint degraded'. Timing-dependent measured PASS=5 FAIL=2; the UNMO HEAD script against the same live stack later measured PASS=7 FAIL=0 SKIP=. Acceptance criteria: The operator has the classification question below, the an recorded as consumer DATA, and assert implements it. The challenge then prod same verdict across 3 consecutive runs rate-limited stack (11.4.50 deterministic consistency), and the fix ships a paired mutation proving assertion (c) still catch genuinely degraded sibling endpoint - n it must not blind it (11.4.201(1)). Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:56:53Z Reported-By: BOB-110 UX-class coverage, discovered new axe-core suite on its first live run V report, verbatim): This is a REAL user defect, not a test-tuning artifact, and it by the automated regime rather than by squinting at the page - which is exactly 11.4.238 posture the project is aiming for static-grep oracle would NOT have found Measured: the served root ships an emp hydration, so any check reading the raw audits a page nobody sees. The violation exists in the hydrated DOM, which is w 11.4.170 requires a rendered oracle and value-equality assertions as the proof a correct. Severity reasoning, stated rather assumed: this is user-visible and affects primary dashboard heading, but it degrades legibility rather than breaking function, surface is operator-facing rather than p Medium, not High. The failing test was FAILING on purpose (11.4.238) instead silenced or marked xfail. tests/ux/ current reports 1 failed, 16 passed; that 1 is this. Anyone reading a red UX suite should r this item, not as flakiness - and when the suite goes fully green, which is the it is closed. HONEST SCOPE LIMIT (11 the dashboard landing view was scanned

ATM ID	Type	Status	Severity	Description
BOB-165	Bug	Queued	High	jackett/* sub-routes and the ng-serve-ho:4200 route set were NOT audited – the commands to close both are recorded in testing/ux_accessibility.md. So this item nodes are a floor, not a total. Affected file-scope manifest: the Angular dashboard served at http://localhost:7187/ (same c SPA as frontend/); production component not test files Reproduction / context: ux/test_live_dashboard_accessibility.py the running merge service. axe-core v4 scanning a real Playwright-rendered JS DOM, reports color-contrast violations on nodes. Measured pairs: .brand / h1 text on background #3c3f41 = 1.43-1.62:1 (large-text floor is 3:1); tagline #808080 #3c3f41 = 2.68:1 (body-text floor is 4.5). Acceptance criteria: axe-core reports color-contrast violations against the live dashboard, with the fix made in product component CSS rather than by relaxing assertion or excluding the rule (11.4.12 reconcile to the correct mechanism, never weaken the check). The existing tests/ux the guard and already fails today, so the captured – closure requires it flipping G against the live surface, which is runtime evidence per 11.4.226. Reported-Via: §11.4.202 reporting direct on 2026-08-21T20:00:08Z Reported-By: surfaced while capturing closure evidence BOB-092; diagnosed rather than worked What (the report, verbatim): This is a ABI-AFTER-INTERPRETER-UPGRADE of a missing package. The package is installed compiled half is built for the previous in That distinction matters because ‘pip in py’ style advice can appear to succeed v leaving the 313 artefact in place. WHY NOT NOTICED EARLIER, stated as fact paths that DO work all avoid system python ci.sh selects an interpreter through _select_python; the long-running suites session ran .venv/bin/python -m pytest and passed. So the automated regime is gre the DOCUMENTED operator command – an 11.4.238-shaped gap, since the escape found by an agent capturing evidence r by the regime. WORKAROUND (measured theorised):

ATM ID	Type	Status	Severity	Description
				PYTEST_DISABLE_PLUGIN_AUTOLOAD the needed plugins passed explicitly (-p pytest_timeout ...) avoids the schemathesis autoloader that drags in jsonschema -> re -> rpds. That is a workaround, NOT the silences the import path rather than re install, and it will surprise the next person runs the documented command verbatim. RELATED, and worth deciding together than twice: BOB-154 wants .venv rebuilt to match the container (which runs CPython 3.12.13 while both system and venv run Python 3.12.13). There are therefore THREE interpreters. Whoever rebuilds the venv should configure resolves for the TARGET interpreter after or this same class reappears one directly. HONEST BOUNDARY (11.4.6): this item NOT claim any test is wrong or any product is broken. The product and the venv-run are unaffected. What is broken is the documented entry point. Affected scope / file-scope manifest: host user site-packages (~/.local/lib/python3/site-packages/rpds/); every CLASH documented 'python3 -m pytest ...' invocation any tooling that uses system python3 rather than .venv/bin/python Reproduction / python3 -m pytest tests/e2e/test_live_stack_evidence.py -q -importmode=importlib -> ModuleNotFoundError: module named 'rpds.rpds', raised during collection via the schemathesis plugin and -> jsonschema -> referencing._core -> MEASURED root cause: system python3 but ~/.local/lib/python3/site-packages/rpds.cpython-313-x86_64-linux-gnu.so - extension built for 3.13. rpds/init.py does 'from .rpds import *', and a 313-tagged importable by 3.14, so the submodule does not exist for this interpreter. The v CORRECT and unaffected: .venv/lib64/python3/site-packages/rpds/ ships rpds.cpython-313-x86_64-linux-gnu.so and '.venv/bin/python3 -m pytest tests/unit/ -v -importmode=importlib - the command CLASH documents - reaches collection without ModuleNotFoundError, OR CLAUDE.m corrected to document the supported run explicitly. Either way the documented command and the working command agree (11.4.6).

ATM ID	Type	Status	Severity	Description
				guide that misleads is the documentation equivalent of a PASS-bluff). WORKAROUND: EFFECTS MEASURED 2026-08-21 — this is not free, and the item previously implemented PYTEST_DISABLE_PLUGIN_AUTOLOAD disables ALL plugin autoloader, not just the one, so anything the suite silently relied on be re-enabled by hand: - It BREAKS 5 plugin env-dependent tests in tests/unit/test_plugin_rutracker.py — TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace need an autoloader plugin. PROVEN CAUSED BY ANY IN-FLIGHT CHANGE: running OWN copy of that file under identical flags produces the same 5 failures. - It makes -timeout=60 (set in pyproject.toml) an UNKNOWN ARGUMENT unless -p pytest is passed explicitly, so pytest.mark.timeout is silently inert under the naive recipe — a believed to be time-bounded is not. - With autoloader ON and only schemathesis disabled (no:schemathesis), that same file is 96/99 CONSEQUENCE FOR ANYONE USING WORKAROUND: -p no:schemathesis alone is strictly better than blanket autoloader-disabled where it suffices, because it removes only the broken plugin. Where the blanket form -p pytest_timeout must be added or timeout is inert, and the 5 env-dependent failures are recognised as workaround artifacts rather than filed as defects. That last point is the reason the workaround manufactures failures that are exactly like product defects. This does not fix the root cause (the venv's rpds ships a cpython-313 ABI tag CPython 3.14 cannot handle) or the fix (rebuild the venv — BOB-154) documents that the interim recipe has a small radius, so nobody reads a workaround as a regression. WHAT. §11.4.19 requires closure migration ATOMIC: a resolving item moves to Fixed and DISAPPEARS from Issues_Summary (open) and APPEARS in Fixed_Summary (closed). Ten rows violate all three properties right now: they carry a terminal status yet current_location='Issues', so they render
BOB-166	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				Issues.md and are listed in docs/Issues_Summary.md with a status that reads '(→ Fixed.md)', while appearing 0 in Fixed_Summary.md. Any reader of the open tracker is told these are open work items. MANIFEST (measured 2026-08-21): BOB-129 (Completed), BOB-129, BOB-131, BOB-131, BOB-145, BOB-146, BOB-148, BOB-153, BOB-155, BOB-157 (all 'Fixed (→ Fixed.md)'). Confirmed rendering: BOB-155/087/157. grep 1 in Issues.md, 0 in Fixed.md, present in Issues_Summary.md, 0 hits in Fixed_Summary.md. ROOT CAUSE (reported at runtime, not inferred). 'close' performs atomic migration and REQUIRES -evidence. 'update -status' sets any §11.4.15 close with NO evidence and NO migration — it knows the location, because it prints it. COPY of the real DB: \$ workable-items BOB-065 -db repro.db -status 'Fixed (→ Fixed.md)' update: BOB-065 updated in (status=Fixed (→ Fixed.md), type=Task repro.db 'SELECT atm_id,status,current...' BOB-065 Fixed (→ Fixed.md) Issues. seam that is supposed to be the ONLY close path (close, evidence-gated per §11.4.14) has a parallel unguarded path that reaches same status while skipping BOTH the evidence requirement and the migration. WHY NOT CAUGHT IT (§11.4.238 coverage-escape). Three standing checks stay GREEN on a holding the forbidden row, verified on the poisoned copy from the repo root so no artifact is involved: (1) 'validate' → 'OK' items, all invariants satisfied' (it has no status↔location coherence invariant); (2) 'in sync' (DB and Markdown agree — or WRONG state; agreement is not correct). CM-CLOSURE-SEAM-BINDS CHECK A flags NON-terminal rows whose id appears work commit; these rows are terminal, are outside its predicate by construction was found by reading a status tally, not automated regime — a §11.4.238 discovery channel escape, which is itself a defect standing to the mis-located rows. ACCEPT (a) 'update' REFUSES a terminal status names 'close' as the correct path, with a §1.1 mutation proving the refusal (removal guard must make the mutation pass). (b)

ATM ID	Type	Status	Severity	Description
				‘validate’ grows a status↔location coherence invariant that FAILS on the forbidden state. a golden-bad fixture and a negative control legitimately terminal row in Fixed must — §11.4.201(1)). (c) The 10 existing rows drained to Fixed with class-matched evidence row, or, where a row’s evidence cannot be produced, honestly re-opened rather than migrated on a bare assertion. (d) Honest boundary: this closes the update-path here the detection gap; it does not claim even historical status write was evidence-backed CLAIMED. No fix is implemented by this. The 10 rows are untouched; draining them acceptance (c) and each needs its own evidence not a bulk UPDATE. WHAT. download-proxy exposes two Server Events routes. They are the same expensive — long-lived connections that hold a worker generator for their lifetime — but only one limit classed: routes.py:801 @router.get('/search/stream/{search_id}') @_rl('sse_stream') classed routes.py:150 @router.get('/theme/stream') async def stream_theme(...) <- limiter decorator @_rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/stream bound to the sse_stream class while /theme/stream falls through to the application class. Measured live on the running stack: /api/theme/stream reports x-ratelimit-limit 1. PROVENANCE + A CORRECTION WORKS. KEEPING (§11.4.6). This was surfaced by BOB-109 scaling agent, whose report from as: ‘sse_stream_limit_decorator is defined imported/applied nowhere ... SSE falls to the 120/minute default: 24x less protection framing is WRONG and was NOT filed as a bug. The agent searched for one symbol NAME wiring uses a different mechanism (@_rl('sse_stream')), and it IS applied — search/stream. Verified by reading routes.py:795-806 and the _rl helper at routes.py:46-51, with a control needle on other *_limit_decorator symbols show results in api/init.py so the zero-hit was not a bug search. The agent’s MEASUREMENT (100 theme/stream) was correct and is what makes this a real finding; its MECHANISM was halves are recorded so the next reader can re-derive the same wrong cause. WHY I
BOB-167	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-168	Task	In progress	Low	MATTERS. An unclassified SSE endpoint is the cheapest way to pin server resources: every connection is held open, and the default rate limit permits 120/min of them. The declared <code>sse_stream</code> class exists precisely because the route shape needs a tighter bound than the GETs. ACCEPTANCE. (a) A decision, recorded, whether <code>/theme/stream</code> belongs in the <code>sse_stream</code> class or genuinely warrants the default rate limit policy question, not automatically a bug. (b) If it belongs in <code>sse_stream</code>, the decorator is applied and a test drives BOTH SSE routes to assert each returns its INTENDED class. (c) A guard that enumerates SSE-shaped routes and fails if one without a class is caught. (d) A guard that carries no explicit rate-limit class — in general form, so the third SSE route does not repeat this. (e) Honest boundary: this does not claim 120/min is exploitable in this deployment with <code>network_mode</code> host and no reverse proxy; every caller shares the 127.0.0.1 bucket. BOB-111 measured and recorded separately. CLAIMED. No change made. The limit was never read from headers, never driven to exhaustion — the limiter is per-IP and shared with concurrent agents on this host (§11.4.201). WHAT. <code>scripts/run_all_challenges.sh:66</code> “<code>scaling_horizontal_challenge.sh</code>” in its <code>set</code>. <code>challenges/scripts/scaling_horizontal_challenge.sh</code> does not exist. <code>sed -n '66p' scripts/run_all_challenges.sh</code> “<code>scaling_horizontal_challenge.sh</code>” <code>\$ ls -l scripts/scaling_horizontal_challenge.sh</code> access ...: No such file or directory VERIFIED independently, not taken on report — see the BOB-109 scaling agent and re-checked by direct invocation. WHY IT MATTERS. this is cosmetic or a §11.4.201 gate-honour defect depends entirely on how the runner handles a missing entry, and that is the first thing to determine: if it SKIPS silently, the challenge advertises coverage it does not have (a claim-vs-reality row with no passing challenge behind it); if it FAILS, the runner is permitted for a reason unrelated to the system test, which trains readers to ignore it. No outcome is acceptable; they need different ACCEPTANCE. (a) Determine and record runner’s actual behaviour on the missing

ATM ID	Type	Status	Severity	Description
				invoking it, not by reading it. (b) EITHER the challenge, OR remove the entry — v §11.4.124 discipline: check git history for whether it once existed and was deleted. A silently-dropped challenge is the more i defect. (c) If the runner silently skips m entries, that is its own finding: a missing challenge must be loud (§11.4.3 SKIP-w at minimum), never absent-and-quiet. S Low as a defect, but it sits on the challe coverage seam, so (c) may deserve its o CORRECTION 2026-08-22 (§11.4.6) — T ITEM’S PREMISE WAS FALSE, and the mine. Recorded openly rather than quiet scaling_horizontal_challenge.sh EXISTS submodules/challenges/challenges/scripts/ scaling_horizontal_challenge.sh, execut bytes, dated Aug 7. So do all 14 other e the meta-runner. It was never deleted: a the roster at 1c959ba, the challenge its in the submodule at 873c0b1, and a pic search across both repositories and all finds ZERO deletions. HOW I GOT IT W run_all_challenges.sh reads submodules challenges/challenges/scripts/. I checke challenges/scripts/ — a DIFFERENT dir belonging to a DIFFERENT, glob-based Two similarly-named directories; I meas wrong one. That is a §11.4.201(9) field-i error. Worse, I wrote that it was “verifie invocation”: I had listed a directory and the aggregator, so the claim overstated evidence class as well as getting the an wrong. THE LESSON, which generalise item. In the SAME commit I correctly c sibling agent making this exact class of BOB-167 — a zero-hit produced by sear symbol name — and applied a control n there. Then I omitted the needle one pa later on my own measurement. The dist that would have caught it: A CONTROL PROVES THE INSTRUMENT CAN SEE, NOT PROVE IT IS POINTED AT THE TH UNDER TEST. §11.4.201(7)(b) as comm applied guards blindness; it does not gu aiming. The needle for THIS query wou been to resolve the path the runner itse not to confirm that some directory listin WHAT IS ACTUALLY WRONG, and it is than what I filed. The missing-entry pat

ATM ID	Type	Status	Severity	Description
				challenges by definition, so it can be ex safely. Run against a checkout where the submodule is absent — the ordinary sta clone made without -recursive — the R runner (byte-identity sha-asserted, not a reports: PASS: 0 FAIL: 0 SKIP: 16 TOTAL OBSERVED EXIT CODE: 0 An ENTIREL RUN BANK REPORTS SUCCESS. Not o dangling entry tolerated — the whole su silently passes without executing anythi is §11.4.201(6) false-null verbatim, and observed-reachable rather than constru (§11.4.115(G)). REVISED ACCEPTANCE ANSWERED, though not as filed: the ru reports a loud, counted SKIP and then r blocks, because the exit expression reac FAIL count. (b) The dangling-entry half nothing is dangling. (c) The real work is contract: a roster entry the runner cann execute must not be indistinguishable f healthy run. See the decision recorded BOB-168/decision_missing_vs_skip.md. REPORTED NOT FIXED: scripts/ pre_build_verification.sh:1792 carries t SKIP/MISSING conflation one level up, other runner. Filed separately. WHAT. The §11.4.65 markdown-export r requires every in-scope doc to ship .htm twins. 286 of the 326 .html files under c (excluding dist/ and node_modules/) are FRAGMENTS — they begin at WHAT. The BOB-109 scaling suite gates timing tests on loadavg_1m/nproc <= 0 this 8-cpu host, a 1-minute load average below 6.0 — and SKIPS loudly otherwise gate is correct and was adopted for a g (§11.4.201(8)): the author validated a 2 growth-exponent threshold, then observ FALSE-FAIL correct code at load 11.06 span 2.136), and at load 15-23 measure spans of 2.358 against injected-cubic m 2.278-2.331 — the signal is smaller than and the two are not separable in either Gating was the honest response. THE P The gate has consequently NEVER RUN own precondition. Every recorded run t measured load 9.15-23; the independent measured 9.66 during review; the cond measured 6.45 at its lowest. GREEN po observed — at loads 9.4-11.6 (span 1.88
BOB-169	Bug	In progress	Medium	
BOB-170	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-171	Bug	Queued	Low	three stability runs (1.759/1.832/1.848) never once beneath the shipped ceiling; author states this honestly as an inference on data rather than a run they can paste, a reviewer confirmed the inference is sound (quiescence is strictly more favourable) inference is still not an observed run. W MATTERS (§11.4.226). A registered, top present guard that never executes is indistinguishable from a guard that would have did. Nothing currently ensures it ever runs a freshness contract, no scheduled quiet-time invocation, no tracked item — which is the unexecuted-standing-guard class §11.4.226 names, where registration gets treated as coverage while execution is nobody's job. forensic precedent in that anchor is two guards that, when finally executed, immediately emitted latent FAILs. ACCEPTANCE. (a) both quiescence-gated timing tests on a run whose 1-minute loadavg is genuinely < 3.0 — no agent fleet, no concurrent build — capture the run as an artifact showing the resolved load reading alongside the mean span exponent, so the precondition is proven from the artifact and not asserted. (b) Record observed GREEN polarity in the item and test source, replacing the current honest inference note. (c) If the gate FAILS under genuine quiescence, that is the finding that exists to surface, and it reopens the threshold calibration question rather than being waved around. (d) Establish a freshness contract §11.4.226: how stale a quiescent verdict is before the gate counts as uncovered against this does not silently return to never-existence. HONEST BOUNDARY. This item does not say the gate is wrong. The evidence points the other way — quiescence is strictly more favourable than the loads where GREEN was already observed. It claims only that the run has happened, and that an unexecuted guard can be cited as coverage. PROVENANCE. Record IMPORTANT-2 by the independent Fable substrate reviewer of the BOB-109 world independently corroborated: the reviewer says the SKIP is loud and prints its resolved load and confirmed the shipped tests do SKIP today's load.

ATM ID	Type	Status	Severity	Description
				WHAT. Both rate limiters key their per-client bucket on the LEFTMOST element of X-Forwarded-For when TRUST_FORWARDER is enabled: download-proxy/src/api/rate_limit.py:117-123 (:7187) and the limiter in plugins/download_proxy.py (:7186) which was deliberately built to exact parity with it. That element is CLIENT-SUPPLIED, an honest reverse proxy APPENDS the peer's IP to the RIGHT of whatever arrived, so the entry is whatever the client sent — attached, not controlled by construction, not by misconfiguration. Consequences with this (1) rotating a forged leftmost value minimizes the unlimited sequence of fresh buckets, with a total bypass of the limit the flag is supposed to make MORE accurate; (2) it also defeats idle-reap cap, since each forged identity is a distinct key — the bucket map is a memory growth surface bounded only by the cap; the cap's eviction then discards the budgets of real clients to make room for forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, and that's why this is Low, not High. TRUST_FORWARDER defaults OFF for external sites, and the deployed stack runs networked on a host with no reverse proxy in front (very common across all five compose services), so peer addresses are real client IPs and no NAT occurs. The defect is latent: it arms the system if someone puts a proxy in front and turns on the flag to recover real client IPs — which is the situation the flag exists for. The failure is therefore 'correct-looking configuration silently disables the limiter', not 'current configuration is broken'. PROVENANCE. Raised as MINOR by the independent reviewer of the BOB-1000; the limiter work and independently confirmed by the implementing agent, which noted it at EBN source sites as a tracked follow-up rather than fixing it in that change. Filed here because neither agent has write access to the trunk. WHY IT COVERS BOTH PORTS. :7186's limiter was built to deliberate policy parity with :7187, including this parsing. Fixing one and not the other would leave the two surfaces disjoint on client identity — worse than the shared one because it becomes surface-dependent and untestable as a single invariant. ACCEPTED. (a) Replace leftmost-element parsing with

ATM ID	Type	Status	Severity	Description
				aware resolution at BOTH sites: either minus-N-trusted-hops, or a trusted-proxy allowlist where only a peer inside the allowlist may assert XFF at all — the choice is a deployment-topology decision and should be recorded, not assumed. (b) A test that rotates the leftmost element and asserts that the key does NOT rotate, per site. (c) Paired mutation: restore leftmost parsing and the test must FAIL. (d) Negative control (§11.4.2) with TRUST_FORWARDED_FOR OFF, the test must be ignored entirely and keying must be back to the peer address — a guard that honours XFF when the flag is off would be a worse defect than the one being fixed. (e) boundary: this closes header-forgery by ensuring that does not make per-IP limiting fair behind a proxy where many real users legitimately share an address. NOT CLAIMED. No change made; both sites still parse leftmost; the flag still does nothing. WHAT. The rutracker SEARCH endpoint is refused by bot protection while the site is not fully reachable. Measured unauthenticated on 2026-08-21 (no cookie file read, no credentials, no value in any artifact — §11.4.10): GET /index.php -> 200, 96283 B, 1.2s. WHAT. download-proxy/src/api/hooks.py:174 def _save_hooks(hooks: list[dict[str, Any]]): if hooks is None: try: os.makedirs(os.path.dirname(HOOKS_FILE), exist_ok=True) with open(HOOKS_FILE, 'w') as f: json.dump(hooks, f, indent=2) except Exception as e: logger.error(f'Failed to save hooks: {e}') catches EVERY exception, logs, and returns None. The signature returns None, so the caller has no channel to learn the write failed. The test sites then report success unconditionally. WHAT. _save_hooks(hooks):174 _save_hooks(hooks):175 logger.info('Created hook: ...'):175 logger.info('Deleted hook: ...'):156 return HookResponse(hook_id=..., ...):176 return {'message': 'Hook deleted', ...} USER-VISIBILITY CONSEQUENCE, which is why this is HIGH. This is not a code-hygiene nit. A user POSTs a hook and receives HTTP 200 and a hook_id, and the hook was never written — their automation script never fires, and the API told them it existed. Symmetrically, a user DELETES a hook, the API returns 'Hook deleted', the file still holds it, and the hook is still there again after the next restart. In both directions.
BOB-172	Bug	Ready for testing	High	
BOB-173	Bug	Ready for testing	High	

ATM ID	Type	Status	Severity	Description
				product reports the opposite of what has happened and the only trace is a log line nobody is watching. THIS IS THE §11.4.252 SHARPSHIELD path combines two dangerous capabilities: that anchor's taxonomy — MUTATION of a resource (a filesystem write) and EXTENDED SIDE EFFECT (hooks are outbound calls that the system will or will not make) — so it is not to FAIL CLOSED: verify the precondition when it cannot be satisfied, and surface the refusal. Instead it fails open, and a bare 'Exception:' that only logs is the exact anomaly §11.4.252 enumerates. At the product level, also a §11.4.201(6) false-null: a successful and a swallowed failure are indistinguishable to the caller. PROVENANCE. Surfaced as a scope observation by the independent review of the BOB-135 test-isolation work — this is what turned that defect into an 'assertion' mystery, because the EACCES on /conf was logged and discarded while the endpoint was returning 200. Verified here directly from the code before filing (the function body and both the read and write above), not taken on report. Record §11.4.238 discovery-channel escape: found in the agent reading code during an unrelated investigation, not by the automated QA. The coverage gap is a defect of equal stature to the defect itself. ACCEPTANCE. (a) _save propagates failure — raise, or return a success; callers must consume. (b) Both endpoints must translate a persistence failure into an HTTP 500-class, never a success body; a creation that did not persist must not return a hook. (c) The test drives each endpoint with the hook for unwritable (read-only dir or a patched directory raising OSError) and asserts a non-2xx response AND that a subsequent GET does not list the phantom hook — assert on the user-observed outcome, not on the log line. (d) Paired mutation: restore the swallow; the test fails. FAIL. (e) Audit the same file for sibling mutations — this is a pattern, and one instance is not alone. (f) Honest boundary: this does not mean the write currently fails in production; it means that WHEN it fails the user is told the outcome and the BOB-135 investigation shows it fails in at least one real environment. NOT CLOSURE. No change made. No assessment of how the write fails in the operator's deployment.

ATM ID	Type	Status	Severity	Description
BOB-174	Bug	In progress	High	RECORDING A SHELL ERROR OF MY C (§11.4.6): the first version of this description was written with the anti-pattern snippet inside backticks in a double-quoted shell argument. The shell ran it as command substitution and the text was replaced by nothing — the stored description read ‘and is the exact anti-pattern’. This is the SECOND time this session the backticks-inside-double-quotes has corrupted content (the first mangled a commit message). Fixed here by editing through a Python script with no shell quoting in the path. Noted as a silently-truncated defect description is the kind of quiet corruption §11.4.201(7) is about — the path is part of the instrumentation failed without erroring. WHAT. A corrupt hooks file is silently indistinguishable from “no hooks configured” and the next create then DESTROYS every existing hook while returning HTTP 200 OK. Defects on one path, filed together because any one alone leaves the data loss reachable — download-proxy/src/api/hooks.py:104. <code>_load_hooks</code> wraps the read in <code>except Exception: logger.error(...)</code> and falls through to read truncated or malformed <code>hooks.json</code> is then reported to every caller as an empty, <code>hook_id</code> list. Confirmed at source. A2 — the console and the reason this is High. <code>create_hook</code> appends, saves. Given a corrupt file that turns into <code>[]</code>, the save writes a one-element over the top. Every previously-configured gone. Measured by the implementing against the ALREADY-FIXED tree, so the the BOB-173 write-failure fix: BEFORE <code>['prod-hook-0', 'prod-hook-1', 'prod-hook-2']</code> reports : 200 <code>{'hooks': [], 'count': 0}</code> <- ZERO hooks configured DELETE report <code>{'detail': 'Hook not found'}</code> <- for a hook in the file POST reports : 200 <code>hook_id=3c8a5930-...</code> AFTER on disk : <code>['3c8a5930-...']</code> Three prod hooks destroyed HTTP 200 throughout, nothing surfaced to user. A5 — <code>_save_hooks</code> writes with a <code>pl</code> <code>open(path, "w")</code>. A crash, <code>ENOSPC</code>, or a write truncates the file in place. That is the corruption A1 then reads as “no hooks configured”. A2 overwrites. The same codebase already has the correct pattern: <code>theme_state.py:115</code> <code>tmp-file + os.replace</code>. WHY THE THREE

ATM ID	Type	Status	Severity	Description
				ONE ITEM. A5 manufactures the corrupt, misreads it as empty, A2 destroys the corrupt. Fixing only A1 leaves truncation reachable; only A5 leaves an existing corrupt file a trigger; fixing only A2 leaves the API lying about what is configured. The chain is the default. DELIBERATELY NOT FIXED UNDER BOB-173 and the reasoning is sound. The implementation agent identified all three while auditing the swallows as that item required, and decided to fix them in the same change because: the change GET semantics on an endpoint to the frontend consumes (200 -> 5xx); they need CORRUPT-file reproduction rather than UNWRITABLE-file one BOB-173 built; and need a design decision that is genuinely obvious — a MISSING hooks file legitimately means “no hooks configured”, while a CORRUPT one does not, and today’s code cannot tell them apart. Expanding BOB-173’s scope to corruption would have meant shipping that decision unexamined. ACCEPTANCE. (a) Distinguish MISSING from CORRUPT: a missing file is an empty list; a corrupt one is an error, reported silently []. (b) Decide and RECORD what the frontend does on corruption — 5xx, or 200 with a degraded marker the frontend can render. (c) the design decision, and it should be stated rather than inferred from whatever the frontend happens to do. (c) create/delete MUST NOT overwrite a file they could not parse — MUST NOT clobber (§11.4.252: mutation plus error side effect must fail closed). (d) Make the operation atomic via tmp + os.replace, reusing theme_state.py’s existing pattern rather than inventing it (§11.4.28). (e) Tests asserting a USER-OBSERVABLE outcome, not log level. (f) a corrupt file, assert GET does not claim to have hooks, assert a create refuses rather than destroying, and assert the file still holds the original hooks afterwards. (f) Paired §11.4.201(1): a mutation per guard. (g) NEGATIVE COMPARISON (§11.4.201(1)): a MISSING file must still be an empty list and a working create; a VALID file must behave exactly as today. A fix that makes every load fail closed by failing always is acceptable. A3, RECORDED SEPARATELY, NOT PART OF THIS CHAIN. VALID_EVENTS and HookEventType are two sources of truth. The closed set. Verified identical today, unguaranteed.

ATM ID	Type	Status	Severity	Description
				against drift: if they diverge a hook registers the direction successfully and then silently never fires. The assertion pinning them would close it. No change made. CLAIMED. No change made. The BOB-166 failure fix is real and orthogonal — it makes the FAILED write honest; it does nothing about the SUCCESSFUL write of wrong data derived from a misread file. WHAT. The workable-items update seams the status-location invariant in one direction. After BOB-166, update --status <terminal> Issues-located row is refused. The mirror image of the open: update --location Fixed --status <terminal> exits 0 and creates a row that the own validator immediately rejects. Verified empirically by the independent reviewer. BOB-166: update -location Fixed -status 'progress' -> exit 0 validate -> FAILS, no row (check (f), fixedLocationNonTerminal). So the command can still mint a state it validate call refuses. That is the §11.4.1 shape at the seam layer: the invariant is CONFIGURED in validate but not ENFORCED every write path that can violate it. WHY DELIBERATELY LEFT OPEN IN BOB-166 why that was right. Three reasons, all correct rather than asserted: (1) the real corpus has ZERO instances of this direction against the direction BOB-166 closed — the asymmetry the data matched the asymmetry in the detective coverage ALREADY existed and demonstrably fires, so unlike BOB-166's this cannot rot silently; (3) and the load one — reopenCmd documents and SANCTUARY transient Fixed-plus-non-terminal state location Fixed operator override. A naive preventive guard at the update seam would collide with that override's semantics. Correcting this therefore requires a design decision on how the two interact, which is a different question from the one BOB-166 was solved by answer, and expanding that item would have meant shipping the decision unexamined ACCEPTANCE. (a) Decide and RECORD the preventive guard on this direction coexists with reopenCmd's sanctioned override — this is actual work, and the answer should be verified down rather than inferred from whatever patch does. Candidates worth weighing the reopen path explicitly, require an over
BOB-175	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				flag on update mirroring reopen's, or leave in the direction detective-only with the rationale recorded so the asymmetry is a decision rather than an oversight. (b) If a guard lands, it's terminalStatuses() — BOB-166 establishes a predicate source specifically so the two cannot drift. (c) Paired §1.1 mutation. (d) NEGATIVE CONTROL (§11.4.201(1)), and a sharp one here: the sanctioned reopen j MUST still work. A guard that breaks reopenCmd's documented override is a positive refusal, and worse than the gap because it breaks a workflow operators to use. HONEST BOUNDARY. This is no data defect. Zero corpus instances, and a detective gate catches it at the next value filed so a scoped, deliberate omission stands as a tracked decision instead of living on code comment where the next reader will find it for an oversight (§11.4.197: a started reaches a documented terminal state, not abandonment). PROVENANCE. Raised by implementing agent of BOB-166 as a known asymmetry it deliberately did not close, independently verified and endorsed as an acceptable scope boundary by that item's reviewer, on the condition that it become a tracked item rather than a comment. This is the item. WHAT. _get_enabled_trackers() in down proxy/src/merge_service/search.py gate the rutracker on RUTRACKER_USERNAME and RUTRACKER_PASSWORD only. It never checks RUTRACKER_COOKIES. So an operator configured with cookies alone — no user password — never has rutracker enabled and the tracker is silently absent from evidence out rather than failing visibly. WHY THIS IS INCOHERENT WITH THE REST OF THE CODE. _search_rutracker treats COOKIES as the PREFERRED auth path, not a fallback. A sibling nnmclub DOES check its own _cookies variable. So the same codebase holds the positions at once: cookies preferred at the site, cookies ignored at the enablement, cookies honoured for a sibling tracker. WHY THIS MATTERS MORE THAN IT LOOKS. CLARIFICATION documents a standing operator mandate (2026-08-15) that per-tracker Netscape files at \${TRACKER_COOKIE_DIR}/
BOB-176	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				cookies_<tracker>.txt are auto-loaded in the configuration as <TRACKER>_COOKIES before every update, restart, install Stage 6 and start.sh is run, and rutracker is named in that set. So the configuration is SANCTIONED, documented configuration for this tracker produces exactly the state that the gate ignores. An operator who follows the documented instructions gets a tracker that never runs, with no error to explain it. It is in PROVENANCE MODE. Silent absence, not a visible failure. The same §11.4.201(6) false-null shape as BOB-172 one layer earlier. BOB-172 fixes a tracker that never RAN and was REFUSED being reported. This is a tracker that never ran at all, and the merged result simply has one fewer component with nothing to indicate it. PROVENANCE is handled by the BOB-172 implementing agent which is tracing the enablement path to locate the defect handling defect. Recorded here rather than in that change: it is a separate defect on a separate seam with its own oracle, and it is in would have widened BOB-172 past it with evidence. ACCEPTANCE. (a) The enablement gate honours RUTRACKER_COOKIES as an independently sufficient credential, matching what _search_rutracker already prefers to the nnmclub sibling already does. (b) A fix asserting that a cookies-only configuration ENABLES rutracker, driving the real _get_enabled_trackers rather than a replacement. Paired §1.1 mutation restoring the user password-only gate — the test must go to NEGATIVE CONTROL (§11.4.201(1)): a configuration with NO credentials at all leave rutracker DISABLED. A fix that enables an unauthenticated tracker is worse than the one because it produces failing searches instead of absent ones. (e) Audit the other trackers for the same asymmetry rather than fixing the one that was noticed — three positions in the codebase suggests nobody has checked them together (§11.4.118). HONEST BOUNDARY verified against a live cookies-only deployment read from source by the BOB-172 agent and recorded on its report. The reading is specific and checkable, but whoever takes this should be by invocation before relying on it. BC Bug Queued High WHAT: BOB-172 identical 4-line HTTP-refusal guard at file private-tracker fetch sites in download-

ATM ID	Type	Status	Severity	Description
				merge_service/search.py (rutracker cookie, rutracker credential :1468, kinozal :1641, nnmclub :1761, iptorrents :1934). Only rutracker COOKIE path is exercised by EVIDENCE (reviewer-authored mutation 11.4.194(6)(d), during the BOB-172 independent review): deleting ONLY the kinozal guard (search.py:1655-1658) while leaving the others intact left the full merge_service suite as passed, ZERO failures. The suite cannot pass all of the five sites. FAILING SCENARIO: a drop or subtly breaks the wiring at kinozal, nnmclub, iptorrents, or rutracker-credentials. 403 at that site silently folds back to status=empty with error=None -- the expected BOB-172 signature -- and no test reddens. DIRECTION (reviewer preferred): collapse the five duplicated wirings into one shared guard, e.g. <code>_check_search_response(tracker_name, status, body) -> bool</code>, so there is ONE condition and a sixth site cannot be added unguarded. (11.4.251 byte-identical-fork extraction). Alternative: parametrise the guard tests to cover all five sites with per-site stub sessions. ACCEPTANCE: mutating the wiring at ANY of the five sites reddens at least one test. Provided by running the same R1 deletion at each site. BOB-178 Bug Queued Medium Vulnerability download-proxy/src/merge_service/search.py:1638-1640 handles a failed kinozal LOGIN response with <code>if login_resp.status in (200, 301, 302): logger.error(...); return</code> sets NO diagnostic. FAILING SCENARIO: Cloudflare returns 403 on <code>takelogin.php</code>, kinozal chip reads <code>status=empty, error=</code> That is the BOB-172 signature exactly: reported to the user as an empty result. CONTRAST establishing this is an oversight is a design choice: the rutracker and nnmclub failures DO set diagnostics (<code>upstream_cookie_auth_failure</code>), and the iptorrents login failure through to the search fetch where BOB-172 guard catches it. Kinozal is the one leg out from neither. FIX DIRECTION: stash <code>_classify_upstream_http_status(login_response)</code> before the early return, or an <code>auth_failure</code> for the status-in-(200,301,302)-but-no-cookie case. ACCEPTANCE: a stubbed 403 on the kinozal login endpoint produces a non-Malicious on the kinozal chip, with a RED capture

ATM ID	Type	Status	Severity	Description
				the pre-fix code first (11.4.115). BOB-172 Task Queued Medium WHAT: the BOB-172 independent review demonstrated two failure paths that still report a refusing or unretried tracker as an empty result. Both are PR EXISTING and neither is a regression from BOB-172, but the fix evidence log presents five-site coverage without stating the bo (11.4.194(5) requires un-analysed dime explicit gaps, never silently assumed sa -- exception before resp.status is read. I probe B, captured: a connection-refused rutracker yields status=empty, error=No http_status=None, metadata.errors=[], metadata.status=completed. Each _search method swallows exceptions (except Ex logger.error; return results) before _search error handling can see them, so a track DOWN is indistinguishable from one that genuinely empty. GAP B -- soft refusal a 200. Reviewer probe A, captured: _classify_upstream_http_status(200, <code> body>) returns None, which is CORRECT design (a 2xx is a usable response; body triggering would risk the over-fire the m controls exist to prevent). But all five se use aiohttp's default allow_redirects=True session-expiry 302 -> login-page-200 ch to zero rows and reports empty. FIX DIF Gap A needs the exception to reach a di rather than being swallowed. Gap B nee DISTINCT detector (final-URL check or form marker) with its own RED -- explic widening of the status-code trigger. ACCEPTANCE: both gaps closed with th REDs, or explicitly closed per 11.4.112 evidence. Immediate sub-task: append a gaps paragraph to docs/qa/BOB-172/ fix_evidence_20260822.log. BOB-180 Queued Low WHAT: download-proxy/ routes.py:727-745 sets status=captcha with a hardcoded message naming RuTr 'RuTracker requires CAPTCHA. Use /api/ rutracker/captcha'. FAILING SCENARIO search returns zero results and the cap flavoured diagnostic came from NNMCI Turnstile, not RuTracker. The user is to a RuTracker captcha endpoint for an NI problem. SEVERITY BOUNDED: the ful and tracker_stats travel in the same res

ATM ID	Type	Status	Severity	Description
				payload, so the truth is present and a cl reads them is not misled -- only the head wrong. The branch was revived by BOB error propagation (correctly: suppressing propagation would recreate the same fa one layer up, 11.4.247) and is already c the routes layer by tests/unit/api_layer/ test_routes_coverage.py:782. FIX DIRE derive the message from the tracker(s) actually erred rather than hardcoding o SCOPE NOTE: api/ was owned by a sibl stream during BOB-172; the author was not to touch it. ACCEPTANCE: an NNM captcha diagnostic on a zero-result sear produces a headline naming NNMClub. BOB-181 Bug Queued High WHAT generate_markdown_exports.sh accepts arguments. It unconditionally scans PROJECT_ROOT/*.md, docs/ recursively scripts/ recursively (scan loops at lines 105/111/116 as of 2026-08-23; they wer when this was filed — line numbers shif the BOB-169 fix, the substance is uncha was independently verified by execution the generator with an explicit path proc out-of-scope export and NOT the request and PROJECT_ROOT is derived from BASH_SOURCE (line 17). A caller who v 'bash scripts/generate_markdown_expo docs/qa/SOME/file.md' gets a WHOLE-T with the argument silently discarded. W MATTERS: the caller reasonably believe was scoped to one file. It was not. In a s checkout with parallel work streams thi regenerates any export whose .md is ne its derived files anywhere in the tree -- artifacts in other streams' scopes and, v BOB-068 auto-commit hazard, into anot stream's commit. CLASS: a false afforda script's interface implies a capability it have (11.4.201: an interface must mean claims). FIX DIRECTION: either accept path arguments and honour them (scop run to exactly those files), or REFUSE v clear message when arguments are pas Silently discarding them is the one opti must not remain. ACCEPTANCE: passin either scopes the run to it, or exits non- naming the unsupported argument. A R capturing today's silent-discard behavior

ATM ID	Type	Status	Severity	Description
BOB-182	Task	Queued	Medium	(11.4.115). LIVE INSTANCE, MEASURED 2026-08-23 (this is no longer hypothetical). During the BOB-169 fix the conductor in the generator with an explicit path to re-ONE document. The argument was discarded the whole tree was scanned. It generated qa/BOB-174/DESIGN_DECISION.{html, at 11:50 -- a SIBLING STREAM's document that stream's author was still editing the (.md mtime 12:00). The exports were the stale on arrival, and additionally carried fix BOB-169 charset defect. The sibling independently reported them as stale in hand-off, having no idea another stream created them. SECOND-ORDER LESSON keeping with this item: the conductor's verify the blast radius ALSO failed, and silently. 'git status --porcelain grep -E pdf)\$"' returned only its own files and not clean - but docs/qa/BOB-174/ is an UNT DIRECTORY, which porcelain collapses line, so the grep was structurally incapable seeing the files inside it (11.4.201(7)(a) the wrong surface; the quiet zero read absence). The instrument that DOES see WHAT: CM-EXPORT-CHARSET-VALID (p invariant 50) adopts its 301 pre-existing violations via a monotone-decrease rate than a hard floor. Two things are owed. ADOPTION DECISION IS THE OPERATION NOT THE SCRIPT'S. 11.4.224(E) requires brownfield adoption question be SURFACE the operator and the answer recorded a consumer DATA - never an invented rate. What exists today is a header comment BOBA_EXPORT_CHARSET_BASELINE of that DOCUMENTS a decision the agent does not RECORD an answer the operator. The independent reviewer judged the change defensible and would not reverse it (a hope on 301 violations makes the build unrealistic which 11.4.234 forbids; 11.4.101 favours reversible option) but correctly held that yet compliant. Options for the operator: immediate hard floor (BASELINE=0) / keep monotone ratchet / per-corpus phase-in changed-code-only with a scheduled full deadline. (2) THE RATCHET DOES NOT ACTUALLY RATCHET. BASELINE is a hope constant. On improvement the gate PAS

ATM ID	Type	Status	Severity	Description
				PRINTS the value to lower it to, but not lowers it, so the gate keeps permitting all the way back to 301 after the corpus. The header now says tightening is MAN rather than claiming automatic behavior does not exist (11.4.6), but the honest is a stopgap, not the fix. WHY IT WAS N WITH THE FIX: a gate that writes its own threshold during a pre-build run becomes PRODUCER as well as a GATE (11.4.24 separation), and that is a design change operator should approve rather than real side effect. ACCEPTANCE: the operator adoption answer recorded as consumer and either a persisted baseline the gate and never raises (with the role-separation question answered), or an explicit decision manual tightening is acceptable. Three icon-glyph controls (.bridge-retry toggle, .caret) render text made of non-points. axe-core declines to decide their reporting them under the incomplete change with messageKey nonBmp and resolving NEITHER fgColor NOR bgColor, so the recomputation in docs/qa/BOB-164/axe_contrast_scan.py cannot decide either such nodes across 16 palette x mode combinations. BOB-164 round 2 stopped being scored as silent passes by naming GLYPH_UNMEASURED, printing them on run and fencing them, so any NEW unmeasured glyph node blocks. What remains genuine unproven is their contrast itself: as non-interface components they owe 3:1 under SC 1.4.11, and no oracle in this repo measures that today. Acceptance: a measurement resolves the effective foreground and background for glyph controls (for example reading getComputedStyle colour against the color backdrop, or replacing the glyphs with ones carrying explicit fill tokens), each of the controls shown to clear 3:1 in all 16 palette mode combinations, and the GLYPH_UNMEASURED fence shrunk by the nodes then proven. docs/qa/BOB-164/axe_contrast_scan.py static dist with no backend, so every node existence depends on fetched data is absent the page it measures: the results table remains empty, so .type-badge and .quality-badge
BOB-184	Task	Queued	Medium	
BOB-185	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
BOB-186	Bug	Queued	—	appear, and the hooks list renders empty so .status pills never appear. That is precisely why the BOB-164 round-1 regression was unseen in a rendered DOM even though scanned in all 16 themes, and it is why I had to close the class in the arithmetic class instead. Those nodes are now covered by frontend/src/app/models/style-contrast.js which extracts declared foreground-on-background from the stylesheets, but that is an ARITHMETIC claim about declared colour pairs, not a rendered-pixel one: it cannot see opacity by an ancestor, a composited backdrop, cascade this repo's static extractor does not model. Acceptance: the scan drives the dashboard with seeded fixture data (a seeded backend, a route fixture, or an injected component harness) so the badge and status nodes are present in the DOM, axe measures them directly, and a control needle provides rendered oracle sees a seeded sub-floor before any clean result from it is believed workable-items diff returned exit 0 and verdict "DB and Markdown are in sync" PARTIAL read: invoked with -issues but -fixed it compared 77 Markdown items 184 DB items and called that in sync, precisely both mismatched counts in its own output 107 DB items were never compared against anything. Measured 2026-08-25: (A) no markdown paths -> correctly REFUSES (B) -db-only -> honest, "no Markdown compared"; (C) -issues alone -> exit 0, (compared 77 Markdown item(s) against item(s))". The 11.4.201(6) false-null: a previously blind instrument returned the same quietly genuinely-synced corpus returns. IMPACT CORRECTED 2026-08-25, my original fix OVERSTATED this. I wrote that CHECK 11.4.106(F) commit seam was exposed. NOT. scripts/hooks/docs-sync-commit-seam.sh:277 passes BOTH -issues and -a repo-wide sweep found NO executable that breaks: the two real callers (that see pre_build_verification.sh:549) both pass paths; the single-path hits are prose evidence files plus a static-scan test fixture. So the was real and latent, never live in this previously corrected claim is: any FUTURE single-caller would have been silently mis-told

ATM ID	Type	Status	Severity	Description
BOB-187	Bug	In progress	—	PROVENANCE, twice corrected. A suba attributed the false-null to the no-paths which actually refuses correctly; the co re-measured and relocated it to the par Then the fixing agent corrected the con impact. Both corrections are recorded s wrong version propagates. FIXED (unco constitution submodule, fetched to upst 7f16739 before edit per 11.4.26). Desig by default (11.4.201(4)); the narrow per comparison PRESERVED behind an exp -partial-scope rather than deleted (dele existing documented capability would b 11.4.122 silent removal), mirroring how preserved the no-Markdown shape. Key MEASURED DB content, never on flag p — a DB with zero Fixed rows IS fully ac for by -issues alone and still passes, so keying would have traded the false-null 11.4.201(1) false-positive refusal. Resul property: “in sync” is reachable only wh two printed counts account for each oth fail/5 pass -> GREEN 9/9; paired 1.1 mu the fixs own revert, killing exactly the 4 guarding tests while the 5 preservation green. Case C now refuses naming “107 located in Fixed (supply -fixed)”, the 10 confirmed against an independent sqLite (107 Fixed + 78 Issues = 185). Incident closed an 11.4.238 gap: case As refusal test anywhere; TestDiffCmd_NoPathsSti now guards it. OPEN: -partial-scope is t naming choice, not an operator decision renaming is cheap now and expensive a consumers adopt it. Other consumers o shared submodule were not swept. scripts/ownership_precondition.sh cons SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe it, but it is NOT fenced. T028 remediati ownership_path_fence() to scripts/lib/ ownership.sh and wired it into ownership_repair.sh, closing the path-es class there; the precondition was outsid agents file scope and still accepts whate config/owned_paths.yaml expands to. B is lower than a recursive chown (it writ files rather than mutating ownership of existing tree) but the INPUT is identical equally unreviewed, so the same

ATM ID	Type	Status	Severity	Description
BOB-188	Bug	In progress	—	QBITTORRENT_DATA_DIR value that will drive a filesystem-wide chown will drive file creation at an arbitrary location. No escape is empirically confirmed for the path, not merely reasoned: during T028 remediation a probe declaring the verb shipped entry with QBITTORRENT_DATA_DIR hung the suite, and ps showed the real process executing “find / (! -uid 1000 -o ! -gid 1000 -printf ...” before it was killed by verification. Acceptance: ownership_precondition.sh SAME ownership_path_fence() predicate scripts/lib/ownership.sh (never a second 11.4.251), refuses fail-closed on a rejection and ships a paired 1.1 mutation proving refusal fires PLUS a negative control proving six live shipped entries still ACCEPT so not a 11.4.201(1) false-positive refusal. Discovered out-of-band during T028 remediation and reported by the agent as needing tracking; it is also a 11.4.238 coverage escape. pre_build_verification.sh invariant 17 (CONSTANT-ITEMS-VALIDATE) resolves the binary through the candidate loop at :50. FIRST entry is constitution/scripts/workable-items/bin/workable-items. That file is GTRACKED (md5 17644a248363, identical to workable-items-linux) and executable, so resolution over the current untracked script constitution/scripts/workable-items/workable-items (md5 43376a6d0184). The tracker is STALE: it does not contain the guards it source now has. MEASURED, needle-probe 2026-08-25. String presence in the tracker binary: “refusing to set terminal status” “Issues-location item has TERMINAL status” control needle from the SAME updateControl; at least one mutable field flag is required” the instrument sees through that path and zeros are real absences, not a blind read. untracked sibling returns 1 / 1 / 2 for the three queries. A check whose message is absent from the binary cannot fire. RUN PROOF: the same injected violation (a read of terminal status while current_location=run through both binaries -> stale repository violation (only the body_md desync class reports 2 including the location-status check verbatim. This is exactly why the ten BC rows

ATM ID	Type	Status	Severity	Description
BOB-189	Bug	Queued	—	(BOB-087/129/131/144/145/146/148/150) survived undetected: the gate that was to catch them was running a binary structure incapable of seeing them. The 11.4.108 >ARTIFACT gap: source correct, artifact every gate reading the artifact reports q This is the THIRD instance of that class one session, alongside BOB-169 (a character that never opened a file) and BOB-183 (a bundle five days older than its sources). COMPOUNDING: the comment at :519- directly above the loop asserts “The naïve workable-items path never existed in the checkout (bin/ is a gitignored local build dir nothing ever populated)”. That state is now FALSE - bin/ holds two tracked binaries, a stale comment asserting the absence of a file that now shadows resolution is how it stayed invisible. Also an 11.4.30 question: who owns: bin/workable-items and workable-items-linux are VERSIONED BY ARTIFACTS in a constitution submodule inherited by reference by every consumer; every consumer inherits whichever binary was last committed. Flagged independently by separate agents this session. ACCEPTABLE resolution must never prefer a stale artifact over a current one - either the tracked binary is removed and resolution falls through to demand, or a freshness check refuses anything older than its sources (see the BOB-183 fingerprint gate for a working pattern); (c) the false comment at :519-523 corrected; (d) the 1.1 mutation proving the new refusal fix is a negative control proving a genuinely fresh one still passes so the fix is not an 11.4.2016 positive refusal; (d) the 11.4.30 tracked decision recorded as operator DATA. Discovered out-of-band during BOB-166 remediation 11.4.238 coverage escape in its own right. WHAT: the fail-open scanner’s shape-(A) heuristic cannot distinguish ‘return False’ meaning PROCEED-AS-IF-FINE from ‘return False’ meaning REFUSE. It flags six text CLOSED guards as fail-open defects: <code>_is_safe_fetch_url</code> (x3), <code>_qbit_add_success</code> and the hooks path-boundary guard. INDEPENDENTLY VERIFIED 2026-08-20 conductor rather than taken on the trial word: <code>download-proxy/src/api/routes.py</code>

ATM ID	Type	Status	Severity	Description
BOB-190	Bug	Queued	—	defines <code>_is_safe_fetch_url</code> returning <code>bool</code> only call site at :1476 reads 'if not <code>_is_safe_fetch_url(url): logger.warning(I</code> SSRF-unsafe download URL (non-public skipping); continue' - a False return RE the URL. WHY THIS MATTERS: per §11.4.99 a false-positive refusal is a FAIL-bluff of severity to a false-negative pass, and he consequence is worse than noise - acting finding would mean making the SSRF g returning False, i.e. deleting an SSRF p to satisfy a gate. A gate that instructs y remove a security control is actively dan not merely imprecise. REPRO: run the f scan; observe six hits; read each call sit ACCEPTANCE: the scanner distinguishe shaped from proceed-shaped False retu site-aware, or an audited waiver list wit entry justification), the six drop out, AN golden-FALSE fixture containing a real guard is added so the discrimination is falsifiable per §1.1. WHAT: <code>~/local/lib/python3/site-packages</code> contains binaries built for the <code>cpython-3</code> while the interpreter is Python 3.14, so <code>pydantic_core</code> and <code>rpds</code> fail to import and test importing FastAPI dies at import time. CONFIRMED NOT OURS: an untouched <code>tests</code> fails identically, so this is environmental existing. WHY IT MATTERS: CLAUDE.n documents 'python3 -m pytest tests/unit -import-mode=importlib' as the canonical invocation and that documented command currently cannot run - docs and host dis §11.4.99 misguidance class at the environment layer. <code>scripts/run-tests.sh</code> already sidestep selecting <code>.venv/bin/python</code>, so a working <code>tests</code> exists and simply is not what the docs tell reader to type. IMPACT: anyone following CLAUDE.md literally concludes the suite is broken; worse, an agent could chase gr rewriting tests. ACCEPTANCE: either re host site-packages so the documented o works, or correct CLAUDE.md + docs/ TESTING.md to name the venv interpreter canonical - decided explicitly, not left to hits it next. Discovered out-of-band by t open triage agent, so per §11.4.238 this a coverage-escape note: no automated c

ATM ID	Type	Status	Severity	Description
BOB-191	Bug	Queued	—	asserts the documented test invocation runs. WHAT: the §11.4.252 fail-open scan report for qBitTorrent-go, and that zero is NOT evidence. The triage agent ran control through the scanner's own path per §11.4.252 (b): a Python 'except Exception: pass' needle SEEN, a TypeScript 'catch (e) {}' needle SEEN (so frontend/src's zero IS a real zero). Go empty-'if err != nil {}' needle was NOT. An instrument that cannot see the idiom of the same quiet zero as a clean tree, and one of those is honest. DISTINCT FROM BOB-189 deliberately not merged with it per §11.4.252. BOB-189 is a false MATCH (fail-closed q reported as fail-open); this is a false NU (entire language unscanned). Same scan in opposite failure directions, different fix. Merging them would hide one behind the other. IMPACT: Go is the language of qbittorrent-go and boba-jackett (port 7189, which contains encrypted tracker credentials), so the user surface is exactly where §11.4.252's creation plus-mutation combination is most likely to have been assessed it; the dashboard says clearly. ACCEPTANCE: the scanner grows a Go needle covering the empty-err-block and swallows idioms, its needle is SEEN through the qBitTorrent-go's result is re-derived, and the finding is triaged as this Python/TS pass. Until then qBitTorrent-go's fail-open position is UNKNOWN and must be reported as UNKNOWN never as 0.