

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	8
Bug	Operator-blocked	1
Bug	Queued	18
Bug	Ready for testing	4
Task	In progress	10
Task	Operator-blocked	3
Task	Queued	24
Task	Ready for testing	2
TOTAL		70

Items

ATM ID	Type	Status	Severity	Description
BOB-008	Bug	Operator-blocked	—	RuTracker automated login blocked by
BOB-065	Task	Queued	High	Lava P2: Egress diagnosis and VPN-hos routing (containers pkg/egress)
BOB-066	Task	In progress	High	Lava P3: BOBA_UPSTREAM_PROXY in proxy + qBitTorrent-go + Jackett + com forward
BOB-068	Bug	In progress	High	RD2-00: unattributed, unreviewed Auto mechanism pushing to main
BOB-069	Task	In progress	High	RD2-40: §11.4.238 QA-discovery-channel — ongoing coverage-escape backfill
BOB-074	Task	In progress	Medium	RD2-07: DDoS-class testing fully absent mandated test-type matrix
BOB-077	Task	Operator-blocked	High	

ATM ID	Type	Status	Severity	Description
				RD2-10: Identify second host running the commit rsync/sync mechanism (OPERATOR DECISION)
BOB-078	Task	Queued	High	RD2-11: Once identified, wire Auto-commit mechanism through §11.4.234 dedicated push script OR retire it
BOB-080	Task	Queued	High	RD2-13: Retroactive Fable-xhigh code re the substantive Auto-commit diffs
BOB-082	Task	Queued	High	RD2-15: Create BOB-064..067 workable the four Lava-porting findings (closes G
BOB-085	Task	Queued	Medium	RD2-18: Create top-level Boba proxy/me service v1.0.0 readiness ledger (closes C
BOB-088	Task	Queued	Medium	RD2-21: Complete/verify README Trac + Status Documents table row-completo (GA-07 remainder)
BOB-090	Task	In progress	Medium	RD2-25: HelixQA Challenge entry exerc three start.sh subcommands end-to-end real compose stack
BOB-093	Task	In progress	High	RD2-28: Live compose bring-up + verify ReDoS regex bounds deployed to contain (closes runtime half of GA-12)
BOB-094	Task	In progress	Medium	RD2-29: Author tests/stress/ test_tracker_fetch_stress_chaos.py with fault injection
BOB-095	Task	In progress	Medium	RD2-30: Author tests/stress/ test_scheduler_hooks_sse_stress_chaos. side triangle
BOB-097	Task	Queued	Low	RD2-32: Author DDoS-class coverage for download-proxy/merge endpoints (canonical of RD2-07)
BOB-100	Task	Queued	Low	RD2-39: Bump submodules/jackett one (canonical impl of RD2-09)
BOB-101	Task	Operator-blocked	High	GA-19/RW-09: Is -profile go parity still a goal? (gates RW-10..13) — OPERATOR-I
BOB-102	Task	Operator-blocked	Medium	RW-05: LAN-exposure threat model — b 127.0.0.1 or keep 0.0.0.0? — OPERATOR DECISION
BOB-104	Task	In progress	Medium	§11.4.238 followup: CodeGraph 1.5.0 nested-.gitignore regression challenge
BOB-106	Task	Queued	Medium	§11.4.238 followup: §11.4.84 quiescence helper for the unattributed auto-commit
BOB-107	Task	Queued	Medium	§11.4.238 followup: pre-dispatch existence for subagent task-brief source inputs
BOB-109	Task		Medium	

ATM ID	Type	Status	Severity	Description
		Ready for testing		BOB-074 followup: scaling-class test coverage absent from mandated test-type matrix
BOB-110	Task	Queued	Medium	BOB-074 followup: UX-class test coverage (accessibility/usability) absent
BOB-111	Task	In progress	High	BOB-074 followup: configure real rate limit for boba's 3 public HTTP endpoints
BOB-114	Task	Queued	Medium	BOB-074 followup: self-validation golden fixture for the rate-limit detector
BOB-120	Bug	Queued	Critical	3rd forced-logout incident 2026-08-18 20:00: SIGKILL user@1000 + preventive-timer for user-slice architectural gap
BOB-121	Task	Ready for testing	Important	External watchdog for the forced-logout architectural gap (task #85, incident #3)
BOB-135	Bug	Ready for testing	Low	Test isolation: test_list_hooks_after_creation bulk suite (Permission denied /config)
BOB-137	Bug	In progress	High	Merge service on 7187 wedges while the process still serves 7186 (GIL starvation on spinning thread)
BOB-141	Task	Queued	Low	CLAUDE.md claims the Go profile server serves 7186/7187/7188 but its container binds to 7186 — doc contradicts the Dockerfile
BOB-143	Bug	Queued	Medium	Orphaned .worktrees/ dirs (46M, unresolved gitdir) pollute gate scan scope and manifest false BOB-126-class findings
BOB-149	Bug	Queued	Low	Managed-plugin count diverges 43/42/41 from constitution, CLAUDE.md and the REAL badge; the badge is hand-maintained and unguarded
BOB-150	Task	Queued	Low	pre_build_verification.sh invariant labelled 44 but only 35 invariants are labelled
BOB-151	Task	Queued	Medium	CM-SCRIPT-DOCS-SYNC is a named gate implementation, and 24 of 36 scripts have companion doc
BOB-152	Bug	Queued	Medium	Constitution sweep walks vendored third-party code: 82% of one gate's 38,291 findings come from submodules/
BOB-154	Bug	Ready for testing	Medium	Host venv and production container run different starlette versions (1.4.1 vs 1.6.0)
BOB-156	Bug	Queued	Medium	BOB-145 event-loop regression guard is too sensitive and flaky: 8786ms under host 900-1500ms ceiling calibrated on a quiet machine
BOB-158	Bug	Queued	High	tests/conftest.py cannot run on the production interpreter: binds to asyncio.events._get_event_loop_policy, a private API, while production is 3.12.13

ATM ID	Type	Status	Severity	Description
BOB-159	Bug	Queued	Medium	Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:01:17Z Reported-By: independent review (IMPORTANT-2), partially remediated What (the report, verbatim): -recreate path was fixed this round: run_ownership_precondition -> stack_down, run_ownership_repair -> stack_up, so the walk walks a tree no container can write to. The ordering is now asserted behaviourally with new checks in tests/unit/test_start_reload_recreate.sh, each killed and paired 1.1 mutation (11/11 RED-OK). The path is NOT covered by that fix and this tracks the remainder. On a warm ./start gate runs before THIS invocation brings up, but if the stack is ALREADY running containers write while the repair walks. A container can then create a new non-owned file BEHIND the walk, after which the completion marker records complete ownership that is not, and those stragglers are never repaired without a manual -force. BOU dismissed: after any successful pass the marker is valid for the scope fingerprint and the repair short-circuits without walking (marker_is_valid so the window requires a stale-or-absent AND a live stack together. Declaring a new entry re-arms the marker, which is exact that combination arises in practice. NOT THIS ROUND, with the reason stated rather hidden: the clean guard is a cheap marker probe mode on scripts/ownership_repair.sh answers has-this-scope-been-repaired without walking the tree. The existing -dry-run serve: it deliberately SKIPS the marker (FORCE=0 && DRY_RUN=0 guard) and walks, so using it as a warm-start probe walk the tree twice on every start. Additional mode requires editing ownership_repair.sh another stream was actively editing during round; duplicating marker_is_valid into instead would be a near-identical fork (Deferred to this item rather than done unilaterally or silently. The misleading comment at the warm-start call site (which claimed gate runs before any container writes) has been corrected to state this boundary honestly. Affected scope / file-scope manifest: (warm-start dispatch, run_ownership_gate site); scripts/ownership_repair.sh (marker

ATM ID	Type	Status	Severity	Description
BOB-160	Task	Queued	Medium	path) Reproduction / context: 1. Ensure stack is UP. 2. Change config/owned_path so the scope fingerprint changes (this requires marker by design, data-model E2). 3. Run warm ./start.sh (no -recreate). 4. The ownership repair walks and chowns the declared test containers are still running and able to run it. Acceptance criteria: A warm ./start.sh complete an ownership repair while a container that writes to a declared location is running. Either the repair is deferred with an acceptance refusal naming -recreate, or the stack is up for the walk. Asserted behaviourally in test_start_reload_recreate.sh alongside existing PRECONDITION_BEFORE_DOWN / REPAIR_AFTER_DOWN / REPAIR_BEFORE checks, each with a paired 1.1 mutation to it. Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:02:48Z Reported-By: What (the report, verbatim): Independent §11.4.209 review (IMPORTANT-3) found the feature's strongest automated check never executed by any runner: (1) tests/ownership/test_container_writes_owned_files — the §11.4.115 RED-turned-regression FR-011/FR-007, proven via a docker-container revert mutation — was moved out of test integration/ (commit 58d340a, to escape autouse fixture hang) but no runner (ci.test.sh, run-all-tests.sh) was ever extended to cover tests/ownership/. (2) tests/pre_build/test.sh, the §1.1 paired-mutation meta-test the scripts/pre_build/check_cm.sh gate was executed by NOTHING — self-reported honestly in commit 04742d7's own message ('STILL OPEN (reported, not fixed): test_pre_build/ is executed by NOTHING') but tracked as a workable item (a §11.4.197 requirements gap) and never wired into pre_build_verification.sh invariant 30, which globbed only tests/unit/test.sh. Affected file-scope manifest: ci.sh; scripts/pre_build_verification.sh invariant 30 (CONTAINER_UNIT-TESTS-EXECUTED) Reproduction context: grep -rn test_container_writes_owned_files ci.sh run-all-tests.sh scripts/ returned zero results (only docs/specs mentioned the path); grep scripts/pre_build_verification.sh showed

ATM ID	Type	Status	Severity	Description
				30's for-loop globbing only tests/unit/te never tests/pre_build/test_.sh. Accepta criteria: ci.sh gains a runtime-gated py ownership/ stage (skips honestly, rc=0, podman/docker present); scripts/ pre_build_verification.sh invariant 30's covers both tests/unit/test_.sh and tests pre_build/test_*.sh, verified by an extra standalone run of the invariant's exact l reporting a non-zero RAN count that inc files from both directories. Reported-Via: §11.4.202 reporting dire on 2026-08-21T19:31:12Z Reported-B BOB-092 remediation, residual finding V report, verbatim): §11.4.69 names CM OPEN-SKIP as one of three mandatory p gates: it audits sink-side probe helpers if any code path converts an empty or unreachable response into a PASS-count for a feature class with a sink-side prob project does not have it. Why this matte rather than in the abstract: the BOB-09 remediation just removed TWO fail-open exactly this class from tests/e2e/ test_live_stack_evidence.py — the nnmc on-404 (already removed at 7baef2b) an surviving sibling at test_iporrents_is_authenticated_in_sea skipped on 'not authenticated and statu success' while asserting 'treating as tra outage'. That assertion was false for rej credentials (upstream_http_403) and fo container (plugin_env_missing), both of definitive product failures the product a classifies via error_type. Two fail-opens class, found one at a time, is the §11.4.1 extend-to-all-cases signal and the §11.4. coverage-escape signal together: the re not surface either, an agent reading cod remediation's own guards are AST-struc lethal under three discriminating mutat they LIVE IN THE FILE THEY GUARD, deleting that file evades them entirely. A gate is the durable home because it auc corpus rather than one file. Honest bou (§11.4.6): this item does NOT claim the remediated fail-opens are unguarded — guarded, with runtime evidence (13 pas 97.36s against the live stack). It claims CLASS has no corpus-wide detector, so
BOB-161	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
BOB-162	Task	Queued	Medium	instance in a different file is invisible again. Affected scope / file-scope manifest: pre_build/ (gate absent); tests/e2e/test_live_stack_evidence.py (guards current inside the file they guard) Reproduction context: grep -rl 'CM-NO-FAIL-OPEN-SKIP' scripts/ tests/ returns exactly ONE hit, at tests/e2e/test_live_stack_evidence.py — the guards live in, not a gate. Control not the same query for CM-OWNERSHIP-INVALID (a gate that does exist) returns 3 files, so the instrument can see gate tokens and the gate is a real absence, not a blind zero (§11.4.201(b)). Acceptance criteria: A scripts/pre_build gate named CM-NO-FAIL-OPEN-SKIP exists, wired into scripts/pre_build_verification.py audits sink-side probe helpers for code errors, converting an empty/unreachable/error into a PASS-counting SKIP for a feature that HAS a sink-side probe. It ships a golden fixture (a real fail-open -> gate FIRES) and a golden-FALSE-with-carrier (an honest test geo skip, and a comment merely MENTION the phrase -> gate MUST NOT fire), per §11.4.107(10)/§11.4.201(1). Paired §11.4.202 makes the gate FAIL before the gate is reported. Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:35:02Z Reported-By: BOB-106/BOB-107 remediation, residual (the report, verbatim): The TESTS for guards are now executed: pre-build invalid glob was extended from tests/unit + tests/pre_build to also cover tests/hooks. BEFORE/after: the glob expanded to 27 subjects, 0 under tests/hooks while 3 existed there; expands to 30 with 3 under tests/hooks. pass. But the GUARDS themselves are still invoked by nothing. check-brief-inputs.sh is honestly conductor-run rather than auto-commit brief's required-input list lives in free-form, that the Agent tool's structured input does expose, so an automatic prompt-scraping would itself be an unproven pattern-matching seam is the dispatch procedure, and that documentation + habit change, not a honest unattributed-commit-guard.sh is different why this item exists. Run against real history names 14 violating commits since the last (and 21 bare Auto-commit subjects exist in all refs). Wiring it into the pre-build or c

ATM ID	Type	Status	Severity	Description
				seam AS-IS would therefore refuse even immediately, on pre-existing debt nobody in the moment. That is precisely the 11.4.135 brownfield-adoption shape, and 11.4.66 put the choice with the operator, not with agent inventing a ratchet. The options, the decision is a decision and not a default immediate hard floor – the seam refuses 14 are attributed; (b) one-time monotonic decrease ratchet (the 11.4.135 pattern) snapshot 14 as the baseline, the count may and may never rise, so day one is green disease cannot spread; (c) changed-commit – gate only commits created from now on scheduled deadline for the historical 14 report-only for a stated period, then escape Recommendation, with the reasoning rather just the pick: (b). It is the pattern this commit already uses for exactly this situation, if the debt visible and bounded immediately cannot be satisfied by deleting the guard (11.4.227 closes that gaming channel). the operator’s call and this item is not closed until that answer is recorded. Honest Bob (11.4.6): this item does NOT claim the 14 commits are harmful in themselves – it is their ATTRIBUTION is absent and that is currently prevents the 15th. Affected s file-scope manifest: scripts/hooks/unattributed-commit-guard.sh; scripts/hooks/check-bob-inputs.sh; scripts/pre_build_verification.sh; scripts/commit-push-all.sh Reproduction context: Both guards run correctly by hand are covered by passing tests (9/9 and 13 proven non-bluff against a no-op stub). 14 invoked by scripts/pre_build_verification.sh; scripts/commit-push-all.sh, so neither escape standing detection pressure (11.4.226: 14 with no execution seam is not coverage registration is not coverage). Acceptance criteria: Each guard is invoked by a new seam, OR carries a registered deferral point this item. For unattributed-commit-guard specifically, the operator has answered adoption question below and the answer is recorded as consumer DATA – never an ratchet. Reported-Via: §11.4.202 reporting directly on 2026-08-21T19:41:08Z Reported-By: BOB-114 remediation, pre-existing defect
BOB-163	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				surfaced by BOB-111 landing a real limit. (the report, verbatim): PRE-EXISTING INTRODUCED – and proven so rather than asserted: git diff shows assertion (c) is IDENTICAL to HEAD, and the failure reverts against the unmodified HEAD script. What changed is not the challenge but the SYNOPSIS. BOB-111 appears at least partially delivered because :7187 now returns real 429s where previously returned none. :7185 and :7189 produce zero 429s. This is a 11.4.248-class corrosion risk and that is why it is filed rather than left as a footnote: run_all_challenges now fail INTERMITTENTLY, and an intermediate red trains everyone to re-run until green, which point a real regression is dismissed as ‘probably the flaky rate-limit one’. THE DECISION, which is an operator call under 11.4.66 and which I deliberately did NOT make. Does a 429 from a WORKING limiter count as a sibling endpoint being RESPONSIVE? (a) YES – 429 proves the endpoint is alive and correct, protecting itself. Assertion (c) accepts 2xx, 429, and only a connection failure / 5xx counts as degraded. Risk: a genuinely working endpoint that happens to answer 429 with (b) NO – keep requiring 2xx, but make the challenge limiter-aware: drain or wait out the window before the sibling probe (retry-served, so the budget is knowable rather than guessed). (c) Probe siblings on a path that exempts (a healthz-class route), so the intermediate question is asked without spending the budget. Recommendation with reasoning, not just a decision. (a) composed with a bounded liveness of 429 carrying a well-formed Retry-After header as evidence of a live, correctly-behaving service. (b) makes the challenge slower and could increase the window value that will drift. But the service is ‘responsive’ here are a product judgement. The answer is recorded, not assumed. RELATIONSHIP stated as fact rather than folded in: the challenge counts 429 only, not 503, though the script header says ‘429 (or equivalent)’. Counting 503 would collide with the crash detector’s logic. Worth deciding when BOB-111 lands line :7185 and :7189. Affected scope / file-manifest: challenges/scripts/ddos_resilience_challenge.sh assertion endpoint isolation; run_all_challenges.sh

ATM ID	Type	Status	Severity	Description
BOB-164	Bug	In progress	Medium	invokes it Reproduction / context: As requires a sibling-endpoint probe to ans 2xx). :7187 now enforces a real limiter (live: x-ratelimit-limit: 120, x-ratelimit-re 86, retry-after: 45, server: uvicorn). A full challenge run sends roughly 250 requests :7187, so sibling probes fired during the endpoints' tiers land on an exhausted backend receive 429 - which assertion (c) scores 'endpoint degraded'. Timing-dependent measured PASS=5 FAIL=2; the UNMOOD HEAD script against the same live stack later measured PASS=7 FAIL=0 SKIP=. Acceptance criteria: The operator has the classification question below, the answer recorded as consumer DATA, and assertion implements it. The challenge then produces the same verdict across 3 consecutive runs on a rate-limited stack (11.4.50 deterministic consistency), and the fix ships a paired mutation proving assertion (c) still catches a genuinely degraded sibling endpoint - not it must not blind it (11.4.201(1)). Reported-Via: §11.4.202 reporting direct on 2026-08-21T19:56:53Z Reported-By: BOB-110 UX-class coverage, discovered a new axe-core suite on its first live run Verdict, report, verbatim): This is a REAL user defect, not a test-tuning artifact, and it is by the automated regime rather than by squinting at the page - which is exactly 11.4.238 posture the project is aiming for. static-grep oracle would NOT have found it. Measured: the served root ships an empty hydration, so any check reading the raw audits a page nobody sees. The violation exists in the hydrated DOM, which is what 11.4.170 requires a rendered oracle and value-equality assertions as the proof of correct. Severity reasoning, stated rather assumed: this is user-visible and affects primary dashboard heading, but it degrades legibility rather than breaking function, surface is operator-facing rather than public Medium, not High. The failing test was FAILING on purpose (11.4.238) instead silenced or marked xfail. tests/ux/ current reports 1 failed, 16 passed; that 1 is this. Anyone reading a red UX suite should recognize this item, not as flakiness - and when the

ATM ID	Type	Status	Severity	Description
				the suite goes fully green, which is the s it is closed. HONEST SCOPE LIMIT (11 the dashboard landing view was scanned jackett/* sub-routes and the ng-serve-ho :4200 route set were NOT audited – the commands to close both are recorded in testing/ux_accessibility.md. So this item nodes are a floor, not a total. Affected s file-scope manifest: the Angular dashb served at http://localhost:7187/ (same c SPA as frontend/); production componen not test files Reproduction / context: ux/test_live_dashboard_accessibility.py the running merge service. axe-core v4. scanning a real Playwright-rendered JS DOM, reports color-contrast violations o nodes. Measured pairs: .brand / h1 text on background #3c3f41 = 1.43-1.62:1 (l large-text floor is 3:1); tagline #808080 #3c3f41 = 2.68:1 (body-text floor is 4.5 Acceptance criteria: axe-core reports color-contrast violations against the live dashboard, with the fix made in product component CSS rather than by relaxing assertion or excluding the rule (11.4.12 reconcile to the correct mechanism, nev weaken the check). The existing tests/u the guard and already fails today, so the captured – closure requires it flipping G against the live surface, which is runtim evidence per 11.4.226. Reported-Via: §11.4.202 reporting dire on 2026-08-21T20:00:08Z Reported-By surfaced while capturing closure eviden BOB-092; diagnosed rather than worke What (the report, verbatim): This is a ABI-AFTER-INTERPRETER-UPGRADE c a missing package. The package is insta compiled half is built for the previous in That distinction matters because ‘pip in py’ style advice can appear to succeed v leaving the 313 artefact in place. WHY NOT NOTICED EARLIER, stated as fact paths that DO work all avoid system pyt ci.sh selects an interpreter through _select_python; the long-running suites session ran .venv/bin/python -m pytest a passed. So the automated regime is gre the DOCUMENTED operator command – an 11.4.238-shaped gap, since the esc
BOB-165	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				found by an agent capturing evidence r by the regime. WORKAROUND (measur theorised): PYTEST_DISABLE_PLUGIN_AUTOLOAD the needed plugins passed explicitly (-p pytest_timeout ...) avoids the schemath autoload that drags in jsonschema -> re -> rpds. That is a workaround, NOT the silences the import path rather than rep install, and it will surprise the next pers runs the documented command verbatim RELATED, and worth deciding together than twice: BOB-154 wants .venv rebuild to match the container (which runs CPy 3.12.13 while both system and venv run There are therefore THREE interpreters Whoever rebuilds the venv should confi resolves for the TARGET interpreter aft or this same class reappears one direct HONEST BOUNDARY (11.4.6): this item NOT claim any test is wrong or any pro is broken. The product and the venv-run are unaffected. What is broken is the do entry point. Affected scope / file-scope manifest: host user site-packages (~/.l python3/site-packages/rpds/); every CL documented 'python3 -m pytest ...' invo any tooling that uses system python3 ra than .venv/bin/python Reproduction / python3 -m pytest tests/e2e/ test_live_stack_evidence.py -q -import- mode=importlib -> ModuleNotFoundError module named 'rpds.rpds', raised durin collection via the schemathesis plugin a -> jsonschema -> referencing._core -> r MEASURED root cause: system python3 but ~/.local/lib/python3/site-packages/r rpds.cpython-313-x86_64-linux-gnu.so - extension built for 3.13. rpds/init.py do 'from .rpds import *', and a 313-tagged importable by 3.14, so the submodule g does not exist for this interpreter. The v CORRECT and unaffected: .venv/lib64/p site-packages/rpds/ ships rpds.cpython- x86_64-linux-gnu.so and '.venv/bin/pytho import rpds' succeeds. Acceptance cri python3 -m pytest tests/unit/ -v -import- mode=importlib - the command CLAUD documents - reaches collection without ModuleNotFoundError, OR CLAUDE.mo

ATM ID	Type	Status	Severity	Description
				corrected to document the supported ru explicitly. Either way the documented c and the working command agree (11.4. guide that misleads is the documentation equivalent of a PASS-bluff). WORKAROUND EFFECTS MEASURED 2026-08-21 — th is not free, and the item previously impl PYTEST_DISABLE_PLUGIN_AUTOLOAD disables ALL plugin autoload, not just th one, so anything the suite silently relied be re-enabled by hand: - It BREAKS 5 p env-dependent tests in tests/unit/ test_plugin_rutracker.py — TestConfig::test_default_mirrors, test_env_mirrors_override, test_get_env_with_default, test_get_mirrors_from_env_empty, test_get_mirrors_from_env_whitespace need an autoloaded env plugin. PROVE caused by any in-flight change: running OWN copy of that file under identical fl produces the same 5 failures. - It makes -timeout=60 (set in pyproject.toml) an UNKNOWN ARGUMENT unless -p pyte is passed explicitly, so pytest.mark.time silently inert under the naive recipe — a believed to be time-bounded is not. - W autoloading ON and only schemathesis dis (no:schemathesis), that same file is 96/9 CONSEQUENCE FOR ANYONE USING WORKAROUND: -p no:schemathesis alon strictly better than blanket autoload-dis where it suffices, because it removes on broken plugin. Where the blanket form -p pytest_timeout must be added or timo inert, and the 5 env-dependent failures recognised as workaround artifacts rathe filed as defects. That last point is the re the workaround manufactures failures t exactly like product defects. This does n the root cause (the venv's rpds ships a cpython-313 ABI tag CPython 3.14 cann or the fix (rebuild the venv — BOB-154) documents that the interim recipe has a radius, so nobody reads a workaround a a regression. WHAT. §11.4.19 requires closure migrat ATOMIC: a resolving item moves to Fixe DISAPPEARS from Issues_Summary (op and APPEARS in Fixed_Summary (close
BOB-166	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				Ten rows violate all three properties right now, but they carry a terminal status yet <code>current_location='Issues'</code>, so they render in <code>Issues.md</code> and are listed in <code>docs/Issues_Summary.md</code> with a status that <code>location</code> reads <code>'(→ Fixed.md)'</code>, while appearing 0 in <code>Fixed_Summary.md</code>. Any reader of the open tracker is told these are open work items. MANIFEST (measured 2026-08-21): BOB-065 (Completed), BOB-129, BOB-131, BOB-132, BOB-145, BOB-146, BOB-148, BOB-153, BOB-155, BOB-157 (all 'Fixed (→ Fixed.md)'). Confirmed rendering: BOB-155/087/157/158/159/160/161/162/163/164/165/166/167/168/169/170/171/172/173/174/175/176/177/178/179/180/181/182/183/184/185/186/187/188/189/190/191/192/193/194/195/196/197/198/199/200/201/202/203/204/205/206/207/208/209/210/211/212/213/214/215/216/217/218/219/220/221/222/223/224/225/226/227/228/229/230/231/232/233/234/235/236/237/238/239/240/241/242/243/244/245/246/247/248/249/250/251/252/253/254/255/256/257/258/259/260/261/262/263/264/265/266/267/268/269/270/271/272/273/274/275/276/277/278/279/280/281/282/283/284/285/286/287/288/289/290/291/292/293/294/295/296/297/298/299/300/301/302/303/304/305/306/307/308/309/310/311/312/313/314/315/316/317/318/319/320/321/322/323/324/325/326/327/328/329/330/331/332/333/334/335/336/337/338/339/340/341/342/343/344/345/346/347/348/349/350/351/352/353/354/355/356/357/358/359/360/361/362/363/364/365/366/367/368/369/370/371/372/373/374/375/376/377/378/379/380/381/382/383/384/385/386/387/388/389/390/391/392/393/394/395/396/397/398/399/400/401/402/403/404/405/406/407/408/409/410/411/412/413/414/415/416/417/418/419/420/421/422/423/424/425/426/427/428/429/430/431/432/433/434/435/436/437/438/439/440/441/442/443/444/445/446/447/448/449/450/451/452/453/454/455/456/457/458/459/460/461/462/463/464/465/466/467/468/469/470/471/472/473/474/475/476/477/478/479/480/481/482/483/484/485/486/487/488/489/490/491/492/493/494/495/496/497/498/499/500/501/502/503/504/505/506/507/508/509/510/511/512/513/514/515/516/517/518/519/520/521/522/523/524/525/526/527/528/529/530/531/532/533/534/535/536/537/538/539/540/541/542/543/544/545/546/547/548/549/550/551/552/553/554/555/556/557/558/559/560/561/562/563/564/565/566/567/568/569/570/571/572/573/574/575/576/577/578/579/580/581/582/583/584/585/586/587/588/589/590/591/592/593/594/595/596/597/598/599/600/601/602/603/604/605/606/607/608/609/610/611/612/613/614/615/616/617/618/619/620/621/622/623/624/625/626/627/628/629/630/631/632/633/634/635/636/637/638/639/640/641/642/643/644/645/646/647/648/649/650/651/652/653/654/655/656/657/658/659/660/661/662/663/664/665/666/667/668/669/670/671/672/673/674/675/676/677/678/679/680/681/682/683/684/685/686/687/688/689/690/691/692/693/694/695/696/697/698/699/700/701/702/703/704/705/706/707/708/709/710/711/712/713/714/715/716/717/718/719/720/721/722/723/724/725/726/727/728/729/730/731/732/733/734/735/736/737/738/739/740/741/742/743/744/745/746/747/748/749/750/751/752/753/754/755/756/757/758/759/760/761/762/763/764/765/766/767/768/769/770/771/772/773/774/775/776/777/778/779/780/781/782/783/784/785/786/787/788/789/790/791/792/793/794/795/796/797/798/799/800/801/802/803/804/805/806/807/808/809/810/811/812/813/814/815/816/817/818/819/820/821/822/823/824/825/826/827/828/829/830/831/832/833/834/835/836/837/838/839/840/841/842/843/844/845/846/847/848/849/850/851/852/853/854/855/856/857/858/859/860/861/862/863/864/865/866/867/868/869/870/871/872/873/874/875/876/877/878/879/880/881/882/883/884/885/886/887/888/889/890/891/892/893/894/895/896/897/898/899/900/901/902/903/904/905/906/907/908/909/910/911/912/913/914/915/916/917/918/919/920/921/922/923/924/925/926/927/928/929/930/931/932/933/934/935/936/937/938/939/940/941/942/943/944/945/946/947/948/949/950/951/952/953/954/955/956/957/958/959/960/961/962/963/964/965/966/967/968/969/970/971/972/973/974/975/976/977/978/979/980/981/982/983/984/985/986/987/988/989/990/991/992/993/994/995/996/997/998/999/1000/1001/1002/1003/1004/1005/1006/1007/1008/1009/1010/1011/1012/1013/1014/1015/1016/1017/1018/1019/1020/1021/1022/1023/1024/1025/1026/1027/1028/1029/1030/1031/1032/1033/1034/1035/1036/1037/1038/1039/1040/1041/1042/1043/1044/1045/1046/1047/1048/1049/1050/1051/1052/1053/1054/1055/1056/1057/1058/1059/1060/1061/1062/1063/1064/1065/1066/1067/1068/1069/1070/1071/1072/1073/1074/1075/1076/1077/1078/1079/1080/1081/1082/1083/1084/1085/1086/1087/1088/1089/1090/1091/1092/1093/1094/1095/1096/1097/1098/1099/1100/1101/1102/1103/1104/1105/1106/1107/1108/1109/1110/1111/1112/1113/1114/1115/1116/1117/1118/1119/1120/1121/1122/1123/1124/1125/1126/1127/1128/1129/1130/1131/1132/1133/1134/1135/1136/1137/1138/1139/1140/1141/1142/1143/1144/1145/1146/1147/1148/1149/1150/1151/1152/1153/1154/1155/1156/1157/1158/1159/1160/1161/1162/1163/1164/1165/1166/1167/1168/1169/1170/1171/1172/1173/1174/1175/1176/1177/1178/1179/1180/1181/1182/1183/1184/1185/1186/1187/1188/1189/1190/1191/1192/1193/1194/1195/1196/1197/1198/1199/1200/1201/1202/1203/1204/1205/1206/1207/1208/1209/1210/1211/1212/1213/1214/1215/1216/1217/1218/1219/1220/1221/1222/1223/1224/1225/1226/1227/1228/1229/1230/1231/1232/1233/1234/1235/1236/1237/1238/1239/1240/1241/1242/1243/1244/1245/1246/1247/1248/1249/1250/1251/1252/1253/1254/1255/1256/1257/1258/1259/1260/1261/1262/1263/1264/1265/1266/1267/1268/1269/1270/1271/1272/1273/1274/1275/1276/1277/1278/1279/1280/1281/1282/1283/1284/1285/1286/1287/1288/1289/1290/1291/1292/1293/1294/1295/1296/1297/1298/1299/1300/1301/1302/1303/1304/1305/1306/1307/1308/1309/1310/1311/1312/1313/1314/1315/1316/1317/1318/1319/1320/1321/1322/1323/1324/1325/1326/1327/1328/1329/1330/1331/1332/1333/1334/1335/1336/1337/1338/1339/1340/1341/1342/1343/1344/1345/1346/1347/1348/1349/1350/1351/1352/1353/1354/1355/1356/1357/1358/1359/1360/1361/1362/1363/1364/1365/1366/1367/1368/1369/1370/1371/1372/1373/1374/1375/1376/1377/1378/1379/1380/1381/1382/1383/1384/1385/1386/1387/1388/1389/1390/1391/1392/1393/1394/1395/1396/1397/1398/1399/1400/1401/1402/1403/1404/1405/1406/1407/1408/1409/1410/1411/1412/1413/1414/1415/1416/1417/1418/1419/1420/1421/1422/1423/1424/1425/1426/1427/1428/1429/1430/1431/1432/1433/1434/1435/1436/1437/1438/1439/1440/1441/1442/1443/1444/1445/1446/1447/1448/1449/1450/1451/1452/1453/1454/1455/1456/1457/1458/1459/1460/1461/1462/1463/1464/1465/1466/1467/1468/1469/1470/1471/1472/1473/1474/1475/1476/1477/1478/1479/1480/1481/1482/1483/1484/1485/1486/1487/1488/1489/1490/1491/1492/1493/1494/1495/1496/1497/1498/1499/1500/1501/1502/1503/1504/1505/1506/1507/1508/1509/1510/1511/1512/1513/1514/1515/1516/1517/1518/1519/1520/1521/1522/1523/1524/1525/1526/1527/1528/1529/1530/1531/1532/1533/1534/1535/1536/1537/1538/1539/1540/1541/1542/1543/1544/1545/1546/1547/1548/1549/1550/1551/1552/1553/1554/1555/1556/1557/1558/1559/1560/1561/1562/1563/1564/1565/1566/1567/1568/1569/1570/1571/1572/1573/1574/1575/1576/1577/1578/1579/1580/1581/1582/1583/1584/1585/1586/1587/1588/1589/1590/1591/1592/1593/1594/1595/1596/1597/1598/1599/1600/1601/1602/1603/1604/1605/1606/1607/1608/1609/1610/1611/1612/1613/1614/1615/1616/1617/1618/1619/1620/1621/1622/1623/1624/1625/1626/1627/1628/1629/1630/1631/1632/1633/1634/1635/1636/1637/1638/1639/1640/1641/1642/1643/1644/1645/1646/1647/1648/1649/1650/1651/1652/1653/1654/1655/1656/1657/1658/1659/1660/1661/1662/1663/1664/1665/1666/1667/1668/1669/1670/1671/1672/1673/1674/1675/1676/1677/1678/1679/1680/1681/1682/1683/1684/1685/1686/1687/1688/1689/1690/1691/1692/1693/1694/1695/1696/1697/1698/1699/1700/1701/1702/1703/1704/1705/1706/1707/1708/1709/1710/1711/1712/1713/1714/1715/1716/1717/1718/1719/1720/1721/1722/1723/1724/1725/1726/1727/1728/1729/1730/1731/1732/1733/1734/1735/1736/1737/1738/1739/1740/1741/1742/1743/1744/1745/1746/1747/1748/1749/1750/1751/1752/1753/1754/1755/1756/1757/1758/1759/1760/1761/1762/1763/1764/1765/1766/1767/1768/1769/1770/1771/1772/1773/1774/1775/1776/1777/1778/1779/1780/1781/1782/1783/1784/1785/1786/1787/1788/1789/1790/1791/1792/1793/1794/1795/1796/1797/1798/1799/1800/1801/1802/1803/1804/1805/1806/1807/1808/1809/1810/1811/1812/1813/1814/1815/1816/1817/1818/1819/1820/1821/1822/1823/1824/1825/1826/1827/1828/1829/1830/1831/1832/1833/1834/1835/1836/1837/1838/1839/1840/1841/1842/1843/1844/1845/1846/1847/1848/1849/1850/1851/1852/1853/1854/1855/1856/1857/1858/1859/1860/1861/1862/1863/1864/1865/1866/1867/1868/1869/1870/1871/1872/1873/1874/1875/1876/1877/1878/1879/1880/1881/1882/1883/1884/1885/1886/1887/1888/1889/1890/1891/1892/1893/1894/1895/1896/1897/1898/1899/1900/1901/1902/1903/1904/1905/1906/1907/1908/1909/1910/1911/1912/1913/1914/1915/1916/1917/1918/1919/1920/1921/1922/1923/1924/1925/1926/1927/1928/1929/1930/1931/1932/1933/1934/1935/1936/1937/1938/1939/1940/1941/1942/1943/1944/1945/1946/1947/1948/1949/1950/1951/1952/1953/1954/1955/1956/1957/1958/1959/1960/1961/1962/1963/1964/1965/1966/1967/1968/1969/1970/1971/1972/1973/1974/1975/1976/1977/1978/1979/1980/1981/1982/1983/1984/1985/1986/1987/1988/1989/1990/1991/1992/1993/1994/1995/1996/1997/1998/1999/2000/2001/2002/2003/2004/2005/2006/2007/2008/2009/2010/2011/2012/2013/2014/2015/2016/2017/2018/2019/2020/2021/2022/2023/2024/2025/2026/2027/2028/2029/2030/2031/2032/2033/2034/2035/2036/2037/2038/2039/2040/2041/2042/2043/2044/2045/2046/2047/2048/2049/2050/2051/2052/2053/2054/2055/2056/2057/2058/2059/2060/2061/2062/2063/2064/2065/2066/2067/2068/2069/2070/2071/2072/2073/2074/2075/2076/2077/2078/2079/2080/2081/2082/2083/2084/2085/2086/2087/2088/2089/2090/2091/2092/2093/2094/2095/2096/2097/2098/2099/2100/2101/2102/2103/2104/2105/2106/2107/2108/2109/2110/2111/2112/2113/2114/2115/2116/2117/2118/2119/2120/2121/2122/2123/2124/2125/2126/2127/2128/2129/2130/2131/2132/2133/2134/2135/2136/2137/2138/2139/2140/2141/2142/2143/2144/2145/2146/2147/2148/2149/2150/2151/2152/2153/2154/2155/2156/2157/2158/2159/2160/2161/2162/2163/2164/2165/2166/2167/2168/2169/2170/2171/2172/2173/2174/2175/2176/2177/2178/2179/2180/2181/2182/2183/2184/2185/2186/2187/2188/2189/2190/2191/2192/2193/2194/2195/2196/2197/2198/2199/2200/2201/2202/2203/2204/2205/2206/2207/2208/2209/2210/2211/2212/2213/2214/2215/2216/2217/2218/2219/2220/2221/2222/2223/2224/2225/2226/2227/2228/2229/2230/2231/2232/2233/2234/2235/2236/2237/2238/2239/2240/2241/2242/2243/2244/2245/2246/2247/2248/2249/2250/2251/2252/2253/2254/2255/2256/2257/2258/2259/2260/2261/2262/2263/2264/2265/2266/2267/2268/2269/2270/2271/2272/2273/2274/2275/2276/2277/2278/2279/2280/2281/2282/2283/2284/2285/2286/2287/2288/2289/2290/2291/2292/2293/2294/2295/2296/2297/2298/2299/2300/2301/2302/2303/2304/2305/2306/2307/2308/2309/2310/2311/2312/2313/2314/2315/2316/2317/2318/2319/2320/2321/2322/2323/2324/2325/2326/2327/2328/2329/2330/2331/2332/2333/2334/2335/2336/2337/2338/2339/2340/2341/2342/2343/2344/2345/2346/2347/2348/2349/2350/2351/2352/2353/2354/2355/2356/2357/2358/2359/2360/2361/2362/2363/2364/2365/2366/2367/2368/2369/2370/2371/2372/2373/2374/2375/2376/2377/2378/2379/2380/2381/2382/2383/2384/2385/2386/2387/2388/2389/2390/2391/2392/2393/2394/2395/2396/2397/2398/2399/2400/2401/2402/2403/2404/2405/2406/2407/2408/2409/2410/2411/2412/2413/2414/2415/2416/2417/2418/2419/2420/2421/2422/2423/2424/2425/2426/2427/2428/2429/2430/2431/2432/2433/2434/2435/2436/2437/2438/2439/2440/2441/2442/2443/2444/2445/2446/2447/2448/2449/2450/2451/2452/2453/2454/2455/2456/2457/2458/2459/2460/2461/2462/2463/2464/2465/2466/2467/2468/2469/2470/2471/2472/2473/2474/2475/2476/2477/2478/2479/2480/2481/2482/2483/2484/2485/2486/2487/2488/2489/2490/2491/2492/2493/2494/2495/2496/2497/2498/2499/2500/2501/2502/2503/2504/2505/2506/2507/2508/2509/2510/2511/2512/2513/2514/2515/2516/2517/2518/2519/2520/2521/2522/2523/2524/2525/2526/2527/2528/2529/2530/2531/2532/2533/2534/2535/2536/2537/2538/2539/2540/2541/2542/2543/2544/2545/2546/2547/2548/2549/2550/2551/2552/2553/2554/2555/2556/2557/2558/2559/2560/2561/2562/2563/2564/2565/2566/2567/2568/2569/2570/2571/2572/2573/2574/2575/2576/2577/2578/2579/2580/2581/2582/2583/2584/2585/2586/2587/2588/2589/2590/2591/2592/2593/2594/2595/2596/2597/2598/2599/2600/2601/2602/2603/2604/2605/2606/2607/2608/2609/2610/2611/2612/2613/2614/2615/2616/2617/2618/2619/2620/2621/2622/2623/2624/2625/2626/2627/2628/2629/2630/2631/2632/2633/2634/2635/2636/2637/2638/2639/2640/2641/2642/2643/2644/2645/2646/2647/2648/2649/2650/2651/2652/2653/2654/2655/2656/2657/2658/2659/2660/2661/2662/2663/2664/2665/2666/2667/2668/2669/2670/2671/2672/2673/2674/2675/2676/2677/2678/2679/2680/2681/2682/2683/2684/2685/2686/2687/2688/2689/2690/2691/2692/2693/2694/2695/2696/2697/2698/2699/2700/2701/2702/2703/2704/2705/2706/2707/2708/2709/2710/2711/2712/2713/2714/2715/2716/2717/2718/2719/2720/2721/2722/2723/2724/2725/2726/2727/2728/2729/2730/2731/2732/2733/2734/2735/2736/2737/2738/2739/274

ATM ID	Type	Status	Severity	Description
BOB-167	Bug	Queued	Medium	names 'close' as the correct path, with a §1.1 mutation proving the refusal (removal of guard must make the mutation pass). (b) 'validate' grows a status↔location coherence invariant that FAILS on the forbidden state; a golden-bad fixture and a negative control legitimately terminal row in Fixed must — §11.4.201(1)). (c) The 10 existing rows drained to Fixed with class-matched evidence row, or, where a row's evidence cannot be produced, honestly re-opened rather than migrated on a bare assertion. (d) Honest boundary: this closes the update-path hence the detection gap; it does not claim even historical status write was evidence-backed CLAIMED. No fix is implemented by this. The 10 rows are untouched; draining them acceptance (c) and each needs its own evidence not a bulk UPDATE. WHAT. download-proxy exposes two Server Events routes. They are the same expensive — long-lived connections that hold a word generator for their lifetime — but only one limit classed: routes.py:801 @router.get('/search/{search_id}') @_rl('sse_stream') classed routes.py:150 @router.get('/theme/stream') async def stream_theme(...) <- limiter decorator _rl(cls) resolves to limiter.limit(limit_for(cls)), so /search/stream bound to the sse_stream class while /theme/stream falls through to the application class. Measured live on the running stack: /api/theme/stream reports x-ratelimit-limit 1. PROVENANCE + A CORRECTION WORKING KEEPING (§11.4.6). This was surfaced by BOB-109 scaling agent, whose report from as: 'sse_stream_limit_decorator is defined but not imported/applied nowhere ... SSE falls to the 120/minute default: 24x less protection framing is WRONG and was NOT filed as such. The agent searched for one symbol NAME wiring uses a different mechanism (@_rl('sse_stream')), and it IS applied — search/stream. Verified by reading routes.py:795-806 and the _rl helper at routes.py:46-51, with a control needle on other *_limit_decorator symbols show results in api/init.py so the zero-hit was not a bug search. The agent's MEASUREMENT (100 theme/stream) was correct and is what

ATM ID	Type	Status	Severity	Description
BOB-168	Task	In progress	Low	this a real finding; its MECHANISM was halved, and the next reader will re-derive the same wrong cause. WHY IT MATTERS. An unclassified SSE endpoint is the cheapest way to pin server resources: each connection is held open, and the default limit permits 120/min of them. The declared sse_stream class exists precisely because the route shape needs a tighter bound than GETs. ACCEPTANCE. (a) A decision, recorded whether /theme/stream belongs in the sse_stream class or genuinely warrants the default policy question, not automatically a bug. (b) If it belongs in sse_stream, the decorator is applied and a test drives BOTH SSE routes to assert each returns its INTENDED class from x-ratelimit-limit, so a future route without a class is caught. (c) A guard that enumerates SSE-shaped routes and fails if one that carries no explicit rate-limit class is found in general form, so the third SSE route does not repeat this. (d) Honest boundary: this does not claim 120/min is exploitable in this deployment with network_mode host and no reverse proxy; every caller shares the 127.0.0.1 bucket. BOB-111 measured and recorded separately. CLAIMED. No change made. The limit was never read from headers, never driven to exhaustion — the limiter is per-IP and shared with concurrent agents on this host (§11.4.201). WHAT. scripts/run_all_challenges.sh:66 “scaling_horizontal_challenge.sh” in its set. challenges/scripts/scaling_horizontal_challenge.sh does not set -n ‘66p’ scripts/run_all_challenges.sh “scaling_horizontal_challenge.sh” \$ ls -l scripts/scaling_horizontal_challenge.sh access ...: No such file or directory VERIFIED independently, not taken on report — see the BOB-109 scaling agent and re-checked by direct invocation. WHY IT MATTERS. this is cosmetic or a §11.4.201 gate-honour defect depends entirely on how the runner has a missing entry, and that is the first thing to determine: if it SKIPS silently, the challenge advertises coverage it does not have (a claim-vs-reality row with no passing challenge behind it); if it FAILS, the runner is permitted for a reason unrelated to the system test, which trains readers to ignore it. N

ATM ID	Type	Status	Severity	Description
				outcome is acceptable; they need differ ACCEPTANCE. (a) Determine and recor runner's actual behaviour on the missin invoking it, not by reading it. (b) EITHE the challenge, OR remove the entry — v §11.4.124 discipline: check git history fo whether it once existed and was deleted silently-dropped challenge is the more i defect. (c) If the runner silently skips m entries, that is its own finding: a missin challenge must be loud (§11.4.3 SKIP-w at minimum), never absent-and-quiet. S Low as a defect, but it sits on the challe coverage seam, so (c) may deserve its o CORRECTION 2026-08-22 (§11.4.6) — T ITEM'S PREMISE WAS FALSE, and the mine. Recorded openly rather than quie scaling_horizontal_challenge.sh EXISTS submodules/challenges/challenges/scrip scaling_horizontal_challenge.sh, execut bytes, dated Aug 7. So do all 14 other e the meta-runner. It was never deleted: a the roster at 1c959ba, the challenge its in the submodule at 873c0b1, and a pic search across both repositories and all finds ZERO deletions. HOW I GOT IT W run_all_challenges.sh reads submodules challenges/challenges/scripts/. I checke challenges/scripts/ — a DIFFERENT dir belonging to a DIFFERENT, glob-based Two similarly-named directories; I meas wrong one. That is a §11.4.201(9) field-i error. Worse, I wrote that it was "verifie invocation": I had listed a directory and the aggregator, so the claim overstated evidence class as well as getting the an wrong. THE LESSON, which generalise item. In the SAME commit I correctly c sibling agent making this exact class of BOB-167 — a zero-hit produced by sear symbol name — and applied a control n there. Then I omitted the needle one pa later on my own measurement. The dist that would have caught it: A CONTROL PROVES THE INSTRUMENT CAN SEE, NOT PROVE IT IS POINTED AT THE TH UNDER TEST. §11.4.201(7)(b) as comm applied guards blindness; it does not gu aiming. The needle for THIS query wou been to resolve the path the runner itse

ATM ID	Type	Status	Severity	Description
				not to confirm that some directory listing WHAT IS ACTUALLY WRONG, and it is than what I filed. The missing-entry pat challenges by definition, so it can be ex safely. Run against a checkout where th submodule is absent — the ordinary sta clone made without -recursive — the R runner (byte-identity sha-asserted, not a reports: PASS: 0 FAIL: 0 SKIP: 16 TOTAL OBSERVED EXIT CODE: 0 An ENTIREL RUN BANK REPORTS SUCCESS. Not o dangling entry tolerated — the whole su silently passes without executing anyth is §11.4.201(6) false-null verbatim, and observed-reachable rather than constru (§11.4.115(G)). REVISED ACCEPTANCE ANSWERED, though not as filed: the ru reports a loud, counted SKIP and then r blocks, because the exit expression rea FAIL count. (b) The dangling-entry half nothing is dangling. (c) The real work is contract: a roster entry the runner cann execute must not be indistinguishable f healthy run. See the decision recorded BOB-168/decision_missing_vs_skip.md. REPORTED NOT FIXED: scripts/ pre_build_verification.sh:1792 carries t SKIP/MISSING conflation one level up, other runner. Filed separately. WHAT. The §11.4.65 markdown-export r requires every in-scope doc to ship .html twins. 286 of the 326 .html files under c (excluding dist/ and node_modules/) are FRAGMENTS — they begin at WHAT. The BOB-109 scaling suite gates timing tests on loadavg_1m/nproc <= 0 this 8-cpu host, a 1-minute load average below 6.0 — and SKIPS loudly otherwise gate is correct and was adopted for a g (§11.4.201(8)): the author validated a 2 growth-exponent threshold, then observ FALSE-FAIL correct code at load 11.06 span 2.136), and at load 15-23 measure spans of 2.358 against injected-cubic m 2.278-2.331 — the signal is smaller than and the two are not separable in either Gating was the honest response. THE P The gate has consequently NEVER RUN own precondition. Every recorded run t measured load 9.15-23; the independent
BOB-169	Bug	In progress	Medium	
BOB-170	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				measured 9.66 during review; the cond measured 6.45 at its lowest. GREEN po observed — at loads 9.4-11.6 (span 1.88 three stability runs (1.759/1.832/1.848) never once beneath the shipped ceiling. author states this honestly as an inferen data rather than a run they can paste, a reviewer confirmed the inference is sou (quiescence is strictly more favourable) inference is still not an observed run. W MATTERS (§11.4.226). A registered, top present guard that never executes is indistinguishable from a guard that wou did. Nothing currently ensures it ever r freshness contract, no scheduled quiet- invocation, no tracked item — which is the unexecuted-standing-guard class §1 names, where registration gets treated coverage while execution is nobody's jo forensic precedent in that anchor is two guards that, when finally executed, imm emitted latent FAILs. ACCEPTANCE. (a both quiescence-gated timing tests on a whose 1-minute loadavg is genuinely < — no agent fleet, no concurrent build — capture the run as an artifact showing t resolved load reading alongside the me span exponent, so the precondition is pr from the artifact and not asserted. (b) R observed GREEN polarity in the item ar test source, replacing the current hones inference note. (c) If the gate FAILS unc genuine quiescence, that is the finding exists to surface, and it reopens the thr calibration question rather than being v around. (d) Establish a freshness contra §11.4.226: how stale a quiescent verdict before the gate counts as uncovered ag this does not silently return to never-ex HONEST BOUNDARY. This item does n the gate is wrong. The evidence points t way — quiescence is strictly more favou than the loads where GREEN was alrea observed. It claims only that the run ha happened, and that an unexecuted guar be cited as coverage. PROVENANCE. R IMPORTANT-2 by the independent Fabl substrate reviewer of the BOB-109 worl independently corroborated: the review the SKIP is loud and prints its resolved

ATM ID	Type	Status	Severity	Description
				and confirmed the shipped tests do SKI today's load. WHAT. Both rate limiters key their per-bucket on the LEFTMOST element of X-Forwarded-For when TRUST_FORWARDER is enabled: download-proxy/src/api/rate_limit.py:117-123 (:7187) and the limiter in plugins/download_proxy.py (:7186) which was deliberately built to exact parity with it. That element is CLIENT-SUPPLIED honest reverse proxy APPENDS the peer to the RIGHT of whatever arrived, so the entry is whatever the client sent — attacked by construction, not by misconfiguration. Consequences with the (1) rotating a forged leftmost value min unlimited sequence of fresh buckets, with total bypass of the limit the flag is supposed make MORE accurate; (2) it also defeats idle-reap cap, since each forged identity has a distinct key — the bucket map is a memory growth surface bounded only by the cap; cap's eviction then discards the budgets of clients to make room for forged ones. NOT EXPLOITABLE AS DEPLOYED TODAY, and why this is Low, not High. TRUST_FORWARDED_FOR defaults OFF on source sites, and the deployed stack runs network host with no reverse proxy in front (verified across all five compose services), so peer addresses are real client IPs and no NAT occurs. The defect is latent: it arms the someone puts a proxy in front and turns on to recover real client IPs — which is the situation the flag exists for. The failure is therefore 'correct-looking configuration silently disables the limiter', not 'current broken'. PROVENANCE. Raised as MINOR the independent reviewer of the BOB-1 limiter work and independently confirmed implementing agent, which noted it at E source sites as a tracked follow-up rather than fixing it in that change. Filed here because neither agent has write access to the trust WHY IT COVERS BOTH PORTS. :7186's was built to deliberate policy parity with including this parsing. Fixing one and not other would leave the two surfaces disjoint on client identity — worse than the shared because it becomes surface-dependent
BOB-171	Bug	Queued	Low	

ATM ID	Type	Status	Severity	Description
				untestable as a single invariant. ACCEPT (a) Replace leftmost-element parsing with aware resolution at BOTH sites: either minus-N-trusted-hops, or a trusted-proxy allowlist where only a peer inside the allowlist may assert XFF at all — the choice is a deployment-topology decision and should be recorded, not assumed. (b) A test that for rotating leftmost element and asserts that the key does NOT rotate, per site. (c) Paired mutation: restore leftmost parsing and the test must FAIL. (d) Negative control (§11.4.2) with TRUST_FORWARDED_FOR OFF, the test must be ignored entirely and keying must be back to the peer address — a guard that is honouring XFF when the flag is off would be a worse defect than the one being fixed. (e) boundary: this closes header-forgery by ensuring it does not make per-IP limiting fair behind a proxy where many real users legitimately share an address. NOT CLAIMED. No change made; both sites still parse leftmost; the flag still does nothing. WHAT. The rutracker SEARCH endpoint is not refused by bot protection while the site is not fully reachable. Measured unauthenticated on 2026-08-21 (no cookie file read, no credentials, no value in any artifact — §11.4.10): GET /index.php -> 200, 96283 B, 1.2s. WHAT. download-proxy/src/api/hooks.py:174 def _save_hooks(hooks: list[dict[str, Any]]): None: try: os.makedirs(os.path.dirname(HOOKS_FILE), exist_ok=True) with open(HOOKS_FILE, 'w') as f: json.dump(hooks, f, indent=2) except Exception as e: logger.error(f'Failed to save hooks: {e}') catches EVERY exception, logs, and returns None. The signature returns None, so the caller has no channel to learn the write failed. The test sites then report success unconditionally. _save_hooks(hooks):174 _save_hooks(hooks):175 logger.info('Created hook: ...'):175 logger.info('Deleted hook: ...'):156 return HookResponse(hook_id=..., ...):176 return {'message': 'Hook deleted', ...} USER-VOLUNTARILY CONSEQUENCE, which is why this is HIGH and not a code-hygiene nit. A user POSTs a hook and receives HTTP 200 and a hook_id, and the hook was never written — their automation script never fires, and the API told them it existed. Symmetrically, a user DELETES a hook,
BOB-172	Bug	Ready for testing	High	
BOB-173	Bug	Ready for testing	High	

ATM ID	Type	Status	Severity	Description
				‘Hook deleted’, the file still holds it, and again after the next restart. In both direct product reports the opposite of what happens and the only trace is a log line nobody is watching. THIS IS THE §11.4.252 SHAR path combines two dangerous capabilities that anchor’s taxonomy — MUTATION of resource (a filesystem write) and EXTENSION SIDE EFFECT (hooks are outbound calls the system will or will not make) — so it is not to FAIL CLOSED: verify the precondition when it cannot be satisfied, and surface the refusal. Instead it fails open, and a bare Exception:’ that only logs is the exact and §11.4.252 enumerates. At the product level also a §11.4.201(6) false-null: a success and a swallowed failure are indistinguishable to the caller. PROVENANCE. Surfaced as a scope observation by the independent review the BOB-135 test-isolation work — this is what turned that defect into an ‘assert (a mystery, because the EACCES on /conf is logged and discarded while the endpoint is returning 200. Verified here directly from before filing (the function body and both read above), not taken on report. Record §11.4.238 discovery-channel escape: found agent reading code during an unrelated investigation, not by the automated QA the coverage gap is a defect of equal stature the defect itself. ACCEPTANCE. (a) _save propagates failure — raise, or return a success callers must consume. (b) Both endpoints translate a persistence failure into an HTTP (500-class), never a success body; a creation did not persist must not return a hook in test drives each endpoint with the hook unwritable (read-only dir or a patched code raising OSError) and asserts a non-2xx AND that a subsequent GET does not list phantom hook — assert on the user-observed outcome, not on the log line. (d) Paired mutation: restore the swallow; the test fails FAIL. (e) Audit the same file for sibling — this is a pattern, and one instance is alone. (f) Honest boundary: this does not the write currently fails in production; in that WHEN it fails the user is told the outcome and the BOB-135 investigation shows it in at least one real environment. NOT C

ATM ID	Type	Status	Severity	Description
BOB-174	Bug	In progress	High	No change made. No assessment of how write fails in the operator's deployment RECORDING A SHELL ERROR OF MY C (§11.4.6): the first version of this description written with the anti-pattern snippet inside backticks in a double-quoted shell argument the shell ran it as command substitution text was replaced by nothing — the stored description read 'and is the exact anti-pattern. This is the SECOND time this session the backticks-inside-double-quotes has corrupted content (the first mangled a commit message). Fixed here by editing through a Python script with no shell quoting in the path. Noted a silently-truncated defect description is the kind of quiet corruption §11.4.201(7) about — the path is part of the instrument failed without erroring. WHAT. A corrupt hooks file is silently indistinguishable from "no hooks configured" and the next create then DESTROYS every existing hook while returning HTTP 200 OK. Defects on one path, filed together because any one alone leaves the data loss reachable — download-proxy/src/api/hooks.py:104. <code>_load_hooks</code> wraps the read in <code>except Exception as e: logger.error(...)</code> and falls through to report truncated or malformed <code>hooks.json</code> is then reported to every caller as an empty, <code>hooks</code> list. Confirmed at source. A2 — the console and the reason this is High. <code>create_hook</code> appends, saves. Given a corrupt file that turns into <code>[]</code>, the save writes a one-element over the top. Every previously-configured gone. Measured by the implementing against the ALREADY-FIXED tree, so the BOB-173 write-failure fix: BEFORE <code>['prod-hook-0', 'prod-hook-1', 'prod-hook-2']</code> reports : 200 <code>{'hooks': [], 'count': 0}</code> <- ZERO hooks configured DELETE report <code>{'detail': 'Hook not found'}</code> <- for a hook in the file POST reports : 200 <code>hook_id=3c8a5930-...</code> AFTER on disk : <code>['3c8a5930-...']</code> Three prod hooks destroyed HTTP 200 throughout, nothing surfaced to user. A5 — <code>_save_hooks</code> writes with a <code>pl</code> <code>open(path, "w")</code>. A crash, <code>ENOSPC</code>, or a write truncates the file in place. That is the corruption A1 then reads as "no hooks" A2 overwrites. The same codebase already

ATM ID	Type	Status	Severity	Description
				the correct pattern: theme_state.py:115 tmp-file + os.replace. WHY THE THREE ONE ITEM. A5 manufactures the corrupt misreads it as empty, A2 destroys the co Fixing only A1 leaves truncation reacha only A5 leaves an existing corrupt file a trigger; fixing only A2 leaves the API lyi what is configured. The chain is the def DELIBERATELY NOT FIXED UNDER BO and the reasoning is sound. The implem agent identified all three while auditing swallows as that item required, and dec fix them in the same change because: th change GET semantics on an endpoint t frontend consumes (200 -> 5xx); they n CORRUPT-file reproduction rather than UNWRITABLE-file one BOB-173 built; a need a design decision that is genuinely obvious — a MISSING hooks file legitim means “no hooks configured”, while a C one does not, and today’s code cannot t apart. Expanding BOB-173’s scope to co would have meant shipping that decisio unexamined. ACCEPTANCE. (a) Disting MISSING from CORRUPT: a missing file an empty list; a corrupt one is an error, silently []. (b) Decide and RECORD wha does on corruption — 5xx, or 200 with a degraded marker the frontend can rend the design decision, and it should be sta rather than inferred from whatever the happens to do. (c) create/delete MUST overwrite a file they could not parse — not clobber (§11.4.252: mutation plus e side effect must fail closed). (d) Make th atomic via tmp + os.replace, reusing theme_state.py’s existing pattern rather inventing it (§11.4.28). (e) Tests assertin USER-OBSERVABLE outcome, not log l a corrupt file, assert GET does not claim hooks, assert a create refuses rather th destroying, and assert the file still holds original hooks afterwards. (f) Paired §1. mutation per guard. (g) NEGATIVE COM (§11.4.201(1)): a MISSING file must stil empty list and a working create; a VALI must behave exactly as today. A fix that every load fail closed by failing always i A3, RECORDED SEPARATELY, NOT PAR THIS CHAIN. VALID_EVENTS and

ATM ID	Type	Status	Severity	Description
BOB-175	Task	Queued	Low	HookEventType are two sources of truth. The closed set. Verified identical today, ungu against drift: if they diverge a hook regi successfully and then silently never fire assertion pinning them would close it. M CLAIMED. No change made. The BOB-1 failure fix is real and orthogonal — it m FAILED write honest; it does nothing ab SUCCESSFUL write of wrong data deri misread file. WHAT. The workable-items update seam the status-location invariant in one dire After BOB-166, update --status <termina Issues-located row is refused. The mirro open: update --location Fixed --status < terminal> exits 0 and creates a row that own validator immediately rejects. Verifi empirically by the independent reviewe BOB-166: update -location Fixed -statu progress' -> exit 0 validate -> FAILS, na row (check (f), fixedLocationNonTermin So the command can still mint a state it validate call refuses. That is the §11.4.1 shape at the seam layer: the invariant is CONFIGURED in validate but not ENFC every write path that can violate it. WH DELIBERATELY LEFT OPEN IN BOB-16 why that was right. Three reasons, all c rather than asserted: (1) the real corpus ZERO instances of this direction against the direction BOB-166 closed — the asy the data matched the asymmetry in the detective coverage ALREADY existed and demonstrably fires, so unlike BOB-166's this cannot rot silently; (3) and the load one — reopenCmd documents and SANCT transient Fixed-plus-non-terminal state location Fixed operator override. A naiv preventive guard at the update seam w collide with that override's semantics. C this therefore requires a design decision how the two interact, which is a differen question from the one BOB-166 was sco answer, and expanding that item would meant shipping the decision unexamined ACCEPTANCE. (a) Decide and RECORD preventive guard on this direction coexi reopenCmd's sanctioned override — thi actual work, and the answer should be v down rather than inferred from whatever

ATM ID	Type	Status	Severity	Description
				patch does. Candidates worth weighing the reopen path explicitly, require an ov flag on update mirroring reopen's, or lea direction detective-only with the rational recorded so the asymmetry is a decision than an oversight. (b) If a guard lands, i terminalStatuses() — BOB-166 establish predicate source specifically so the two cannot drift. (c) Paired §1.1 mutation. (c) NEGATIVE CONTROL (§11.4.201(1)), and sharp one here: the sanctioned reopen MUST still work. A guard that breaks reopenCmd's documented override is a positive refusal, and worse than the gap because it breaks a workflow operators to use. HONEST BOUNDARY. This is no data defect. Zero corpus instances, and detective gate catches it at the next val. filed so a scoped, deliberate omission st as a tracked decision instead of living o code comment where the next reader w it for an oversight (§11.4.197: a started reaches a documented terminal state, n abandonment). PROVENANCE. Raised i implementing agent of BOB-166 as a kn asymmetry it deliberately did not close, independently verified and endorsed as acceptable scope boundary by that item reviewer, on the condition that it becom tracked item rather than a comment. Th item. WHAT. <code>_get_enabled_trackers()</code> in down proxy/src/merge_service/search.py gate rutracker on <code>RUTRACKER_USERNAME</code> <code>RUTRACKER_PASSWORD</code> only. It never <code>RUTRACKER_COOKIES</code>. So an operator configured with cookies alone — no use password — never has rutracker enable and the tracker is silently absent from e out rather than failing visibly. WHY THA INCOHERENT WITH THE REST OF TH <code>_search_rutracker</code> treats <code>COOKIES</code> as th PREFERRED auth path, not a fallback. A nnmclub sibling DOES check its own * variable. So the same codebase holds th positions at once: cookies preferred at t site, cookies ignored at the enablement cookies honoured for a sibling tracker. V MATTERS MORE THAN IT LOOKS. CLA documents a standing operator mandat
BOB-176	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				(2026-08-15) that per-tracker Netscape files at <code>\${TRACKER_COOKIE_DIR}/cookies_<tracker>.txt</code> are auto-loaded in as <code><TRACKER>_COOKIES</code> before every up, restart, install Stage 6 and start.sh and rutracker is named in that set. So t SANCTIONED, documented configuration for this tracker produces exactly the sta gate ignores. An operator who follows t documented instructions gets a tracker never runs, with no error to explain it. I MODE. Silent absence, not a visible fail same §11.4.201(6) false-null shape as B one layer earlier. BOB-172 fixes a track RAN and was REFUSED being reported this is a tracker that never ran at all, an merged result simply has one fewer cor with nothing to indicate it. PROVENANC by the BOB-172 implementing agent wh tracing the enablement path to locate th handling defect. Recorded here rather t in that change: it is a separate defect or separate seam with its own oracle, and in would have widened BOB-172 past it evidence. ACCEPTANCE. (a) The enable gate honours RUTRACKER_COOKIES a independently sufficient credential, mat what _search_rutracker already prefers the nnmclub sibling already does. (b) A asserting that a cookies-only configurat ENABLES rutracker, driving the real _get_enabled_trackers rather than a rep Paired §1.1 mutation restoring the user password-only gate — the test must go NEGATIVE CONTROL (§11.4.201(1)): a configuration with NO credentials at all leave rutracker DISABLED. A fix that en unauthenticated tracker is worse than t because it produces failing searches ins absent ones. (e) Audit the other tracker for the same asymmetry rather than fixi the one that was noticed — three positio codebase suggests nobody has checked together (§11.4.118). HONEST BOUND verified against a live cookies-only depl read from source by the BOB-172 agent recorded on its report. The reading is sp checkable, but whoever takes this shou by invocation before relying on it. BC Bug Queued High WHAT: BOB-172

ATM ID	Type	Status	Severity	Description
				identical 4-line HTTP-refusal guard at first private-tracker fetch sites in download-proxy/src/merge_service/search.py (rutracker cookie:1468, kinozal:1640, nnmclub:1761, iptorrents:1934). Only rutracker COOKIE path is exercised by BOB-172 EVIDENCE (reviewer-authored mutation 11.4.194(6)(d), during the BOB-172 independent review): deleting ONLY the kinozal guard (search.py:1655-1658) while leaving the others intact left the full merge_service suite as passed, ZERO failures. The suite cannot pass all of the five sites. FAILING SCENARIO: a guard that drops or subtly breaks the wiring at kinozal, nnmclub, iptorrents, or rutracker-credentials. BOB-172 403 at that site silently folds back to status=empty with error=None -- the expected BOB-172 signature -- and no test reddens. DIRECTION (reviewer preferred): collapse the five duplicated wirings into one shared function, e.g. _check_search_response(tracker_name, status, body) -> bool, so there is ONE case and a sixth site cannot be added unguarded (11.4.251 byte-identical-fork extraction). Alternative: parametrise the guard tests for all five sites with per-site stub sessions. ACCEPTANCE: mutating the wiring at all five sites reddens at least one test. Provided running the same R1 deletion at each site. BOB-178 Bug Queued Medium V download-proxy/src/merge_service/search.py:1638-1640 handles a failed kinozal LOGIN response with if login_response.status in (200, 301, 302): logger.error(...); return sets NO diagnostic. FAILING SCENARIO: Cloudflare returns 403 on takelogin.php, kinozal chip reads status=empty, error=empty. That is the BOB-172 signature exactly: a 403 reported to the user as an empty result. CONTRAST establishing this is an oversight, a design choice: the rutracker and nnmclub failures DO set diagnostics (upstream_cookie_auth_failure), and the iptorrents login failure through to the search fetch where BOB-172 guard catches it. Kinozal is the one leg that is neither. FIX DIRECTION: stash _classify_upstream_http_status(login_response) before the early return, or an auth_failure for the status-in-(200,301,302)-but-no-cookie case. ACCEPTANCE: a stubbed 403 on the

ATM ID	Type	Status	Severity	Description
				kinozal login endpoint produces a non-M on the kinozal chip, with a RED capture the pre-fix code first (11.4.115). BOB Task Queued Medium WHAT: the BO independent review demonstrated two f paths that still report a refusing or unre tracker as an empty result. Both are PR EXISTING and neither is a regression fr BOB-172, but the fix evidence log prese five-site coverage without stating the bo (11.4.194(5) requires un-analysed dime explicit gaps, never silently assumed sa -- exception before resp.status is read. I probe B, captured: a connection-refused rutracker yields status=empty, error=N http_status=None, metadata.errors=[], metadata.status=completed. Each _sear method swallows exceptions (except Ex logger.error; return results) before _se error handling can see them, so a track DOWN is indistinguishable from one tha genuinely empty. GAP B -- soft refusal a 200. Reviewer probe A, captured: _classify_upstream_http_status(200, <c body>) returns None, which is CORREC design (a 2xx is a usable response; body triggering would risk the over-fire the n controls exist to prevent). But all five se use aiohttp's default allow_redirects=Tr session-expiry 302 -> login-page-200 ch to zero rows and reports empty. FIX DIF Gap A needs the exception to reach a di rather than being swallowed. Gap B nee DISTINCT detector (final-URL check or form marker) with its own RED -- explic widening of the status-code trigger. ACCEPTANCE: both gaps closed with th REDs, or explicitly closed per 11.4.112 evidence. Immediate sub-task: append a gaps paragraph to docs/qa/BOB-172/ fix_evidence_20260822.log. BOB-180 Queued Low WHAT: download-proxy/ routes.py:727-745 sets status=captcha_ with a hardcoded message naming RuT 'RuTracker requires CAPTCHA. Use /ap rutracker/captcha'. FAILING SCENARIO search returns zero results and the cap flavoured diagnostic came from NNMCI Turnstile, not RuTracker. The user is to a RuTracker captcha endpoint for an NI

ATM ID	Type	Status	Severity	Description
				problem. SEVERITY BOUNDED: the full and tracker_stats travel in the same res payload, so the truth is present and a cl reads them is not misled -- only the head wrong. The branch was revived by BOB error propagation (correctly: suppressing propagation would recreate the same fa one layer up, 11.4.247) and is already c the routes layer by tests/unit/api_layer/ test_routes_coverage.py:782. FIX DIRE derive the message from the tracker(s) actually erred rather than hardcoding o SCOPE NOTE: api/ was owned by a sibl stream during BOB-172; the author was not to touch it. ACCEPTANCE: an NNM captcha diagnostic on a zero-result sear produces a headline naming NNMClub. BOB-181 Bug Queued High WHAT generate_markdown_exports.sh accepts arguments. It unconditionally scans PROJECT_ROOT/*.md, docs/ recursively scripts/ recursively (scan loops at lines 105/111/116 as of 2026-08-23; they were when this was filed — line numbers shifted the BOB-169 fix, the substance is unchanged was independently verified by execution the generator with an explicit path produced out-of-scope export and NOT the request and PROJECT_ROOT is derived from BASH_SOURCE (line 17). A caller who v 'bash scripts/generate_markdown_exports docs/qa/SOME/file.md' gets a WHOLE-T with the argument silently discarded. W MATTERS: the caller reasonably believed was scoped to one file. It was not. In a checkout with parallel work streams this regenerates any export whose .md is new its derived files anywhere in the tree -- artifacts in other streams' scopes and, v BOB-068 auto-commit hazard, into another stream's commit. CLASS: a false affordance script's interface implies a capability it have (11.4.201: an interface must mean claims). FIX DIRECTION: either accept path arguments and honour them (scope run to exactly those files), or REFUSE v clear message when arguments are passed Silently discarding them is the one option must not remain. ACCEPTANCE: passing either scopes the run to it, or exits non-

ATM ID	Type	Status	Severity	Description
BOB-182	Task	Queued	Medium	naming the unsupported argument. A R capturing today's silent-discard behavior (11.4.115). LIVE INSTANCE, MEASURE 2026-08-23 (this is no longer hypothetical). During the BOB-169 fix the conductor in the generator with an explicit path to re ONE document. The argument was disc the whole tree was scanned. It generate qa/BOB-174/DESIGN_DECISION.{html, at 11:50 -- a SIBLING_STREAM's docum that stream's author was still editing th (.md mtime 12:00). The exports were th stale on arrival, and additionally carried fix BOB-169 charset defect. The sibling independently reported them as stale in hand-off, having no idea another stream created them. SECOND-ORDER LESSO keeping with this item: the conductor's verify the blast radius ALSO failed, and silently. 'git status --porcelain \\\ grep -E pdf)\$"' returned only its own files and r clean - but docs/qa/BOB-174/ is an UNT DIRECTORY, which porcelain collapses line, so the grep was structurally incap seeing the files inside it (11.4.201(7)(a) the wrong surface; the quiet zero read a absence). The instrument that DOES se WHAT: CM-EXPORT-CHARSET-VALID (p invariant 50) adopts its 301 pre-existing violations via a monotone-decrease ratc than a hard floor. Two things are owed. ADOPTION DECISION IS THE OPERAT NOT THE SCRIPT'S. 11.4.224(E) requir brownfield adoption question be SURFA the operator and the answer recorded a consumer DATA - never an invented rat What exists today is a header comment BOBA_EXPORT_CHARSET_BASELINE c that DOCUMENTS a decision the agent does not RECORD an answer the operat The independent reviewer judged the cl defensible and would not reverse it (a h on 301 violations makes the build unrea which 11.4.234 forbids; 11.4.101 favour reversible option) but correctly held tha yet compliant. Options for the operator: immediate hard floor (BASELINE=0) / k monotone ratchet / per-corpus phase-in changed-code-only with a scheduled ful deadline. (2) THE RATCHET DOES NOT

ATM ID	Type	Status	Severity	Description
BOB-183	Bug	In progress	High	ACTUALLY RATCHET. BASELINE is a h constant. On improvement the gate PAS PRINTS the value to lower it to, but not lowers it, so the gate keeps permitting r all the way back to 301 after the corpus The header now says tightening is MAN rather than claiming automatic behavior does not exist (11.4.6), but the honest s is a stopgap, not the fix. WHY IT WAS N WITH THE FIX: a gate that writes its ov threshold during a pre-build run becom PRODUCER as well as a GATE (11.4.24 separation), and that is a design change operator should approve rather than re side effect. ACCEPTANCE: the operator adoption answer recorded as consumer and either a persisted baseline the gate and never raises (with the role-separati question answered), or an explicit decis manual tightening is acceptable. The compiled Angular bundle shipped a download-proxy/src/ui/dist/frontend/bro carries the BOB-164 colour-contrast de verbatim: scanned 2026-08-23 it report: violation nodes on darcula/dark and 1 o light, while the same sources built to a path report 0 across all 16 palette x mo combinations. The fix is therefore corre SOURCE layer and absent at the ARTIF so §11.4.108 layer 2 is NOT closed and still see the low-contrast dashboard. Two compounding facts make this silent rath obvious: dist/ is gitignored, so the diver never shows in a diff; and scripts/install asserts only that the directory EXISTS, it is newer than the sources it was built stale bundle passes install unchallenged Acceptance: install (or an equivalent ga on a bundle older than its sources, the l rebuilt, and docs/qa/BOB-164/ axe_contrast_scan.py run against downl proxy/src/ui/dist/frontend/browser exits zero violation nodes and zero blocking i nodes.
				Three icon-glyph controls (.bridge-retry toggle, .caret) render text made of non-points. axe-core declines to decide their reporting them under the incomplete ch with messageKey nonBmp and resolving NEITHER fgColor NOR bgColor, so the
BOB-184	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				recomputation in docs/qa/BOB-164/axe_contrast_scan.py cannot decide either such nodes across 16 palette x mode combinations. BOB-164 round 2 stopped being scored as silent passes by naming GLYPH_UNMEASURED, printing them, and run and fencing them, so any NEW unmeasured glyph node blocks. What remains genuine and unproven is their contrast itself: as non-contrast interface components they owe 3:1 under SC 1.4.11, and no oracle in this repo measures that today. Acceptance: a measurement that resolves the effective foreground and background for glyph controls (for example reading getComputedStyle colour against the computed backdrop, or replacing the glyphs with controls carrying explicit fill tokens), each of the controls shown to clear 3:1 in all 16 palette mode combinations, and the GLYPH_UNMEASURED fence shrunk by the nodes then proven. docs/qa/BOB-164/axe_contrast_scan.py is a static dist with no backend, so every node's existence depends on fetched data is about the page it measures: the results table is empty, so .type-badge and .quality-badge never appear, and the hooks list renders empty so .status pills never appear. That is precisely why the BOB-164 round-1 regression was unseen in a rendered DOM even though scanned in all 16 themes, and it is why I had to close the class in the arithmetic model instead. Those nodes are now covered by frontend/src/app/models/style-contrast.py which extracts declared foreground-on-background from the stylesheets, but that is an ARIF claim about declared colour pairs, not a rendered-pixel one: it cannot see opacity by an ancestor, a composited backdrop, or cascade this repo's static extractor does not model. Acceptance: the scan drives the dashboard with seeded fixture data (a seeded backend, a route fixture, or an injected component harness) so the badge and status nodes are present in the DOM, axe measures them directly, and a control needle proves rendered oracle sees a seeded sub-floor before any clean result from it is believed.
BOB-185	Task	Queued	Medium	
BOB-186	Bug	Queued	—	workable-items diff returned exit 0 and verdict "DB and Markdown are in sync"

ATM ID	Type	Status	Severity	Description
				PARTIAL read: invoked with -issues but -fixed it compared 77 Markdown items 184 DB items and called that in sync, pr both mismatched counts in its own outp 107 DB items were never compared aga anything. Measured 2026-08-25: (A) no markdown paths -> correctly REFUSES (B) -db-only -> honest, "no Markdown compared"; (C) -issues alone -> exit 0, (compared 77 Markdown item(s) against item(s))". The 11.4.201(6) false-null: a p blind instrument returned the same qui genuinely-synced corpus returns. IMPA CORRECTED 2026-08-25, my original fi OVERSTATED this. I wrote that CHECK 11.4.106(F) commit seam was exposed. NOT. scripts/hooks/docs-sync-commit-seam.sh:277 passes BOTH -issues and - a repo-wide sweep found NO executable that breaks: the two real callers (that se pre_build_verification.sh:549) both pass paths; the single-path hits are prose evi files plus a static-scan test fixture. So th was real and latent, never live in this pr corrected claim is: any FUTURE single-caller would have been silently mis-told PROVENANCE, twice corrected. A suba attributed the false-null to the no-paths which actually refuses correctly; the co re-measured and relocated it to the par Then the fixing agent corrected the con impact. Both corrections are recorded s wrong version propagates. FIXED (unco constitution submodule, fetched to upst 7f16739 before edit per 11.4.26). Desig by default (11.4.201(4)); the narrow per comparison PRESERVED behind an exp -partial-scope rather than deleted (dele existing documented capability would b 11.4.122 silent removal), mirroring how preserved the no-Markdown shape. Key MEASURED DB content, never on flag p — a DB with zero Fixed rows IS fully ac for by -issues alone and still passes, so keying would have traded the false-null 11.4.201(1) false-positive refusal. Resul property: "in sync" is reachable only wh two printed counts account for each oth fail/5 pass -> GREEN 9/9; paired 1.1 mu the fixs own revert, killing exactly the 4

ATM ID	Type	Status	Severity	Description
BOB-187	Bug	In progress	—	guarding tests while the 5 preservation green. Case C now refuses naming “107 located in Fixed (supply -fixed)”, the 10 confirmed against an independent sqlite (107 Fixed + 78 Issues = 185). Incident closed an 11.4.238 gap: case As refusal test anywhere; TestDiffCmd_NoPathsSt now guards it. OPEN: -partial-scope is naming choice, not an operator decision renaming is cheap now and expensive as consumers adopt it. Other consumers of shared submodule were not swept. scripts/ownership_precondition.sh cons SAME unreviewed .env-driven scope as ownership_repair.sh and creates probe it, but it is NOT fenced. T028 remediation ownership_path_fence() to scripts/lib/ownership.sh and wired it into ownership_repair.sh, closing the path-es class there; the precondition was outside agents file scope and still accepts what config/owned_paths.yaml expands to. B is lower than a recursive chown (it writes files rather than mutating ownership of existing tree) but the INPUT is identical. equally unreviewed, so the same QBITTORRENT_DATA_DIR value that will drive a filesystem-wide chown will drive file creation at an arbitrary location. No escape is empirically confirmed for the path, not merely reasoned: during T028 remediation a probe declaring the verb shipped entry with QBITTORRENT_DATA hung the suite, and ps showed the real executing “find / (! -uid 1000 -o ! -gid 1 -printf ...” before it was killed by verification. Acceptance: ownership_precondition.sh SAME ownership_path_fence() predicate scripts/lib/ownership.sh (never a second 11.4.251), refuses fail-closed on a rejection and ships a paired 1.1 mutation proving refusal fires PLUS a negative control proving six live shipped entries still ACCEPT so not a 11.4.201(1) false-positive refusal. Discovered out-of-band during T028 remediation and reported by the agent as needing triage it is also a 11.4.238 coverage escape. pre_build_verification.sh invariant 17 (C WORKABLE-ITEMS-VALIDATE) resolves binary through the candidate loop at :5
BOB-188	Bug	In progress	—	

ATM ID	Type	Status	Severity	Description
				FIRST entry is constitution/scripts/workable-items/bin/workable-items. That file is GIT TRACKED (md5 17644a248363, identical to workable-items-linux) and executable, so resolution over the current untracked source constitution/scripts/workable-items/workable-items (md5 43376a6d0184). The tracker is STALE: it does not contain the guards it the source now has. MEASURED, needle-prime, 2026-08-25. String presence in the tracked binary: "refusing to set terminal status". "Issues-location item has TERMINAL status control needle from the SAME updateControl, at least one mutable field flag is required". The instrument sees through that path and that the zeros are real absences, not a blind read. The untracked sibling returns 1 / 1 / 2 for the three queries. A check whose message is absent from the binary cannot fire. RUNTIME PROOF: the same injected violation (a runtime terminal status while current_location=run through both binaries -> stale repository violation (only the body_md desync class) reports 2 including the location-status class verbatim. This is exactly why the ten BOB rows (BOB-087/129/131/144/145/146/148/150) survived undetected: the gate that was supposed to catch them was running a binary structure incapable of seeing them. The 11.4.108 >ARTIFACT gap: source correct, artifact incorrect. Every gate reading the artifact reports correct. This is the THIRD instance of that class in one session, alongside BOB-169 (a character that never opened a file) and BOB-183 (a bundle five days older than its sources). COMPOUNDING: the comment at :519- directly above the loop asserts "The naïve workable-items path never existed in the checkout (bin/ is a gitignored local build dir nothing ever populated)". That state is now FALSE - bin/ holds two tracked binaries, a stale comment asserting the absence of a file that now shadows resolution is how it stayed invisible. Also an 11.4.30 question: who owns: bin/workable-items and workable-items-linux are VERSIONED BY ARTIFACTS in a constitution submodule inherited by reference by every consumer. Every consumer inherits whichever binary

ATM ID	Type	Status	Severity	Description
				last committed. Flagged independently separate agents this session. ACCEPTA resolution must never prefer a stale art a current one - either the tracked binary removed and resolution falls through to demand, or a freshness check refuses a older than its sources (see the BOB-183 fingerprint gate for a working pattern); false comment at :519-523 corrected; (c 1.1 mutation proving the new refusal fir negative control proving a genuinely fre still passes so the fix is not an 11.4.201 positive refusal; (d) the 11.4.30 tracked decision recorded as operator DATA. Di out-of-band during BOB-166 remediation 11.4.238 coverage escape in its own rig