

README

BOB-075 followup — concurrency guard anti-bluff evidence

Revision: 1 Last modified: 2026-08-18T15:00:00Z

Followup from BOB-075 SDD task review (subagent a584f8cc): the reviewer inadvertently corrupted docs/codegraph/Status.md by running two concurrent status_docs_freshness_challenge.sh invocations (the backdate+trap+restore was not thread-safe). This directory carries the §11.4.115 polarity proof for the fix: real corruption BEFORE the fix, real lock rejection AFTER the fix, and a §1.1 mutation proving the lock is load-bearing.

All reproductions below run against a **scratch harness** (a throwaway project-root directory containing only the challenge script, a stub docs_chain engine binary that sleeps to simulate the real engine's wall-clock, and synthetic docs/codegraph/Status.md / docs/features/Status.md fixtures) — never against this repository's real, committed docs/codegraph/Status.md. The scratch harness exercises the exact same code path (same script, same directory layout relative to PROJECT_ROOT, same backdate+trap+restore logic) without any risk of leaving the real committed doc corrupted if a repro run goes wrong.

Methodology

1. Build a scratch project root: <scratch>/challenges/scripts/status_docs_freshness_challenge.sh (a copy of the script under test), <scratch>/constitution/submodules/docs_chain/docs_chain (a stub engine — sleep "\${STUB_DOCS_CHAIN_SLEEP:-N}"; echo ...; exit 0 — standing in for the real engine's genuine 70-90s verify --all wall-clock, which is exactly what gives two manually-started shells in real usage enough time to overlap), and fresh docs/codegraph/Status.md / docs/features/Status.md fixtures with an in-SLA ****Last modified:**** header.

2. Launch `RED_MODE=1 bash challenges/scripts/status_docs_freshness_challenge.sh` twice in the background, staggered by a small, swept delay (0s-2.5s), from that scratch root.
3. Wait for both to exit; diff the final `docs/codegraph/Status.md` against the pre-run original. A byte-identical result is `CLEAN`; any difference (in every observed case, the `**Last modified:**` header stuck at the backdated “100 days ago” value) is `CORRUPTED`.

No synchronization hooks, sleeps, or barriers were injected into the script under test anywhere in this evidence — every reproduction is genuine, unforced concurrent execution (“run 2 shells simultaneously”, per the task brief), with only the launch stagger swept to find a hit.

Results summary

Scenario	Script variant	Lock available	Trials	Corrupted
before_fix/	HEAD (pre-fix)	n/a — no lock exists	15 across two sweeps (10-stagger exploratory + 5-trial repeat at the hit stagger)	8/15 (see before_fix_reliability_sweep_s
after_fix/	fixed	flock	6 (staggers 0-1.0s)	0/6
mkdir_fallback/	fixed	mkdir (flock hidden via PATH)	3 attempts (2 hit unrelated fakebin PATH-harness gaps — see mkdir_fallback/notes.txt — but still correctly rejected the 2nd instance in both; 1 fully clean end-to-end pass saved here)	0
signal_handling/	fixed	flock	1 (SIGTERM sent mid-backdate)	0 (file restored, locl in that order)
mutation_test/				

Scenario	Script variant	Lock available	Trials	Corrupted
	fixed with acquire_lock call commented out	none (mutated)	20 (dense stagger sweep)	2/20 — see mutation notes.txt

Polarity confirmed (§11.4.115): unguarded concurrent execution corrupts the file (reproducible, before_fix/); guarded concurrent execution never does across 26 combined trials (after_fix/ + mkdir_fallback/); stripping *only* the lock invocation from the fixed script (the §1.1 mutation) restores the ability to corrupt (mutation_test/) — proving the lock, not some incidental side effect, is what prevents the corruption.

Honesty note on hit rate (§11.4.6): this is a genuine OS-scheduler race, not a deterministic bug — the exact hit rate varies with host load and with incidental script-length/parse-timing changes (the fixed script is longer, which measurably shifts the race window; see mutation_test/notes.txt for detail). What is invariant across every sweep run for this evidence package is the *polarity*: 0 corruptions observed with the lock engaged, and repeated corruptions observed with it absent (either pre-fix or via the mutation).

Directory contents

- before_fix/ — unmodified (git HEAD, pre-fix) script. instance_A.log / instance_B.log are the two concurrent runs' full transcripts for one representative corrupted trial (stagger=0.3s). original_codegraph_status.md vs final_codegraph_status_CORRUPTED.md / diff_original_vs_final.diff show the corruption directly. reliability_sweep_summary.txt records 5 repeated trials at the same stagger (3/5 corrupted) plus an earlier 10-stagger exploratory sweep (3/10 corrupted at staggers 0s/0.05s/0.3s).
- after_fix/ — fixed script, flock available. instance_A_winner.log completes the full baseline+RED+GREEN cycle and exits 0. instance_B_rejected.log fails fast with the actionable lock-held message and exits 1 *before touching the file*. stagger_sweep_summary.txt covers 6 staggers, all clean.
- mkdir_fallback/ — same test with PATH restricted to a directory with no flock binary, forcing the portable mkdir-based lock. Same outcome (one winner, one fast-rejected loser, clean final state, no stray lock artifacts).
- signal_handling/ — a single long-running instance (8s stub sleep to widen the window), polled until its backdate was confirmed live on disk, then sent SIGTERM. victim.log + notes.txt show the file was restored to the true original *before*

the process exited, and no lock artifact was left behind — the required “restore first, then release the lock” ordering.

- `mutation_test/` — the fixed script with the single `acquire_lock` invocation commented out (`mutation.diff`). A 20-trial dense stagger sweep reproduces corruption twice (`instance_A.log / instance_B.log / diff_original_vs_final.diff` are the `idx=5` hit); `notes.txt` has the full sweep breakdown and the honesty note on hit-rate variance.

Fix reference

The concurrency guard lives in `challenges/scripts/status_docs_freshness_challenge.sh`: `acquire_lock()` / `release_lock()` (`flock` on the script’s own file, `mkdir` fallback with §11.4.180 stale-holder-PID reaping for portability), invoked once near the top of the script before any file mutation, and a unified `cleanup()` trap (registered `EXIT INT TERM`) that restores the real file first (if a backdate is in flight) and only then releases the lock — see the script’s own top-of-file doc block for the full design rationale.