

forensics

BOB-137 — merge service (7187) wedge: captured forensics

Revision: 3 Last modified: 2026-08-20T16:10:00Z

Captured 2026-08-20 on the live stack. Evidence class: runtime (§11.4.226) — live process, socket and thread state plus in-process stack dumps, not source inspection.

STATUS: ROOT CAUSE ESTABLISHED. The spinning frame is named in §7 and proven by seven all-thread stack dumps taken while the defect was active (`stall_dump_1.txt`). Two inherited premises from Revision 2 are **REFUTED** under control conditions in §6 — read that section before citing Rev-2.

No fix is proposed or applied here (§11.4.102 Iron Law). The only code change is an observe-only instrument (§9).

1. The wedge, measured (Revision 1 — unchanged, still accurate)

Container qbittorrent-proxy, up 3h53m, reported Up 4 hours (healthy):

```
curl --max-time 6 http://localhost:7186/      -> HTTP 200 in 0.096s
curl --max-time 6 http://localhost:7187/      -> HTTP 000 after 6.004s
```

Both ports are served by the SAME process — `ss -ltnp` shows one pid holding fd 3 (7186) and fd 7 (7187). Not a crashed worker: one loop inside a live process stopped servicing requests while another in the same process kept working. See `sockets.txt`, `threads.txt`, `container_log_tail.txt`.

2. Reproduced deterministically (NEW, Revision 3)

scripts/diagnostics/bob137_soak.sh 2400 6 drives sustained concurrent POST /api/v1/search fan-out while an independent probe samples both ports every 5s. Measured 2026-08-20T15:39Z-15:46Z:

```
probes total: 26
7187 outages: 22      <- merge service dead
7186 outages:  0      <- download proxy healthy throughout
```

The asymmetry is total: **22 of 26 probes found 7187 unreachable and 0 of 26 found 7186 unreachable**, in the same process, at the same instants.

The reproduced wedge carries the SAME socket fingerprint as the originally reported one (repro_socket_state.txt), which is what binds this repro to the reported defect rather than to a look-alike (§11.4.199, precondition provenance = **observed**, not constructed):

Signal	Originally reported (§1)	Reproduced (§2)
7187 LISTEN accept backlog	Recv-Q 6	Recv-Q 88
7187 conns in CLOSE-WAIT with unread bytes	162/162/162/79/1/1/1	85/168/194/220/201/...
7186 LISTEN backlog	0	0
fd exhaustion	no (48/16384)	no (118)
thread count	16	19

3. Why 7186 survives while 7187 dies

They are different servers on different threads of one process:

- **7186** is `download_proxy.run_server()` — a blocking socketserver on its own thread (main.py:242). CPython releases the GIL every switch interval (default 5 ms), so it keeps getting scheduled even while another thread runs Python continuously.
- **7187** is uvicorn on a single asyncio event loop (main.py:288). It has no second thread to fall back on: if any callback on that loop does not return, **no** other callback runs — not accept, not read, not the response write.

That is why the accept queue backs up (Recv-Q 88 on the LISTEN socket) and why established connections sit in CLOSE-WAIT with the client's request bytes still unread: the kernel completed the handshakes and buffered the bytes, the client gave up and sent FIN, and the application never ran to read or close either.

4. Log silence and self-clearing, explained — and measured

Both follow from §3. Nothing is logged because nothing on the loop progresses; the service recovers with no restart because the callback eventually returns.

Revision 3 measured two complete episodes end to end, with NO restart and no intervention (watchdog log):

```
15:40:03 SILENT for 22s (episode 1) -- dumping all thread stacks
15:49:18 RECOVERED (episode 1)          -> ~9 min 37 s of event-
loop silence
15:49:38 SILENT for 21s (episode 2) -- dumping all thread stacks
15:55:13 RECOVERED (episode 2)          -> ~5 min 56 s of event-
loop silence
```

Load was stopped at 15:44Z, yet the loop kept spinning until 15:55Z — **eleven minutes of stall after the last request was sent**, draining merges already queued. Longest single stretch of continuous silence recorded by the watchdog: loop_silent_for=572.7s. This is the mechanism behind the originally reported ~2 h of log silence: the stall outlives the traffic that caused it, and concurrent searches queue their merges behind one another.

Immediately after recovery the service answered normally (HTTP 200 in 3.05s, then sub-second), with no restart — which is exactly the Revision 2 §6 observation, now explained rather than merely noted.

Episode 3: it is NOT soak-only

At 15:56:17Z a THIRD episode began with **zero load from this investigation** (verified by enumerating real pids and excluding the probing shell's own command line — the §11.4.196(D)/§12.12 pgrep -f self-match footgun, which did initially produce a false hit here). Its dump shows the identical path:

```
deduplicator.py, line 446 in _detect_content_type
deduplicator.py, line 201 in _extract_identity_from_result
deduplicator.py, line 224 in _check_match
deduplicator.py, line 81 in merge_results
```

```
search.py,          line 914 in _run_search
routes.py,          line 447 in _background
```

So the defect fires under ordinary traffic reaching POST /api/v1/search, not only under the synthetic soak. Any search large enough to make merge_results run long will stall the merge service for other callers.

5. SIGTERM slowness is a SEPARATE, benign issue (NEW — decoupled)

podman restart reported StopSignal SIGTERM failed to stop container in 10 seconds, resorting to SIGKILL on a **healthy, freshly-started** container. It is therefore NOT a symptom of the wedge. Cause is in main(): the handler sets _shutdown_event, then main() does proxy_thread.join(timeout=5) + fastapi_thread.join(timeout=5) on two threads that never exit — 10 s total, which equals podman's stop grace period. Tracked separately; not part of this root cause.

6. REFUTED inherited premises (§11.4.205(4), §11.4.201)

Revision 2 diagnosed “GIL starvation” from two readings. Both are refuted:

(a) “1 thread in state R with wchan=0” is the OBSERVER ITSELF. Demonstrated under control conditions on this host: a thread that reads /proc/self/task is by definition running while it reads, so it reports state=R, wchan=0 — and d_utime=0, i.e. it consumed no CPU:

```
observer tid: 312722
  tid=312721 state=R wchan=futex_wait_queue
  tid=312722 state=R wchan=0 <== THE OBSERVER ITSELF
```

The same R/wchan=0/d_utime=0 row appears in the self-test's *golden-bad* case where the loop is blocked in time.sleep and **nothing is spinning**. Observer-induced signal (§11.4.201(10)); it is not evidence of a spinner.

(b) ps %CPU 20.6 is a LIFETIME AVERAGE, not an instantaneous rate. %CPU is total CPU time ÷ elapsed time since process start. Over 3h53m it says roughly 48 minutes of CPU were consumed at some point; it says nothing about what the process was doing at the moment of measurement (§11.4.201(9), field identity and semantic class — a cumulative ratio read as an occupancy).

(c) “14 threads in futex_wait_queue” is the normal idle state of any condition-variable waiter — here the anyio worker threads and the default ThreadPoolExecutor (visible by name in stall_dump_1.txt). Not a signature.

The Rev-2 conclusion (“a thread busy-looping holds the GIL and everything else queues on it”) was directionally close to the truth but rested on three non-load-bearing readings. The actual mechanism (§7) was established from per-thread CPU deltas and in-process stack dumps instead.

7. ROOT CAUSE — the spinning frame

Deduplicator.merge_results() performs an $O(N^2)$ CPU-bound, regex-heavy deduplication synchronously on the asyncio event-loop thread.

Path: routes.py:447 _background (an asyncio task) → await orch._run_search → search.py:914 calls self.deduplicator.merge_results(all_results) as a plain synchronous call. It is never handed to an executor, so for its entire duration the merge service’s only event loop runs no other callback.

Seven dumps across one continuous 5.5-minute stall (stall_dump_1.txt):

```
===== BOB-137 STALL DUMP #1 loop_silent_for=21.7s  loop_tid=11
=====
  File ".../re/__init__.py", line 186 in sub
  File ".../merge_service/deduplicator.py", line 294 in
_normalize_name
  File ".../merge_service/deduplicator.py", line 357 in
_compare_identities
  File ".../merge_service/deduplicator.py", line 226 in
_check_match
  File ".../merge_service/deduplicator.py", line 81 in
merge_results
  File ".../merge_service/search.py",          line 914 in
_run_search
  File ".../api/routes.py",                    line 447 in
_background
  File ".../asyncio/events.py",                line 88 in
_run      <-- ON THE EVENT LOOP
  File ".../asyncio/base_events.py",          line 1999 in _run_once
```

Per-thread CPU delta, sampled over 1.0 s inside each dump:

```
DUMP #1 loop_silent_for= 21.7s  tid=11 state=R d_utime=97
d_stime=1  <-- ASYNCIO LOOP
DUMP #2 loop_silent_for= 82.9s  tid=11 state=R d_utime=97
```

```

d_stime=1 <-- ASYNCIO LOOP
DUMP #3 loop_silent_for=144.1s   tid=11 state=R d_utime=98
d_stime=1 <-- ASYNCIO LOOP
DUMP #4 loop_silent_for=205.3s   tid=11 state=R d_utime=81
d_stime=1 <-- ASYNCIO LOOP
DUMP #5 loop_silent_for=266.5s   tid=11 state=R d_utime=90
d_stime=1 <-- ASYNCIO LOOP

```

Every other thread in every dump: `d_utime=0`. The loop thread alone burns 81-98 ticks per second — 81-98 % of one core, in **user** time, continuously.

It is executing, not blocked. The outer frames are identical in all seven dumps (`merge_results` → `_check_match`) while the inner frame MOVES between them — `_normalize_name/re.sub` (line 294), `_detect_content_type/re.search` (474), `_extract_identity_from_result/re.search` (209), `_compare_identities` (357, 358). A stationary outer frame with a moving inner frame and sustained user-space CPU is a CPU-bound loop, and is exactly what distinguishes it from a blocking syscall (which would show a fixed frame and `d_utime=0`).

Why it grows super-linearly and is expensive per step

`merge_results` (`deduplicator.py:73-95`) pops a seed and compares it against every remaining candidate — $N(N-1)/2$ comparisons in the worst case, when few results match. Measured growth on synthetic data is super-linear but not yet fully quadratic at $N \leq 400$; see §8 for the numbers and the honest boundary.

Each comparison re-derives everything from scratch; nothing is cached (`grep -n "lru_cache" merge_service/deduplicator.py` → no matches):

- `_check_match` calls `_extract_identity_from_result` **twice** — once for the seed and once for the candidate — so the seed's identity is recomputed for every one of its candidates. Each call runs a season/episode regex, `_detect_content_type` (a chain of up to ~15 sequential `re.search` calls), plus resolution and codec regexes.
- `_compare_identities` then calls `_normalize_name` on both titles; each call runs 5 `re.sub` passes.

The module contains 30 `re.*` call sites on this path. Order ~25 regex operations per comparison × $N^2/2$ comparisons ≈ $O(10 N^2)$ regex operations. Single trackers were observed returning up to 189 results (`container_log_tail.txt`); summed across the 43 managed plugins, N in the low thousands puts this in the 10^7 - 10^8 regex-operation range for ONE search.

Why stalls reach hours

Concurrent searches do not overlap — each queues its own `merge_results` on the same single loop, so their costs are **additive**. Observed here: load was stopped at 15:44Z and the loop was still spinning at 15:46Z, draining merges already queued. The originally reported ~2 h of log silence is consistent with several large fan-outs queued behind one another.

The last lines before the original silence (`container_log_tail.txt`) are a security-suite fan-out across ~43 trackers with XSS/SQLi payload queries, followed by ~11 repeats of “Loaded configuration for qBittorrent” — the memoised `get_config()` logging once per concurrent caller that raced its `_config` is `None` check. That is the concurrency burst which fed the merges.

Every observation, accounted for

Observation	Explained by
7187 dead, 7186 alive, same process	§3 — one loop vs a separate thread
CLOSE-WAIT with unread request bytes	loop never runs <code>accept/read</code> callbacks
7187 LISTEN accept backlog (Recv-Q 6 → 88)	same
container log silent for ~2 h self-clears with no restart	nothing on the loop progresses the merge finally returns
16 threads wedged vs 5 idle	fan-out’s anyio + executor pool
high cumulative process CPU	~1 core of regex work per merge
stall duration unbounded/variable	$O(N^2)$ in results, additive across searches

8. Scale, measured (contention-immune)

A first attempt to TIME `merge_results` ran inside the wedged container and was contaminated by CPU contention with the live spin (N=100 took 8.5 s while N=200 took 3.7 s — an ordering only contention explains). Those numbers are discarded, not reported (§11.4.201(10)).

Instead the REGEX OPERATION COUNT was measured. A count is exact regardless of CPU contention, so it is valid under the conditions actually available (§11.4.201(8)). `re.sub/search/match/findall` were wrapped with a counter and `merge_results` run over synthetic result sets:

dataset A -- names drawn from a small vocabulary (many matches, groups collapse)

N	regex_ops	ops/N ²	growth per doubling	
50	40,256	16.1		
100	138,077	13.8	x3.4	(linear=x2, quadratic=x4)
200	388,142	9.7	x2.8	
400	1,046,531	6.5	x2.7	

dataset B -- highly distinct names (few matches, groups do not collapse)

N	regex_ops	ops/N ²	growth per doubling	
50	20,395	8.2		
100	49,888	5.0	x2.4	
200	129,764	3.2	x2.6	
400	387,224	2.4	x3.0	

Honest reading (§11.4.6): this is SUPER-LINEAR but, on these datasets, NOT yet fully quadratic — growth per doubling is 2.4–3.4 against 2.0 for linear and 4.0 for quadratic, and in dataset B it RISES with N (2.4 → 2.6 → 3.0). The claim supported by this measurement is “grows substantially faster than the number of results”, not “measured N²”. The *worst case* is quadratic by code structure — `merge_results` (`deduplicator.py:73-95`) compares each seed against every remaining candidate, i.e. $N(N-1)/2$ comparisons when nothing matches — and dataset B’s rising trend is consistent with approaching that bound as N grows.

Per-comparison cost is the other half, and it is pure recomputation: `_normalize_name` is called at **four** sites per comparison (lines 275, 276 via `_compare_name_and_size`, and 357, 358 via `_compare_identities`), each running 5 `re.sub` passes; `_extract_identity_from_result` runs twice more (lines 224-225), each including `_detect_content_type`’s chain of up to ~15 sequential `re.search` calls. Nothing is memoised — `grep -n "lru_cache" merge_service/deduplicator.py` returns no matches — so the seed’s own normalisation and identity are recomputed for every candidate it is compared against.

Scale actually observed on the live service: **572.7 s of continuous event-loop silence in a single episode**, and the loop was still inside `merge_results` more than 15 minutes after all load had stopped.

9. The instrument (download-proxy/src/main.py)

py-spy cannot attach on this host (kernel.yama.pttrace_scope=1 denies non-child attach), so the process was made to dump its own stacks:

- an asyncio heartbeat task stamps a monotonic clock once a second from inside the merge-service loop;
- a plain watchdog thread notices when that stamp stops advancing and writes an all-thread traceback plus a per-thread CPU delta to /config/download-proxy/diagnostics/stall_dumps.log and to stderr;
- SIGUSR1 (dump to log) and SIGQUIT (dump to file) give an on-demand dump with no ptrace: podman exec qbittorrent-proxy kill -USR1 1.

faulthandler.dump_traceback() is C-level and does not need the GIL, which is why it was used instead of sys._current_frames() — the latter needs the GIL and would hang in the same way as the thing it is measuring.

It adds **no** timeout, retry, restart or change to request handling.

The instrument is self-validated (§11.4.107(10), §11.4.201(7)(b))

A watchdog that never fires proves nothing — it could equally mean “no stall” or “blind instrument” (§11.4.201(6) false-null). Both poles are asserted by scripts/diagnostics/bob137_watchdog_selftest.py, run against the container’s own Python 3.12:

```
== GOLDEN-BAD: event loop blocked -> watchdog MUST dump ==
    PASS: dumped and named the blocking frame
    PASS: per-thread CPU delta present (spin-vs-block
discriminator)
== GOLDEN-GOOD: event loop healthy -> watchdog MUST stay quiet ==
    PASS: stayed quiet on a healthy loop
```

SELFTTEST PASS - watchdog fires on a real stall and not otherwise.

Is it safe to leave in production?

Yes, with one disclosed behaviour change. Each property was checked, not assumed:

Concern	Finding
Steady-state cost	

Concern	Finding
	heartbeat wakes 1×/s, watchdog thread wakes 1×/2 s. No work while healthy.
Can it hang the thing it measures?	No. <code>faulthandler.dump_traceback()</code> is C-level and does not require the GIL.
Disk growth	Capped: <code>BOBA_STALL_MAX_DUMPS</code> (default 200) × ~9.5 KB observed ≈ 1.9 MB per process lifetime. Re-dump interval 60 s.
Repo pollution (§11.4.30)	None. Dumps land in <code>download-proxy/diagnostics/*.log</code> , already covered by <code>.gitignore:64 (*.log)</code> — verified with <code>git check-ignore</code> .
Signal conflicts	None. No <code>SIGUSR1/SIGQUIT</code> handler exists in app source; <code>gunicorn</code> is not installed and <code>uvicorn.workers</code> (the only dependency referencing <code>SIGQUIT</code>) is never imported — verified in the live interpreter.
Shutdown path	Untouched. <code>SIGTERM/SIGINT</code> handlers are unchanged.
Failure mode	Fails open and loud: <code>install</code> is wrapped and logs <code>install FAILED</code> ; internal dump errors are appended to <code>_diag_errors</code> and surfaced by the watchdog, never silently swallowed.
Kill switch	<code>BOBA_STALL_WATCHDOG=0</code> disables it entirely; thresholds tunable via <code>env</code> .
§11.4.263	The instrument sends no signals at all. It only <i>registers</i> handlers.

Disclosed behaviour change: `faulthandler.register(SIGQUIT, chain=False)` replaces `SIGQUIT`'s default disposition, so `kill -QUIT` now dumps stacks and the process `CONTINUES` instead of terminating with a core dump. For a service this is the safer disposition, but an operator relying on `kill -QUIT` to stop the container would find it ineffective — `SIGTERM` is unaffected and remains the correct stop signal. If that trade is unwanted, drop the `SIGQUIT` registration and keep `SIGUSR1`, which is unused by anything.

Recommendation: keep it. The defect is transient and self-clearing, so without an always-armed in-process observer the next occurrence is again only catchable by luck — which is how it escaped automated QA in the first place.

10. Discovery channel (§11.4.238)

Found by an agent hand-probing the host during an unrelated investigation — NOT by automated QA. Coverage escape. Two escapes are now visible:

1. the health check could not observe this (sibling item BOB-138);
2. no automated check exercises concurrent search fan-out and asserts the merge service stays responsive during it — which is precisely what bob137_soak.sh does, and it reproduced the defect in under two minutes.

11. Artifacts

File	Contents
stall_dump_1.txt	7 all-thread stack dumps + per-thread CPU deltas
repro_socket_state.txt	live socket/thread state during the reproduced wedge
soak_probe.log	per-probe 7186/7187 latency and HTTP code
soak_driver.log	fan-out request outcomes
sockets.txt, threads.txt, container_log_tail.txt	Revision 1 captures
scripts/diagnostics/bob137_watchdog_selftest.py	golden-good/golden-bad self-test
scripts/diagnostics/bob137_soak.sh	soak reproducer