

BOB-168 “ acceptance (a) answered: what the runner does with a missing entry

Revision: 2 Last modified: 2026-08-22T10:45:58Z **Evidence class:** SOURCE as originally written (Â§11.4.226); the exit-semantics claim has since been confirmed at RUNTIME “ see the supersession note below.

(Revision 1 shipped without a Â§11.4.44 header; added here on first modification, and the counter starts at the revision this edit produces.)

SUPERSEDED IN PART, 2026-08-22. The exit-semantics analysis below is CORRECT and has since been confirmed at RUNTIME “ see `runtime_confirmation_20260822.md`, which observes exit 0 with 16 SKIPS and 0 challenges executed. Two things below have changed status:

- The “Why not runtime evidence” section is discharged. The missing-entry path was exercised on the real runner at zero cost, because that path runs no challenges by definition.
- The item’s remaining premise did not survive. The entry is **not** dangling: `scaling_horizontal_challenge.sh` exists and is executable, and all 15 entries plus the meta-runner are present. See `premise_correction_20260822.md`. Acceptance (b) is void “ there is nothing to author and nothing to remove.

The decision this document called for is recorded in `decision_missing_vs_skip.md` and is implemented.

Captured 2026-08-21. **Evidence class: SOURCE (Â§11.4.226).** Read from `scripts/run_all_challenges.sh`, not observed at runtime “ see “Why not runtime” below, which is a real limitation and not a formality.

The answer is a SPLIT verdict, and both halves matter

The filed item posed two possibilities “ silent skip (advertises coverage it lacks) or hard fail (permanently red for an unrelated reason) “ and said they need different fixes. Neither is what happens.

Half 1 “ the SKIP is LOUD and counted. This half is correct (lines 88-92):

```
if [[ ! -f "${script_path}" ]]; then
  echo "  SKIP: ${script} “ not found"
  SKIP=$((SKIP + 1))
```

```
    continue
fi
```

It names the missing script, gives the reason, and increments a counter that is printed in the summary. That is Â§11.4.3 SKIP-with-reason done properly, and it rules out the “silent” hypothesis outright.

Half 2 “ the SKIP NEVER BLOCKS (lines 145-152):

```
echo "FAIL: ${FAIL}"
echo "SKIP: ${SKIP}"
...
if [[ "${FAIL}" -gt 0 ]]; then
    exit 1
fi
exit 0
```

The exit code reads ONLY the FAIL count. Any number of SKIPS “ including a challenge that has never existed “ exits 0.

Why that combination is the actual defect

A loud message nobody is required to act on decays into wallpaper. The roster can rot indefinitely: an entry may be listed, absent, and reported, while the suite stays green forever and no gate ever forces the question. That is the Â§11.4.266 shape “ a capability advertised in the bank with no passing challenge behind it “ and it is the Â§11.4.135 principle (ABSENCE blocks exactly as a FAIL does) simply not applied at this seam.

Note the contrast with a seam in this same repo that gets it right: the release verdict-coverage rule treats a missing verdict as blocking precisely because a quiet zero and a clean result are indistinguishable to an exit code.

What this means for the itemâ€™s acceptance

Acceptance (a) is ANSWERED. Acceptance (c) “ “œif the runner silently skips missing entries, that is its own findingâ€™ “ needs restating, because the premise was wrong: it does NOT skip silently. The finding is narrower and more precise: **a listed-but-absent challenge is reported and then tolerated forever**. The fix is not to make the SKIP louder; it is to decide whether an absent entry should participate in the exit code, and that is a policy choice worth stating explicitly rather than patching.

Acceptance (b) is unchanged: author the challenge or remove the entry, with Â§11.4.124 git-history discipline first, since a silently-dropped challenge is the more serious possibility.

Why not runtime evidence (Â§11.4.226 honest boundary)

Confirming this by INVOCATION would mean running the real `run_all_challenges.sh`, which is the right instrument (Â§11.4.201(11): probe the artifact through its real path, never a replica). It was not run, for two reasons that are about safety and not convenience:

1. Four subagents are live on this host and the bank includes DDoS-class and sustained-load challenges. The constitution caps test execution at 30-40% of host resources; firing that bank now would breach it and corrupt the siblingsâ€™TM measurements (Â§11.4.119).
2. At a 180s timeout per entry across ~20 entries the run is up to an hour.

A hand-built replica of the loop was deliberately NOT used as a substitute: it would prove a property of the replica, not of the runner â€” the Â§11.4.201(11) prerequisite-vs-artifact trap.

So this reads the exit logic directly and says so. Runtime confirmation on a quiescent host remains OWED, and the claim above is scoped to source class accordingly.