

rw08-latency-diagnosis-20260616

RW-08 — Search-latency / / search/sync reset over tunnel — READ-ONLY diagnosis

Revision: 1 **Last modified:** 2026-06-16T10:00:00Z **Status:** diagnosis only — NO code edits (another agent owns download-proxy/) **Author:** background diagnostic subagent. Anti-bluff (§11.4.6): every claim cites a real file:line, the running tunnel log, or the source. Hypotheses are labelled as hypotheses; the confirmed facts are stated as facts. **Scope:** characterize RW-08 (docs/REMAINING_WORK_PLAN.md §RW-08) and propose a fix direction for the main agent to implement later.

1. The symptom (from the plan + live evidence)

- Broad ubuntu search $\approx$ **67 s** in-container wall-clock (plan RW-08).
 - Over the SSH tunnel, POST /search/sync **resets** (ConnectionResetError) at ~13–40 s; browser console shows ERR_INCOMPLETE_CHUNKED_ENCODING.
 - The SSE path (POST /search + GET /search/stream/{id}) **survives** and the dashboard uses it; scripted/curl callers of /search/sync hang/reset.
-

2. Root cause — there are TWO distinct, independent problems

2a. The reset is an SSH-tunnel idle-timeout artifact, NOT a server fault (FACT)

`/search/sync` (`download-proxy/src/api/routes.py:377-406`) is **fully blocking**: it does `metadata = await orch.search(...)` (line 400), and `SearchOrchestrator.search` (`search.py:903-929`) calls `start_search + await self._run_search(...)` and only returns after the **entire** 29-tracker fan-out + dedup completes. FastAPI buffers the whole `SearchResponse` and sends the first bytes **only at the end**. So the HTTP response sends **zero bytes for the full ~67 s**.

An SSH -L tunnel (and any intermediate proxy/NAT) drops a connection that is idle — no bytes in either direction — past its keepalive window. With nothing flowing for tens of seconds the tunnel tears the socket down; the client sees `ConnectionReset / ERR_INCOMPLETE_CHUNKED_ENCODING` while the **server side still returns 200** once it finishes (it just has no socket left to write to). This is corroborated by the live keepalive supervisor log `qa-results/tunnel-keepalive-testsession.log`, which shows the tunnel itself repeatedly going “health check FAILED on 127.0.0.1:7187 — re-establishing” and recovering — i.e. the transport layer is intermittently dropping, independent of any single request. **Conclusion: the reset is transport-layer (idle tunnel), not an application bug. The previously-established “console ERR_INCOMPLETE_CHUNKED_ENCODING is an SSH-tunnel artifact, server-side 200” finding holds.**

Why the SSE path survives: `GET /search/stream/{id}` streams frames as trackers complete AND emits : keepalive comment lines (the same keepalive idiom is visible at `routes.py:99` for the theme SSE stream). Continuous bytes keep the tunnel’s idle timer from firing — so the SSE socket never goes quiet long enough to be reset. This is the structural reason the dashboard (SSE) works while `/search/sync` (one big blocking response) resets.

2b. The ~67 s wall-clock is the 29-provider fan-out cost (FACT + arithmetic)

Default fan-out is **29 providers**: 24 live public trackers (`PUBLIC_TRACKERS = 38`, minus `DEAD_PUBLIC_TRACKERS = 14`, default-excluded) + up to 4 private + Jaxxnet (`_get_enabled_trackers`,

search.py:977-998). They run under a **concurrency cap of 5** (MAX_CONCURRENT_TRACKERS, search.py:611) with a **per-public-tracker deadline of 60 s** (PUBLIC_TRACKER_DEADLINE_SECONDS, search.py:1027, 1082, clamped 5–120).

So worst-case wall-clock $\approx \text{ceil}(29 / 5) \approx \mathbf{6 \text{ waves}}$, each bounded by the slowest tracker in the wave ($\leq 60 \text{ s} + 10 \text{ s}$ outer `asyncio.wait_for` margin at search.py:1031-1034). A single slow/borderline-alive tracker in each wave dominates that wave. The observed $\sim 67 \text{ s}$ is consistent with a handful of slow-but-not-dead public trackers each chewing toward the deadline across several waves. The dead-tracker exclusion already removed the 14 worst offenders; the remaining cost is the long tail of live-but-slow trackers running 6 waves deep at a 60 s ceiling.

3. What is NOT the cause (ruled out)

- **Not** an SSRF/auth/parse bug — the fan-out completes and returns 200 server-side.
 - **Not** the rutracker ReDoS (RW-06) — that fix is deployed (the installed engine config/qBittorrent/nova3/engines/rutracker.py contains the `{0,512}` bound; verified this session) and its parse is sub-second; it is not the latency driver.
 - **Not** unbounded fan-out — the semaphore (5) and the per-search cap (MAX_CONCURRENT_SEARCHES, search.py:623) are in place; the issue is total wall-clock under a 60 s per-tracker ceiling $\times 6$ waves, plus the blocking-response transport interaction.
-

4. Proposed fix direction (for the main agent — NOT implemented here)

Two independent levers; do BOTH for the cleanest result.

Fix A — keep the /search/sync socket warm (addresses the reset; highest value)

Make /search/sync stop being a single silent blocking response. Lowest-risk option: emit a periodic heartbeat so the tunnel never sees an idle socket. Either: - (preferred) keep /search/sync returning the final SearchResponse JSON, but stream it as a StreamingResponse that flushes a whitespace/newline (" `/\n`") heartbeat every $\sim 10 \text{ s}$ while the

fan-out runs, then the JSON body at the end — mirrors the `: keepalive` idiom already used for the theme SSE at `routes.py:99`; OR - **steer non-browser callers to the streaming endpoint** — document `POST /search + GET /search/stream/{id}` as the canonical path for slow/tunnelled clients, and treat `/search/sync` as a `localhost/test` convenience only. Note: the operator's own `tunnel-keepalive` supervisor only keeps the *tunnel process* alive; it does not keep an *individual idle request* from being reset — so a per-request heartbeat (or SSE) is still required.

Fix B — cut the wall-clock (addresses the 67 s)

- Lower the default `PUBLIC_TRACKER_DEADLINE_SECONDS` (e.g. 60 → 20–25) so a slow tracker can't dominate a wave; the dashboard already shows a `deadline_hit` chip (`search.py:1187`) so truncation stays honest/visible.
- Optionally raise `MAX_CONCURRENT_TRACKERS` (5 → 8–10) to cut the wave count from ~6 to ~3–4 — bounded by §12.6 memory budget (each public tracker is a subprocess) and the existing event-loop-starvation guardrails noted at `search.py:617-623`.
- Capture before/after timing of a broad search (e.g. `ubuntu`) per RW-08 acceptance.

Fix C — regression guard (§11.4.85)

Add a chaos test for sync-over-slow-link: simulate a server that produces no bytes for > tunnel-idle-window and assert the heartbeat path keeps the socket alive (and/or assert the SSE path is unaffected). Add a deterministic latency test asserting the wave-count / deadline math so a future deadline-default regression is caught.

5. Root-cause hypothesis (one line)

The reset is an **SSH-tunnel idle-timeout** on `/search/sync`'s fully-blocking, zero-bytes-until-done response (server returns 200; the tunnel drops the idle socket mid-flight); the ~67 s wall-clock is the **29-provider fan-out** at 5-concurrent × 60 s-per-tracker (~6 waves). Fix = heartbeat/stream the sync path (or steer callers to SSE) **and** tune the per-tracker deadline / concurrency. The SSE path is structurally robust because its continuous frames + `: keepalive` keep the tunnel warm.