

scripts/hooks/check-brief-inputs.sh “ pre-dispatch input existence check

Revision: 1 **Last modified:** 2026-08-21T19:55:00Z **Purpose:** Â§11.4.201(6) / Â§11.4.252 fail-closed check that every input a task brief declares actually exists and is non-empty BEFORE a subagent is dispatched against it. **Last verified:** 2026-08-21

Overview

A task brief names its source materials. If one of those files is missing when the subagent starts, the agent does not stop “ it proceeds against what it can read and produces confident output grounded in less evidence than the brief promised. That is the Â§11.4.201(6) FALSE-NULL at the dispatch layer: absent input and empty input are indistinguishable from “œnothing to say here”.

MEASURED, the case this guard exists for: the forensic brief .superpowers/sdd/task-phase1a-brief.md still exists, but the scratchpad session directory holding its inputs is **entirely gone**. Run against that brief today, this guard separates the 6 genuinely-missing files (naming each) from the 2 that are present, and exits 1 “ reproducing exactly what should have blocked that dispatch.

Prerequisites

- A readable brief file, or an explicit --file list. No network, no runtime.

Usage

```
bash scripts/hooks/check-brief-inputs.sh --brief <path-to-brief>
bash scripts/hooks/check-brief-inputs.sh --file <p1> --file <p2>
bash scripts/hooks/check-brief-inputs.sh --self-test
```

Exit status: 0 when every declared input exists and is non-empty; 1 otherwise, with a per-path report naming each missing file and the remediation (“œrespawn the producer”).

Edge cases

- **Present but EMPTY is a failure, not a pass.** A zero-byte file is the clearest form of the false-null this guard exists to catch.

- **Path shorthand.** Briefs in this repo abbreviate long paths with `|`; the extractor handles that convention rather than silently skipping such lines – a skipped line would be a false-null of its own.
- **A brief with no `## Source materials` section** yields nothing to check; that is reported honestly rather than counted as a pass.

Internal behaviour

Fails CLOSED (Â§11.4.252): when an input's status cannot be resolved, the guard refuses rather than proceeding, and prints the resolved evidence for every refusal (Â§11.4.201(5)) so a false positive is diagnosable in one step.

Paired Â§1.1 mutation (verified): neutering the existence/non-empty checks makes the self-test FAIL and makes the real scan report "OK" all 8 present against the forensic brief – the exact fabricate-on-absent-evidence failure mode this guard prevents. Restored, sha256-identical, both green.

Known gap – stated rather than implied

This is **conductor-run tooling, not an automatic PreToolUse hook**. A brief's required-input list lives in free-form prose that the Agent tool's structured input does not expose, so an automatic prompt-scraping gate would itself be an unproven pattern-match (Â§11.4.201(7)(a)). Its seam is therefore the dispatch procedure – a habit backed by a runnable check – and that limitation is tracked in **BOB-162** rather than papered over.

Related

- `tests/hooks/test_check_brief_inputs.sh` – 15 real-invocation assertions
- `scripts/hooks/unattributed-commit-guard.sh` – sibling hygiene guard