

compute-badges

scripts/compute-badges.sh — §11.4.259 machine-derived README badge regenerator

Revision: 1 **Last modified:** 2026-08-18T22:29:59Z **Status:** active
Item: BOB-118 (§11.4.6 bluff audit — task-ad3205a9)

Overview

scripts/compute-badges.sh regenerates the “python tests” and “frontend tests” shields.io badges in README.md, plus the matching Contributing- section bullets and a corroborating ## Test counts table in docs/TESTING.md, from REAL numbers produced by a live invocation of the tools that actually count those tests — never a hand-typed or carried-forward number. It is the fix for BOB-118 and the standing mechanism that keeps the same class of bug from recurring.

Why this exists — BOB-118

docs/qa/task-ad3205a9-* (a §11.4.6 no-guessing bluff audit) found that README.md’s badge row read python tests-585 passing and frontend tests-182 passing, while a real pytest --collect-only over tests/ collected **5248** tests — roughly 9x more than the badge claimed. The cited “authoritative source”, docs/TESTING.md, contained **zero** occurrences of either number anywhere in the file — the badge could not be corroborated by any tracked document. This is exactly the class of defect §11.4.259 exists to make structurally impossible: every README badge **MUST** be machine-derived, never hand-picked.

What it computes

Badge / bullet	Real command	Notes
python tests badge (whole tree)	pytest --collect-only -q tests/ --import-mode=importlib	Parses pytest's own N tests collected summary line.
Contributing bullet "unit + e2e + contract"	sum of pytest --collect-only -q tests/ {unit,e2e,contract}/	Scoped to match the bullet's own stated suite scope — never the whole-tree number under a narrower label.
frontend tests badge	frontend/ node_modules/.bin/ vitest list --run	Exact per-test-case count. Falls back to a labelled grep proxy over it(/test(declarations when node_modules is absent (never silently presented as exact).
## Test counts table in docs/TESTING.md	all of the above, plus tests/ integration/	The corroborating source the badge audit found missing.

Two badges (challenges, pre-build invariants) were verified fresh in a prior session and are cross-checked (not rewritten) by this script — see the counts printed to stdout on every run.

Usage

```
# Compute + rewrite README.md and docs/TESTING.md in place
scripts/compute-badges.sh

# Compute only – print whether the badges are stale, write
# nothing.
# Exit 0 = in sync, exit 2 = stale (advisory; callers must not
# treat
# this as fatal per §11.4.234 always-unblocked).
scripts/compute-badges.sh --check

# Point at alternate files (used by the paired §1.1 mutation test)
scripts/compute-badges.sh --check --readme /tmp/fixture-README.md
--testing-md /tmp/fixture-TESTING.md
```

Wording: “collected” not “passing”

The regenerated badges say “**N collected**”, not “N passing”. Running the full 5248-test tree (many of which need live services, Docker/Podman, or hardware) end-to-end is outside what a badge-refresh script can honestly assert on every invocation. “Collected” is the claim this script can always verify in seconds with zero infrastructure dependencies: the test exists and pytest/vitest can find it. Asserting “passing” without having actually run every one of those tests in the same invocation would just be a differently-shaped version of the same bluff BOB-118 reports.

Edge cases (honest failure mode)

If the python interpreter cannot be resolved (mirrors `ci.sh`’s fallback chain: `$PYTHON → .venv/bin/python → python3.13 → python3.12 → python3`) or `pytest --collect-only` produces no parseable summary line, the badge is emitted as N/A (<reason>) in shields.io’s lightgrey color — never a fabricated number (§11.4.6). Same discipline for the frontend count when neither vitest nor any `*.spec.ts` file can be found.

Related scripts

- `scripts/pre_build_verification.sh` invariant CM-BADGE-FRESHNESS-CHECK calls `compute-badges.sh --check` as a non-blocking WARN (never FAIL, per §11.4.234) so a drifted badge surfaces on every pre-build sweep instead of silently aging for months.
- `docs/TESTING.md` — the corroborating authoritative source this script keeps in sync.

Last verified

2026-08-18, this session: `--check` correctly reported the pre-fix README as stale (5248 vs the stale 585, 371 vs the stale 182), then a clean run regenerated both files and a follow-up `--check` reported in-sync. See `docs/qa/task-bob118-fix/` for the captured before/after evidence.