

generate_markdown_exports

scripts/ generate_markdown_exports.sh — the §11.4.65 document-twin exporter

Revision: 2 **Last modified:** 2026-09-26T12:54:43Z **Status:** active

Overview

scripts/generate_markdown_exports.sh is the project's canonical §11.4.65 exporter. For every in-scope *.md source it produces the sibling document twins:

Twin	Produced by	Present when
.html	pandoc --standalone, else python-markdown	always (one leg or the other must exist)
.pdf	weasyprint rendering the .html	the weasyprint CLI is on PATH
.docx	pandoc -f markdown -t docx (direct from the .md, BOB-011)	pandoc is on PATH

It is the script the pre-build sync invariant, scripts/compute-badges.sh, and scripts/testing/update_readme_doc_links.sh all sit downstream of.

Prerequisites

Every leg is **optional and independently degradable**, and a leg that cannot run writes **nothing** — the script never emits a blank or placeholder file (§11.4.3 honest SKIP, §11.4.6 no fabricated output).

- pandoc — HTML + DOCX. The highest-fidelity path.
- weasyprint (the **CLI**, not the Python module) — PDF.
- a Python interpreter that can import markdown — the HTML fallback used only when pandoc is absent.
- If neither pandoc nor a markdown-capable interpreter is found, the script **exits 1** and names both the interpreters it probed and the remedy. It does not half-run.

All three install without root:

```
uv tool install weasyprint #
    weasyprint CLI
uv pip install --python .venv/bin/python markdown # python-
    markdown fallback
# pandoc: a user-local tarball unpacked under ~/.local/opt +
    symlinked into ~/.local/bin
```

Usage

```
# Full in-scope sweep: project root *.md + docs/**/*.md + scripts/
    **.md
bash scripts/generate_markdown_exports.sh

# Explicit scope – regenerate ONLY these sources' twins
bash scripts/generate_markdown_exports.sh CLAUDE.md docs/
    USER_MANUAL.md
```

Inputs

Input	Meaning
\$@ (optional)	Explicit .md paths. When present, only those files are converted and the discovery sweep is skipped entirely.
\$BOBA_EXPORT_PYTHON (optional)	Interpreter to use for the python-markdown leg. Absent → auto-resolved.

Outputs

.html, .pdf, and .docx siblings written next to each source .md. Counts of generated files are printed per twin type.

Exit codes

Code	Meaning
0	Completed (including “nothing was stale”).
1	No HTML converter available at all, or an explicitly-named argument was not an existing .md file.

Why the explicit-scope argument exists

The full sweep touches ~355 exports. A session that edits nine documents should resync exactly those nine twins: re-rendering 346 unrelated exports costs minutes of `weasyprint` time and needlessly churns the corpus that `scripts/pre_build/check_cm_export_charset_valid.sh` measures against its monotone ratchet baseline. Passing the changed files keeps the blast radius equal to the change.

An argument that is not an existing .md file is a **hard error**, never a silent skip — a typo’d path must not be able to masquerade as “nothing to do” (§11.4.201(6): a quiet zero from a blind instrument is indistinguishable from a clean corpus).

Idempotency and the staleness rule

Re-running is a no-op: a second run over unchanged sources rewrites nothing and leaves every twin byte-identical.

Staleness uses the same oracle as the pre-build gate (BOB-249)

The writer decides whether a twin is stale with `scripts/lib/export_staleness.sh` — the **same** content-history oracle CM-MARKDOWN-EXPORT-SYNC (pre-build invariant 16) uses, so writer and gate agree on what “stale” means:

Situation	Verdict
twin missing	stale
.md (or twin) locally edited / untracked	stale when the twin's mtime is older (mtime is meaningful here)
both committed and clean	stale only when the .md's last-touching commit is more recent than the twin's (git history, never mtime)
file outside any git work tree, or the lib missing next to the script	plain mtime fallback, announced on stderr when the lib is missing

Why not mtime: a fresh `git checkout` writes `x.docx` and `x.html` **before** `x.md` (path order), so on a scratch clone 226 of 456 `.md` files were strictly newer than their `.docx` and 42 than their `.html`. The old mtime-only

writer rewrote all 310 of those twins with zero content change. With the oracle the same fresh checkout rewrites **0**, while a real content edit of a `.md` still regenerates all three twins.

Byte-stable output (SOURCE_DATE_EPOCH)

pandoc stamps the wall-clock time into every `.docx` (`docProps/core.xml`), so each regeneration used to be a byte-different file. The script now exports `SOURCE_DATE_EPOCH` per file before rendering:

- the `.md`'s last-commit time (`git log -1 --format=%ct -- <md>`), so the stamp is a property of the source history, identical in every clone;
- 1785674948 (2026-08-02T12:49:08Z, the value `constitution/scripts/render/render-governance-twins.sh` pins) for a never-committed `.md`;
- a `SOURCE_DATE_EPOCH` the caller already exported wins.

Regenerating identical content twice now yields a byte-identical `.docx`. weasyprint 69.0 PDFs were already byte-stable (measured); HTML from `pandoc --standalone` carries no timestamp. **Limit:** a twin rendered while its `.md` is still uncommitted carries the *previous* commit's time; a forced re-render after that `.md` is committed carries the new one. Output is deterministic per (content, history), not per content alone.

The charset self-heal rule

Staleness is **not** history alone. An `.html` that exists and is mtime-fresh but declares **no charset** is treated as STALE and regenerated, because such a file is a pre-BOB-169 charset-less fragment that would otherwise persist forever and keep re-baking mojibake into any PDF rendered from it. PDF staleness additionally keys on the `.html` (the PDF is derived from it), so healing the HTML half cannot leave a corrupt PDF behind. An HTML regenerated in the **current** run always re-derives its PDF explicitly: the oracle caches the working-tree state once per run, so to it a just-rewritten HTML still looks clean.

Internal behaviour worth knowing

Capability probes test the artifact, not a prerequisite

The weasyprint probe runs `weasyprint --version` — the real entry point the render path uses. It previously also required `python3 -c "from weasyprint import HTML"` to succeed, which is a **prerequisite** probe, not an artifact probe: a weasyprint CLI installed into its own isolated environment (`uv tool install`, `pipx`, a `venv` on `PATH`) renders PDFs perfectly while being invisible to a system-interpreter import. On this host that made the script announce *“PDF support: not available”* and produce **zero PDFs while exiting 0** — a false negative of exactly the shape §11.4.201(11) names. Do not reintroduce an import-based probe for a CLI-invoked tool.

The python-markdown interpreter is resolved by trying the import

`PY_MD` is chosen by actually attempting `import markdown` against, in order: `$BOBA_EXPORT_PYTHON`, the repo’s `.venv/bin/python`, then `python3`. It is never assumed from a path’s existence (§11.4.6). Probing only system `python3` is insufficient here: this host’s `/usr/bin/python3` has no `pip` and is PEP-668 EXTERNALLY-MANAGED, so `markdown` cannot live there without `root` — while the repo `.venv` holds it with no privilege at all.

Encoding is pinned on both legs

The corpus carries Cyrillic (Боба) and §. The python-markdown leg opens and writes with an explicit `encoding='utf-8'` rather than inheriting the ambient locale, so a C/POSIX locale cannot mangle the read and leave the emitted `<meta charset="utf-8">` advertising a lie.

Verifying the output

```
# charset declared on every generated export (ratcheted gate)
```

```
bash scripts/pre_build/check_cm_export_charset_valid.sh
```

```
# the PDF text layer really preserves non-ASCII
```

```
bash tests/unit/test_export_pdf_charset_integrity.sh
```

```
# DOCX shape + idempotency
```

```
bash tests/unit/test_docx_export.sh
```

```
# checkout-order touch rewrites nothing; content edit / missing
    twin /
# charset fragment still regenerate; docx + pdf byte-stable
bash tests/unit/
    test_generate_markdown_exports_content_staleness.sh

# the shared staleness oracle (incl. .docx history)
bash tests/unit/test_export_staleness_oracle.sh
```

When checking a heading survived into the HTML, match against a **whitespace-normalized** stream, not lines: pandoc hard-wraps its HTML at ~72 columns, so a heading routinely spans a newline and a line-oriented `grep -F` returns a confident, wrong zero (\$11.4.201(7)(c)). Remember also that `&` is correctly emitted as `&`;

Related scripts

- `scripts/pre_build/check_cm_export_charset_valid.sh` — the gate over this script's output; reads a monotone-decreasing ratchet baseline.
- `scripts/regenerate-continuation-exports.sh` — deliberately narrow CONTINUATION-only regen; see its own guide for why it does not call this.
- `scripts/lib/export_staleness.sh` — the staleness oracle shared by this script and the pre-build gate.
- `tests/unit/test_export_pdf_charset_integrity.sh` — the paired §1.1 guard.

Last verified: 2026-09-26 — fresh scratch-clone checkout of `docs/ + README.*` (456 .md): 310 twins rewritten before the fix, 0 after; a control edit still regenerated its three twins. Content-staleness, oracle, DOCX, PDF charset-integrity and path-argument guards all PASS.