

scripts/ownership_repair.sh

Ownership Repair

Revision: 1 **Last modified:** 2026-08-21T15:00:00Z **Purpose:** Operator guide for the tool that brings pre-existing content back under the ownership of the person who started the system. **Last verified:** 2026-08-21

Overview

scripts/ownership_repair.sh walks every location declared in config/owned_paths.yaml and chowns every item that is **not** owned by the operator back to the operator's uid:gid.

It exists because of a measured state, not a hypothetical one: on 2026-08-21 the download root contained an item owned by uid 100999 and config/contained **51** of them a rootless-container sub-uid with no host account. The operator could not edit their own configuration, and config/boba.db (mode 600, owner unresolvable) could not be read by the operator at all, which made the backup procedure that docs/BOBA_DATABASE.md *mandates* impossible to perform.

The precondition (scripts/ownership_precondition.sh) **detects** that state. This script **fixes** it.

Four behaviours a reader will otherwise get wrong

These are the non-obvious parts. Each is a deliberate design decision with a failure mode behind it.

1. The completion marker is written ONLY on success so an interrupted run resumes rather than being skipped forever. logs/ownership/repair-marker.json is the **last** act of a fully successful pass. Writing it at *start* would let a single interruption mark the repair *done* permanently and silently skip everything it had not reached yet the run-once optimisation would defeat the repair it optimises. Interrupted marker absent the next run walks the whole scope again. This is also the operator's escape hatch from a long run: stopping it mid-repair is safe.

2. The marker carries a scope FINGERPRINT so changing config/owned_paths.yaml re-arms the repair. A marker keyed on mere existence would report *already done* about work that was never performed the moment a **new** path is declared: the newly-declared location would be silently never repaired. The marker instead stores the sha256 of the sorted declared scope (ownership_scope_fingerprint), and a normal run is a no-op only when the marker's fingerprint matches the scope as

it is *right now*. Add, remove, or edit an entry and the next run repairs the new scope.

3. `preserve_mode` guards MODE BITS, not ownership – such an entry is still chowned. `preserve_mode: true` on `config/boba.db` does **not** exempt the credential store from ownership repair; it is repaired like everything else. What `preserve_mode` protects is its **permission bits**. `chown(2)` performed by a non-root user clears `setuid/setgid`, so the script reads each item's original mode before the change and restores it afterwards. Bringing a credential store under the operator's ownership while widening who can read it would trade a usability defect for a security one – strictly worse than the defect being fixed (FR-015). Mode restoration also fires, `preserve_mode` or not, for any item that actually carried special bits (a 4-digit octal mode).

4. `preserve_mode` entries are processed FIRST – and that ordering is load-bearing. The real scope nests: `config/boba.db` (`preserve_mode: true`) lives inside `config/` (`preserve_mode: false`). Whichever entry reaches an item first is the one whose rules apply, and once repaired the item no longer matches the “wrongly-owned” filter, so the later entry never sees it. Ordering the strictest constraint first makes the strictest rule win, and yields exactly one change-record entry per altered item with no cross-entry deduplication pass.

Prerequisites

- bash 4+ and coreutils (find, chown, chmod, mktemp, sha256sum, date, tr, wc)
- python3 with **PyYAML** – used by the shared scope parser
- `scripts/lib/ownership.sh` – sourced, never re-implemented
- `config/owned_paths.yaml` – the declared scope (data-model **E1**)
- Write access to `logs/ownership/` under the project root (created on demand)
- **Optional:** podman or docker, only for the namespace fallback described under *Internal behaviour*
- **Optional:** nice and ionice, for the host-safety re-exec

Usage examples

Example 1 – the normal invocation

```
scripts/ownership_repair.sh
```

Repairs the declared scope. If a valid marker already exists for the current scope fingerprint, this is a no-op and prints `already complete` for this scope (marker: `logs/ownership/repair-marker.json`) – nothing to do.

Example 2 “ preview without changing anything

```
scripts/ownership_repair.sh --dry-run
```

Prints one would chown <uid>:<gid> (was <uid>:<gid>) <path> line per item that would change. Creates **nothing** “ not the state directory, not the change record, not the marker. A dry run that left a new directory behind would already have changed the tree it claimed only to report on (measured 2026-08-21: an earlier revision created logs/ownership/ under --dry-run, which the unit suite does not observe).

--dry-run also **bypasses the marker check**, so it previews the scope even when a valid marker says the repair is complete.

Example 3 “ re-walk despite a valid marker

```
scripts/ownership_repair.sh --force
```

Ignores the marker and walks the scope again. Idempotent: items already correct are never matched by the walk, so a forced run on a clean tree changes nothing and writes an empty change record.

Example 4 “ repair a different scope

```
scripts/ownership_repair.sh --scope /path/to/other_owned_paths.yaml
scripts/ownership_repair.sh --scope=/path/to/other_owned_paths.yaml # eq
```

Both forms are accepted; the value is exported as OWNED_PATHS_FILE for the shared library.

Example 5 “ combine preview with an alternate scope

```
scripts/ownership_repair.sh --dry-run --scope config/owned_paths.yaml
```

Example 6 “ suppress the namespace fallback

```
CONTAINER_RUNTIME= scripts/ownership_repair.sh --dry-run
```

An explicitly **empty** CONTAINER_RUNTIME means “no container runtime is available” and disables the <runtime> unshare fallback entirely. This is how gates and test suites run the script without touching a runtime.

Flags

Flag	Effect
--force	Ignore a valid completion marker and re-walk the declared scope.
--dry-run	Report what would change; change nothing, record nothing, create

Flag

--scope <path> / --scope=<path>

-h, --help

(anything else)

Effect

nothing. Also bypasses the marker check.

Override the declared scope file. A missing argument to the space-separated form is exit 2.

Print usage and exit 0.

Unrecognised: error + usage to stderr, exit 2.

Env vars

Variable

OWNED_PATHS_FILE

CONTAINER_RUNTIME

OWNERSHIP_REPAIR_RENIED

PYTHON_BIN

TMPDIR

scope placeholders

Meaning

Scope file path when --scope is absent. Default config/owned_paths.yaml under the project root.

Set-but-empty = no runtime available; the unshare fallback is disabled. **Unset** = detect podman, then docker. **Non-empty** = use that command verbatim. Set-but-empty is honoured rather than re-detected: probing behind a value the caller explicitly set would do work the caller said not to do (Â§11.4.201).

Internal. Set by the script's own host-safety re-exec so it happens exactly once; not intended to be set by hand.

Consulted first when the shared library chooses a PyYAML-capable interpreter.

Where the run's scratch directory is created (removed on exit).

`${VAR:-default}` inside a declared path is expanded against the live environment

QBITTORRENT_DATA_DIR resolves the host-specific download root at run time.

Exit codes

Code

0

Meaning

Code	Meaning
1	Every in-scope item is operator-owned " repaired now, or already correct. The marker is written. At least one item could not be repaired. Each is named individually. No marker is written, so the next run retries the whole declared scope.
2	Could not run: the scope is missing/unparseable, the fingerprint cannot be computed, the state directory cannot be created, or the change record cannot be opened for append. Nothing was touched.

Two signal exits are also possible and are deliberate: 130 on SIGINT and 143 on SIGTERM, both printing no marker written; the next run resumes.

A --dry-run exits with the same RC it would have produced for real " so a preview that names a declared path which cannot be repaired exits 1, not 0. Exiting 0 there would make --dry-run useless as a pre-flight check.

Edge cases

- **Absent + optional: true** " logged as absent, declared optional " skipped, and recorded with outcome: "skipped" and every ownership field null. config/boba.db is exactly this shape before first boot.
- **Absent + not optional** " FAILED " declared path does not exist and is not optional, recorded with outcome: "failed", exit 1. Per E1 an absent non-optional path is an error, not a skip.
- **A directory the walk cannot fully read** " find reports both what it could read *and* a non-zero status; both halves are honoured. The readable part is repaired, the unreadable part is a named failure line, and the run exits 1. It is never a silent gap.
- **A symlink inside the scope** " chown -h acts on the link itself, never on its target. Without -h, a symlink pointing anywhere on the filesystem would let the repair mutate a file **outside** the declared scope; FR-005 requires out-of-scope reach to be impossible by construction, not merely unintended.
- **Content at a rootless sub-uid (e.g. 100999)** " an unprivileged chown cannot touch it. The script falls back to <runtime> unshare chown -h 0:0, which re-enters the same user namespace where the host operator's uid maps to uid 0. The fallback runs **only after** a plain chown has already failed, so a run that does not need it never invokes the runtime at all.
- **Neither chown nor the fallback works for an item** " the item is named on stderr, a corrective outcome: "failed" record is appended for it, and the run exits 1. It is never reported as success (FR-006).

- **A batch fails as a whole** â†’ the script degrades to per-path repair for that batch, because a batch failure does not say *which* path failed and FR-006 requires items to be listed individually.
- **Paths containing tabs or newlines** â†’ survive intact: find emits three fixed numeric fields first and the path last, NUL-terminated, and read assigns the remainder verbatim to the last variable.
- **A crash between record and mutation** â†’ the record already exists. Records are flushed to disk *before* the chown (FR-004b), so the recorded outcome: "changed" is the **intent**; if the mutation then fails, a corrective failed entry is appended for that path.
- **Concurrent downloads while repairing** â†’ out of scope by construction: the repair blocks every download-writing service (FR-004g). Recorded in the contract so a later reader sees it was decided, not overlooked.
- **Host load** â†’ the script re-execs itself under `nice -n 19 ionice -c 3` on first entry (guarded by `OWNERSHIP_REPAIR_RENICKED` so it happens once, skipped entirely if either tool is absent). `exec` keeps the same pid, so a supervisor that captured the pid can still signal it.
- **Wiring, measured 2026-08-21T15:10Z:** `start.sh` calls this script from `run_ownership_gate()`, immediately after the ownership precondition passes, on the normal start path and on `--recreate`. A non-zero exit from the repair refuses the start; a missing script file also refuses the start (FR-004d â€” the repair blocks before any download-writing service accepts work). It is invoked under `nice -n 19 ionice -c 3` when both tools are present, which is belt and braces: this script also re-execs itself that way. The systemd unit `scripts/systemd/user/boba-stack.service` deliberately declares no `ExecStartPre` of its own â€” it runs `./start.sh --no-build` and inherits the single gate, so the two lifecycle paths cannot disagree about whether the repair ran. This wiring landed while this document was being written â€” it was measured directly in the working tree, not taken from the contract (Â§11.4.6). **Note the consequence:** because the marker makes a completed run a no-op, the per-start cost after the first successful pass is one scope read and one fingerprint comparison.

Internal behaviour

Output artifacts

Both live in `logs/ownership/` at the **project root** â€” operator-readable, on the host (never only inside a container, an E3 rule), and inside the already gitignored `logs/` tree so an operational log never becomes a Â§11.4.30 versioned-artifact violation. `docs/qa/` was deliberately not used: that tree is curated QA evidence (Â§11.4.83) and is the wrong home for an operational log.

`logs/ownership/repair-changes.ndjson` â€” the **E3 change record**. Append-only, one JSON object per line, greppable by path:

```
{"path":"/â€¦/config/some.file","previous_uid":100999,"previous_gid":10099
```

outcome is one of changed | skipped | failed. For a declared path that does not exist there is no previous ownership state, so every ownership field is JSON null rather than a sentinel “-1” -1 is a plausible-looking uid and a later reader could not tell a sentinel from a real value.

This record is the operator’s recovery trail for a repair they did not approve item by item. It holds **paths, uids, gids and modes only** “never file contents, never credential values (§11.4.10). `boba.db` appears as a path; nothing of its contents ever does.

logs/ownership/repair-marker.json “the E2 completion marker:

```
{
  "completed_at": "2026-08-21T14:00:00Z",
  "scope_fingerprint": "<sha256 of the sorted declared scope>",
  "items_changed": 52
}
```

Written via `mktemp + mv -f`, so a half-written marker can never be read by the next start. Validity is a literal match on the current fingerprint inside the file “a stale fingerprint is treated exactly as an absent marker.

The pass

1. **Re-exec** under `nice/ionice` (once).
2. **Parse arguments**, resolve the operator `uid:gid` (`id -u / id -g` “the identity running the script, deliberately, because the feature’s definition of “correct owner” is *whoever started the system*, not a configured constant that could drift).
3. **Read the scope**. Unreadable “exit 2: reporting an unreadable scope as an empty scope would be the §11.4.201(6) false-null, where a blind instrument and a clean tree return the same quiet zero. Compute the fingerprint; failure “exit 2.
4. **Order entries**, preserve `_mode` first (see *Overview*, point 4).
5. **Marker check** “skipped under `--force` or `--dry-run`.
6. **Open the change record** for append (skipped entirely under `--dry-run`).
7. **Per entry**: find the location for items whose uid **or** gid differs from the operator’s, emitting `%U\t%G\t%m\t%p\0`. Non-recursive entries and regular files get `-maxdepth 0`. **This filter is the scope fence**: only items under a declared path are ever named, so an out-of-scope item cannot be reached even by accident (FR-005).
8. **Batch of 256**: append every record to the record file, **then** `chown` the batch, **then** restore mode bits where required. 256 is not a tuning knob “it is the width of the window in which records exist for items not yet mutated, traded against one fork per batch instead of one per item.
9. **Progress** is emitted as *items processed against items discovered*, throttled to once per second (FR-004e). A spinner or fixed-step estimate would not satisfy the requirement: it exists so a long run on a large library is distinguishable from a hang.
10. **Verdict**. All good “write the marker and exit 0. Any failure “**no marker**, list the failures, exit 1.

Related scripts

- [ownership_precondition.md](#) “ scripts/ownership_precondition.sh, the startup check whose failure message points here. The precondition **detects**; this script **fixes**.
- scripts/lib/ownership.sh “ the shared library both scripts source (ownership_scope_entries, ownership_operator_uid/_gid, ownership_scope_fingerprint, probe_location). Sourced rather than copied: a second copy would be the near-identical fork Â§11.4.251 forbids and would drift from the precondition and the gate that read the same scope.
- scripts/pre_build/check_cm_ownership_invariants.sh “ invariant 33 (CM-OWNERSHIP-INVARIANTS), the standing regression gate that keeps the ownership routes from being reverted.
- config/owned_paths.yaml “ the declared scope (E1). **Editing it changes the fingerprint and re-arms this repair.**
- tests/unit/test_ownership_repair.sh “ the paired contract suite (golden-bad / golden-good / out-of-scope negative control / interrupt / preserve-mode cases).
- tests/ownership/test_container_writes_owned_files.py “ the Â§11.4.115 RED that captured the original defect at uid 100999.
- [../guides/file-ownership.md](#) “ the operator-facing narrative for the whole ownership feature.
- [../BOBA_DATABASE.md](#) “ Â§ 3, the backup procedure that this repair made performable again.

Cross-references

- specs/002-user-owned-downloads/contracts/repair-cli.md “ the contract this script implements (FR-004, FR-004a“g, FR-005, FR-006, FR-015).
- specs/002-user-owned-downloads/data-model.md “ **E1** scope, **E2** marker, **E3** change record.
- Â§11.4.201 “ assert the real condition; an unreadable scope is not an empty one, and a guard that refuses a healthy tree is as broken as one that passes a broken one.
- Â§11.4.10 “ the change record carries no credential values.
- Â§11.4.30 “ the record lives under the already-gitignored logs/tree.
- Â§11.4.251 “ why the scope/probe logic is sourced from one library rather than duplicated across the three consumers.
- **Principle XIII / Â§12.6** “ the host runs mission-critical work, which is why the script re-execs itself under nice -n 19 ionice -c 3.

Anti-bluff & Â§11.4 discipline

- **Records precede mutation.** A record written *after* the chown is lost exactly when it matters “ a crash mid-repair.

- **Failure is never rounded up to success.** Items that could not be repaired are listed individually, recorded with outcome: "failed", and the run exits 1 with no marker (FR-006).
- **The marker is earned, not assumed.** It is written only after a fully successful pass, so "œalready done" can never be claimed about work that was not performed " including work that a *scope change* newly declared.
- **Permission bits are never widened.** `preserve_mode` and special-bit items have their original mode restored after `chown`.
- **The scope fence is structural.** `find` is rooted at declared paths and `chown -h` refuses to follow symlinks out, so out-of-scope reach is impossible by construction rather than by intention.
- **11.4.263:** this script signals no processes " there is no `kill`, `pkill`, or `killpg` anywhere in it, so the `pgid <= 1` broadcast-kill hazard cannot arise.

Last verified

2026-08-21 " every flag, exit code, environment variable, record field, and edge case in this document was read from `scripts/ownership_repair.sh` and `scripts/lib/ownership.sh`, not from the script's usage text alone. The automatic first-start invocation is described from the call actually present in the working tree at 2026-08-21T15:10Z (`start.sh` 'run_ownership_gate()), not from the contract's word for it.