

boba user@1000 out-of-scope watchdog (BOB-116 / BOB-120 / BOB-123)

Revision: 2 **Last modified:** 2026-08-19T02:30:00Z **Class:** operator-run system-slice service “captures forensics on user@1000 SIGKILL”
Anchor: §11.4.4 four-layer, §11.4.108 runtime-signature, §11.4.115 RED-first, §11.4.6 anti-guessing, §11.4.10 credentials, §11.4.128 recording, §11.4.201 guard-real-condition, §11.4.161 rootless (no; system-slice by intent), §6.U (operator runs su, agent never invokes it).

Purpose

4 forced-logout incidents (BOB-116 / BOB-120 / BOB-123 + the 2026-07-07 one) share this mechanism: user@1000.service is SIGKILLED → cascade kill of every child process → GDM greeter respawns → operator perceives sudden logout.

Every in-scope monitor (user systemd timer, user cron, rootless podman container) DIES IN THE SAME CASCADE and cannot preserve forensics “it is architecturally incapable of observing the kill of the scope it runs in.

This watchdog is the ONLY mechanism that can attribute the SIGKILL because it lives in system.slice (a peer of user.slice), so it survives the kill of user@1000.service and captures the surrounding journal + cgroup + audit + process-tree state DURABLY.

What gets installed

Path	Purpose
/usr/local/bin/boba-user1000-watchdog	The monitor script (bash)
/etc/systemd/system/boba-user1000-watchdog.service	Systemd unit “Slice=system.slice, runs as root, Restart=always
/var/log/boba-watchdog/YYYY-MM-DDTHH-MM-SSZ/	Per-incident evidence dir, chmod 700

How it works

1. On boot, systemd starts the service in system.slice as root.
2. The script runs a pre-flight (§11.4.201): verifies uid=0 AND its own cgroup is (a) NOT under /user.slice/ AND (b) POSITIVELY under /system.slice/ (fail-closed rejection of cgroup-namespaced contexts

- where the negative check alone would pass while actually being inside user.slice). REFUSES to run in any other case.
- NOTE: because the preflight requires the system.slice cgroup, a manual root-terminal debug invocation (`su -c '/usr/local/bin/boba-user1000-watchdog'` from a tty) will refuse " for that, use `su -c 'systemd-run --slice=system.slice --pty /usr/local/bin/boba-user1000-watchdog'` which drops the shell into system.slice for the duration (this host has no sudo, only su).
 - Enters a `journalctl -f -u user@1000.service` loop.
 - On matching the prefix `<TARGET_UNIT>`: Main process exited AND either `status=9` OR `code=killed`, checks the `CAPTURE_COOLDOWN_SEC` dedup gate; if outside cooldown, triggers `capture_forensics` immediately.
 - Captures: journal window (`PRE_KILL_WINDOW_SEC` before + `POST_KILL_WINDOW_SEC` after), cgroup state (memory + PSI + pids), full process tree, kernel audit log (if Path-1 audit rules installed), uptime + last, `loginctl show-user`.
 - Rotates evidence dirs to keep the last `RETAIN_LAST_N` (default 20).
 - Keeps tailing " captures every subsequent incident too.

Tunables (environment)

- `BOBA_WATCHDOG_TARGET` " target unit to monitor (default `user@1000.service`; set to a scratch unit for a 11.4.108 layer-4 install-time drill without touching production)
- `BOBA_WATCHDOG_ROOT` " evidence root (default `/var/log/boba-watchdog`)
- `BOBA_WATCHDOG_RETAIN` " retain last N incident dirs (default 20)
- `BOBA_WATCHDOG_PRE_SEC` " pre-kill journal window (default 60)
- `BOBA_WATCHDOG_POST_SEC` " post-kill journal window (default 15)
- `BOBA_WATCHDOG_COOLDOWN_SEC` " dedup window for cascade triggers (default 60; I1 fix " cascade lines within 1s must not evict real incident dirs)

Set via drop-in `/etc/systemd/system/boba-user1000-watchdog.service.d/env.conf` with `[Service] + Environment="BOBA_WATCHDOG_ROOT=/path"`.

! **BOBA_WATCHDOG_ROOT misconfiguration is destructive:** rotation `rm -rf` targets ANY dir under the configured root. Never point it at `/`, `/home`, etc.

Install / uninstall

Per 6.U + 11.4.109, no agent tool invokes su/sudo. Install and uninstall scripts print `su -c 'â€|'` commands you copy/paste at your terminal:

```
bash scripts/system-slice-watchdog/install.sh    # prints commands to run
bash scripts/system-slice-watchdog/uninstall.sh  # prints commands to run
```

Anti-bluff verification (Â§11.4.108 + Â§11.4.115)

Run the challenge WITHOUT root â€” verifies structural correctness:

```
bash challenges/scripts/user1000-watchdog-challenge.sh # GREEN
POLARITY_MODE=red bash challenges/scripts/user1000-watchdog-challenge.sh
```

The challenge asserts:

1. All source files exist + parseable (bash -n, systemd unit shape)
2. All installer scripts are executable
3. **Load-bearing**: service declares Slice=system.slice (not user.slice)
4. Watchdog pre-flight refuses to run in user.slice
5. Evidence root is durable (/var/log, not tmpfs)
6. Journal trigger matches the exact BOB-116/BOB-120/BOB-123 signature
7. Rotation implemented (bounded disk)
8. Resource limits set (MemoryMax, CPUQuota, TasksMax)

After an incident

The forensics are in /var/log/boba-watchdog/<timestamp>/:

```
state.log          â€” uptime, last, loginctl, meminfo, PSI, cgroup snapshots
journal-pre.log    â€” journal for PRE_KILL_WINDOW_SEC before kill
journal-post.log   â€” journal captured POST_KILL_WINDOW_SEC after kill
proctree.log       â€” full ps auxf + top-20-by-RSS
audit.log          â€” kernel audit logs (requires Path-1 sudo audit rules)
```

Compose with Path 1 (kernel audit rules)

For **initiator attribution** (PID + UID + cmdline of the SIGKILL sender), run the audit-rule install from docs/incidents/2026-08-19-sudo-audit-rules-for-operator.md BEFORE the next incident. The watchdog captures the audit log entries; without the rules, audit.log will say â€œ(ausearch failed â€” audit rules may not be installed)â€.

Honest boundaries (Â§11.4.6)

- The watchdog DIAGNOSES; it does NOT prevent the kill.
- **M3 fix**: â€œsurvives the killâ€ is DESIGN INTENT proven at Â§11.4.108 layer-1 (source declares system.slice) + layer-2 (post-install systemctl show boba-user1000-watchdog.service -p Slice returns system.slice). Layers 3-4 (survives an actual kill event; captures durable evidence during real incident) are **PENDING** operator install + next incident. A cross-referable positive fact: GDM and auditd both survived all 4 forced-logout incidents on this host and both live in

- system.slice “ this is external evidence of the category-survival, not proof of THIS unit’s per-incident behavior.
- The watchdog CANNOT observe events pre-kill that happen outside user@1000.service’s journal “ e.g. a mutter/gnome-shell segfault would be in the user’s journal which may already be flushed by the time the watchdog captures.
 - The watchdog runs as root (needed to survive user@1000 kill + read all journals); it is bounded by MemoryMax=200M and CPUQuota=15% per §12.6.
 - **I3 fix “ credential leak honesty (§11.4.10):** evidence dirs contain ps auxf and ps -eo cmd output. Those capture **full process cmdlines**, which on this host CAN contain credentials (curl -u, env-prefixed invocations, some tracker plugins). Mitigations: (a) evidence dirs are root:root 0700, (b) the watchdog itself logs ONLY operational lines to journal (trigger lines, filenames, preflight status) “ never evidence contents, and no secret is passed through log(), (c) never commit raw captures unredacted per §11.4.128(4). If you export a capture for sharing, redact cmdline fields first.
 - **M4 fix “ KillUserProcesses claim was speculative:** prior text said “systemd’s baked-in default is yes“. Distro-patched behavior varies; the authoritative source is busctl get-property org.freedesktop.login1 /org/freedesktop/login1 org.freedesktop.login1.Manager KillUserProcesses at runtime “ check it, don’t assume.
 - **Restart-gap-of-blindness:** Restart=always RestartSec=5s means there is a ~5 second window after any watchdog exit during which no capture will happen. If an incident hits inside that window, evidence will be sparse.